

mEm Inc

3ND EDITION

Lua

Programming

**The Ultimate Beginner's Guide to
Learn Lua Step by Step**

CLAUDIA ALVES & ALEXANDER ARONOWITZ



Lua Programming

The Ultimate Beginner's Guide to Learn Lua Step by Step

Third edition
2021

By
Claudia Alves
& Alexander Aronowitz

TABLE OF CONTENTS

Introduction

The audience

About the third edition

Other resources

Certain Typographical Conventions

Running examples

Acknowledgments

PART I

Language

Chapter 1. Getting Started

1.1. Blocks

1.2. Some lexical conventions

1.3. Global variables

1.4. Separate interpreter

Exercises

Chapter 2. Types and values

2.1. Nil

2.2. Boolean

2.3. Numbers

2.4. Strings

Literals

Long lines

Type conversions

2.5. Tables

2.6. Functions

2.7. userdata and threads

Exercises

Chapter 3. Expressions

3.1. Arithmetic Operators

3.2. Comparison operators

3.3. Logical operators

3.4. Concatenation

3.5. Length operator

3.6. Operator Priorities

3.7. Table constructors

Exercises

Chapter 4. Operators

4.1. Assignment operators

4.2. Local variables and blocks

4.3. Control structures

if then else

while

repeat

Numeric for statement

General for statement

4.4. break, return and goto

Exercises

Chapter 5. Functions

5.1. Multiple Results

5.2. Functions with a variable number of arguments

5.3. Named arguments

Exercises

Chapter 6. More about functions

6.1. Closures

6.2. Non-global functions

6.3. Optimization of tail calls

Exercises

Chapter 7. Iterators and the generic for

7.1. Iterators and Closures

7.2. Generic for semantics

7.3. Stateless iterators

7.4. Complex State Iterators

7.5. True iterators

Exercises

Chapter 8. Compilation, Execution, and Errors

8.1. Compilation

8.2. Precompiled Code

8.3. C code

8.4. Errors

8.5. Error and Exception Handling

8.6. Error messages and call stack

Exercises

Chapter 9. Coroutines

9.1. Basics of coroutines

9.2. Channels and Filters

9.3. Coroutines as iterators

9.4. Non-displacing multi-threading

Exercises

Chapter 10. Completed Examples

10.1. The problem of eight queens

10.2. The most common words

10.3. Markov chain

Exercises

PART II

Tables and Objects

Chapter 11. Data Structures

11.1. Arrays

11.2. Matrices and multidimensional arrays

11.3. Linked Lists

11.4. Queues and double queues

11.5. Sets and Sets

11.6. Line buffers

11.7. Counts

Exercises

Chapter 12. Data Files and Persistence

12.1. Data files

12.2. Serialization

Saving tables without loops

Saving tables with loops

Exercises

Chapter 13. Metatables and Metamethods

13.1. Arithmetic metamethods

13.2. Comparison Methods

13.3. Library Metamethods

13.4. Methods for accessing the table

__Index metamethod

__Newindex metamethod

Default tables

Tracking access to a table

Read-only tables

Exercises

Chapter 14. Environment

14.1. Global variables with dynamic names

14.2. Descriptions of global variables

14.3. Non-global environments

14.4. Using _ENV

14.5. _ENV and load

Exercises

Chapter 15. Modules and Packages

15.1. Require function

Renaming a module

Search along the path

File crawlers

15.2. The Standard Approach for Writing Lua Modules

15.3. Using environments

15.4. Submodules and Packages

Exercises

Chapter 16. Object Oriented programming

16.1. Classes

16.2. Inheritance
16.3. Multiple inheritance
16.4. Hiding
16.5. Single Method Approach
Exercises

Chapter 17. Weak Tables and Finalizers

17.1. Weak tables
17.2. Functions with caching
17.3. Object Attributes
17.4. Again tables with default values
17.5. Ephemeral tables
17.6. Finalizers

Exercises

PART III

Standard Libraries

Chapter 18. Math Library

Exercises

Chapter 19. Library for Bitwise Operations

Exercises

Chapter 20. Library for Working with Tables

20.1. Insert and remove functions
20.2. Sorting
20.3. Concatenation

Exercises

Chapter 21. Library for working with strings

21.1. Basic functions for working with strings
21.2. Functions for working with templates

String.find function

String.match function
String.gsub function
String.gmatch function

21.3. Templates

21.4. Grips

21.5. Substitutions

URL encoding
Replacing tabs

21.6. Tricky tricks

21.7. Unicode

Exercises

Chapter 22. Library I / O

22.1. Simple I / O Model

22.2. Full I / O model

A small trick to increase performance
Binaries

22.3. Other operations with files

Exercises

Chapter 23. Library of operating room functions systems

23.1. Date and time

23.2. Other system calls

Exercises

Chapter 24. Debug Library

24.1. Accessibility (introspection)

Accessing Local Variables
Access to non-local variables
Access to other coroutines

24.2. Hooks

24.3. Profiling

Exercises

PART IV

WITH API

Chapter 25. Overview of the C API

25.1. First example

25.2. Stack

Putting items on the stack

Referring to items

Other stack operations

25.3. Error handling in the C API

Handling errors in the application code

Error handling in the library code

Exercises

Chapter 26. Extending Your Application

26.1. Basics

26.2. Working with tables

26.3. Lua function calls

26.4. Generalized function call

Exercises

Chapter 27. Calling C from Lua

27.1. Functions on C

27.2. Continuations

27.3. C modules

Exercises

Chapter 28. Techniques for writing functions in C

28.1. Working with arrays

28.2. Working with strings

28.3. Saving state in functions on C

Register

Function related values

Function related values used

several functions

Exercises

Chapter 29. User-Defined Types in C

29.1. User data (userdata)

29.2. Metatables

29.3. Object Oriented Access

29.4. Access as an ordinary array

29.5. Light objects of type userdata (light userdata)

Exercises

Chapter 30. Resource Management

30.1. Directory iterator

30.2. XML Parser

Exercises

Chapter 31. Threads and States

31.1. Numerous threads

31.2. Lua States

Exercises

Chapter 32. Memory Management

32.1. Memory allocation function

32.2. Garbage collector

Garbage collector API

Exercises

INTRODUCTION

When Waldemar, Louis, and I started developing Lua in 1993, we could hardly imagine that Lua would spread like that. On- started as home language for two specific projects, now Lua is widely used in all areas one can get benefits from simple, extensible, portable and efficient scripting language such as embedded systems, mobile devices swarms and, of course, games. We designed Lua from the beginning to integrate with software software written in C / C ++ and other common strange languages. There are many benefits to this integration. Lua is a tiny and simple language, partly because it doesn't tries to do what C is already good at, such as speed, low-level operations and interaction with third-party programs parties. Lua relies on C for these tasks. Lua offers something for which C is not good enough: sufficient distance from the hardware go support, dynamic structures, no redundancy and ease of testing and debugging. For these purposes, Lua has safe environment, automatic memory management and good possibilities for working with strings and other types resizable data. Some of Lua's strength comes from its libraries. And this is no coincidence. In the end Finally, one of the main strengths of Lua is extensibility. Many language features contribute to this. Dynamic typing tion provides a large degree of polymorphism. Automatic memory management simplifies interfaces because there is no need to the ability to decide who is responsible for allocating and freeing memory or how to handle overflows. Higher-order functions and anonymous functions allow a high degree of parameterization, making functions more versatile.

More than an extensible language, Lua is a “ *glue vayushim* ”(*glue*) language . Lua supports a component-based approach to software development when we create an application by gluing together existing high-level components. These components are written in a compiled language with a static typing such as C / C ++; Lua is the glue we use use to arrange and connect these components. Usually components (or objects) represent more specific low- tier entities (such as widgets and data structures), which which hardly change during the development of the program and which take up the bulk of the final program execution time. Lua gives the final form to the application, which is most likely changes a lot during the life of a given software product. However, unlike other "glue" technologies, Lua is is a complete programming language. Therefore we can use Lua not only to "glue" components, but also to adaptation and customization of these components, as well as to create a floor new components. Of course, Lua isn't the only scripting

language. Exists other languages you can use for roughly the same goals. Nevertheless, Lua provides a whole set of possibilities, which make him the best choice for many tasks and gives him your unique profile:

- *Extensibility* . Lua's extensibility is so great that many consider Lua not as a language, but as a set for DSL structures (domain-specific language, a language created for specific area, application). We developed Lua from the beginning so that it is extensible both through code on Lua as well as through C code. How the Lua proof implements most of its basic functionality via external library. Interaction with C / C ++ really simple and Lua has been successfully integrated with many others languages such as Fortran, Java, Smalltalk, Ada, C #, and even with scripting languages like Perl and Python.
- *Simplicity* . Lua is a simple and small language. It is founded on a small number of concepts. This simplicity makes it easier to study nie. Lua contributes to its very small size. Floor- distribution kit (source code, manual, binaries ly for some platforms) is quietly placed on one floppy disk.
- *Efficiency* . Lua has a very efficient implementation tions. Independent tests show that Lua is one of the most most fast languages among scripting languages.
- *Portability* . When we talk about portability, we say Rome on running Lua on all platforms you are only talking about heard: all versions of Unix and Windows, PlayStation, Xbox, Mac OS X and iOS, Android, Kindle Fire, NOOK, Haiku, QUALCOMM Brew, big servers from IBM, RISC OS, Symbian OS, process Sora Rabbit, Raspberry Pi, Arduino and more. Original the code for each of these platforms is almost the same. Lua does not use conditional compilation to adapt its code for different machines, instead it keeps standard ANSI (ISO) C. Thus, you usually do not you need to adapt it to the new environment: if you have a com- piler with ANSI C, then you just need to compile Lua.

The audience

Lua users generally fall into one of three broad groups: those who use Lua already built into the application, those who use use Lua separately from any application (standalone), and those who use Lua and C together. Many people use Lua built into some application, like Adobe Lightroom, Nmap, or World of Warcraft. These pri- applications use the Lua-C API to register new functions,

creating new types and changing the behavior of some operations language by configuring Lua for its area. Often users are whom applications don't even know that Lua is an independent language, adapted for the given area. For example, many developers Lightroom plugins are unaware of other uses of this language; Nmap users generally view Lua as scripting Nmap language; World of Warcraft players can view Lua as a language exclusively for this game.

Lua is also useful as just an independent language, not only for word processing and one-time small programs, but also for a variety of medium to large sized projects. For the In this way, the main functionality of Lua comes from its libraries. Standard libraries, for example, provide basic new functionality for working with templates and other functions for working with strings. As Lua improves its support libraries, a large number of external packages appeared. Lua Rocks, a system for building and managing modules for Lua, now has over 150 packages. Finally, there are programmers who use Lua as a library flow for C. Such people write more in C than in Lua, although they need requires a good understanding of Lua to create interfaces that are simple, easy to use and well integrated tied with the tongue.

This book can be useful to all of these people. The first part covers the tongue itself, showing how the whole its potential. We focus on various language constructs and use numerous examples and exercises to show know how to use them for practical tasks. Some chap- this part covers basic concepts such as managing structures, while the rest of the chapters cover more advanced tricky topics like iterators and coroutines.

The second part is completely devoted to tables, the only structure data tour in Lua. The chapters in this part discuss data structures, persistence, packages and object-oriented programming. It is there that we will show the full power of language. The third part introduces the standard libraries. This part especially useful for those who use Lua on their own language, although many applications include part or all standard libraries. In this part, each library is dedicated separate chapter: math library, bitwise library, library for working with tables, library for working with strings mi, I / O library, operating system library and debug library. Finally, the last part of the book covers the API between Lua and C.

This section is markedly different from the rest of the book. In this part we will be programming in C, not Lua. For some, this part may be uninteresting, but for someone - on the contrary, the most useful part of the book.

Running examples

You will need a Lua interpreter to run the examples from this books. Ideally, you should use Lua 5.2, however most in the examples will work on Lua 5.1 without any changes. The Lua site (<http://www.lua.org>) stores all the source code for interpreter. If you have a C compiler and know how compile the C code on your computer, you better ask try to install Lua from source; it's really easy. The *Lua Binaries* site (look for luabinaries) already offers compiled native interpreters for all major platforms. If you using Linux or another UNIX-like system, you can check the repository of your distribution; many distributions already offer ready-made packages with Lua. For Windows, a good choice is rum is *Lua for Windows* (look for luaforwindows), which is a convenient set for working with Lua. It includes interpreting torus, integrated editor and many libraries. If you are using Lua embedded in an application like WoW or Nmap, then you may need a manual for this. placement (or the help of a "local guru") in order to understand, how to run your programs. Nevertheless, Lua remains all the same the same language; most of the examples we will look at in this book are applicable regardless of how you use Lua. But I recommend that you start learning Lua with an interpreter to run you examples.

Part I

I am the language

CHAPTER 1

Begin

Continuing the tradition, our first Lua program simply printed em “Hello World” :
`print (“Hello World”)` If you are using a separate Lua interpreter, then all you need is you need to run your first program - this is to run the interpreter tator - usually called `lua` or `lua5.2` - with the name of the text file containing your program. If you saved the above the above program in `hello.lua` file , then you should run following command:

```
% lua hello.lua
```

As a more complex example, our next program defines There is a function for calculating the factorial of a given number, asking gives the user a number and prints its factorial:

- defines a factorial function

```
function fact (n)
if n == 0 then
return 1
else
return n * fact (n-1)
end
end
print (“enter a number:”)
a = io.read (“* n”) - reads a number
print (fact (a))
```

1.1. Blocks

Every piece of code that Lua executes, such as a file or from- a smart string in interactive mode is called *a* chunk. A block is simply a sequence of commands (or statements). Lua does not need a separator between consecutive operators, but you can use semicolon if you like. I personally use I use a semicolon only to separate statements written in one line. Line splitting does not play any role in the syntax system Lua; so, the following four blocks are valid and equivalent:

```
a = 1
b = a * 2
a = 1;
b = a * 2;
```

```
a = 1; b = a * 2
a = 1 b = a * 2 - ugly, but valid
```

A block can consist of just one statement, as in the example "Hello World", or consist of a set of operators and definitions functions (which are actually just assignments, as we will see later), as in the factorial example. Block can be as great as you want. Since Lua is also used as language for describing data, blocks of several megabytes are not a rarity. The Lua interpreter does not have any problems with bot with large blocks.

Instead of writing your programs to a file, you can run the interpreter interactively. If you run- those lua without arguments, then you will see its prompt for input:

```
% lua
Lua 5.2 Copyright (C) 1994-2012 Lua.org, PUC-Rio
>
```

Accordingly, each command that you enter (like, for example measure, print "Hello World") is executed immediately after how you enter it. To exit the interpreter, just type end-of-file character (ctrl-D in UNI, ctrl-Z in Windows) or call the `exit` function from the operating system library - you need type `os.exit ()` .

In interactive mode, Lua usually interprets each the line you enter as a separate block. However, if he is detects that the line is not a complete block, then it waits continue typing until a complete block is obtained. This way you can enter multi-line definitions, so as a factorial function , directly interactively. One- but it is usually more convenient to place such definitions to a file and then call Lua to execute that file. You can use the `-i` option to force Lua switch to interactive mode after executing the given block ka:

```
% lua -i prog
```

A command like this will execute the block in the `prog` file and then go into interactive mode. This is especially useful for debugging and manual th testing. At the end of this chapter, we will look at other options. command line for the Lua interpreter. Another way to trigger blocks is the `dofile` function , which paradise executes the file immediately. For example, let's say you have `lib1.lua` file with the following code:

```
function norm (x, y)
return (x ^ 2 + y ^ 2) ^ 0.5
end
```

```
function twice (x)
return 2 * x
end
```

Then interactively you can type

```
> dofile ("lib1.lua") - load your library
> n = norm (3.4, 1.0)
> print (twice (n)) -> 7.0880180586677
```

The `dofile` function is also useful when you are testing a piece of code. You can work with two windows: one contains the text editor with your program (for example, in the `prog.lua` file), and in other window is a console running the Lua interpreter in the interactive mode. After you have saved the changes to your program, you do `dofile ("prog.lua")` in the console to load the new code; then you can start using the new code, calling functions and printing the results.

1.2. Some lexical agreements

Identifiers (or names) in Lua are strings from Latin letters, numbers and underscores that do not start with a number;

eg:

```
ij i10 _ij
aSomewhatLongName _INPUT
```

You are better off avoiding identifiers consisting of underscores followed by capital Latin letters (for example, `_VERSION`); they are reserved for special purposes in Lua. I usually use `id_` (single underscore) for dummy variables. In older versions of Lua, the concept of what a letter is depended on from the locale. However, these letters make your program unsuitable to run on systems that do not support this locale. Therefore, Lua 5.2 considers only letters as letters from the following ranges: `AZ` and `az`. The following words are reserved, you cannot use them as identifiers:

```
and break do
else elseif
end false goto for function
if in
local nil not
or repeat return then true
```

until while

Lua is case sensitive: **and** is a reserved word, however, And and AND are two different identifiers. The comment starts with two minus signs (-) and continues with- until the end of the line. Lua also supports block comments, which starts with -[[and goes to the next]] ¹ . Standard the way to comment out a piece of code is to put it between

-[[and -]] as shown below:

```
-[[
print (10) - no action (commented out)
-]]
```

To make this code active again, just add one minus to the first line:

```
--- [[
print (10) -> 10
-]]
```

In the first example -[[in the first line starts a block commentary, and the double minus in the last line is also in this comment. In the second example --- [starts the usual single line comment, so the first and last lines become regular independent comments. In this case print is outside of comments.

1.3. Global Variables

Global variables do not need descriptions; you just use them you whine. It is not an error to refer to uninitialized variable; you just get nil as a result- that:

```
print (b) -> nil
b = 10
print (b) -> 10
```

If you assign nil to a global variable, Lua will itself as if this variable has never been used:

```
b = nil
print (b) -> nil
```

After this assignment, Lua may eventually regain its password. the space occupied by this variable.

1.4. Separate interpreter

Stand-alone interpreter (also called `lua.c` in relation to the name of its source file or just `lua` by name executable file) is a small program that allows direct use of Lua. This section presents her basic options.

When the interpreter loads a file, it skips the first string if it starts with `#` character . This allows you to use make Lua a scripting interpreter on UNIX systems. If you start your script with something like

```
#!/usr/local/bin/lua
```

(assuming the interpreter is in `/usr/local/bin`) or

```
#!/usr/bin/env lua ,
```

then you can directly run your script without explicitly start the Lua interpreter. The interpreter is called like this:

```
lua [options] [script [args]]
```

All parameters are optional. As we have seen, when we start If `lua` has no arguments, it goes into interactive mode.

The `-e` option allows you to directly specify the code directly in the command-line like below:

```
% lua -e "print (math.sin (12))" -> -0.53657291800043
```

(UNIX requires double quotes so that the command interpreter the `torus` (shell) did not parse parentheses).

The `-l` option loads the library. As we have seen previously, `-i` pe- puts the interpreter into interactive mode after processing the OS arguments. So the next call will load the library `lioteku lib` , then execute the assignment `x = 10` and finally pass into interactive mode.

```
% lua -i -llib -e "x = 10"
```

Interactively, you can print the value of the expression by simply typing a line that starts with an equal sign for followed by the expression:

```
> = math.sin (3) -> 0.14112000805987
> a = 30
> = a -> 30
```

This feature allows Lua to be used as a calculator. Before executing its arguments, the interpreter looks for environment variable named `LUA_INIT_5_2` or, if such a change no, no, `LUA_INIT` . If one of these variables is present and its has the form *@filename* , the interpreter launches this file. If `LUA_INIT_5_2` (or `LUA_INIT`) is defined but not start- is started with the '@' character , then the interpreter assumes that it contains executable Lua code and executes it. `LUA_INIT` gives huge the ability to configure the interpreter, since when configuration, all the power of Lua is available to us. We can download packages, change the current path, define your own functions change, rename or delete functions, etc. The script can get its arguments in a global variable

`arg` . If we have a call like `% lua script abc` , then the interpreter creates `arg` table with all command line arguments before by executing the script. The script name is located at index 0, per- the first argument (in the example it is “a”) is located at index 1, and so on. The preceding options are arranged in negative indices, as they appear before the script name. For example, ra- look at the following call:

```
% lua -e “sin = math.sin” script ab
```

The interpreter collects arguments as follows:

```
arg [-3] = "lua"
arg [-2] = "-e"
arg [-1] = "sin = math.sin"
arg [0] = "script"
arg [1] = "a"
arg [2] = "b"
```

Most often the script uses only positive indices (in the example, these are `arg [1]` and `arg [2]`).

Since Lua 5.1 the script can also receive its arguments using the expression ... (three dots). In the main part of the script, this is expression gives all the arguments of the script (we will discuss similar expressions (see section 5.2).

Exercises

Exercise 1.1 . Run the factorial example. What will happen with your program if you enter a negative number? Modify the example to avoid this problem.

Exercise 1.2 . Run the `twice` example by loading once file using the `-l` option , and another time using `dofile` . what faster?

Exercise 1.3 . Can you name another language using `(-)` for comments?

Exercise 1.4 . Which of the following lines are valid by our identifiers?

___ `_end End end until? nil NULL`

Exercise 1.5 . Write a simple script that prints your name without knowing it in advance.

CHAPTER 2

Types and values

Lua is a dynamically typed language. The language has no definitions of `ty-pov`,

each value carries its own type.

There are eight basic types in Lua: *nil* , *boolean* , *number* , *string* , *userdata* , *function* , *thread* and *table* . The `type` function returns the type for any passed value:

```
print (type ("Hello world"))
-> string
print (type (10.4 * 3))
-> number
print (type (print))
-> function
print (type (type))
-> function
print (type (true))
-> boolean
print (type (nil))
-> nil
print (type (type (X)))
-> string
```

The last line will always return `string` regardless of the value `Niya` the `X` , as the result of a function `type` is always a string.

Variables have no predefined types; any variable can contain values of any type:

```
print (type (a)) -> nil ('a' not yet defined)
a = 10
print (type (a)) -> number
a = "a string !!"
print (type (a)) -> string
a = print
-- Yes it is possible!
a (type (a))
-> function
```

Notice the last two lines: functions are first class values in Lua; they can be manipulated like any other values. (More on this in Chapter 6.) Usually when you use the same variable for value different types, you get disgusting code. However, sometimes judicious use of this opportunity is beneficial, for example using `nil` to distinguish normal return value from any error.

2.1. Nil

Nil is a type with only one value, `nil` , the main which is different from all other

values. Lua

uses `nil` to indicate a missing value. Like us already seen, globals default to `nil` before its first assignment, you can also assign `nil` a global variable to delete it.

2.2. Boolean (boolean values)

The boolean type has two values, ***true*** and ***false***, which serve to pre-setting traditional logical values. However, these values do not monopolize all conditional values: in Lua, any value can represent a condition. Relevant checks (checking conditions in various control structures) interpret both **nil** and **false** as false and all other values as valid muddy. In particular, Lua treats zero and an empty string as true. in logical conditions. Throughout the book, false will mean **nil**. and **false**. In the case when it is about boolean values, there will be explicitly set to **false**.

2.3. Numbers

The number type represents floating point values specified with double precision. Lua does not have a built-in integer type.

Some fear that even simple operations such as increase by one (increment) and comparison, may be incorrect work with floating point numbers. However, in fact it is not this way. Almost all platforms now support the standard IEEE 754 for floating point representation. According to this standard, the only possible source of errors is There is a case where the number cannot be accurately represented. Opera- walkie-talkie rounds its result only if the result cannot be accurately represented as the corresponding float value point. Any operation whose result can be accurately predicted put, will have the exact meaning. In fact, any integer up to 2^{53} (approximately 10^{16}) has an exact floating point representation with double precision. When you use a float value double-precision dot to represent integers, no rounding errors unless the value is greater than 2^{53} in absolute value. In particular, Lua is capable of representing any 32-bit integer values without rounding problems. Of course, fractional numbers will have rounding problems. This the situation is no different than when you have paper and a pen. If we want to write $1/7$ in decimal form, then we must somewhere us stop. If we use ten digits to represent number, then $1/7$ becomes 0.142857142. If we calculate $1/7 * 7$ s ten digits, we get 0.999999994, which is different from 1. Moreover, numbers that have a finite representation in the form

decimal fractions can have an infinite representation in the form binary fractions. So, $12.7 \cdot 20 + 7.3$ is not zero, since both numbers 12.7 and 7.3 do not have an exact binary representation (see Figure 2.3). Before we continue, remember that integers have an exact representation and therefore have no rounding errors. Most modern CPUs perform floating operations as fast (or even faster) than with integers. However, it is easy to compile Lua so that for numeric values a different type was used, for example long integers single-precision floating-point values or numbers. It is especially useful for platforms without hardware support for numbers with floating point such as embedded systems. For this, refer to the `luaconf.h` file in the Lua source files. We can write numbers, if necessary, specifying a fraction part and decimal degree. Examples of valid numeric constants are:

```
4 0.4 4.57e-3 0.3e12 5E + 20
```

Moreover, we can also use hexadecimal constants starting at `0x`. Since Lua 5.2 hexadecimal constants can also have a fractional part and a binary power (ne- 'p' or 'P' is used before the degree), as in the following examples:

```
0xff (255) 0x1A3 (419) 0x0.2 (0.125) 0x1p-1 (0.5)  
0xa.bp2 (42.75)
```

(We've added a decimal representation for each constant.)

2.4. Strings

Lua strings have the usual meaning: a sequence of characters. Lua supports all 8-bit characters, and strings can keep characters with any codes, including zeros. It means that you can store any binary data as strings. you also you can store unicode strings in any representation (UTF-8, UTF-16, etc.). The standard library that comes with Lua is does not contain built-in support for these views. However less, you may well be working with UTF-8 strings, which we will consider in section 21.7. Lua strings are immutable values. You can not change the character inside the string, as you can in C; instead of this you create a new line with the desired changes as shown in following example:

```
a = "one string"  
b = string.gsub(a, "one", "another") - change part of the string  
print(a) -> one string
```

print (b) -> another string

Lua strings are subject to automatic memory management, just like other Lua objects (tables, functions, etc.). This is a sign cheat that you don't have to worry about allocating and freeing lines; Lua will do it for you. The string can be one character or a whole book. Programs that work with lines of 100K or 10M characters are not uncommon in Lua. You can get the length of a string using as a prefix operator `#` (called length operator):

```
a = "hello"  
print (#a) -> 5  
print (# "good \ 0bye") -> 8
```

Literals

We can put strings inside single or double quotes- check:

```
a = "a line"  
b = 'another line'
```

These types of records are equivalent; the only difference is what is inside a string, limited to one type of quotation marks, you can directly insert quotation marks of a different type. Usually most programmers use quotation marks of one type for the same string type. For example, a library that works with XML, can use single quotes for strings, containing XML fragments, since these fragments often contain double quotes. Lua strings can contain the following escape sequence nos:

```
\ a  
bell  
\ b  
Back space  
\ f  
page translation (form feed)  
\ n  
newline  
\ r  
carriage return  
\ t  
tab (horizontal tab)  
\ v  
vertical tab  
\\
```

backslash
\"
double quote
\
single quote

The following example illustrates their use:

```
> print ("one line \nnext line \n\" in quotes \", 'in quotes'")
one line
next line
"In quotes", 'in quotes'
> print ('a backslash inside quotes: \' \\\ "')
a backslash inside quotes: '\'
> print ("a simpler way: '\\")
a simpler way: '\'
```

We can specify a character in a string using its numeric value using constructions like `\ddd` and `\x\hh`, where *ddd* is a sequence of no more than three decimal digits, and *hh* is a sequence of exactly two hexadecimal digits. As complex example two lines `"alo \n123\" "` and `"\ 97lo \ 10 \ 04923"` have the same meaning on a system using ASCII: 97 is the ASCII code for 'a', 10 is the code for the line feed character, and 49 is the code for the digit '1' (in this example we should write the value 49 using three decimal digits `\049`, because it is followed by another number; otherwise Lua interpreted it as code 492). We we can also write the same string as `' \x61 \x6c \x6f \x0a \x31 \x32 \x33 \x22 '`, representing each character in its sixteen-normal value.

Long lines

We can delimit character strings with double square brackets, as we did with the comments. Line in this form can span many lines, and control sequences the numbers on these lines will not be interpreted. Moreover, this the form ignores the first character of the line if it is a jump character next line. This form is especially useful for writing lines, containing large snippets of code as shown below:

```
page = [[
<html>
<head>
<title> An HTML Page </title>
</head>
```

```

<body>
<a href=32http://www.lua.org> Lua </a>
</body>
</html>
]]
write (page)

```

Sometimes you may want to put something like

`a = b [c [i]]` (note the `]]` in this code) or you can want to put in a line a piece of code where some piece is already commented out. To deal with such cases, you can use place any number of equal signs between two opening in square brackets, for example `[=== [` . After that the line will terminate only on a pair of closing square brackets with the same the most equal signs (`]===]` for our example). The scanner will ignore pairs of brackets with a different number of values. kov equality. By choosing a suitable number of equal signs You can wrap any fragment in a string. The same is true for comments. For example, if you start those long comment with `- [= [` , then it will continue all the way before `] =]` . This feature allows you to comment out any fragment a code snippet containing already commented out fragments.

Long lines are very handy for including text in your code, but you shouldn't use them for non-text strings. Although the lines in Lua can contain any characters, this is not a good idea - use these symbols in your code: you may run into problems with your text editor; moreover, lines of the form `"\ R \ n"` can become `"\ n"` . Therefore, to represent the arbitrary binary data, it is better to use control software sequences starting with a `\` , such as `\ x13 \ x01 \ xA1 \ xBB "` . However, this presents a problem for long lines because of the resulting length. For situations like this, Lua 5.2 offers a control after- the `\ z` sequence : it skips all characters in the string up to the first non-whitespace character. The following example illustrates its use use:

```

data = "\ x00 \ x01 \ x02 \ x03 \ x04 \ x05 \ x06 \ x07 \ z
\ x08 \ x09 \ x0A \ x0B \ x0C \ x0D \ x0E \ x0F "

```

`\ z` at the end of the first line skips the following end of line and indentation of next line so that byte `\ x07` immediately followed by a byte `\ x08` in the resulting row.

Type casts

Lua provides automatic conversion of values between

strings and numbers at run time. Any numeric operation applied to a string, tries to convert the string to a number:

```
print ("10" + 1) -> 11
print ("10 + 1") -> 10 + 1
print ("- 5.3e-10" * "2") -> -1.06e-09
print ("hello" + 1) - ERROR (cannot convert "hello")
```

Lua applies similar conversions not only in arithmetic operators, but also in other places where the expected number, for example for the `math.sin` argument . Similarly, when Lua expects to receive a string, but receives a number, it converts the number to string:

```
print (10 .. 20) -> 1020
```

(The `..` operator is used in Lua to concatenate strings. When you use it write immediately after the number, then you must separate them from each a friend with a space; otherwise Lua will assume that the first point is decimal point of the number.) Today we are not sure if these automatic conversions types were a good idea in Lua design. It is generally best on do not count them. They are handy in some places; but adding- There are complexities in both the language and the programs that use them. After all, strings and numbers are different types despite everything these transformations. A comparison like `10 = "10"` gives a false value, because `10` is a number and `"10"` is a string. If you need to explicitly convert a string to a number, then you can use the `tonumber` function , which returns **nil** if string ka does not contain a number:

```
line = io.read ()
- read the line
n = tonumber (line)
- try to translate it into a number
if n == nil then
error (line .. "is not a valid number")
else
print (n * 2)
end
```

To convert a number to a string, you can use the function `tostring` or concatenate the number with an empty string:

```
print (tostring (10) == "10")
-> true
print (10 .. "" == "10")
```

-> true

These transformations always work.

2.5. Tables

The table type corresponds to an associative array. An associative array is an array that can be indexed not only numbers, but also strings or any other value from the language, except for **nil**.

Tables are the main (actually the only) mechanism for structuring data in Lua, and very powerful. We use tables to represent regular arrays, sets, records and other data structures simple, homogeneous and in an efficient way. Lua also uses tables to pre- placing packages and objects. When we write `io.read`, we think- This is about the “`read` function from the `io` module”. For Lua, this expression means "Take the value from the `io` table by the `read` key".

Lua tables are neither values nor variables; they *objects*. If you are familiar with arrays in Java or Scheme, then you should know what I mean. You can think of a table as dynamically allocated object; your program only works with a link (pointer) to it. Lua never uses covert copying or creating new tables. Moreover, you even no need to declare a table in Lua; in fact there is not even a way declare a table. You create tables using a special expression, which in the simplest case looks like `{}` ;

```
a = {}  
- create a table and remember the link to it in 'a'  
k = "x"  
a[k] = 10  
- new record with key "x" and value 10  
a[20] = "great" -- new record with key 20 and value "great"  
print (a["x"]) -> 10  
k = 20  
print (a[k]) -> "great"  
a["x"] = a["x"] + 1 -- increase the record "x"  
print (a["x"]) -> 11
```

The table is always anonymous. There is no permanent connection between the variable that contains the table and the table itself:

```
a["x"] = 10  
b = a  
- 'b' refers to the same table as 'a'  
print (b["x"]) -> 10  
b["x"] = 20
```

```
print (a ["x"]) -> 20
a = nil - only 'b' still refers to the table
b = nil - there are no references to the table
```

When there are no more references to the table left in the program, the collector garbage in Lua will eventually destroy the table and reuse its memory. Each table can contain values with different types of data, and the table grows as new records are added:

```
a = {}
- empty table
- create 1000 new entries
for i = 1, 1000 do a [i] = i * 2 end
print (a [9]) -> 18
a ["x"] = 10
print (a ["x"]) -> 10
print (a ["y"]) -> nil
```

Pay attention to the last line: as with the global variables, uninitialized table fields return **nil**. As with global variables, you can assign the table field is **nil** to destroy it. This is not a coincidence: Lua stores global variables in regular tables. We'll take a closer look at this in Chapter 14. To represent records, you use the field name as an index. Lua supports this view by offering the following syntax sugar: instead of `a [" name "]` you can write `a.name`. So Thus, we can rewrite the last few lines of the previous example in a cleaner way:

```
ax = 10 - same as a ["x"] = 10
print (ax) - same as print (a ["x"])
print (ay) - same as print (a ["y"])
```

For Lua, these two forms are completely equivalent and can be freely used. For the reader, however, each form can communicate a specific intent. Dot notation clearly shows that we are using the table as a record (structure), where we have a certain set of given, predefined keys.

Another entry suggests that the table can be used use any string as a key and for some reason in this place we work with a specific key.

A common newbie mistake is that they confuse `ax` and `a [x]`. The first form actually matches `a ["x"]`, that is, referring to the table with the key "x". In the second case the key is the value of the variable `x`. Shown below

difference:

```
a = {}  
x = "y"  
a[x] = 10 - write 10 in the "y" field  
print (a[x]) -> 10 - value of the "y" field  
print (ax) -> nil - value of field "x" (undefined)  
print (ay) -> 10 - value of the field "y"
```

To represent a traditional array or list, simply use
use a table with integer keys. There is no way, no
the need to declare size; you just initialize those elements
the cops you need:

```
- read 10 lines, memorizing them in the table  
a = {}  
for i = 1, 10 do  
  a[i] = io.read ()  
end
```

Since you can index the table by any value-
you can start indices in an array with any number that
you like. However, it is customary in Lua to start arrays with one.
(not from scratch as in C) and some Lua tools stick to this-
th agreement.

Usually, when you are working with a list, you need to know its length.
It can be constant or it can be written somewhere. Usually
we write the length of the list in a field with a non-numeric key; by history
For logical reasons, some programs use for these purposes
Leave the field "n" .

Often, however, the length is not explicitly specified. Since anyone who does not
the value of **nil** corresponds to the initialized field , then we can use
use this value to determine the end of the list. For instance,
if you've read ten lines into a list, it's easy to remember that its
the length is 10, since its keys are the numbers 1, 2,..., 10.
This approach only works with lists that have no *holes* , which
they contain the value **nil** . We call such lists of *sequences*
sequences .

For sequences, Lua offers the '#' length operator. is he
returns the last index or length of a sequence. On-
example, you can print the lines read in the previous
example using the following code:

```
- print the lines  
for i = 1, #a do  
  print (a[i])  
end
```

Since we can index the table with values of any type, then when indexing a table, the same subtleties arise as and when checking for equality. Although we can index the table and using the integer 0 , and using the string "0" , these two mean The values are different and correspond to different elements of the table. Analogically, the lines "+1" , "01" and "1" also correspond to different elements table cops. When you are unsure about the type of your indexes, use explicit casting:

```
i = 10; j = "10"; k = "+10"
```

```
a = {}
```

```
a[i] = "one value"
```

```
a[j] = "another value"
```

```
a[k] = "yet another value"
```

```
print(a[i])
```

```
-> one value
```

```
print(a[j])
```

```
-> another value
```

```
print(a[k])
```

```
-> yet another value
```

```
print(a[t tonumber(j)]) -> one value
```

```
print(a[t tonumber(k)]) -> one value
```

If you do not pay attention to these subtleties, then it is easy to add to the program hard-to-find errors.

2.6. Functions

Functions are first class values in Lua: programs can write functions to variables, pass functions as arguments to other functions and return functions as a result. This capability lends tremendous flexibility to the language; software ma can override the function to add new functionality nality, or simply remove the function to create a safe environment for executing a piece of untrusted code (for example, code received over the network). Moreover, Lua provides a good support for functional programming, including nested functions with the appropriate lexical environment; simple then wait until chapter 6. Finally, first class functions play an important role in the object-oriented features of Lua as we will see in chapter 16.

Lua can call functions written in Lua and functions that

written in C. Usually we use functions written in C. C, in order to get high performance and access to the features not available directly from Lua, such as accessing operating system tools. All standard libraries in Lua are written in C. They include functions for working with strings, working with tables, input / output, access to basic operating system, math functions and debugging. We will discuss Lua functions in Chapter 5 and C functions in Chapter 27.

2.7. userdata and threads

The userdata type allows you to store arbitrary C language data in Lua variables. This type has no built-in operations, except assignment and equality testing. Values of this type are used to represent new types created by applications or a library written in C; for example, the standard library the I / O library uses them to represent open files. We will discuss this type in more detail later when we move on to the C API.

The *thread* type will be covered in Chapter 9, where we will look at coroutines.

Exercises

Exercise 2.1. What is the meaning of an expression `type(nil) == nil`? (You can use Lua to check your answer.) Can you explain the result?

Exercise 2.2. Which of the following is acceptable by numbers? What are their meanings?

`.0e12 .e12 0.0e 0x12 0xABFG 0xA FFFF 0xFFFFFFFF
0x 0x1P10 0.1e1 0x0.1p1`

Exercise 2.3. 12.7 is equal to 127/10, where all numbers are decimal. Can you imagine it as the value of a binary fraction? And the number 5.5?

Exercise 2.4. How would you write the following XML snippet in a Lua string?

```
<![CDATA [  
Hello world  
]]>
```

Use at least two different methods.

Exercise 2.5. Let's say you need to write a long after-a sequence of arbitrary bytes as a string constant in Lua. How do you do it? Pay attention to readability, maximum line length and performance.

Exercise 2.6. Consider the following code:

```
a = {}; aa = a
```

What will be the value of `aaaa` ? Any `a` in this sequence is somehow different from the rest?

Now add the following line to the previous code:

```
aaaa = 3
```

What will be the value of `aaaa` now ?

CHAPTER 3

Expressions

Expressions represent values. Expressions in Lua include numeric constants and string literals, variables, unary and binary operations and function calls. Expressions also include- contain non-standard function definitions and constructors for tables.

3.1. Arithmetic operators

Lua supports the standard arithmetic operators: binary- nye '+' (addition), '-' (subtraction), '*' (multiplication), '/' (div- number), '^' (exponentiation), '%' (remainder of division), and unary '-' (change sign). All of them work with floating-point numbers.

point. For example, $x^{0.5}$ calculates the square root of x , and $x^{(-1/3)}$ computes the inverse of the cube root of x .

The following rule defines the modulus operator:

```
a% b == a - math.floor (a / b) * b
```

For integer operands, it has a standard value, and re-

The result has the same sign as the second operand. For real operands, it has some additional features. On-

for example, `x% 1` gives the fractional part of `x`, and `xx% 1` gives the integer part.

Ana-

logical `xx% 0.01` gives `x` with exactly two decimal places after comma:

```
x = math.pi
```

```
print (x - x% 0.01) -> 3.14
```

As another example of using the remainder operator

division, consider the following example: let's say you want to know will the vehicle after turning at a given angle

move in the opposite direction. If the angle is given in degrees, then you can use the following formula:

```
local tolerance = 10
```

```
function isturnback (angle)
```

```
angle = angle% 360
```

```
return (math.abs (angle - 180) < tolerance)
```

```
end
```

This definition works even for negative angles:

```
print (isturnback (-180)) -> true
```

If you want to work in radians instead of degrees, we just let's change the constants in the functions:

```
local tolerance = 0.17
```

```
function isturnback (angle)
```

```
angle = angle% (2 * math.pi)
```

```
return (math.abs (angle - math.pi) < tolerance)
```

```
end
```

All we need is an `angle% (2 * math.pi)` operation to convert any angle to the interval $[0, 2\pi)$.

3.2. Comparison Operators

Lua provides the following comparison operators:

```
<> <= > == ~=
```

All of these operators are always boolean.

The `==` operator tests for equality; operator `~=` is negation

equality. We can use both of these operators to any two

values. If the values are of different types, then Lua assumes that they are not equal. Otherwise Lua compares them accordingly

their type. The **nil** value is equal only to itself.

Lua compares tables and objects of type userdata by reference, i.e. two such values are considered equal only if they are the same object. For example, after doing the following code:

```
a = {}; ax = 1; ay = 0
b = {}; bx = 1; by = 0
c = a
```

we get `a == c` , but `a ~ b` .

We can only apply order operators to a pair of numbers, or a couple of lines. Lua compares strings alphabetically, following the convention The locale set for Lua. For example, for the Portuguese locale Latin-1 we get `"acai" < "açai" < "acorde"` . Values of types from-personal from strings and numbers, can only be compared for equality (and inequality).

When comparing values of different types, you need to be careful nym: remember that `"0"` is different from `0` . Moreover, `2 < 15` is obvious true, but `"2" < "15"` is false. In case you are trying to compare a string and a number, for example `2 < "15"` , an error occurs.

3.3. Logical operators

The logical operators are **and** , **or** and **not** . As well as managing constructs, logical operators treat **false** and **nil** as false and all others as true values. **And** return operator gives its first operand if it is false, otherwise it returns its second operand. The **or** operator returns its first operand if it is not false; otherwise it returns its second operand:

```
print (4 and 5)
-> 5
print (nil and 13)
-> nil
print (false and 13)
-> false
print (4 or 5)
-> 4
print (false or 5)
-> 5
```

Both operators (**and** and **or**) use shorthand evaluation, then there they calculate their second operand only when it is necessary dimo. This ensures that expressions like `(type(v) == "table" and`

`v.tag == "h1")` will not cause errors in their computation: Lua will not try to compute `v.tag` when `v` is not a table.

A useful construct in Lua is `x = x or v`, which is equivalent to the following code:

```
if not x then x = v end
```

That is, the `x` value is set equal to the default `v` if `x` is undefined (assuming `x` is not **false**).

Another useful construct is `(a and b) or c`, or simply the `a and b or c`, as the operator **and** has a higher preference than **or**. It is equivalent to the expression `a ? B: c` in C, when provided that `b` is not false. For example, we can choose the maximum from two numbers `x` and `y` using the following operator:

```
max = (x > y) and x or y
```

When `x > y`, then the first expression in the **and** operator is true, to this it returns its second value (`x`), which is always tiny (since it's a number), and then the **or** operator returns its first operand, `x`. If the expression `x > y` is false, then the result of the operator **and** is also false, and so the **or** operator returns its second operand, `y`.

The `not` operator always returns a boolean value:

```
print (not nil)
-> true
print (not false)
-> true
print (not 0)
-> false
print (not not 1)
-> true
print (not not nil)
-> false
```

3.4. Concatenation

Lua refers to the concatenation operator as `..` (dots). If operand is a number, Lua will convert it to a string. (Some languages use the `+` operator for concatenation, but Lua `3 + 5` differs from `3..5`.)

```
print ("Hello" .. "World")
-> Hello World
print (0 .. 1)
-> 01
print (000 .. 01)
```

-> 01

Remember that strings in Lua are immutable values.

The concatenation operator always creates a new string without changing its operands:

```
a = "Hello"
print (a .. "World")
-> Hello World
print (a)
-> Hello
```

3.5. Length operator

The length operator works with strings and tables. With strings it gives the number of bytes per line. With tables, it returns the length of the *pos-research* presented in the table.

There are several common idioms associated with the length operator to work with sequences.

```
print (a [#a]) - prints the last element of the sequence 'a'
a [#a] = nil - removes the last element
a [#a + 1] = v - adds 'v' to the end of the list
```

As we saw in the previous chapter, the length operator is not let's say for lists with holes (nil). It only works for software sequences, which we defined as lists without holes.

More precisely, a *sequence* is a table where the keys are give a sequence $1, \dots, n$ for some n . (Remember, that any **nil** key is not actually in the table.)

In particular, a table without numeric keys is a sequential length zero.

Over the years, there have been many proposals to expand the meaning of lengths to lists with holes, but this is easier said than done.

The problem is that since a list is a table, the concept "Length" is somewhat vague. For example, consider a list produced by the following code snippet:

```
a = {}
a [1] = 1
a [2] = nil - does nothing, since a [2] is already nil
a [3] = 1
a [4] = 1
```

It is easy to say that the length of this list is four and it has a hole in index 2. However, what about the following example?

```
a = {}
```


a[1] = 1

a[10000] = 1

Should we treat this a as a list with 10,000 elements where 9998 elements are **nil** ? Now let the program do following:

a[10000] = nil

What happened to the length of the list? Should it be 9999, since the program removed the last item? Or maybe 10,000, because the program just changed the value of the last element to **nil** ? Or should the length be 1?

Another common suggestion is to do so that the # operator returns the number of elements in the table. This semantic is clear and well defined, but of no benefit.

Let's consider all the previous examples and imagine how much

A similar operator would be useful for algorithms that work with lists or arrays.

Even more problematic are the **nil** values at the end of the list. Ka-

What should be the length of the following list?

a = {10, 20, 30, nil, nil}

Recall that for Lua, a **nil** field is no different from accompanying field. Thus, the previous table is indistinguishable from {10,20,30} ; its length is 3, not 5.

You might think that **nil** at the end of the list is a special case.

However, many lists are built by adding elements, one after another. Any list with holes built like this thus, simply obtained by adding **nil** to its end.

Many of the lists we use in our programs are sequences (for example, a line in a file cannot be **nil**), and so most of the time the length operator is not passive for use. If you really have to work with lists with holes, then you'd better clearly remember the length somewhere list.

3.6. Operator Priorities

The operator precedence in Lua is given in the table below, from the oldest to the lowest:

^

not # - (unary)

* /%
+ -
..
<> <=> = ~ ==

and

or

All binary operators are left associative, except
'^' (exponentiation) and '.' (concatenation), which are associated
tive to the right. Therefore the following expressions on the left are equivalent
expressions on the right:

$a + i < b / 2 + 1$

$<-->$

$(a + i) < ((b / 2) + 1)$

$5 + x ^ 2 * 8$

$<-->$

$5 + ((x ^ 2) * 8)$

$a < y \text{ and } y <= z <-->$

$(a < y) \text{ and } (y <= z)$

$-x ^ 2$

$<-->$

$-(x ^ 2)$

$x ^ y ^ z$

$<-->$

$x ^ (y ^ z)$

When in doubt, always use parentheses. It's easier than
look in the manual, and most likely later when you read
this code, you will have doubts again.

3.7. Table constructors

Constructors are expressions that create and initialize
tables. They are the hallmark of Lua and one of its
most useful and versatile mechanisms.

The simplest constructor is an empty constructor, `{}` , which
creates an empty table; we've seen this before. Constructors
also initialize lists. For example, the operator

`days = {"Sunday", "Monday", "Tuesday", "Wednesday",
"Thursday", "Friday", "Saturday"}`

will initialize `days [1]` to "Sunday" (the first element
constructor has index 1, not 0), `days [2]` is "Monday"
etc.:

`print (days [4]) -> Wednesday`

Lua also offers special syntax for initialization
tables by fields, as in the following example:

`a = {x = 10, y = 20}`

This line is equivalent to the following commands:

```
a = {}; ax = 10; ay = 20
```

The original expression is simpler and faster because Lua immediately creates
It has a table with the correct size.

Regardless of which constructor we used
to create a table, we can always add and remove fields
from her:

```
w = {x = 0, y = 0, label = "console"}
x = {math.sin (0), math.sin (1), math.sin (2)}
w [1] = "another field" - add key 1 to table 'w'
xf = w
- add key "f" to table 'x'
print (w ["x"])
-> 0
print (w [1])
-> another field
print (xf [1])
-> another field
wx = nil
- remove field "x"
```

However, creating the table right away with the correct boot-
more efficiently and clearly.

We can mix these two initialization styles (list and by
fields) in the same constructor:

```
polyline = {color = "blue",
thickness = 2,
npoints = 4,
{x = 0, y = 0},
- polyline [1]
{x = -10, y = 0},
- polyline [2]
{x = -10, y = 1},
- polyline [3]
{x = 0, y = 1}
- polyline [4]
}
```

The example above also shows how contributions can be-
nest constructors one into another to represent more complex
data structures. Each of the `polyline [i]` elements is a table,
representing a record:

```
print (polyline [2] .x)
-> -10
print (polyline [4] .y)
-> 1
```

These two forms of constructor have their limitations. For instance,

you cannot initialize fields with negative indices or with indices that are not identifiers. For such goals have a different, more general format. In this format, we clearly write we use the index as an expression between square brackets:

```
opnames = {[["+"] = "add", [{"-"} = "sub",
[["*"] = "Mul", [{"/"} = "div"]}
i = 20; s = "-"
a = {[i + 0] = s, [i + 1] = s..s, [i + 2] = s..s..s}
print (opnames [s])
-> sub
print (a [22])
-> ---
```

This syntax is more awkward, but also more general: considered previously constructor forms are special cases of this more general syntax. The constructor $\{x = 0, y = 0\}$ is equivalent to $\{["X"] = 0, [{"y"} = 0\}$, and the constructor $\{"r", "g", "b"\}$ is equivalent to $\{[1] = "r", [2] = "g", [3] = "b"\}$.

You can always put a comma after the last entry in constructor. These commas are optional:

```
a = {[1] = "red", [2] = "green", [3] = "blue",}
```

This frees programs that generate Lua constructors from the need to process the last element in a special way.

Finally, you can always use dot c in the constructor comma instead of comma. I usually use semicolons to de-division of various sections in the constructor, for example, the division of the hour

from the part formatted as a list:

```
{x = 10, y = 45; "One", "two", "three"}
```

Exercises

Exercise 3.1. What will the following program print?

```
for i = -10, 10 do
  print (i, i% 3)
end
```

Exercise 3.2. What is the result of the expression $2 \wedge 3 \wedge 4$?

What about $2 \wedge -3 \wedge 4$?

Exercise 3.3. We can represent the polynomial

$$a_n X^n + a_{n-1} X^{n-1} + \dots + a_1 X^1 + a_0$$

in Lua as a list of its coefficients $\{a_0, a_1, \dots, a_n\}$.

Write a function that receives a polynomial (represent

table) and `x` returns a polynomial value
ma in `x`.

Exercise 3.4. Can you write a function from the previous present exercises so as to use n additions and n multiplied scaling (and not using exponentiation)?

Exercise 3.5. How can you check if the value is boolean without resorting to the `type` function ?

Exercise 3.6. Consider the following expression:

`(x and y and (not z)) or ((not y) and x)`

Are parentheses needed in this expression? How would you go advised to use them in this expression?

Exercise 3.7. What will the following code snippet print? Explain.

```
sunday = "monday"; monday = "sunday"  
t = {sunday = "monday", [sunday] = monday}  
print (t.sunday, t[sunday], t[t.sunday])
```

Exercise 3.8. Suppose you want to create a table that-paradise with each escape sequence (escape sequence) for strings binds its value. How would you write a constructor for such a table?

CHAPTER 4

Operators

Lua supports an almost traditional set of operators like to a set used in C or Pascal. Traditional operators include assignment, control constructs and calls procedures. Lua also supports the less common operators such as multiple assignment and locale definition variables.

4.1. Assignment operators

Assignment is the basic means of changing the values of variables noah and table fields:

```
a = "hello" .. "world"  
tn = tn + 1
```

Lua allows *multiple assignments*, which swarm assigns a list of values to a list of variables in one step.

For example, in the operator

```
a, b = 10, 2 * x
```

variable `a` gets the value 10, and variable `b` gets the value `2 * x`.

In multiple assignment, Lua first evaluates all values.

and only then performs the assignments. Therefore, we can use multiple assignments to change

in places two values, as in the following examples:

```
x, y = y, x
```

```
- swap 'x' and 'y'
```

```
a[i], a[j] = a[j], a[i] - swap 'a[i]' and 'a[j]'
```

Lua always converts the number of values to the number of variables:

when the number of values is less than the number of variables, then

Lua pads the list of values with the appropriate number of **nils**,

and when the number of values is greater, then the extra values are simply thrown:

```
a, b, c = 0, 1
```

```
print(a, b, c)
```

```
-> 0 1 nil
```

```
a, b = a + 1, b + 1, b + 2 - the value b + 2 is discarded
```

```
print(a, b)
```

```
-> 1 2
```

```
a, b, c = 0
```

```
print(a, b, c)
```

```
-> 0 nil nil
```

The last assignment in the example above shows the distribution

an error. In order to initialize the list, re-

variables, you must provide a value for each variable:

```
a, b, c = 0, 0, 0
```

```
print(a, b, c)
```

```
-> 0 0 0
```

In fact, most of the previous examples are somewhat

then artificial. I rarely use multiple assignments

just to connect several unrelated

about of assignments in one line. In particular, the multiple assignment

an assignment is no faster than a set of corresponding single assignments.

vanii. However, often we really need a plural

assignment. We have already seen an example that changes two variables

values. More frequent use is to obtain

several values returned by the function at once. How do we discuss

dim in section 5.1, a function can return several values at once.

In such cases, multiple assignment is usually used -

to get all these values. For example, in the assignment

a, b = f () call f gives two values: the first one is written to a ,
and the second is at b .

4.2. Local variables and blocks

In addition to global variables, Lua also supports local variable variables. We create local variables with operator **local** :

```
j = 10
```

- global variable

```
local i = 1 - local variable
```

Unlike global variables, the scope of local the variable is limited to the block where it was declared. *Block* is control structure body, function body and code block (file or the line where the variable was declared):

```
x = 10
```

```
local i = 1 - local in the block
```

```
while i <= x do
```

```
local x = i * 2 - local inside the while block
```

```
print (x) -> 2, 4, 6, 8, ...
```

```
i = i + 1
```

```
end
```

```
if i > 20 then
```

```
local x
```

- local inside “then”

```
x = 20
```

```
print (x + 2) - (will print 22 if the condition is met)
```

```
else
```

```
print (x) -> 10 (global)
```

```
end
```

```
print (x)
```

```
-> 10 (global)
```

Please note that this example will not work like expected if you enter it interactively. In the interactive mode, each line is an independent block (for except for the case when the string is not a complete const manual). Once you enter the second line of the example (`local i = 1`), Lua will execute it and start a new block of code (next line ka). By that time, the scope of local variable `i` is already will end. To solve this problem, we can explicitly conclude this whole block between the **do - end keywords** . When you enter **do** , the block will end only when you enter the corresponding **end** , so Lua won't try to execute each line as separate block.

These **do** -blocks prove useful when we need BO
More precise control over the scope of local variables:

```
do
local a2 = 2 * a
local d = (b ^ 2 - 4 * a * c) ^ (1/2)
x1 = (-b + d) / a2
x2 = (-b - d) / a2
end -- scope 'a2' and 'd' ends here
print (x1, x2)
```

It is good style to use local variables
wherever possible. Local variables help
you avoid clogging the global environment with unnecessary names
mi. Moreover, accessing a local variable is faster than accessing
to the global. Finally, the local variable ceases to exist.
as soon as its scope ends, allowing the collection
the garbage bin to free the memory occupied by its value.
Lua treats local variable declarations simply as
operators. Therefore, you can insert a description of the local re-
wherever you can insert a statement. Scope
of the described variables begins immediately after the description and ends
end of the block. Each description may include an assignment
initial value, which acts the same as the operator for
piles: extra values are discarded, extra variables
get **nil** . If there is no assignment in the variable description
nil , the corresponding variable is set to **nil** :

```
local a, b = 1, 10
if a < b then
print (a) -> 1
local a - implies '= nil'
print (a) -> nil
end
- ends the block started by 'then'
print (a, b) -> 1 10
```

A common idiom in Lua is the following:

```
local foo = foo
```

This code creates a local variable `foo` and initializes it
the value of the global variable `foo` . (Local variable `foo`
becomes visible only after this announcement.) This idiom
useful when a block needs to store a value
the original variable, if it changes somewhere later in the code;
it also speeds up access to this variable.

Since many languages are forced to declare all local

variables at the beginning of a block (or procedure), some people think that it is bad practice to declare variables in the middle of a block.

In fact, the opposite is true: declaring a variable only when you really need it, you rarely need to declare it without initial value (and therefore you are unlikely to forget its initial lyse). Moreover, you reduce the scope of change - noah, which makes the code easier to read.

4.3. Control constructs

Lua provides a small and fairly traditional set of control constructs using **if** for conditional execution and **while** , **repeat** and **for** to iterate. All control structures have an explicit ending: **end** terminates **if** , **for** and **while** , while how **until** ends **repeat** .

The conditional execution of the control structure can give any value. Remember that Lua considers all values other than nye from **false** and **nil** , as true. (In particular, Lua considers zero and an empty string as true values.)

if then else

The **if statement** checks the condition and execute its *then-part* or its *else-part* respectively. The **else** part is optional.

```
if a < 0 then a = 0 end
if a < b then return a else return b end
if line > MAXLINES then
  showpage ()
  line = 0
end
```

For writing nested **if statements**, you can use **elseif** . This is similar to **else** followed by an **if** , but it does not no need for many **end** :

```
if op == "+" then
  r = a + b
elseif op == "-" then
  r = a - b
elseif op == "*" then
  r = a * b
elseif op == "/" then
  r = a / b
else
  error ("invalid operation")
```

end

Since there is no *switch statement* in Lua , such constructions can be are freely common.

while

As the name suggests, this operator repeats its body while the condition is true. As usual, Lua checks the condition first; if a it is false, then the cycle ends; otherwise Lua does the body of the loop and repeats the given process.

```
local i = 1
while a [i] do
  print (a [i])
  i = i + 1
end
```

repeat

As the name suggests, the **repeat - until** statement repeats its body until the condition becomes true. Checking the condition is added after the loop body is executed, so the loop body will be executed at least once.

```
- print the first non-empty line
repeat
  line = io.read ()
until line ~= ""
print (line)
```

Unlike many languages, in Lua the scope of local variables include the loop condition:

```
local sqr = x / 2
repeat
  sqr = (sqr + x / sqr) / 2
local error = math.abs (sqr ^ 2 - x)
until error < x / 10000 - local variable 'error' is visible here
```

Numeric for statement

The **for** statement **comes** in two flavors - numeric **for** and generic **for** .

The numeric **for** statement looks like this:

```
for var = exp1, exp2, exp3 do
```

<something>

end

This loop will execute *something* for every `var` value from `exp1` to `exp2`, using `exp3` as a *step* to increase `var`. This is a your expression (`exp3`) is optional; when absent, Lua uses 1 as a step. As typical examples of such cycles can be considered

```
for i = 1, f(x) do print(i) end
```

```
for i = 10, 1, -1 do print(i) end
```

If you want to get a loop without an upper limit, then you can

use constant `math.huge` :

```
for i = 1, math.huge do
```

```
if (0.3 * i ^ 3 - 20 * i ^ 2 - 500 >= 0) then
```

```
print(i)
```

```
break
```

```
end
```

```
end
```

The **for** loop has some subtleties that you know best, to use it well. First, all three expressions are counted only once, before the start of the cycle. For example, in our first example, Lua will execute `f(x)` just once. Secondly, the control variable is a local variable, automatically declared operator **for**, and it is only within the visible cycle. It is a common mistake to think that this variable is all still exists after the end of the loop:

```
for i = 1, 10 do print(i) end
```

```
max = i - possibly wrong! Here 'i' is global
```

If you need the value of the control variable after the loop (usually when you exit the loop ahead of time), then you should store its value in another variable:

- find a value in the list

```
local found = nil
```

```
for i = 1, #a do
```

```
if a[i] < 0 then
```

```
found = i - save value of 'i'
```

```
break
```

```
end
```

```
end
```

```
print(found)
```

Third, you should never change the value of the control variable: the effect of such changes is unpredictable. If you want to end the **for** loop to terminate it normally, use **break** (as we did in the previous example).

General for statement

The generic **for** statement iterates over all the values returned by iterating function:

```
- print all values in table 't'  
for k, v in pairs (t) do print (k, v) end
```

This example uses `pairs`, a convenient iteration function to traverse the entire table provided by the underlying library Lua. At each step of this cycle, `k` gets an index, and `v` gets a value, associated with this index.

Despite its seeming simplicity, the general **for** operator - it is a very powerful language construct. With suitable iterators you can bypass just about anything you want in an easy-to-read form. The standard libraries provide multiple iterators, allowing us to iterate over the lines of the file (`io.lines`), pairs from the tables (`pairs`), elements of sequence (`ipairs`), words inside strings (`string.gmatch`), etc.

Of course, we can write our own iterators as well. Although using the **for** statement in its general form is easy, the task of writing an iterator function has its own subtleties; we will cover this topic later in chapter 7.

The general loop operator has two features in common with numeric loop operator: loop variables are local to loop bodies, and you should never write any values.

Let's consider a more specific example of using the operator **for** general view. Let's say you have a table with the names of non-Delhi:

```
days = { "Sunday", "Monday", "Tuesday", "Wednesday",  
         "Thursday", "Friday", "Saturday" }
```

Now you want to translate the name of the day to its position in the week.

You can traverse the entire table looking for a given name. But, as you'll soon find out, you rarely need to search in Lua. More efficient an effective approach would be to build a reverse table, for example `revDays`, where day names are indices and values are `mi` are day numbers. This table will look like this way:

```
revDays = { ["Sunday"] = 1, ["Monday"] = 2,  
           ["Tuesday"] = 3, ["Wednesday"] = 4,
```

```
["Thursday"] = 5, ["Friday"] = 6,  
["Saturday"] = 7}
```

Then all you need to find the day number is
refer to this reverse table:

```
x = "Tuesday"  
print (revDays [x]) -> 3
```

Of course, you don't have to explicitly set this reverse table. We can
build it automatically from the original:

```
revDays = {}  
for k, v in pairs (days) do  
  break, return and goto  
revDays [v] = k  
end
```

This loop will perform an assignment for each element of `days` ,
where the variable `k` gets the key (1, 2, ...) and `v` gets the value
("Sunday", "Monday", ...).

4.4. break, return and goto

The **break** and **return statements** allow us to jump right out of the block.
ka. The **goto statement** allows us to jump to almost any
place of function.

We use the **break** statement to end the loop. This opera-
the generator interrupts the inner loop (**for** , **repeat** or **while**) containing
shying him. Also it can be used to return from a function,
so you don't have to use the **return statement** unless you
return no value.

For syntactic reasons, the **return statement** can only be
to the last statement of the block: in other words, or the last
operator, or right before **end** , **else**, or **until** . In the following case-
as the **return statement** - the last statement block **the then** .

```
local i = 1  
while a [i] do  
  if a [i] == v then return i end  
  i = i + 1  
end
```

This is usually the place where we use **return** , because
any other operators following it would never
filled up. Sometimes it's actually useful to write **return** in
middle of the block; for example you can debug a function and want
avoid doing it. In cases like this, you can use

put an explicit **do** block around the **return statement** :

```
function foo ()  
  return  
- << SYNTAX ERROR  
- 'return' is the last statement in the next block  
do return end  
- OK  
<other statements>  
end
```

The **goto statement** translates program execution to the appropriate label. There were long discussions about **goto** , some people even now they believe that they are harmful to programming and should us to be excluded from programming languages. However, many languages offer a similar operator, and they have a reason for it. These operators are a powerful mechanism that will used carefully, can improve the quality of your th code.

In Lua, the syntax for the **goto statement** is quite traditional: it is a **goto** reserved word followed by a label name which can be any valid identifier. Syntax for the label, however, is more complex: it consists of two colons, followed by followed by the label name followed by two more colons, for example `:: name ::` . This complexity is intentional, its purpose is to force programmer think twice before using **goto** .

Lua puts some restrictions on where you can jump with **goto** . First, the labels follow the usual rules of visibility, so you cannot jump directly inside the block (since the label inside the block is invisible outside of it). In-second, you cannot jump out of the function. (Note, that the first rule excludes the possibility of jumping *inside the* function tion.) Third, you cannot jump inside the area of effect of the local variable.

A typical and well-used use of the **goto statement** is an emulation of some construction you learned from in another language, but which is absent in Lua, such as **continue** , many go-level **break** , **redo** , etc. The **continue** statement is just a transition to the label at the end of the loop, the **redo** statement jumps to the beginning of the block:

```
while some_condition do  
  :: redo ::  
  if some_other_condition then goto continue
```

```

else if yet_another_condition then goto redo
end
<some code>
:: continue ::
end

```

A useful nuance in the Lua specs is that the scope

```

while some_condition do
if some_other_condition then goto continue end
local var = something

```

of a local variable ends with the last *neppure* operator of the block where the variable is defined; tags are counted empty operators. In order to see the usefulness of this, look at the following code snippet:

```

break, return and goto
<some code>
:: continue ::
end

```

You might think this **goto statement** jumps straight to the scope of `var`. However, the `continue` label finds after the last non-empty block operator, and therefore not in the region. These are the actions of `var`.

The **goto statement** is also useful when writing state machines.

As an example, Listing 4.1 is an example program that which believes whether its input contains an even number of zeros. Exists better ways of writing this program, but this approach very useful if you want to automatically translate the ending automaton to Lua code (think about automatic code generation).

As another example, consider a simple maze game.

The maze contains several rooms, each with up to four doors: north, south, east and west. At each step, the user enters movement board. If there is a door in this direction, then the user the provider enters the appropriate room; otherwise the print program there is a warning. The goal is to get from the starting room to the ultimate room.

This game is a typical machine where the current room is is a state. We can implement this game using one a block for each room and a **goto statement** to jump from one rooms to another. Listing 4.2 shows how you can write a simple the smallest labyrinth of four rooms.

For this simple game, you might decide that the program that controls data when you describe rooms and movements when

using tables is a better solution. However, if in each room has its own characteristics, then this approach provides is quite successful.

Listing 4.1. State machine example using **goto**

```
:: s1 :: do
local c = io.read (1)
if c == '0' then goto s2
elseif c == nil then print'ok '; return
else goto s1
end
end
```



```

:: s2 :: do
local c = io.read (1)
if c == '0' then goto s1
elseif c == nil then print'not ok '; return
else goto s2
end
end
goto s1

```

Listing 4.2. Maze game

```

goto room1 - starting room
:: room1 :: do
local move = io.read ()
if move == "south" then goto room3
elseif move == "east" then goto room2
else
print ("illegal move")
goto room1 - stay in the same room
end
end
:: room2 :: do
local move = io.read ()
if move == "south" then goto room4
elseif move == "west" then goto room1
else
print ("illegal move")
goto room2
end
end
:: room3 :: do
local move = io.read ()
if move == "north" then goto room1
elseif move == "east" then goto room4
else
print ("illegal move")
goto room3
end
end
:: room4 :: do
print ("Congratulations, you won!")
end

```

Exercises

Exercise 4.1. Most languages with C-like syntax catfish doesn't offer an **elseif** construct . Why is this construct Is it more needed in Lua than in other languages?

Exercise 4.2. Write four different ways to implement

to create an unconditional loop in Lua. Which one do you prefer curls?

Exercise 4.3. Many people think that **repeat - until** is used rarely and therefore should not be present in minimalist languages like Lua. What do you think about it?

Exercise 4.4. Rewrite the state machine in Listing 4.2. without using **goto** .

Exercise 4.5. Can you explain why Lua has Is there a restriction that you can't jump out of the function? (Hint: How would you implement this feature?)

Exercise 4.6. Assuming goto can jump out of functions, explain that the program in Listing 4.3 should make. (Try to reason about the label using the same rules that are used to describe the the scope of local variables.)

Listing 4.3. Strange (and incorrect) **goto** usage

```
function getlabel ()
return function () goto L1 end
:: L1 ::
return 0
end
function f (n)
if n == 0 then return getlabel ()
else
local res = f (n - 1)
print (n)
return res
end
end
x = f (10)
x ()
```

CHAPTER 5

Functions

Functions are the main mechanism for abstraction of operators and expressions in Lua. Functions can perform a specific task (in other languages this is often called a *procedure* or *subroutine*) or calculate and return values. In the first case, we use you function call as operator; in the second case we use it as expression:

```
print (8 * 9, 9/8)
a = math.sin (3) + math.cos (10)
print (os.date ())
```

In both cases, the argument list is enclosed in parentheses, indicating a call; if the function has no arguments, then we are all equal but must write `()` to indicate a function call. Existence

There is a special exception to this rule: if a function has all one argument and this argument, either a literal (character string) or table constructor, the parentheses are optional:

```
print "Hello World"
<--> print ("Hello World")
dofile 'a.lua'
<--> dofile ('a.lua')
print [[a multi-line
<--> print ([[a multi-line
message]] message]])
f {x = 10, y = 20}
<--> f ({x = 10, y = 20})
type {}
<--> type ({})
```

Lua also offers special syntax for object-directed calls, colon operator. An expression like

`o:foo(x)` is just a way to write `o.foo(o, x)`, that is, to call `o.foo` by adding `o` as an optional argument. In Chapter 16, we discuss we make similar calls (and object-oriented programming in more detail.

A Lua program can use functions written as in Lua and C (or any other language used by the application niy). For example, all functions from the Lua standard library are written sana in C. However, when calling a function, there is no difference between functions written in Lua and functions written in C.

As we saw in other examples, the function definition follows traditional syntax like below:

```
- add the elements of the sequence 'a'
function add (a)
local sum = 0
for i = 1, #a do
sum = sum + a [i]
end
return sum
end
```

In this syntax, the function definition contains a *name* (in the example `add`), a list of *parameters*, and a *body* that is a list of operators

moat.

Parameters work as local variables, initials values of arguments passed when calling the function. You can call a function with a number of arguments, different from its parameter list. Lua will cast the number of arguments to the number of parameters in the same way as it is done in the plural piling: unnecessary arguments are discarded instead of missing ones **nil** is added. For example, consider the following function:

```
function f (a, b) print (a, b) end
```

She has this behavior:

```
f (3) -> 3 nil
```

```
f (3, 4) -> 3 4
```

```
f (3, 4, 5) -> 3 4 (5 is discarded)
```

Although this behavior can lead to errors (easy at runtime), it is also useful, especially for default arguments. For example, consider the following a function that increments the global counter:

```
function incCount (n)
```

```
  n = n or 1
```

```
  count = count + n
```

```
end
```

This function has one default parameter; if we call we put it `incCount ()` without arguments, then it will increment `count` by one. When you call `incCount ()`, Lua first initializes `n` values by **nil**; the `or` operator returns its second argument, and as a result Lua sets the variable `n` to 1.

5.1. Multiple results

A little common, but nevertheless very convenient feature

The thing about Lua is that a function can return multiple values.

Some of the predefined functions in Lua return multiple values. As an example, take the `string.find` function,

which looks for a pattern in a string. This function returns two indices `s` when it finds a pattern: the index of the start of the pattern in the string and the index

end of the pattern. Multiple assignment allows the program get both results:

```
s, e = string.find ("hello Lua users", "Lua")
```

```
print (s, e)
```

-> 7 9

(Note that the index of the first character of the string is 1.)

Functions we write ourselves can also return immediately multiple values by simply listing them after the word **return** . On-example, a function that searches for the maximum element in a sequence value, can return both the maximum element itself and its index:

```
function maximum (a)
local mi = 1
- index of the maximum element
local m = a [mi]
- maximum value
for i = 1, #a do
if a [i]> m then
mi = i; m = a [i]
end
end
return m, mi
end
print (maximum ({8,10,23,12,5})) -> 23 3
```

Lua always lists the number of values returned by a function.

her, to the circumstances of her call. When we call the function like operator, then Lua discards all return values. When we use a function call in an expression, then Lua only stores first value. We get all the values only when the function call is the last (or only) expression in the list expressions. In Lua, these lists appear in four constructs: multiple assignment, function call arguments, const table handler and **return statement** . To illustrate all these cases, we consider the following function definitions:

Multiple results

```
function foo0 () end
- returns nothing
function foo1 () return "a" end - returns 1 value
function foo2 () return "a", "b" end - returns 2 values
```

In multiple assignment, function call as last (or a single) expression uses as many results as how much is needed to match the list of variables:

```
x, y = foo2 ()
- x = "a", y = "b"
x = foo2 ()
- x = "a", "b" is discarded
x, y, z = 10, foo2 ()
```

- x = 10, y = "a", z = "b"

If the function returns no value or returns, but not so much as required, then use

is nil :

x, y = foo0 ()

- x = nil, y = nil

x, y = foo1 ()

- x = "a", y = nil

x, y, z = foo2 ()

- x = "a", y = "b", z = nil

A function call that is not the last item in list, gives exactly one value:

x, y = foo2 (), 20

- x = "a", y = 20

x, y = foo0 (), 20, 30

- x = nil, y = 20, 30 is discarded

When the function call is the last (or only) one argument of another call, then all results of the first call go as arguments to the input of the second call. We have already seen examples of this

constructs with the `print` function . Since the `print` function can receive a variable number of arguments, the operator `print (g ())` prints melts all the values returned by `g` .

```
print (foo0 ())
```

```
->
```

```
print (foo1 ())
```

```
-> a
```

```
print (foo2 ())
```

```
-> ab
```

```
print (foo2 (), 1)
```

```
-> a 1
```

```
print (foo2 () .. "x")
```

```
-> ax (see below)
```

When a call to `foo2` is inside an expression, Lua converts the number of returned values to one; therefore in the last on this string, concatenation only uses "a" .

If we write `f (g (x))` and `f` has a fixed number of arguments, then Lua converts the number of values returned to the number of arguments `f` , as we saw earlier.

The table constructor also uses all values returned by function, without any changes:

```
t = {foo0 ()}
```

```
- t = {} (empty table)
```

```
t = {foo1 ()}
```

```
- t = {"a"}  
t = {foo2 ()}  
- t = {"a", "b"}
```

As always, this behavior occurs only if the call is the last expression in the list; calls anywhere else give exactly one value:

```
t = {foo0 (), foo2 (), 4} - t [1] = nil, t [2] = "a", t [3] = 4
```

Finally, the `return f ()` statement returns all values that returned `f` :

```
function foo (i)  
if i == 0 then return foo0 ()  
elseif i == 1 then return foo1 ()  
elseif i == 2 then return foo2 ()  
end  
end  
print (foo (1))  
-> a  
print (foo (2))  
-> ab  
print (foo (0))  
- (no values)  
print (foo (3))  
- (no values)
```

You can "force" the call to return only one value by enclosing it in an extra pair of parentheses:

```
print ((foo0 ()))  
-> nil  
print ((foo1 ()))  
-> a  
print ((foo2 ()))  
-> a
```

Be careful: the **return statement** does not require parentheses around return value. So an expression like `return (f (x))` always returns exactly one value, regardless of whether how many values the function `f` returns . Sometimes that's exactly what you need; sometimes not.

A special function that returns multiple values is

I wish to set up `table.unpack` . It takes an array as input and returns everything elements of this array, starting at 1:

```
print (table.unpack {10,20,30}) -> 10 20 30  
a, b = table.unpack {10,20,30} - a = 10, b = 20, 30 is discarded
```

An important use of `unpack` is a generic mechanism function call. The generalized mechanism allows you to call any `buoy` function with any arguments dynamically. In ANSI C, for example

measures, there is no way to build a generalized call. You can declare a function that takes a variable number of arguments (with `stdarg.h`) and you can call various functions, using function pointers. However, you cannot call a function with a variable number of arguments: every time you write those in C, you have a fixed number of arguments, and each argument is of a fixed type. In Lua, if you want to call a function `f` with a variable number of arguments from array `a`, you can simply write the following:

```
f (table.unpack (a))
```

The `unpack` function call returns all values from `a` that become arguments to `f`. For example consider the following call:

```
print (string.find ("hello", "ll"))
```

You can dynamically build an equivalent call when the power of the following code:

```
f = string.find
a = {"hello", "ll"}
print (f (table.unpack (a)))
```

Usually `unpack` uses the length operator to find out if how many items should be returned, so it only works with sequences. If necessary, then it can be explicitly limited:

```
print (table.unpack ({ "Sun", "Mon", "Tue", "Wed"}, 2, 3))
```

```
-> Mon Tue
```

Although the `unpack` function is written in C, we can write it in Lua, using recursion:

```
function unpack (t, i, n)
  i = i or 1
  n = n or #t
  if i <= n then
    return t [i], unpack (t, i + 1, n)
  end
end
```

The first time we call it with a single argument `volume`, `i` writes 1, and `n` writes the length sequentially. The function then returns `t [1]` along with all the results `unpack (t, 2, n)`, which in turn returns `t [2]` and all the results `call unpack (t, 3, n)` and so on, stopping after `n` items

Comrade

5.2. Variable functions number of arguments

A Lua function can have an arbitrary number of arguments (*variadic*). For example, we've already called the `print` function with one, two name and a large number of arguments. Although `print` is defined in C, we and in Lua, we can write functions with a variable number of arguments.

As the next example, the function below returns the sum all its arguments:

```
function add (...)  
  local s = 0  
  for i, v in ipairs { ... } do  
    s = s + v  
  end  
  return s  
end  
print (add (3, 4, 10, 25, 12)) -> 54
```

Three dots (`...`) in the parameter list indicate that this function has a variable number of arguments. When we call such a function, Lua collects all of its arguments into a list; we call these assembled arguments *additional arguments* function. Function can access their additional ones again with three dots, now as an expression. In our example-re expression `{...}` gives an array with all the arguments collected. The function iterates over the elements of this array in order to find their amount.

We call expression *... variable expression arguments* (vararg expression). It behaves like a function that returns giving many values, returning all additional arguments current function. For example, the `print (...)` command will print everything before additional arguments to the current function. Similarly, the following the command will create two local variables with the values of the first two additional arguments (or **nil** if there are no such arguments).

```
local a, b = ...
```

In fact, we can mimic the standard mechanism for passing parameters to Lua, translating the following construction

```
function foo (a, b, c)
in
```

```
function foo (...)
local a, b, c = ...
```

For those who like Perl's parameter passing mechanism, this is a like.

The function shown below simply returns all passed arguments:

```
function id (...) return ... end
```

The following function behaves the same as `foo`, except the fact that before calling her, she prints a message with all transmitted arguments:

```
function foo1 (...)
print ("calling foo:", ...)
return foo (...)
end
```

This is a pretty useful trick for keeping track of all calls to this function.

Let's take a look at another useful example. Lua provides

There are separate functions for formatting text (`string.format`) and for writing text (`io.write`). It's quite easy to combine the two functions into one function with a variable number of arguments:

```
function fwrite (fmt, ...)
return io.write (string.format (fmt, ...))
end
```

Note the presence of the `fmt` parameter in front of the periods.

Functions with a variable number of arguments can have any number of number of fixed parameters before the part with variable number scrap parameters. Lua assigns the first values to these variables; the rest (if any) go as additional parameters. Below we we will show several examples of calls and corresponding parameters moat:

Call

Options

`fwrite ()`

`fmt = nil`, no additional parameters

`fwrite ("a")`

`fmt = "a"`, no additional parameters

`fwrite ("% d% d", 4, 5)` `fmt = "% d% d"`, optional 4 and 5

(Note that calling `fwrite ()` will result in an error,

since `string.format` requires a string as its first argument.)

The function can be used to bypass all additional parameters.

use the expression `{...}` to collect them all in a table `tsu`, as we did in the `add` function definition .

In rare cases where the passed arguments may take value of **nil** , table created using `{...}` will not *nastoya*-sequence. For example, there is no way for in order to find out from this table whether there were any arguments at the end of the list.

to **nil** 's. For these cases, Lua offers the `table.pack` function ¹ . This the function takes an arbitrary number of arguments and returns a new a table containing all its arguments, like `{...}` , but this the table will have an additional field `n` containing the total number of its arguments. The following function uses `table.pack` to so that its arguments include **nil** values .

```
function nonils (...)  
  local arg = table.pack (...)  
  for i = 1, arg.n do  
    if arg[i] == nil then return false end  
  end  
  return true  
end  
print (nonils (2,3, nil))  
-> false  
print (nonils (2,3))  
-> true  
print (nonils ())  
-> true  
print (nonils (nil))  
-> false
```

Remember, however, that `{...}` is faster and cleaner than `table.pack` .

5.3. Named arguments

The mechanism for passing parameters in Lua is *positional* : when we call the function, then the correspondence between arguments and form- the minimum parameters is carried out according to their position. First the argument gives the value to the first parameter, and so on. Sometimes, however,

it is useful to specify the parameter by name. To illustrate

To complete this, let's look at the `os.rename` function (from the library `os`), which renames the file. Quite often we forget which name comes first, new or old; so we might want

override this function so that it receives two named parameters:

- wrong

```
rename (old = "temp.lua", new = "temp1.lua")
```

Lua has no direct support for this syntax, but we can achieve the desired effect with a small syntax change. The idea is to collect everything arguments to a table and use this table as the only one function argument. The special syntax that Lua provides is used to call a function, with the table constructor as the only one argument will help us achieve this:

```
rename {old = "temp.lua", new = "temp1.lua"}
```

Accordingly, we override the `rename` function with only one parameter and we get the real arguments from this parameters:

```
function rename (arg)
return os.rename (arg.old, arg.new)
end
```

This method of passing parameters is especially useful when functions There are many arguments and most of them are optional. For example measures, the function that creates a new window in the GUI library can have dozens of arguments, most of which are optional, and it is best to pass them using names:

Listing 5.1. Function with named optional parameters

```
function Window (options)
- check mandatory options
if type (options.title) ~= "string" then
error ("no title")
elseif type (options.width) ~= "number" then
error ("no width")
elseif type (options.height) ~= "number" then
error ("no height")
end
- everything else is optional
_Window (options.title,
options.x or 0, - default value
options.y or 0, - default value
options.width, options.height,
options.background or "white", - default value
options.border - default value false (nil)
)
end
w = Window {x = 0, y = 0, width = 300, height = 200,
title = "Lua", background = "blue",
border = true
```

}

Exercises

Exercise 5.1. Write a function that gets produced any number of strings and returns them concatenated together.

Exercise 5.2. Write a function that receives an array and prints all the elements of this array. Consider the benefits advantages and disadvantages of using `table.unpack` in this function.

Exercise 5.3. Write a function that gets produced any number of values and returns all but the first.

Exercise 5.4. Write a function that receives an array and prints all combinations of elements in this array.

(*Hint* : you can use a recursive formula

for the number of combinations: $C(n, m) = C(n - 1, m - 1) + C(n - 1, m)$).

To get all $C(n, m)$ combinations of n elements in group of size m , you first add the first element to the result and then generate all $C(n - 1, m - 1)$ combinations from the remaining elements in the remaining places. When n is less, than m , there are no more combinations. When m is zero, there is there is only one combination, and it does not use any items.)

CHAPTER 6

More about functions

Functions in Lua are first class values with appropriate lexical scope.

What does it mean that functions are "first class values"?

This means that in Lua, a function is a value that has the same rights as standard values for numbers and strings. We can save functions in variables (local and global) and in tables, we can pass functions as arguments and return them from other functions.

What does "lexical scope" mean for functions? This is a sign

cheat that functions can access variables containing them functions ¹. As we will see in this chapter, this is a kind of harmless property. `v` gives tremendous power to the language because it allows Lua to use many powerful tricks from the world of functional programming. Even if you are not at all interested in functional programming, it is still worth learning a little about how to use take advantage of these opportunities as they can make your Rammu is smaller and simpler.

A somewhat confusing concept in Lua is that functions, like other values, are anonymous; they have no names. When we are talking about a function name like `print`, we mean a belt that contains this function. As with any other variable containing any other value, we can manipulate manipulate these variables in many different ways. Following the example, although somewhat contrived, shows possible examples:

```
a = {p = print}
ap ("Hello World") -> Hello World
print = math.sin - 'print' now refers to sine
ap (print (1)) -> 0.841470
sin = ap
- 'sin' now refers to the print function
sin (10, 20)
-> 10 20
```

(We'll see useful uses of this feature later.)

If functions are values, then do there exist expressions which create functions? Yes. In particular, the standard way create a function in Lua like for example

```
function foo (x) return 2 * x end ,
this is just an example of what we call syntactic sugar ;
it's just a nicer way to write the following code:
```

```
foo = function (x) return 2 * x end
```

Therefore, the definition of a function is actually the operator (assignment), which creates a value of type “function” and assigns it to a variable. We can consider the expression

`function (x) body end` as a function constructor, just like

`{}` is a table constructor. We call the result

the definition of such constructors with *an anonymous function*. Although we are

then we assign functions to global variables, giving them something like a name, there are times when functions remain anonymous.

Let's take a look at some examples.

The `table` library provides the `table.sort` function, which paradise gets a table and sorts its elements. Similar function should provide infinite sort order variations: ascending and descending, numeric or alphabetically, as my key, etc. Instead of trying to provide all possible options `sort` provides an additional parameter which is *an ordering function*: a function that takes two arguments and determines if the first element should be sorted before the second tied list. For example, let's say we have the following record table:

```
network = {  
  {name = "grauna", IP = "210.26.30.34"},  
  {name = "arraial", IP = "210.26.30.23"},  
  {name = "lua", IP = "210.26.23.12"},  
  {name = "derain", IP = "210.26.23.20"},  
}
```

If you want to sort the table by `name` field in reverse alphabetically, then you can simply write:

```
table.sort(network, function(a, b) return (a.name > b.name) end)
```

See how convenient it was to use anonymous function in this statement.

A function that takes another function as an argument is what we call *a higher-order function*. High functions of the following order are a convenient software mechanism, and using anonymous functions to create their functional arguments is a great source of flexibility. However, remember the thread that higher-order functions are not something special, they are simply a consequence of Lua's ability to handle functions like values of the first class.

In order to illustrate the use of the functions of high orders, we will write a simplified definition of higher order function, derivative. Following informal definition, the derivative of the function f at the point x is the value $(f(x + d) - f(x)) / d$ when d becomes infinitesimal. According to and with this definition, we can write an approximate value derivative as follows:

```
function derivative(f, delta)  
  delta = delta or 1e-4  
  return function(x)  
    return (f(x + delta) - f(x)) / delta  
  end
```

end

Having received the function `f`, the call to `derivative (f)` will return an approximate the value of its derivative, which is another function:

```
c = derivative (math.sin)
> print (math.cos (5.2), c (5.2))
-> 0.46851667130038
0.46856084325086
print (math.cos (10), c (10))
-> -0.83907152907645
-0.83904432662041
```

Since functions are first class values in Lua, we can remember them not only in global variables, but also in local variables and table fields. As we will see later, the using functions in table fields is a key component some of the advanced features of Lua, such as modules and object-oriented programming.

6.1. Closures

When we write a function that is enclosed within another function, then it has full access to the local variables of the environment its functions; we call this *lexical scope* (lexical scoping). While this rule of visibility may seem obvious, in fact it is not. Lexical scope together with functions that are first class objects is a very powerful concept in a programming language, but many languages don't support this.

Let's start with a simple example. Suppose you have a list of names students and a table comparing their grades to them; you want to sort list students according to their grades, students with higher grades should come before. You can achieve this by following-in a way:

```
names = {"Peter", "Paul", "Mary"}
grades = {Mary = 10, Paul = 7, Peter = 8}
table.sort (names, function (n1, n2)
return grades [n1]> grades [n2]
— compare grades
end)
```

Now let's say that you want to create a function to solve the given noah tasks:

```
function sortbygrade (names, grades)
table.sort (names, function (n1, n2)
return grades [n1]> grades [n2]
— compare grades
```



```
end)
end
```

An interesting feature in this example is that anonymous the function passed to the `sort` function accesses the parameter `grades`, which is local to the enclosing function `sortBygrade`. Within this anonymous function, `grades` is not neither a global variable nor a local variable, but the fact that we called a *nonlocal variable*. (For historical reasons, for notation for non-local variables in Lua is also used term *upvalue*.)

Why is this so interesting? Because functions are important first class, and so they can *leave* primary education the power of the action of its variables. Consider the following example:

```
function newCounter ()
  local i = 0
  return function ()
    — anonymous function
    i = i + 1
    return i
  end
end
c1 = newCounter ()
print (c1 ()) -> 1
print (c1 ()) -> 2
```

In this code, the anonymous function refers to the non-local re-variable `i` to account for the value. However, by the time we call anonymous function, variable `i` will already leave its scope by validity, since the function that created this variable (`newCounter`) has already finished. However, Lua correctly handles this situation using the concept of a *closure*. Simply put, a closure is a function plus whatever it needs to access non-local variables. If we call `newCounter` again, then it will create a new local variable `i`, so we get a new one a closure working on this new variable:

```
c2 = newCounter ()
print (c2 ()) -> 1
print (c1 ()) -> 3
print (c2 ()) -> 2
```

Thus, `c1` and `c2` are different closures of the same functions, and each uses its own independently instantiated local variable `i`.

In fact, in Lua, the meaning is a closure, not a function.

tion. A function is just a prototype for a closure. However, we will use the term "function" to denote the closure whenever it will not lead to confusion.

Closures prove to be a very handy tool in many cases. As we have already seen, they turn out to be convenient as arguments to higher-order functions such as `sort`. Closures are also useful for functions that build other functions like the `newCounter` function in our example, or a function for finding a derivative; this mechanism allows Lua programs to use advanced techniques from the world of functional programming. Closures are also handy for various *callable functions* (callback). A typical example occurs when you create various buttons in its library to create a GUI. Each button has its own function, which should be called when the user clicks on this button; usually needed for different buttons led to different actions. For example, a calculator you need ten buttons, one for each number. You can co-build with a similar function:

```
function digitButton (digit)
return Button {label = tostring (digit),
action = function ()
add_to_display (digit)
end
}
end
```

In this example, we will assume that `Button` is a function from a library that creates new buttons; `label` is the label of the button; `action` is the closure to be called when the button is pressed. The closure can be noticeable after a long time after `digitButton` is executed, and after local variable `digit` went out of their field of sight, but it is not less, the closure can still access it.

Closures are also useful in a very different case.

Since functions are stored in regular variables, we can it is easy to override functions in Lua, including even standard ones. This possibility is one of the reasons why Lua is so flexible. However when you override the function you still need the old one function. For example, you want to override the `sin` function to work in degrees instead of radians. This new function converts its argument and then calls the original `sin` function to produce

completing work. Your code might look like this below:

```
oldSin = math.sin
math.sin = function (x)
return oldSin (x * math.pi / 180)
end
```

The following is a slightly neater way to accomplish this override:

```
do
local oldSin = math.sin
local k = math.pi / 180
math.sin = function (x)
return oldSin (x * k)
end
end
```

We now store the old version in a local variable; the only way to access it is through a new function.

You can use this same approach to create secure environments, also called *sandboxes*. Safe environments are extremely important when executing code from untrusted sources.

sources such as the Internet. For example, to restrict files, which the program can access, we can override the `io.open` function using closures:

```
do
local oldOpen = io.open
local access_OK = function (filename, mode)
<check access>
end
io.open = function (filename, mode)
if access_OK (filename, mode) then
return oldOpen (filename, mode)
else
return nil, "access denied"
end
end
end
```

What makes this example especially enjoyable is that after that overrides there is absolutely no way for the program to call the original `open`, except through a new version with control. Nebezo- the passable version is stored in a local variable inside the closure, not attainable from the outside in any way. With this approach, you can build sandboxes for Lua on Lua itself, while receiving as advantages simplicity and flexibility. Instead of some kind of one-stop solution for

all problems Lua provides a meta mechanism, so you can adjust your environment to your goals.

6.2. Non-global functions

The obvious consequence of the fact that functions are values first class, is that we can save functions not only in global variables, but also in local variables-data and fields of the table.

We have already seen various examples of functions stored in fields. tables: most Lua libraries use this mechanism (by example, `io.read` , `math.sin`). To create similar functions in Lua we just need to combine the standard syntax for functions with syntax for tables:

```
Lib = {}  
Lib.foo = function (x, y) return x + y end  
Lib.goo = function (x, y) return x - y end  
print (Lib.foo (2, 3), Lib.goo (2, 3)) -> 5 -1
```

Of course, we can also use constructors:

```
Lib = {  
  foo = function (x, y) return x + y end,  
  goo = function (x, y) return x - y end  
}
```

Moreover, Lua also provides another syntax for similar functions:

```
Lib = {}  
function Lib.foo (x, y) return x + y end  
function Lib.goo (x, y) return x - y end
```

When we store a function in a local variable, we get we tea a *local function* , that is, a function with a limited scope visibility. Definitions like this are especially handy for packages: since Lua treats each block as a function, a block can Can define local functions that are only visible from the block. Lexical scoping ensures that other functions from package can use these local functions:

```
local f = function ( <params> )  
  <body>  
end  
local g = function ( <params> )  
  <some code>  
  f () - 'f' is visible here  
  <some code>
```

end

Lua also supports the following syntactic sugar for local functions:

```
local function f ( <params> )  
<body>  
end
```

When defining recursive local functions, one gets subtlety. The naive approach doesn't work here. Consider the following definition:

```
local fact = function (n)  
  if n == 0 then return 1  
  else return n * fact (n-1) - error  
end  
end
```

When Lua compiles the `fact (n-1)` call in the body of a function, then the local function `fact` has not yet been defined. Therefore, this definition will try to call the global function `fact`, not the local one.

We can solve this problem by first defining a local variable and then already defining the function itself:

```
local fact  
fact = function (n)  
  if n == 0 then return 1  
  else return n * fact (n-1)  
end  
end
```

Now `fact` inside the function refers to a local variable.

Its value at the moment the function is defined does not mean anything; by the time the function is executed, it will already receive correct value.

When Lua "reveals" its syntactic sugar for the local function, it does not use the "naive" way. Instead, define division as shown below:

```
local function foo ( <params> ) <body> end ,  
goes into
```

```
local foo; foo = function ( <params> ) <body> end
```

Therefore, we can safely use this syntax to recursive functions.

Of course, this trick won't work unless you have direct recursion - the two functions call each other. In such cases, you need to explicitly write put the appropriate descriptions of local variables:

```
local f, g - described local variables  
function g ()
```

```

<some code> f () <some code>
end
function f ()
  <some code> g () <some code>
end

```

In this example, the function `f` can not write a local function `f` , because in such a case Lua will create a new local variable- the new `f` , leaving the old (referenced by `g`) uninitialized- Noah.

6.3. Tail optimization calls

Another interesting feature of functions in Lua is that Lua performs tail call optimizations. (It means that Lua supports *tail recursion* optimization , although it is not directly related to recursion here, see exercise- (see section 6.3.)

The tail call is actually a **goto** that looks like function call. A tail call happens when a function is called has another function as its last action. For example, in the following In the following code, the call to the function `g` is tail:

```
function f (x) return g (x) end
```

After `f` calls `g` , it has nothing else to do. In similar situations, the program does not need to return to the calling function when the nested call has completed. Therefore after tail call the program does not need to store any information about the calling function on the stack. When the call to `g` ends it seems that control goes directly to the point where it was named `f` . Some language implementations, such as the Lua interpreter, use this fact and do not allocate additional stack space for tail call. We say that these implementations are supported by *eliminating tail calls* (tail-call elimination).

Since tail calls do not use stack space, number of nested tail calls that the program can execute, it is simply not limited by anything. For example, we can you-call the next function, passing any number as an argument cop:

```
function foo (n)
if n > 0 then return foo (n - 1) end
end
```

This call will never result in a stack overflow.

The subtle point in eliminating tail calls is the question of what is a tail call. For some quite obvious candidates require that a defiant the function does nothing else after the call, it is not executed. On-example, in the following code, the function call `g` is not tail.

```
function f (x) g (x) end
```

The problem in this example is that after calling `g`, the function `f` should discard the results of `g` before returning. Similarly, all the following the next calls also do not satisfy the condition:

```
return g (x) + 1 - addition is required
return x or g (x) - must be converted to 1 value
return (g (x))
- must be converted to 1 value
```

In Lua, only a call like `return func (args)` is tail.

However, both *func* and its arguments can be complex expressions, as Lua will execute them before calling. For example the following the call is tailored:

```
return x [i] .foo (x [j] + a * b, i + j)
```

Exercises

Exercise 6.1. Write a function `integral` that gets function `f` and returns an approximate value of its integral `rala`.

Exercise 6.2. In Exercise 3.3, you had to write a function the function that receives the polynomial (represented by the table) and the value of the variable and returns the value of the polynomial for this variable. Write a function that receives a lot of `goclen` and returns a function that, when called for any `x` value, will return the polynomial value for this `x`. For instance:

```
f = newpoly ({3, 0, 1})
print (f (0))
-> 1
print (f (5))
-> 76
```

```
print (f (10))  
-> 301
```

Exercise 6.3. Sometimes the language supporting optimization tail calls is called *supporting tail recursion* (properly tail recursive), if the optimization of tailings out calls is only supported for recursive calls.

wwii (Without recursive calls, the maximum call depth is wow is statically defined.)

Show that this is not true in a language like Lua: write a program that implements unlimited call depth wow without using recursion. (*Hint* : see section 8.1.)

Exercise 6.4. As we have seen, the tail call is a a clipped **goto** . Using this idea, rewrite the code for the maze games from section 4.4 using tail calls. Each block should become a new function and each **goto** becomes a tail call.

CHAPTER 7

Iterators

and the generalized for

In this chapter, we will show you how to write iterators for generalized **for** (general **for** operator). Starting with simple iterators, we learn how to use the full power of generic for to write simpler and more efficient iterators.

7.1. Iterators and Closures

An *iterator* is any construct that allows you to iterate over host elements of the set. In Lua, we usually represent iterators by *help functions*: every time we call a function, it will rotate the "next" item from the set.

Any iterator has to save its state somewhere between challenges to know where he is and how to proceed.

Closures are an excellent mechanism for this task.

Recall that a closure is a function that refers to one or several local variables from your environment.

These variables retain their values between successive calls to the closure, thereby helping the closure understand where it is on its way. Of course, to create a new closure we also have to create non-local variables for it. Therefore, the structure of a closure usually includes two functions at once: itself closure and *factory*, a function that creates a closure along with surrounding variables.

As an example, let's write a simple iterator for a list.

Unlike `ipairs`, this iterator will not return the index of every the same element, but only its value:

```
function values (t)
  local i = 0
  return function () i = i + 1; return t [i] end
end
```

In this example, `values` is a factory. Every time we call

We use this factory, it creates a new closure (iterator). This is

The *kaney* stores its state in its external variables `t` and `i`. Each

Every time we call this iterator, it returns the following

value from list `t`. After the last element, the iterator will return **nil**, which marks the end of iterations.

We can use this iterator in a **while loop** :

```
t = {10, 20, 30}
iter = values (t)
- create an iterator
while true do
  local element = iter ()
  - call the iterator
  if element == nil then break end
```

```
print (element)
end
```

However, it is much easier to use the generic **for** statement .

After all, it was created for exactly this kind of iteration-
niya:

```
t = {10, 20, 30}
for element in values (t) do
  print (element)
end
```

The generic **for** does all the behind-the-scenes work for the iteration rationing: it stores the iterating function inside, so we don't needs a variable `iter` , it calls the iterator for each new iteration walkie-talkie, and it completes the iteration when the iterator returns **nil** . (We'll see in the next section that the generic **for** does even more than that.)

As a more advanced example, see Listing 7.1 for re-fetch all words from the current input file. For such a search, we two values are needed: the content of the current line (variable `line`) and where we are inside this line (variable `pos`). With these data, we can always generate the next word. Basic

The most important part of the iterating function is the call to `string.find` . This call searches for a word in the current line, starting at the current position. He described

matches a 'word' using the pattern `'% w +'` , which one or more alphanumeric characters. If this call finds word, then the function updates the current position by the first character after a word and returns that word ₁ . Otherwise, the iterator reads the following-string and repeats the search. If there are no more rows, it returns **nil** to signal the end of the traversal.

Despite its complexity, using `allwords` is extremely easy then:

```
for word in allwords () do
  print (word)
end
```

This is a typical situation with iterators: they may not be so easy to write, but they are easy to use. This is not a problem, much more often, end users who program in Lua do not write their own iterators, but use the iterators provided by application.

Listing 7.1. Iterator to iterate over all words from the input file
function `allwords ()`

```

local line = io.read ()
- current line
local pos = 1
- current position in the line
return function ()
- iterator function
while line do
— repeat as long as there are lines
local s, e = string.find (line, "% w +", pos)
if s then
— have you found the word?
pos = e + 1
— next position after word
return string.sub (line, s, e) — return word
else
line = io.read () - word not found; trying the trail. string
pos = 1
- start at the beginning of the line
end
end
return nil - no more lines, end of traversal
end
end

```

7.2. Generic for semantics

One of the drawbacks of the iterators discussed above is that that we need to create a new closure for initialization each new cycle. For most cases, this is not problem. For example, in the case of the `allwords` iterator, the creation price is a single closure is incomparable to the cost of reading an entire file. but in some situations this can be significant. Such cases we can use the generic **for itself** to store states. In this section, we will see what possibilities for storing state is offered by the generic **for** .

We have seen that a generic **for** loop stores iterables during a loop. function within itself. It actually stores three values: iterated function, *unchangeable state* (invariant state) and *control variable* . Now let's get down to the details.

The syntax **for a generic for is** shown below:

```

for <var-list> in <exp-list> do
<body>
end

```

Here *var-list* is a list of one or more re-variables separated by commas, and *exp-list* is a list of one or multiple expressions, also separated by commas. Often a list expression consists of a single element, a call to the iterate factory rators. In the following code, for example, the list of variables is *k* , *v* , and the list of expressions consists of a single element, *pairs (t)* :

```
for k, v in pairs (t) do print (k, v) end
```

Often, a variable list also consists of just one variable - *noah*, as in the following loop:

```
for line in io.lines () do  
io.write (line, "\n")  
end
```

We name the first variable in the *control variable list-noah* . During the whole cycle, its value is not equal to **nil** , because when it becomes **nil** , the loop ends.

The first thing a **for** loop does is compute the values expressed by following **in** . These expressions must give three meanings, used by the **for** statement : iterating function, immutable the state and initial value of the manipulated variable. As in multiple assignment, only the last (or the only nth) list item can give more than one value; and the number of these values are reduced to three, excess values are discarded, together then the missing **nil** 's are added . (When we use simple iterators, the factory only returns the iterating function, this invariant state and control variable get the value is **nil** .)

After this initialization, **for** calls the iterating function with two arguments: an invariant state and a control transfer change. (From the point of view of the **for** statement , this is an invariant state doesn't make any sense at all. The **for** statement only passes the value reading the state from the initialization step to calling the iterating function **for**) then assigns the values returned by the iterating function, variables declared in the variable list. If a the first value (assigned to the manipulated variable) is **nil** , then the loop ends. Otherwise, **for** executes its body and calls again an iterating function, repeating the process.

More precisely the construction of the view

```
for var_1, ..., var_n in <explist> do <block> end
```

equivalent to the following code:

```
do
```

```

local _f, _s, _var = <explist>
while true do
local var_1, ..., var_n = _f(_s, _var)
_var = var_1
if _var == nil then break end
<block>
end
end

```

So if our iterating function is f , the immutable state s and the initial state for the control variable is a_0 , then the control variable will run through the following values $a_1 = f(s, a_0)$, $a_2 = f(s, a_1)$, etc., until a_i equals **nil**. If **for** has other variables, then they just get an extra solid values returned by f .

7.3. Stateless iterators

As its name implies, such an iterator does not store any state. Therefore, we can use the same a stateless iterator in many loops, thus avoiding creating new closures.

As we have seen, the **for** statement calls the iterating function with two arguments: immutable state and control transfer changeable. A stateless iterator builds the next element of the loop, using only these two values. A typical example of this iteration is the `ipairs` , which iterates over all the elements of the array:

```

a = {"one", "two", "three"}
for i, v in ipairs(a) do
print(i, v)
end

```

The state of this iterator is the table we are iterating over (an immutable state that retains its value throughout loop), and the current index (control variable). And `ipairs` (factory) and the iterator itself is very simple, we could write them to Lua as follows:

```

local function iter(a, i)
i = i + 1
local v = a[i]
if v then
return i, v
end
end

```

```
function ipairs (a)
return iter, a, 0
end
```

When Lua is `ipairs (a)` for a cycle **for** , she gets three values: an iterating function `iter` , `a` as an invariant state value and zero as the initial value for the control change. Lua then calls `iter (a, 0)` , which gives `1, a [1]` (unless `a [1]` is no longer **nil**). The next iteration calls `iter (a, 1)` , which returns `2, a [2]` , and so on up to the first **nil** element .

The `pairs` function , which iterates over all the elements in a table, looks like `ms`, except that the iterating function is a function `next` , which is a standard Lua function:

```
function pairs (t)
return next, t, nil
end
```

Calling `next (t, k)` , where `k` is the key of table `t` , returns the following - key in the table in no particular order, and also associated with this key the value as the second return value. Call `next (t, nil)` returns the first pair. When there are no more pairs, then `next` returns **nil** .

Some people prefer to use `next` explicitly , avoiding calling

```
pairs :
for k, v in next, t do
<loop body>
end
```

Recall that `for` casts its list of expressions to three values. `niyam`, which are `next` , `t` and **nil** ; this is exactly what what happens when calling `pairs` .

An iterator for traversing a linked list is another interesting one an example of a stateless iterator. (As we have already mentioned, the connected lists are not common in Lua, but sometimes we need them.)

```
local function getnext (list, node)
if not node then
return list
else
return node.next
end
end
function traverse (list)
return getnext, list, nil
end
```

Here we use the beginning of the list as an invariant state (second value returned by `traverse`) and the current node as

ve control variable. When the iterated function `getnext` will be called the first time, `node` will be **nil**, and therefore the function will return `list` as the first node. In subsequent calls, `node` will already be is not **nil** and therefore the iterator will return `node.next` as expected.

As usual, using this iterator is extremely simple:

```
list = nil
for line in io.lines () do
  list = {val = line, next = list}
end
for node in traverse (list) do
  print (node.val)
end
```

7.4. Iterators with complex condition

Often an iterator needs to store more state than is placed in the variables of the invariant state and the control variable. The simplest solution is to use a kaniy. An alternative solution would be to pack everything you need an iterator into a table and use that table as invariant state for the loop. Using a table, an iterator can store as much data as he needs. Moreover, he can change these data the way he wants. Although the state is the same all the time table (therefore it is invariant), the contents of the table can be run throughout the cycle. Since such iterators store everything their data is in a state, they usually ignore the second argument, provided by the generic **for** loop (loop variable).

As an example of this approach, we will rewrite the iterator `allwords`, which bypasses all words in the input file. This time we will store its state in a table with two fields: `line` and `pos`.

The function that starts the loop is pretty simple. She must ver-
Introduce an iterating function and an initial state:

```
local iterator
- to be determined later
function allwords ()
  local state = {line = io.read (), pos = 1}
  return iterator, state
end
```

The iterator function does most of the work :

```
function iterator (state)
```

```

while state.line do
- repeat while there are lines
- looking for the next word
local s, e = string.find (state.line, "% w +", state.pos)
if s then
- found the word?
- update position
state.pos = e + 1
return string.sub (state.line, s, e)
else - word not found
state.line = io.read () - trying the next line ...
state.pos = 1
-- ... first
end
end
return nil
- there are no more lines, we end the loop
end

```

Whenever possible, you should try to write an iteration-stateless **for** loop, those that keep their entire state in the **for** loop variables. With them, you do not create new objects when start the cycle. If this model does not fit, then you should ask to create closures. Also, it's prettier, closure is usually is more efficient as an iterator than a table:

First, it is cheaper to create a closure than a table; secondly, accessing non-local variables is faster than accessing table fields.

We will see another way to write iterators later, using by using coroutines. This is the most powerful solution, but it somewhat more expensive.

7.5. Genuine iterators (true iterarators)

The term "iterator" is somewhat inaccurate, since it actually iterates. It iterates not an iterator, but a **for** loop. Iterators only provide consecutive values to iterate over. Maybe more a good term would be "generator", but the term "iterator" is already became widespread in languages like Java.

However, there is another way to construct iterators, where iterators actually do iteration. When we use we use such iterators, we don't write a loop; instead we just then we call an iterator with an argument describing that the iterator should to do at each iteration. More precisely, the iterator receives as ve argument is a function that it calls inside its loop.

As an example, let's rewrite the `allwords` iterator again

using this approach:

```
function allwords (f)
for line in io.lines () do
for word in string.gmatch (line, "%w+") do
f (word)
- call function
end
end
end
```

To use this iterator, we just have to provide
inject the body of the loop as a function. If we just want to print each
after a word, then we use `print` :

```
allwords (print)
```

Often an anonymous function is used as the loop body. On-
example, the following code snippet counts how many times the word "hello"
occurs in the file:

```
local count = 0
allwords (function (w)
if w == "hello" then count = count + 1 end
end)
print (count)
```

The same task written using iterators previously
the style looked at is not much different:

```
local count = 0
for w in allwords () do
if w == "hello" then count = count + 1 end
end
print (count)
```

Similar iterators were very popular in older versions

Lua, when the language did not yet have a **for** statement . How do they compare
with iterators of the previously discussed style? Both styles have a
roughly the same overhead: one function call per iteration
walkie-talkie. On the other hand, it is easier to write an iterator using
genuine iterators (although we can get the same ease
using coroutines). On the other hand, previously considered
the new style is more flexible. First, it allows two or more pa-
rallel iteration. (For example, consider the case of bypassing
two files at once, comparing them word by word.) Second, it
allows the use of **break** and **return** inside a loop. With authentic
return iterators **return** from anonymous function, but not
from the loop. Therefore, I usually use traditional (i.e.
seen earlier) iterators.

Exercises

Exercise 7.1. Write an iterator `fromto` such that you follow the two loops are equivalent:

```
for i in fromto (n, m)
  <body>
end
for i = n, m
  <body>
end
```

Can you implement this with an iterator without co-standing?

Exercise 7.2. Add a step parameter to the previous exercise. `neniyu`. Can you still implement this with soup of a stateless iterator?

Exercise 7.3. Write a `uniqewords` iterator that can rotates all words from the given file without repetitions. (*Hint* : start with the `allwords` code in Listing 7.1; use use a spreadsheet to store all the words you have already zeros.)

Exercise 7.4. Write an iterator that returns everything non-empty substrings of the given string. (You will need a function `ration string.sub .`)

CHAPTER 8

Compilation, execution and errors

Although we call Lua an interpreted language, Lua is always pre-compiles the source code into an intermediate form before ex-filling. (In fact, many interpreted languages do ditto.) Having a compile file may sound strange in relation to an interpreted language like Lua. However, the main the feature of interpreted languages is not that they are not compiled. are copied, but what is possible (and easy) to execute the generated

code on the fly. We can say that having a function like `dofile` is this is what allows Lua to be called an interpreted language.

8.1. Compilation

Earlier, we introduced `dofile` as a kind of primitive operation for executing blocks of Lua code, but `dofile` is actually an aid gateway function: all the hard work is done `LoadFile`. Like `dofile`, `loadfile` loads a block of Lua code from a file, but it doesn't complete this block. Instead, it only compiles this block and returns the compiled block as a function. Moreover, in unlike `dofile`, `loadfile` does not throw an error, it just returns error code so that we can handle these errors ourselves. We can Let's define `dofile` like this:

```
function dofile (filename)
local f = assert (loadfile (filename))
return f ()
end
```

Note the use of the `assert` function in order to *cause errors* (raise error), if `loadfile` fulfills a mistake.

For simple tasks, `dofile` is convenient because it does all the work. but in one call. However, `loadfile` is more flexible. In case of error `loadfile` returns **nil** and an error message, which allows us to handle the error in a convenient way. Moreover, if we need but execute the file several times, then we can call once `loadfile` and call the function it returns several times. This the approach is much cheaper than calling `dofile` multiple times, since the file is compiled only once.

The `load` function is similar to `loadfile`, except that it takes a block of code not from a file, but from line 1. For example, consider the following

next line:

```
f = load ("i = i + 1")
```

After executing this code, `f` will be a function that executes-
em `i = i + 1` when called:

```
i = 0
```

```
f (); print (i) -> 1
```

```
f (); print (i) -> 2
```

The `load` function is quite powerful and we have to use it with caution. It is also an expensive feature (compared to some

alternatives) and can lead to code that is very heavy to understand. Before you use it, make sure there is no more simple way to solve the problem.

If you want a quick and dirty `dostring` (i.e. load and execute block), you can directly use the result

```
tat load :
```

```
load (s) ()
```

However, if there is at least one syntax error, then `load` will return **nil** and the final error will be something like “*attempt to call a nil value*”. For clearer error handling use

```
assert :
```

```
assert (load (s)) ()
```

There is usually no point in calling the `load` function on a literal (that is, an explicitly quoted string). For example, the following two lines are roughly equivalent:

```
f = load (“i = i + 1”)
```

```
f = function () i = i + 1 end
```

However, the second line is much faster since Lua compiled functions together with the surrounding block. On the first line, the `load` call includes a separate compilation.

Since `load` does not compile lexically aware actions, then the two lines considered may not be entirely equivalent. To see the difference, let's change it slightly example:

```
i = 32
```

```
local i = 0
```

```
f = load (“i = i + 1; print (i)”)
```

```
g = function () i = i + 1; print (i) end
```

```
f () -> 33
```

```
g () -> 1
```

The `g` function works on the local variable `i` as expected, however the function `f` works with the global `i`, since `load` is always compiles its blocks in the global environment.

The most typical use of `load` is to execute external code, that is, snippets of code coming from outside your programs. For example, you may want to plot the function the option specified by the user; the user enters the function code, and you then use `load` to execute it. Pay

note that `load` expects to receive a block, that is, statements. If a you want to evaluate the expression, then you can append to the beginning return expressions, which will give you a statement that returns a value

given expression. Let's look at an example:

```
print "enter your expression:"
local l = io.read ()
local func = assert (load ("return" .. l))
print ("the value of your expression is" .. func ())
```

Since the function returned by `load` is a normal function, you can call her many times:

```
print "enter function to be plotted (with variable 'x'):"
local l = io.read ()
local f = assert (load ("return" .. l))
for i = 1, 20 do
  x = i - global 'x' (to be visible from outside the block)
  print (string.rep ("*", f ()))
end
```

(The `string.rep` function repeats a string a specified number of times.)

We can also call the `load` function by passing it as argument *reading function* (reader function). Reading function can return a block of code in parts; `load` calls this function before it's time until it returns **nil**, indicating the end of the block. For example measures, the following call is equivalent to `loadfile` :

```
f = load (io.lines (filename, "*" L"))
```

As we will see in chapter 22, the challenge `io.lines (filename, "*" L")` WHO-rotates a function that, when called, returns the following line from file `2` . So `load` will read the block from file line by line. The next option is similar, but more efficient effective:

```
f = load (io.lines (filename, 1024))
```

Here the iterator returned by `io.lines` reads in blocks at 1024 bytes.

Lua treats each independent block as the body of an anonymous variable function with a variable number of arguments. For instance, `load ("a = 1")` returns an analog of the following function:

```
function (...) a = 1 end
```

Like any other function, blocks can define their own locale. variable variables:

```
f = load ("local a = 10; print (a + 20)")
f () -> 30
```

Using these capabilities, we can rewrite our example with plotting the graph so as not to use the global re-variable `x` :

```
print "enter function to be plotted (with variable 'x'):"
local l = io.read ()
local f = assert (load ("local x = ...; return" .. l))
```

```
for i = 1, 20 do
print (string.rep ("*", f (i)))
end
```

We put the description “local x = ...” at the beginning of the block to define divide x as a local variable. When we call f with an argument i , this argument becomes the value of expression

The load function never raises an error, in case of an error it just returns **nil** and an error message:

```
print (load ("ii"))
-> nil [string "ii"]: 1: '=' expected near 'i'
```

Moreover, these functions have no side effect. They only compile the block to an internal representation and return-give the result as an anonymous function. A common mistake is that it is assumed that block loading determines the function tions (defined in this block). In Lua, function definitions are assignment; and as such they happen at runtime, not at compile time. For example, let's say we have a file foo.lua with the following content:

```
- file 'foo.lua'
function foo (x)
print (x)
end
```

Then we run the command

```
f = loadfile ("foo.lua")
```

After this command, foo has been compiled but not yet defined.

To define it, we must execute a block:

```
print (foo) -> nil
f () - defines 'foo'
foo ("ok") -> ok
```

In serious programs that need to execute external code, you have to handle all errors that occur when loading block. Moreover, you may want to start a new unit in the protection environment to avoid unpleasant side effects
Comrade We'll discuss environments in detail in Chapter 14.

8.2. Precompiled code

As I mentioned at the beginning of this chapter, Lua precompiles the running code before executing it. Lua also allows distribution take the code in precompiled form.

The simplest way to get a precompiled file is

also called a *binary block* in Lua - is the use of the `luac` program included in the standard delivery. For instance, next call creates `prog.lc` file with precompiled version of the `prog.lua` file :

```
$ luac -o prog.lc prog.lua
```

The interpreter can then execute this file as a normal Lua code, working in the same way as with the original file:

```
$ lua prog.lc
```

Lua allows for precompiled code pretty much wherever it allows source code. In particular, `loadfile` and `load` take the input is also precompiled code.

We can write a simple `luac` replacement directly to Lua:

```
p = loadfile (arg [1])
f = io.open (arg [2], "wb")
f: write (string.dump (p))
f: close ()
```

The main function here is `string.dump` : it gets the function `lua` and returns its precompiled code as a string, properly formatted to load it into Lua.

The `luac` program also introduces some interesting options.

In particular, the `-l` option prints a list of all opcodes that the compiler generates for the given block. As an example, `foxes Thing 8.1` contains the output of the `luac` program launched with the `-l` option , for the following one line file:

```
a = x + y - z
```

(We will not discuss Lua internals in this book; if you interested in information about these opcodes, then search the Internet the words "lua opcode" will give you reasonably accurate information.)

Precompiled code is not always smaller than the original code, but it loads faster. Another plus is that this gives you protection against accidental changes to the source. However, in deviation from source code, maliciously modified binary code can crash the Lua interpreter or even execute user-supplied machine code. At startup, the usual There is nothing to worry about with this code. However, you should avoid using running untrusted code in precompiled form. The function `load` has a special option for just this task.

Listing 8.1. Sample `luac -l` output

```
main <stdin: 0,0> (7 instructions, 28 bytes at 0x988cb30)
```

0+ params, 2 slots, 0 upvalues, 0 locals, 4 constants, 0 functions

1 [1] GETGLOBAL 0 -2; x

2 [1] GETGLOBAL 1 -3; y

3 [1] ADD 0 0 1

4 [1] GETGLOBAL 1 -4; z

5 [1] SUB 0 0 1

6 [1] SETGLOBAL 0 -1; a

7 [1] RETURN 0 1

In addition to the required first argument, `load` has three more optional arguments. The second argument is the name of the block that the swarm is only used in error messages. The fourth argument `cop` is the environment, we will look at it in detail in chapter 14. Third the argument is exactly what we are interested in now; he controls what types of blocks can be loaded. If this argument is present, then it must be a string: the string `"t"` allows loading only text blocks, line `"b"` allows loading only binary (precompiled) blocks, and the string `"bt"` (the value of the default) allows you to load blocks of both types.

8.3. C code

Unlike Lua code, C code must be forged with the app before using it. In a number of popular operating systems, the easiest way to do this is using the dynamic linking feature. However, given This feature is not part of the ANSI C specification; so no real wearable way to do this.

Lua usually does not include features that cannot be implemented in ANSI C. However, with dynamic linking, tion is different. We can consider it as the basis of all others opportunities: as soon as we have it, we can immediately load reap any feature that is not currently in Lua. Therefore, in this Otherwise, Lua abandons portability rules and implements dynamic linking for a number of platforms. Standard implementation zation offers this feature for Windows, Mac OS X, Linux, FreeBSD, Solaris and most other UNIX implementations. Transfer this capability to other platforms doesn't have to be difficult; refer to your distribution. (To check this, run `print (package.loadlib ("a", "b"))` from the Lua command line and then look at the result. If it reports a file that does not exist, then

you have dynamic linking support. Otherwise an error message will tell you that this feature is not supported or not installed.)

Lua provides all the dynamic linking capabilities through one function, `package.loadlib`. This function receives two string arguments: full library path and function name from this library. Therefore, her typical call looks like but below:

```
local path = "/usr/local/lib/lua/5.1/socket.so"
local f = package.loadlib (path, "luaopen_socket")
```

The `loadlib` function loads the specified library and connects her to Lua. However, it does not call the specified function. Instead he returns a C function as a Lua function. In case of an error when loading a library or finding an initializing function `loadlib` returns **nil** and an error message.

The `loadlib` function is very low level. We must re-Pass the full path to the library and the correct function name (including teas, beginning underscores added by the compiler). Often we load libraries to C using `require`. This function looks for library and uses `loadlib` to load the initialization functions for the library. When called, this initializing function ration builds and returns a table with functions from this library, like does the usual Lua library. We will discuss `require` in section 15.1 and See section 27.3 for more information on C libraries.

8.4. Errors

Errare humanum est ³. Therefore, we must handle errors like this good as we can. Since Lua is an extension language, often embedded in an application, we cannot just fall or exit if an error occurs. Instead, when yes an error occurs, Lua interrupts the execution of the current block and returns control to the application.

Any unexpected situation Lua encounters will cause leads to *call errors* (raises an error). Errors occur when you (more precisely, your program) cannot add values that are not numbers, do not index a table, etc. (You can want to change this behavior using *metatables* as we will see

later.) You can also explicitly cause errors by calling `error` function with an error message as an argument. Usually this function is the correct way to report a bug in your code:

```
print "enter a number:"
n = io.read ("* n")
if not n then error ("invalid input") end
3
```

Humans tend to make mistakes. (*lat* .)

This way of calling `error` is so common that for this

Lua has a built-in `assert` function :

```
print "enter a number:"
n = assert (io.read ("* n"), "invalid input")
```

The `assert` function checks if its first argument is valid is not false, and simply returns this argument; if the argument is false, then `assert` raises an error. Its second argument, the error message, not required. However, keep in mind that `assert` is a normal function. As with all functions, Lua always computes before calling it. passes her arguments. So if you write something like

```
n = io.read ()
assert (tonumber (n), "invalid input:" .. n .. "is not a number")
```

then Lua will always do the concatenation, even if `n` is a number. Therefore- In such cases, it may be better to use an explicit test.

When a function detects an unexpected situation (*excluding reading*), it can go in two ways: return an error code (usually

nil) or raise an error using `error` . There are no tough

rules for choosing between these two options, but we can

offer general advice: an exception that's easy to get around, should throw an error; otherwise, an error code should be returned.

For example, let's consider the `sin` function . How should she behave

lead if its argument is a table? Suppose it is possible

rotates the error code. If we need to check for errors, then we

we can write something like

```
local res = math.sin (x)
if not res then
```

- error?
<error-handling code>

However, we can easily change this exception *before* calling functions:

```
if not tonumber (x) then
- x is not a number?
<error-handling code>
```

Often we do not check for either the argument or the result of calling `sin`, if the argument is not a number, it means that something is wrong in our program. In a similar situation, stop computing and call error is the simplest and most practical way to handle this exception.

On the other hand, let's look at the `io.open` function, which Paradise opens the file. How to behave if asked to open non-existent file? In this case, there is no simple way Soba to check for a call for an exception before calling this function tion. On many systems, the only way to check that file exists, is to try to open it. So if `io.open` cannot open the file for some external reason (for-example, "file does not exist" or "no rights!"), then it just might rotates **nil** along with the error message. In this case, you have a chance to handle the situation in an appropriate way, such as asking with a different filename:

```
local file, msg
repeat
print "enter a file name:"
local name = io.read ()
if not name then return end — nothing was entered
file, msg = io.open (name, "r")
if not file then print (msg) end
until file
```

If you do not want to handle a similar situation, but do not want to be safe anyway, you can simply use

```
use assert :
file = assert (io.open (name, "r"))
```

This is a typical idiom for Lua: if `io.open` completed with The error coy, then `assert` will raise an error.

```
file = assert (io.open ("no-file", "r"))
-> stdin: 1: no-file: No such file or directory
```

Note how the error message that is the second result of `io.open` turns out to be the second argument when call to `assert`.

8.5. Error processing and exceptions

For many applications, you do not need to do any processing. A lot of errors in Lua; all processing is done by the application itself. All work Lua begins by being called by an application, usually consisting of block execution. If an error occurs, then this call returns error code and the application can respond appropriately. In the case of a separate interpreter, its main loop is simply prints an error message and continues.

However, if you need to handle errors in Lua, then you should use `pcall` function (protected call) to encapsulate your code.

Let's say you want to execute a piece of Lua code and catch any error that occurs while executing it. Your first step is to wrap this piece of code in a function; quite an hour then anonymous functions are used for this. Then you call this function using `pcall` :

```
local ok, msg = pcall (function ()
<some code>
if unexpected_condition then error () end
<some code>
print (a [i]) — possible error: 'a' may not be a table
<some code>
end)
if ok then — no error occurred while executing the protected code
<regular code>
else
— the protected code caused an error; process here
<error-handling code>
end
```

The `pcall` call invokes its first argument in *protected mode*, so that all errors during the execution of the function are caught. If there are no errors, then the `pcall` call returns **true** and all values returned by the function. Otherwise, it returns **false** and error message.

Despite its name, the error message need not be by line: calling `pcall` will return whatever Lua value you pass—whether error .

```
local status, err = pcall (function () error ({code = 121}) end)
print (err.code) -> 121
```

These mechanisms provide everything you need to process the data. key in Lua. We throw an exception with `error` and

intercept it using `pcall`. The error message is specifies the type of error.

8.6. Error messages and call stack

Although we can use the sign as an error message any type, usually error messages are strings that describe wondering what went wrong. In the event of an internal errors (e.g. trying to index a non-table) message Lua generates an error; otherwise the error message becomes the value passed to the `error` function. When the error message is a string, then Lua tries to add some information about the place where the error occurred:

```
local status, err = pcall (function () a = "a" +1 end)
print (err)
-> stdin: 1: attempt to perform arithmetic on a string value
local status, err = pcall (function () error ("my error") end)
print (err)
-> stdin: 1: my error
```

The error message contains the file name (in the example it is `stdin`) and the line number in it (in the example it is 1).

The `error` function has a second optional parameter which tells the *level* where to report the error; you are using this as a yardstick to blame someone else for a mistake. For instance, you wrote a function that immediately checks that it was duly called:

```
function foo (str)
if type (str) ~= "string" then
error ("string expected")
end
<regular code>
end
```

Then someone calls your function with the wrong argument-Tom:

```
foo ({x = 1})
```

In this case, Lua points to your function - after all, it was she who caused the error - not the real culprit, the one who called her with the wrong argument. In order to fix this, we can pass `error` that the error you are reporting is Nick at level 2 in the call stack (level 1 is your function):

```
function foo (str)
```

```

if type (str) ~ = "string" then
error ("string expected", 2)
end
<regular code>
end

```

Often when an error occurs, we want to get more more accurate information than just where it originated. How mi- At least we want the call stack that resulted in the error. When `pcall` returns an error, it destroys part of the stack (part from it until the error occurs). Accordingly, if we want to get a call stack, then we must build it before gates from `pcall`. Lua provides the `xpcall` function for this. Cro- the function to be called receives a second argument, *error handling function*. In case of a Lua error calls this error handling function before flushing the stack, therefore, she can use the debug library to obtain any additional information about the error. The two most common error handlers are `debug.debug`, giving you a command line in Lua so you can do it yourself see what was happening when the error occurred; and `debug.traceback` which builds an extended error message ke, including the call stack ⁴. It is the latter function that the independent inter- pretator for printing error messages.

Exercises

Exercise 8.1. It is often necessary to add code to the beginning of the download. pressed block. (We have already seen an example in this chapter when we added code to **return**.) Write a function `loadwithprefix`, which works like `load`, except that it pre- prepends its additional argument to the beginning of the load th block.

Like the original `load` function, `loadwithprefix` should accept blocks represented both as strings and read- functions. Even when the original unit is a string, `loadwithprefix` should not be explicitly concatenated string the passed argument with a block. Instead, she

should call `load` with the appropriate read function - it, which first returns the passed argument, and then - block.

Exercise 8.2. Write a `multiload` function that generalizes `schaft loadwithprefix`, receiving the input list reading functions, as in the following example:

4

```
In Chapter 24, we will learn more about these functions and the debug library.  
f = multiload ("local x = 10;",  
io.lines ("temp", "* L"),  
"Print (x)")
```

For the above example, `multiload` should load block equivalent to concatenating the string `"local ..."` with `co-` by holding the `temp` file and the `"print (x)"` line. Like function `loadwithprefix`, this function itself is nothing concatenated should not.

Exercise 8.3. The `string.rep` function in Listing 8.2 is used *zuet algorithm binary multiplication* (binary multiplication algorithm) to concatenate `n` copies of the given string `s`. For any fixed `n`, we can create a specialized bathroom version of `string.rep`, expanding the loop into a sequential the number of commands `r = r..s` and `s = s..s`. As an example for `n = 5` we get the following function:

```
function stringrep_5 (s)  
local r = ""  
r = r .. s  
s = s .. s  
s = s .. s  
r = r .. s  
return r  
end
```

Write a function that, for a given `n`, returns `stringrep_n` function. Instead of using a closure your function should build the function text in Lua with appropriate commands `r = r..s` and `s = s..s` and then use `load` to get the final function. Compare those are the performance of the `string.rep` function and the one you got functions.

Exercise 8.4. Can you find a value for `p` such that `pcall` expression (`pcall, f`) will return **false** as the first value `nie`?

Chapter 9

Coroutines

A *coroutine* is like a thread (in the multi-threaded sense): it is a thread execution with its stack, its local variables and its *instruction pointer* (instruction pointer); but he shares the global variables and pretty much everything else with other coroutines. The main difference between threads and coroutines is that a multi-threaded program executes all of these threads in parallel but. Coroutines work together: at any given time, a program executes only one of its coroutines, and this one executes. May, a coroutine suspends its execution only when it will explicitly ask for it.

Coroutines are a very powerful concept. And so many of their applications are quite complex. Don't worry if you don't understand some of the examples in this chapter on first reading. You can read to the end of the book and come back later. But please, come back it will be a well spent time.

9.1. Fundamentals of coroutines

Lua keeps all the functions for working with coroutines in a table `coroutine`. The `create` function creates new coroutines. It has everything. Its one argument is the function that the coroutine will execute.

It returns a value of type `thread`, which is the created coroutine. Often the `create` argument is an anonymous function as below:

```
co = coroutine.create (function () print ("hi") end)
print (co) -> thread: 0x8071d98
```

A coroutine can be in one of four states: `suspended`, `running`, `dead`, and `normal`. We can find out the state of the coroutine at using the `status` function :

```
print (coroutine.status (co)) -> suspended
```

When we create a coroutine, it starts in a `suspended` state. `normal` condition; coroutine doesn't start automatically executing

our body when we create it. The `coroutine.resume` function continues starts (starts) execution of the coroutine, changing its state from suspended in running:

```
coroutine.resume (co) -> hi
```

In this first example, the coroutine body simply prints “hi” and terminates execution, leaving the coroutine in destroyed condition:

```
print (coroutine.status (co)) -> dead
```

Until now, coroutines have looked like just a complicated way function calls. The real power of coroutines comes from function `yield`, which allows a running coroutine to pause its execution so that it can be continued later. let's consider a simple example:

```
co = coroutine.create (function ()  
  for i = 1, 10 do  
    print (“co”, i)  
    coroutine.yield ()  
  end  
end)
```

Now, when we continue with this function, it will repairs its execution and executes before the first `yield` :

```
coroutine.resume (co) -> co 1
```

If we now check its status, we will see that this co- the program is paused and, therefore, we can continue its execution:

```
print (coroutine.status (co)) -> suspended
```

From the point of view of the coroutine, all the activity that occurs dit while the coroutine is suspended happens inside the call `yield`. When we continue to execute the coroutine, it returns comes from the `yield` call and continues its execution until the next th call `yield` or the end of the coroutine:

```
coroutine.resume (co) -> co 2
```

```
coroutine.resume (co) -> co 3
```

```
...
```

```
coroutine.resume (co) -> co 10
```

```
coroutine.resume (co) - prints nothing
```

During the last call to `resume`, the loop ends and ends

Execute the function without printing anything. If we try resume its execution again, then `resume` will return **false** and the message about the error:

```
print (coroutine.resume (co))
```

```
-> false cannot resume dead coroutine
```

Note that `resume` executes the coroutine body in protected mode. Therefore, in the event of any errors while executing a coroutine Lua will not show a message error message, but will simply return control to the call to `resume` . When a coroutine continues executing another coroutine, then it is not suspended; in the end we cannot continue live her fulfillment. However, it is not executable, because how many coroutines being executed is another coroutine. Therefore, her status is called *normal* .

A nice feature in Lua is that the `resume` pair - `yield` can exchange data. The first call to `resume` (which there is no `yield` call waiting for it) passes its additional arguments of the main coroutine function:

```
co = coroutine.create (function (a, b, c)
print ("co", a, b, c + 2)
end)
coroutine.resume (co, 1, 2, 3) -> co 1 2 5
```

Calling `resume` returns after **true** , indicating that there is no error-side, all arguments passed to the `yield` call :

```
co = coroutine.create (function (a, b)
coroutine.yield (a + b, a - b)
end)
print (coroutine.resume (co, 20, 10)) -> true 30 10
```

Similarly, `yield` returns all arguments passed according to

The corresponding call to `resume` :

```
co = coroutine.create (function (x)
print ("co1", x)
print ("co2", coroutine.yield ())
end)
coroutine.resume (co, "hi") -> co1 hi
coroutine.resume (co, 4, 5) -> co2 4 5
```

Finally, when the coroutine finishes its execution, everything is values returned by its main function are passed as a result

```
resume :
co = coroutine.create (function ()
return 6, 7
end)
print (coroutine.resume (co)) -> true 6 7
```

We usually rarely use all of these features in one and the same the same coroutine, but they all have their own uses.

For those who already know something about coroutines it is important to clarify some concepts before we move on. Lua offers what is called *asymmetric coroutines* . It means that

it has a function to pause the execution of a coroutine and another function to continue execution of the suspended coroutines. Some languages have *symmetric coroutines* , when there is only one function to transfer control from one coroutines are different.

Asymmetric coroutines are called semi-coroutines by some. mami (not being symmetrical, they are not co-). However, others use the same term semi-routines to denote a limited implementation of coroutines, where a coroutine can be suspended do it only when it does not call any other function, that is, when it has no pending calls. Other words you, only the main body of such a coroutine can `yield` `yield` .

Generators in Python are examples of such semi- routines. Unlike the difference between symmetrical and unbalanced- mi coroutines, the difference between coroutines and generators (as implemented in Python) is much deeper; generators just don't powerful enough to implement some interesting con- instructions that we can do with normal coroutines.

Lua offers complete non-symmetric coroutines. Those who prefer symmetric coroutines, can implement them in based on the asymmetric capabilities of Lua. This is not difficult. (Fact- every control transfer performs a `yield` , followed by `blows resume` .)

9.2. Channels and filters

One of the most important use cases for coroutines is the task of the producer-consumer.

Let's pretend we have a function that is constantly running reads values (for example, reads them from a file), and another function, which constantly consumes these values (for example, writes to another goy file). Typically these two functions look like this:

```
function producer ()
while true do
local x = io.read () - produce new value
send (x)
- send it to the consumer
end
end
function consumer ()
```

```

while true do
local x = receive () - get value from producer
io.write (x, "\n")
- consume it
end
end

```

(In this implementation, both producer and consumer are forever. However, they can be easily changed to stop when more no data.) The challenge here is to connect you-calls `send` and `receive` . This is a typical example of the problem “who has the main cycle”. Both the producer and the consumer are active, each has its own main loops, and each assumes that the other is a call my service. For this particular example, one can easily change structure of one of the functions, expanding its cycle and making it passive noah side. However, in real cases, such a change can to be far from easy.

Coroutines provide an ideal mechanism for connecting producer and consumer, since the `resume` – `yield` pair reverses the usual relationship between caller and caller removable. When a coroutine calls `yield` , it does not call a new one function; instead, it returns control from the current call. `wa (resume)`. Likewise, calling `resume` does not start a new function, and terminates the call to `yield` . This is exactly what we need to connect `neniya send` and `receive` , so that each acts as as if he is the main one, and the second is subordinate. Therefore, `receive` continues to execute the producer, so it can produce a new value; and `send` returns this value back to the consumer:

```

function receive ()
local status, value = coroutine.resume (producer)
return value
end
function send (x)
coroutine.yield (x)
end

```

Of course, the producer must also be a coroutine:

```

producer = coroutine.create (
function ()
while true do
local x = io.read () — produce a new value
send (x)
end
end)

```

With this design, the program starts by calling the consumer. Kog-yes, the consumer needs value, he resumes production body, which is executed until it has a ready value that it conveys to the consumer and does not stop until until the consumer continues executing again. So thus we get what is called *sweet-driven* design *fighter an* (consumer-driven). Another option would be to write program using a *manufacturer-driven* design where sweat-a child is a coroutine.

We can extend this design with filters that are tasks between the manufacturer and the rebel and performing data conversion. *The filter* is producer and consumer at the same time, therefore he owes producer execution to get new value and uses `yield` to pass this value to the consumer. In ka-as a simple example, we can add to our previous code a filter that inserts a line number at the beginning of each line. The code is shown in Listing 9.1. In the end, we just need to create a compo

```
p = producer ()
f = filter (p)
consumer (f)
```

Or even simpler:

```
consumer (filter (producer ()))
```

Listing 9.1. Consumer and manufacturer with filters

```
function receive (prod)
local status, value = coroutine.resume (prod)
return value
end
function send (x)
coroutine.yield (x)
end
function producer ()
return coroutine.create (function ()
while true do
local x = io.read () - produce new value
send (x)
end
end)
end
function filter (prod)
return coroutine.create (function ()
for line = 1, math.huge do
local x = receive (prod) - get new value
```

```

x = string.format ("% 5d% s", line, x)
send (x)
- send it to the consumer
end
end)
end
function consumer (prod)
while true do
local x = receive (prod) - get new value
io.write (x, "\n")
- consume new value
end
end

```

Listing 9.2. Function to get all permutations of n elements a

```

function permgen (a, n)
n = n or #a - default 'n' is the size of 'a'
if n <= 1 then - nothing to do?
printResult (a)
else
for i = 1, n do
- put the i-th element at the end
a [n], a [i] = a [i], a [n]
- create all permutations of other elements
permgen (a, n - 1)
- restore the i-th element
a [n], a [i] = a [i], a [n]
end
end
end

```

If you've thought about pipes in UNIX, then you're not alone. After all, coroutines are a variant of *non-preemptive multi-tasks* (non-preemptive multitasking). With channels every task runs as a separate process; with coroutines each task runs as a separate coroutine. Channels provide buffer between the writer (producer) and the reader (consumer), therefore, some freedom in their relative speeds is possible. This is important for the channel, since the cost of switching between processes themselves are high. With coroutines, the cost of switching between tasks much smaller (like a function call), so writing and the reader can go toe-to-toe.

9.3. Coroutines like iterators

We can consider loop iterators as a separate example producer-consumer: the iterator produces values that consumed by the body of the cycle. Therefore, it is quite natural

use coroutines to write iterators. Really-
but, coroutines are a powerful tool for this purpose.
Again, the key feature is their ability to flip
the relationship between caller and callee. With this feature
we can write iterators without worrying about keeping state between
by successive calls to the iterator.

To illustrate this use case,
let's write an iterator to iterate over all permutations of a given
array. Writing an iterator like this isn't easy, but pretty
just write a recursive function that builds all these overrides
new. The idea is simple: put each element at the end in turn
array and recursively generate any remaining permutations.
The code is shown in Listing 9.2. In order for it to work, we must
we write the corresponding function `printResult` and call
`permgen` with proper arguments:

```
function printResult (a)
  for i = 1, #a do
    io.write (a [i], "")
  end
  io.write ("\n")
end
permgen ({1,2,3,4})
-> 2 3 4 1
-> 3 2 4 1
-> 3 4 2 1
...
-> 2 1 3 4
-> 1 2 3 4
```

Once the generator is ready, it is very easy to convert it to
iterator. First, we'll replace `printResult` with `yield` :

```
function permgen (a, n)
  n = n or #a
  if n <= 1 then
    coroutine.yield (a)
  else
    <as before>
```

Then we define a factory that will run the generator
inside a coroutine, and create an iterating function. For semi-
the next permutation, the iterator simply continues
nenie coroutine:

```
function permutations (a)
  local co = coroutine.create (function () permgen (a) end)
  return function () - iterator
  local code, res = coroutine.resume (co)
```

```
return res
end
end
```

After that we can easily iterate over all permutations of the array.
using the **for** statement :

```
for p in permutations {"a", "b", "c"} do
  printResult (p)
end
-> bca
-> cba
-> cab
-> acb
-> bac
-> abc
```

The `permutations` function uses a typical Lua pattern, which hides the resume of the coroutine inside the function. This pattern is so common that Lua provides special function for it: `coroutine.wrap` . Like `create` , `wrap` creates new coroutine. Unlike `create` , `wrap` does not return itself. `coroutine`; instead, it returns a function that when the call continues execution of this coroutine. Unlike `resume` , it does not return the error code as the first value; together then it causes an error. Using `wrap` , we can write `permutations` as follows:

```
function permutations (a)
  return coroutine.wrap (function () permgen (a) end)
end
```

It is generally easier to use `coroutine.wrap` than `coroutine.create` . It gives us exactly what we need from a coroutine: function to continue its execution. However, it is less flexible than `create` . There is no way to check the status of a coroutine created with `wrap` . Moreover, we cannot check for errors at run time.

9.4. Non-displacing multi-threading

As we saw earlier, coroutines provide an option for noisy multi-threading. Every coroutine is equivalent to a thread. Couple `yield-resume` switches control from one thread to another thread. However, unlike ordinary multi-threading, a coroutine is not preemptive. While the coroutine is running, it cannot be stopped from the outside. She interrupts her execution

only when it explicitly requests it (via a call to `yield`). For In some applications, this is not a problem, rather the opposite. Programming is much easier in the absence of displacement. You do not you need to worry about sync errors because all sync chronization is obvious. You just need to make sure that the coroutine ma calls `yield` outside the critical area of the code.

However, with non-displacing multi-threading, as soon as some the thread calls a blocking operation, the whole program is blocked until this operation completes. For most, this is unacceptable, which leads to the fact that many programmers you don't see coroutines as an alternative to traditional multi-threading. As we will see here, there is an interest in this problem. a new (and obvious, moreover) solution.

Let's look at a typical multi-threaded task: we want download multiple files over HTTP. To download several of these files, we first need to figure out how to download one file. In this example, we will look at the bib-designed by Diego Nehab the LuaSocket *library*. In order to download a file, you must first install connect to the site containing this file, get the file (in blocks) and close the connection. In Lua, we can write this next in a blowing manner. First, we load the LuaSocket library:

```
local socket = require "socket"
```

Then we define the site and the file we want to download. In that for example, we will download the HTML 3.2 reference manual from the site World Wide Web Consortium:

```
host = "www.w3.org"
```

```
file = "/TR/REC-html32.html"
```

Then we open a TCP connection to port 80 (standard port for HTTP connections) of this site:

```
c = assert(socket.connect(host, 80))
```

This operation returns the connection object that we are using we use to send a request to receive a file:

```
c:send("GET" .. file .. "HTTP/1.0\r\n\r\n")
```

We then read the file in 1 KB blocks, writing each block to standard output:

```
while true do
```

```
  local s, status, partial = c:receive(2 ^ 10)
```

```
  io.write(s or partial)
```

```
  if status == "closed" then break end
```

```
end
```

The `receive` function returns either the string it read, or

nil on error; in the latter case, it also returns the code `errors (status)` and what she read before the error (`partial`). When the site closes the connection, we print the remaining data and exit from the loop.

After downloading the file, we close the connection:

```
c: close ()
```

Now that we know how to download one file, let's go back to the problem of downloading multiple files. The simplest approach is to Children download one file at a time. However, this consistent sub-the move when we start reading the file only after we finish with the previous file is too slow. When reading a file by the network, the program spends most of its time waiting for data. More precisely, she spends most of her time blocked in the `receive` call . Therefore, the program can be executed significantly It is much faster if it downloads all files at once. Then when the connection has no ready data, the program can read the data from another connection. It is clear that coroutines provide convenience a good way to organize these simultaneous downloads. we create a new thread for each downloaded file. When the thread there is no ready data, it transfers control to the dispatcher, who calls another thread.

In order to rewrite the program using co-programs, we first need to rewrite the downloading code as function. The result is shown in Listing 9.3. Since we are not interested resche the contents of the file, the function reads and prints the file size instead of writing the file to standard output (when we have several they read several files at once, the output would be complete mishmash).

Listing 9.3. Code for downloading a page from the network

```
function download (host, file)
local c = assert (socket.connect (host, 80))
local count = 0
- counts number of bytes read
c: send ("GET" .. file .. "HTTP / 1.0 \r \n \r \n")
while true do
local s, status = receive (c)
count = count + #s
if status == "closed" then break end
end
c: close ()
print (file, count)
end
```

In the resulting code, we use the helper function
(receive) to receive data from the connection. With successive
However, the code would look like this:

```
function receive (connection)
local s, status, partial = connection: receive (2 ^ 10)
return s or partial, status
end
```

For parallel implementation, this function must receive data
without blocking. Instead, if the required data is not available,
then it calls yield . The new code looks like this:

```
function receive (connection)
connection: settimeout (0)
— do not block
local s, status, partial = connection: receive (2 ^ 10)
if status == “timeout” then
coroutine.yield (connection)
end
return s or partial, status
end
```

The settimeout (0) call does any operation on the connection.
blocking. When the status of the operation is “timeout” , it means that
the operation completed without completing the request. In this case, the thread
is

Give control to another thread. The non- **false** argument passed to
ny yield statement , informs the dispatcher that this thread still performs
your task. Please note that even in the case of the “timeout” status
the partial variable still contains the read data.

Listing 9.4. contains dispatcher code and additional code.

The threads table contains a list of all active threads for the disk.
petcher. The get function ensures that each downloaded file
downloaded in a separate thread. The dispatcher itself is actually
just a loop that iterates over all the threads, starting them to execute-
one by one. It also removes from the list those threads that
have already completed the download. The cycle stops when no more
remains of threads.

Listing 9.4. Dispatcher

```
threads = {}
— list of all running threads
function get (host, file)
- create a coroutine
local co = coroutine.create (function ()
download (host, file)
```

```

end)
- insert it into the list
table.insert (threads, co)
end
function dispatch ()
local i = 1
while true do
if threads [i] == nil then
— no more threads?
if threads [1] == nil then break end — is the list empty?
i = 1 - restart the loop
end
local status, res = coroutine.resume (threads [i])
if not res then
— has the thread finished downloading?
table.remove (threads, i)
else
i = i + 1
— go to the next thread
end
end
end
end

```

Finally, the main routine creates the required threads and calls there is a dispatcher. For example, to download four documents from the site W3C, the main program might look like below:

```

Non-displacing multi-threading
host = "www.w3.org"
get (host, "/TR/html401/html40.txt")
get (host, "/TR/2002/REC-xhtml1-20020801/xhtml1.pdf")
get (host, "/TR/REC-html32.html")
get (host, "/TR/2000/REC-DOM-Level-2-Core-20001113/DOM2-Core.txt")
dispatch () - main loop

```

On my computer, downloading these four files using Coroutine processing takes 6 seconds. With sequential download it takes more than twice (15 seconds).

Despite this optimization, this latest implementation is still far from optimal. Everything works well as long as though one thread would have something to read. However, when no thread has ready to read data, the dispatcher constantly switches from threads to thread just to make sure there are no ready-made data. As a result, this implementation takes almost 30 times more CPU time than the serial version.

In order to avoid this situation, we can use call the `select` function from the `LuaSocket` library: it allows you to

block a program that is pending while changing status in the connection group. Implementation changes are minor us: we only need to change the dispatcher as shown in the listing 9.5. In a loop, the new dispatcher collects connections in the `timedout` table . for which there is no ready-made data. (Remember that `receive` pe- gives similar joins to the `yield` function , so again running them). If none of the connections have ready data, then the dispatcher calls `select` to wait for at least one . One of these connections will change the status. This final re- Alization works as fast as the previous one. However, she uses only slightly more CPU time than the serial naya implementation.

Listing 9.5. Dispatcher using `select` function

```

function dispatch ()
  local i = 1
  local timedout = {}
  while true do
    if threads [i] == nil then — no more threads?
    if threads [1] == nil then break end
    i = 1
    — start the cycle over again
    timedout = {}
    end
    local status, res = coroutine.resume (threads [i])
    if not res then
      — has the thread finished its work?
      table.remove (threads, i)
    else
      — waiting time out
      i = i + 1
      timedout [#timedout + 1] = res
      if #timedout == #threads then — are all threads blocked?
      socket.select (timedout)
      end
    end
  end
end

```

Exercises

Exercise 9.1. Use coroutines to pre- form Exercise 5.4 into a generator for combinations that ry can be used as follows:

```
for c in combinations ({“a”, “b”, “c”}, 2) do
  printResult (c)
end
```

Exercise 9.2. Implement and run the code from the previous section (non-displacing multi-threading).

Exercise 9.3. Implement the `transfer` function in Lua. If a think about calling `resume`-yield the same as calling function and return from it, then this function will be like a `goto` : it interrupts the current coroutine and resumes any other th coroutine passed as an argument.

(*Hint* : use an analog of the `dispatch` procedure to control handling your coroutines. Then `transfer` will transfer control to the dispatcher, informing about which next com- the program needs to be started, and the dispatcher will call `re- sume`).

Chapter 10

Completed examples

To conclude this introduction to the language, we show three simple, but complete programs. The first program solves the problem of eight kings wah. The second program prints the most common words in text. The last example is the implementation of the Markov chain described by Naya Kernighan and Pike in their book " *The Practice of Programming* " (Addison-Wesley, 1999).

10.1. The Eight Queens Problem

Our first example is a very simple program that solves *the problem of eight queens* : you need to arrange the eight queens on chessboard so that none of the queens are under attack. The first step in solving this problem should be noted that each Each solution must have exactly one queen in each line. Thus, we can represent the solution as an array of eight numbers, one for each line; each number tells us in which column is the queen in the corresponding row. On-

example, the array {3,7,2,1,8,6,5,4} means one queen is in row 1 in column 3, another is in row 2 in column 7 and etc. (Note that this is not a valid solution; for example, the queen in row 3 in column 2 attacks the queen in line 4 in column 1). Also note that any solution should be a permutation of numbers from 1 to 8, as the solution should contain one queen in each column.

The complete program is shown in Listing 10.1. The first function is - it is `isplaceok`, which checks that the given position on the board is not falls under the battle of previously placed queens. Remembering that he cannot to be two queens on the same line, this function checks that there are no two queens on one column or one diagonal with a given position.

Listing 10.1. Program for eight queens

```

local N = 8 - board size
- checks that position (n, c) is not under attack
local function isplaceok (a, n, c)
for i = 1, n - 1 do -- for each previously placed queen
if (a[i] == c) or -- the same column?
(a[i] - i == c - n) or -- the same diagonal?
(a[i] + i == c + n) then -- the same diagonal?
return false -- position under attack
end
end
return true -- not under attack
end
- print the board
local function printsolution (a)
for i = 1, N do
for j = 1, N do
io.write (a[i] == j and "X" or "-", "")
end
io.write ("\n")
end
io.write ("\n")
end
- add to the board 'a' all queens from 'n' to 'N'
local function addqueen (a, n)
if n > N then -- have all the queens been placed?
printsolution (a)
else - try to place the nth queen
for c = 1, N do
if isplaceok (a, n, c) then
a[n] = c -- put the n-th queen in column 'c'
addqueen (a, n + 1)

```

```

end
end
end
end
-- run the program
addqueen ({} , 1)

```

Next we have the `printsolution` function , which prints the check mat board. She just goes around the board, typing 'x' in places with queen and '.' elsewhere. Each result looks dit like below:

```

The most common words
X - - - - -
- - - - X - -
- - - - - X
- - - - - X - -
- - X - - - -
- - - - - X -
- X - - - - -
- - - X - - -

```

The last function `addqueen` is the heart of the program. Sleep- then it checks if the solution is complete, and if so, then prints this solution. Otherwise, it iterates over all the pillars. tsy; for each unallocated column, the program puts there a short left and recursively tries to accommodate the remaining queens.

10.2. Most often occurring words

Our next example is a simple program that reads text and prints the most frequent words from that text.

The main data structure of this program is just a table `tsa`, which matches each word with its frequency. Using With this data structure, the program has three main tasks:

- Read the text by counting the number of occurrences of each word.
- Sort the list of words in descending order of frequency of meeting- the bridges of every word.
- Print the first n items from the sorted list.

ka.

To read the text, we can use the `allwords` iterator , which which we discussed in Section 7.1. For every word that we read, we increment the corresponding counter:

```

local counter = {}

```



```

for w in allwords do
counter [w] = (counter [w] or 0) + 1
end

```

The next task is to sort the word list. However, how an attentive reader might have noticed we don't have a word list! but it is easy to create using words that are keys in the table

```

face counter:
local words = {}
for w in pairs (counter) do
words [#words + 1] = w
end

```

Listing 10.2. Program for printing the most common words

```

local function allwords ()
local auxwords = function ()
for line in io.lines () do
for word in string.gmatch (line, "%w +") do
coroutine.yield (word)
end
end
end
return coroutine.wrap (auxwords)
end
local counter = {}
for w in allwords () do
counter [w] = (counter [w] or 0) + 1
end
local words = {}
for w in pairs (counter) do
words [#words + 1] = w
end
table.sort (words, function (w1, w2)
return counter [w1]> counter [w2] or
counter [w1] == counter [w2] and w1 <w2
end)
for i = 1, (tonumber (arg [1]) or 10) do
print (words [i], counter [words [i]])
end

```

Now that we have a list, we can sort it with using the `table.sort` function we discussed in Chapter 6:

```

table.sort (words, function (w1, w2)
return counter [w1]> counter [w2] or
counter [w1] == counter [w2] and w1 <w2
end)

```

The complete program is shown in Listing 10.2. Pay attention- the use of coroutines in the `auxwords` iterator . In the last in the loop that prints the result, the program considers that its first the argument is the number of words to print and uses

value 10 if no arguments were passed.

10.3. Markov chain

Our final example is a *Markov chain* implementation . Program generates pseudo-random text based on which words can follow a sequence of n previous words in the text st. For this implementation, we will assume that n is 2. The first part reads the body text and builds a table that for every two words gives a list of all the words that may be behind them follow in the main text. After building the table, the program uses it to construct a random text, where each word follows the previous two with the same probability as in the base in your text. As a result, we get text that is random, but not absolutely. For example, applying it to the English text of this book, we get texts like “Constructors can also traverse a table constructor, then the parentheses in the following line does the whole file in a field n to store the contents of each function, but to show its only argument. If you want to find the maximum element in an array can return both the maximum value and continues showing the prompt and running the code. The following words are reserved and cannot be used to convert between degrees and radians ”.

We will encode each prefix by connecting two words when using a space:

```
function prefix (w1, w2)
return w1 .. “ ” .. w2
end
```

We will use the string NOWORD (“\n”) to initialize prefix words and end-of-text symbols. For example, for text “The more we try the more we do” the table of the following words will be look like below:

```
{["\ N \ n"] = {"the"},
["\ N the"] = {"more"},
["The more"] = {"we", "we"},
["More we"] = {"try", "do"},
["We try"] = {"the"},
["Try the"] = {"more"},
["We do"] = {"\ n"},
```

```
}
```

The program stores its table in the `statetab` variable . For to insert a new word into the table, we will use the following function:

```
function insert (index, value)
local list = statetab [index]
if list == nil then
statetab [index] = {value}
else
list [#list + 1] = value
end
end
```

It first checks that the given prefix already has a list; if not, it creates a new list with the passed word. Otherwise she inserts the passed word at the end of the existing list.

To build the `statetab` we will use two variables `w1` and `w2` containing the last two words read.

For each new word read, we add it to the list, associated with `w1-w2` , and then update the values for `w1` and `w2` .

After building the table, the program begins to build the text, with consisting of `MAXGEN` words. To begin with, she sets the values of the changes `nym w1` and `w2` . Then, for each prefix, she randomly chooses next word from the list of valid words, prints that word and updates the values of `w1` and `w2` . Listings 10.3 and 10.4 contain the complete program. Unlike our previous example with the most common words, here we use the implementation `allwords` based on closures.

Listing 10.3. Additional program definitions

```
with Markov chain
function allwords ()
local line = io.read () -- current line
local pos = 1
-- current position in the line
return function () -- iterating function
while line do
- repeat until there are lines left
local s, e = string.find (line, "%w+", pos)
if s then
- found the word?
pos = e + 1 - update position
return string.sub (line, s, e) - return word
else
line = io.read () - the word was not found; let's try the trail. string
pos = 1
- start at the beginning of the line
```

```

end
end
return nil
-- no more lines, end of traversal
end
end
function prefix (w1, w2)
return w1 .. "" .. w2
end
local statetab = {}
function insert (index, value)
local list = statetab [index]
if list == nil then
statetab [index] = {value}
else
list [#list + 1] = value
end
end

```

Listing 10.4. Markov chain program

```

local N = 2
local MAXGEN = 10000
local NOWORD = "\n"
- build a table
local w1, w2 = NOWORD, NOWORD
for w in allwords () do
insert (prefix (w1, w2), w)
w1 = w2; w2 = w;
end
insert (prefix (w1, w2), NOWORD)
- generate text
w1 = NOWORD; w2 = NOWORD - initialize
for i = 1, MAXGEN do
local list = statetab [prefix (w1, w2)]
- choose a random word from the list
local r = math.random (#list)
local nextword = list [r]
if nextword == NOWORD then return end
io.write (nextword, "")
w1 = w2; w2 = nextword
end

```

Exercises

Exercise 10.1. Change the program with eight queens, so that it stops after printing the first decision.

Exercise 10.2. An alternative implementation of the problem of eight queens can be the construction of all permutations of numbers from

1 through 8 and check which ones are valid. Change the program mu for using this approach. How fast is the effect of the new program in comparison with the old one? (*Hint* : Compare the total number of permutations with the number of times when the original program calls the `isplaceok` function .)

Exercise 10.3. When we use the program to determine the most common words, then usually the most common the most common words are short, uninteresting words like articles and prepositions. Modify the program so that she skipped words of less than four letters.

Exercise 10.4. Generalize the Markov chain algorithm so that could use any size as length prefix.

Part II

T TABLES AND OBJECTS

CHAPTER 11

Data structures

Lua tables are not just a data structure, they are basic and single data structure. All structures that offer other languages - arrays, records, lists, queues, sets - can

be represented in Lua using tables. Moreover, the tables in Lua effectively implements all of these structures.

In traditional languages such as C and Pascal, we implement more most data structures using arrays and lists (where the list $ki = \text{records} + \text{pointers}$). Although we can implement arrays and lists with tables in Lua (and sometimes we do), tables are much more powerful than arrays and lists; many algorithms with using tables become almost trivial. For instance, we rarely use Lua lookups as tables provide direct access to values of various types.

It takes time to understand how to effectively use tables in Lua. In this chapter, I will show you how to implement typical structures.

rounds of data using tables, and I will give examples of their use.

We won't start with arrays and lists because we need them.

us for other structures, but since most programmers already familiar with them. We have already seen the basis of this material in the pre-

in previous chapters, but I will also repeat it here.

11.1. Arrays

We are implementing arrays in Lua by simply indexing tables with integers numbers. So arrays are not fixed size

and grow as needed. Usually, when initializing an array, we implicitly set its size. For example, after doing the following code, any attempt to access a field outside the range of 1-1000 will return **nil** instead of 0:

```
a = {} - new array
for i = 1, 1000 do
  a[i] = 0
end
```

The length operator ('#') uses this to define the length of the mass-Siwa:

```
print (#a) -> 1000
```

You can start an array from zero or any other value:

- create an array with indices from -5 to 5

```
a = {}
for i = -5, 5 do
  a[i] = 0
```

end

However, it is common in Lua to start arrays at index 1. Bib- the Lua libraries follow this convention; as well as the operator length us. If your arrays don't start at 1, then you won't be able to use call these capabilities of the language.

We can use the constructor to create and initialize- array of one expression:

```
squares = {1, 4, 9, 16, 25, 36, 49, 64, 81}
```

Such constructors can be so large that how much is needed (at least up to several million cops).

11.2. Matrices and multidimensional arrays

There are two main ways of representing matrices in Lua. First - is to use an array of arrays, i.e. a table, each element which is another table. For example, you can create a matrix zu of zeros of size M by N using the following code:

```
mt = {}  
-- create matrix  
for i = 1, N do  
  mt[i] = {} -- create string  
  for j = 1, M do  
    mt[i][j] = 0  
  end  
end
```

Since tables are objects in Lua, to create mat-

You must explicitly create each row of the script. On the one hand, it is is more cumbersome than simply declaring a matrix, as is done in languages C and Pascal. On the other hand, it gives more flexibility. For instance,

you can create a triangular matrix by changing the for loop j = 1, M do ... end in the previous code snippet for j = 1, i do ... end .

With this code, the triangular matrix will only use the posi- memory fault over the original example.

The second way to represent matrices in Lua is to combine converting two indices into one. If both indices are integers lami, then you can simply multiply the first by the corresponding constant and add a second index. With this approach the following

the code will create our matrix of zeros of size M by N :

```
mt = {} -- create matrix
for i = 1, N do
  for j = 1, M do
    mt [(i - 1) * M + j] = 0
  end
end
```

If the indices are strings, then you can create one index by simply concatenating those lines with some character in between.

For example, you can create a matrix m with string indices-
mi s and t using the following code m[s .. ":" .. t] , provided that
that both s and t do not contain colons; otherwise pairs like
(“A:”, “b”) and (“a”, “: b”) will both give the same index “a :: b” .

When in doubt, you can use the control sequence like '\0' for separating indices.

Quite often, applications use a *sparse matrix* ,
that is, a matrix where most of the elements are either 0 or **nil** . On-
example, you can represent a graph using its connectivity matrix-
value in which the value at position m, n is equal to x , if between nodes
m and n is a join at a cost x. When these nodes are not connected, then the value
nil at position m, n is **nil** . In order to represent a graph with ten
with thousands of nodes, where each node has about five neighbors, you
a matrix with one hundred million possible elements is needed, but only
fifty thousand of them will be non- **nil** (five non-zero columns
for each row corresponding to five neighbors). Many books
on data structures discuss in detail how you can implement
similar sparse matrices, without wasting 400 MB of memory on them, but
you rarely need such tricks when programming
in Lua. With our first view (table of tables) you will
ten thousand tables are needed, each of which contains about
five elements, that is, a total of about fifty thousand values. When
the second view we have one table with fifty thousand-
elements. Whichever presentation you use, you
only memory is needed for non- **nil** elements .

When working with sparse matrices, we cannot use
length operator due to holes (**nil** values) between elements. One-
but this is not a big loss; even if we could use it, then
it wouldn't be worth it. For most operations, it was extremely
it is inefficient to iterate over all these empty elements. Instead, we
we can use pairs to traverse only elements other than

from *nil* . For example, in order to multiply a string by a constant, we can use the following code:

```
function mult (a, rowindex, k)
local row = a [rowindex]
for i, v in pairs (row) do
row [i] = v * k
end
end
```

Note, however, that the keys do not have any specific divided order in the table, so iterating with `pairs` does not guarantee that we will visit all columns in ascending order. For some tasks (for example, our previous example) this is no problem. For other purposes, you can use the excellent views, such as linked lists.

11.3. Linked Lists

Since tables are dynamic entities, re-

It's pretty easy to link linked lists in Lua. Each node represented by a table, and the links are just fields of the table, which contain links to other tables. For example, let's re-link the simplest list, where each node contains two fields, `next` and `value` . The root of the list is an ordinary variable:

```
list = nil
```

To insert an element with the value `v` at the beginning of the list, we do:

```
list = {next = list, value = v}
```

To traverse the list, we can write:

```
local l = list
while l do
<visit l.value>
l = l.next
end
```

Other list options, such as bidirectional or circular lists are also easy to implement. However, such structures you will rarely need Lua as there is usually an easier one a way of presenting your data without using related lists. For example, we can represent the stack as (unlimited `ny`) array.

11.4. Queues and doubles queues

The simplest way to implement queues in Lua is to use

The `insert` and `remove` functions from the `table` library . These functions insert and remove elements from an arbitrary position in the array, moving the rest of the array elements. However, such changes Scales can be expensive for large structures. More efficient nth implementation uses two indices, one for the first element and one for the latter:

```
function ListNew ()  
return {first = 0, last = -1}  
end
```

In order not to pollute the global namespace, we define all operations for working with a list within a table, which We'll call it `List` (this way we'll create a *module*). Then we we can rewrite our last example as follows:

```
List = {}  
function List.new ()  
return {first = 0, last = -1}  
end
```

Now we can insert and remove elements from either end beyond constant time:

```
function List.pushfirst (list, value)  
local first = list.first - 1  
list.first = first  
list [first] = value  
end  
function List.pushlast (list, value)  
local last = list.last + 1  
list.last = last  
list [last] = value  
end  
function List.popfirst (list)  
local first = list.first  
if first > list.last then error ("list is empty") end  
local value = list [first]  
list [first] = nil - let the garbage collector remove it  
list.first = first + 1  
return value  
end  
function List.poplast (list)  
local last = list.last  
if list.first > last then error ("list is empty") end  
local value = list [last]  
list [last] = nil - let the garbage collector remove it
```

```
list.last = last - 1
return value
end
```

If you use this structure in the classic way bong, calling only `pushlast` and `popfirst`, then both `first` and `last` will be grow constantly. However, since we represent arrays in Lua with help tables, you can easily index them from 1 to 20 or from 16,777,216 to 16,777,236. Since Lua uses double precision to represent numbers, your program can run for two hundred years, making a million insertions per second, before an overflow problem occurs.

11.5. Sets and sets

Suppose you want to iterate over all ids used by in the program; somehow you need to filter reserved words. Some C programmers can try to use to represent the set of reserved given words an array of strings and then to check if the given word is reserved, search in this sive. You can even use the binary to speed up searches. tree to represent the set.

In Lua, an efficient and easy way to represent sets will use the elements as *indexes* on the table. Then instead of to find if the table contains a given word, you can simply ask try to index the table with this word and see if whether the resulting result is *nil*. For example, we can use the following code:

```
reserved = {
  ["While"] = true, ["end"] = true,
  ["Function"] = true, ["local"] = true,
}
for w in allwords () do
  if not reserved [w] then
    <do something with 'w'>
    - 'w' is an unreserved word
  end
end
```

(Since these words are reserved in Lua, we cannot use use them as identifiers; for example, we cannot

write `while = true` . Instead, we write `["while"] = true` .)

You can also use clearer initialization when using an additional function that builds the set:

```
function Set (list)
  local set = {}
  for _, l in ipairs (list) do set [l] = true end
  return set
end
reserved = Set {"while", "end", "function", "local",}
```

Sets, also called *multisets* , differ from ordinary sets in that each element may not occur how many times. The simple representation of sets in Lua is similar to the previous view for sets, but with each key associated corresponding counter. In order to insert an element, we Liching his counter:

```
function insert (bag, element)
  bag [element] = (bag [element] or 0) + 1
end
```

To remove an element, we decrement its counter:

```
function remove (bag, element)
  local count = bag [element]
  bag [element] = (count and count > 1) and count - 1 or nil
end
```

We store the counter only if it already exists and is not equal zero.

11.6. Line buffers

Suppose you are working with text and reading the file line by line. Then your code might look like this:

```
local buff = ""
for line in io.lines () do
  buff = buff .. line .. "\n"
end
```

Despite its harmless appearance, this code can hit hard for performance for large files: for example, reading a file in 1 MB takes 1.5 minutes on my old computer ¹ .

Why is this so? To understand what is happening, imagine that we are inside a loop; each line is 20 bytes, and we have already read 2500 lines, so `buff` is a 50 KB line. Kog-yes Lua connects `buff..line .. "\n"` ; it allocates a new line in 50,020 bytes and copies 50,000 bytes from `buff` to this newline. Ta-

Thus, for each new line, Lua moves in memory when approximately 50 Kb, and this size is only growing. More precisely this algorithm has quadratic complexity. After reading 100 new lines (2 KB total) Lua has already moved over 2 MB of memory. When Lua is reads 350 KB, more than 50 GB will already be moved in memory (this problem is not unique to Lua: other languages where strings immutable, also face a similar problem, the most Java is a well-known example of such a language). Before we continue, it should be noted that, despite all that said, this is not a typical problem. For small lines the above loop works fine. To read the entire file Lua provides `io.read ("* a")` , this call reads the entire file. However, sometimes we run into this problem. To fight With a similar problem Java uses the `StringBuffer` structure . In Lua, we can use a table as a string buffer.

The key to this approach is the `table.concat` function , which returns the result of concatenating all strings from the given list. With `concat` we can rewrite our previous code to the following in the following way:

```
local t = {}
for line in io.lines () do
  t [#t + 1] = line .. "\n"
end
local s = table.concat (t)
```

This algorithm takes less than 0.5 seconds to read the same self-th file, which took almost a minute with the previously used code. (Regardless, to read the entire file, it is best to use `io.read` with the “* a” option .)

We can do even better. The `concat` function takes as input second optional argument, which is a delimiter, which will be inserted between the lines. Using this separator, we can get rid of the need to insert every time character `'\n'` :

```
local t = {}
for line in io.lines () do
  t [#t + 1] = line
end
```

```
s = table.concat (t, "\n") .. "\n"
```

The `concat` function inserts a separator between lines, but we still need to add one last `\n` character. This last

The new concatenation operation copies the resulting string, which can require a significant amount of time. There is no way to force `concat` insert an extra delimiter, but we can easily achieve this by simply adding an empty line to `t`:

```
t[#t + 1] = ""
```

```
s = table.concat (t, "\n")
```

An extra `\n` character that `concat` will add before the post the ice line is what we need.

11.7. Graphs

Like any sane language, Lua offers various implementations for graphs, each of which is better suited for its own type of algorithm. Here we will look at a simple object-oriented a new implementation in which we will represent the nodes as objects (more precisely, tables, of course) and arcs as links between nodes.

We will represent each node as a table with two fields:

`name`, which is the name of the node, and `adj`, which is the set of nodes.

connected to the data. Since we will be reading the graph from the text file, we need a way to find the node by its name. For

for this we will use an additional table. Function `name2node`,

having received the name of the node, will return this node:

```
local function name2node (graph, name)
```

```
local node = graph[name]
```

```
if not node then
```

```
- there is no node yet, create a new one
```

```
node = {name = name, adj = {}}
```

```
graph[name] = node
```

```
end
```

```
return node
```

```
end
```

Listing 11.1 contains a function that will build a graph. It reads a file where each line contains the names of two nodes, denoting that there is an arc leading from the first node to the second. For each line it uses the `string.match` function to split the string into two names, then finds the corresponding nodes (creating them if necessary) and connects them.

Listing 11.1. Reading a graph from a file

```

function readgraph ()
local graph = {}
for line in io.lines () do
- split the line into two names
local namefrom, nameto = string.match (line, “(% S +)% s + (% S +)”)
- find matching nodes
local from = name2node (graph, namefrom)
local to = name2node (graph, nameto)
- add 'to' to the link list of node 'from'
from.adj [to] = true
end
return graph
end

```

Listing 11.2 illustrates an algorithm using similar graphs. The `findpath` function searches for a path between two nodes using depth-first traversal. Its first parameter is the current node; second for- gives the desired node; the third parameter stores the path from the beginning to the current

mu node; the last parameter is the set of all already visited nodes (to avoid loops). Notice how the algorithm works directly with nodes, avoiding the use of their names. On- example, `visited` is a set of nodes, not node names. Similarly path is a list of nodes.

Listing 11.2. Finding a path between two nodes

```

function findpath (curr, to, path, visited)
path = path or {}
visited = visited or {}
if visited [curr] then — has the node already been visited?
return nil
- there is no way
end
visited [curr] = true
— mark the node as visited
path [#path + 1] = curr — add to the path
if curr == to then
-- target?
return path
end
- try all neighboring nodes
for node in pairs (curr.adj) do
local p = findpath (node, to, path, visited)
if p then return p end
end
path [#path] = nil
— remove a node from the path
end

```

To test this code, we'll add a function that prints path, and additional code to make it work:

```
function printpath (path)
for i = 1, #path do
print (path [i] .name)
end
end
g = readgraph ()
a = name2node (g, "a")
b = name2node (g, "b")
p = findpath (a, b)
if p then printpath (p) end
```

Exercises

Exercise 11.1. Modify the implementation of the queue so that both the index would be zero if the queue is empty.

Exercise 11.2. Repeat exercise 10.3, only instead of in order to use length as a criterion for discarding word, now the program should read from the special th file list of words to skip.

Exercise 11.3. Modify the graph structure so that it matches kept a label for each arc. Each arc must also be represented using an object with two fields: met-coy and knots to which she points. Instead of many neighboring nodes each node must contain a list of arcs going from this node.

Modify the readgraph function so that it is from each line ki file read two node names and a label (assuming that the label this number).

Exercise 11.4. Use the graph representation from the previous the next exercise, where the label of each arc is the distance between the nodes it connects. Write a function tion that finds the shortest path between two nodes. (*Hint* : Use Dijkstra's algorithm.)

CHAPTER 12

Data files and persistence

When working with data files it is usually much easier to write files than to read them. When we write to a file, we are completely in control of everything that happens. On the other hand, when we read from the file, we don't know what to expect. Besides all data types, which the correct data file can contain, the program must also handle bad files intelligently. Therefore writing correct working procedures for reading data is always difficult.

In this chapter, we will see how you can use Lua to do something that would eliminate all the code for reading data from our programs, just writing data in a suitable format.

12.1. Data files

Table constructors provide an interesting alternative to the form data mats. With a little extra work when writing data reading becomes trivial. The approach is to write our data file as a Lua program that, when executed, creates the necessary data. As usual, for the sake of clarity, let's look at the measures. If our data file is in a specific format, for example CSV or XML, our choice is extremely small. However, if we want to create a file for our own use, then we can use Lua constructors as our format. In this format, we represent each entry as a Lua constructor. Instead of writing to our file something like

Donald E. Knuth, *Literate Programming*, CSLI, 1992
Jon Bentley, *More Programming Pearls*, Addison-Wesley, 1990
we're writing:

```
Data files
Entry { "Donald E. Knuth",
  "Literate Programming",
  "CSLI",
  1992 }
```

```
Entry { "Jon Bentley",
"More Programming Pearls",
"Addison-Wesley",
1990}
```

Let's remember? that `Entry { code }` is the same as `Entry ({ code })` , that is, a call to the `Entry` function with the table as the only one argument. Therefore, the above piece of data is by itself actually a Lua program. In order to read such a file, we you just need to execute it with a properly defined the `Entry` function . For example, the following program counts the number entries in the file:

```
local count = 0
function Entry () count = count + 1 end
dofile ("data")
print ("number of records:" .. count)
```

The following program builds a set of all author names, find data in the file, and prints them (not necessarily in the same order, in which they met in the file):

```
local authors = {} -- many authors
function Entry (b) authors [b [1]] = true end
dofile ("data")
for name in pairs (authors) do print (name) end
```

Please note the approach used in these snippets code: `Entry` function acts as a callback function (callback), which is called at runtime `dofile` for each doy record in the file.

When we don't care about file size, we can use our views use name-value pairs ¹ :

```
Entry {
author = "Donald E. Knuth",
title = "Literate Programming",
publisher = "CSLI",
year = 1992
}
```

```
Entry {
author = "Jon Bentley",
1
```

If this format reminds you of BibTeX, then it's no coincidence. BibTeX format was one of the sources that defined the kind of constructors in Lua.

```
title = "More Programming Pearls",
year = 1990,
publisher = "Addison-Wesley",
}
```

This format is what we call a *self - describing format*

data, since each piece of data contains a short description vanishing of its meaning. Self-describing data is more readable (like at least people) than CSV or other compact format; their easy to edit if necessary; and they allow us to contribute small changes to the base format without the need to change data files. For example, if we add a new field, then we you only need to change the reading program by providing a value by default when no field is specified.

Using the name-value format, our program for composing

The author list becomes as shown below:

```
local authors = {} -- set for author names
function Entry (b) authors [b.author] = true end
dofile ("data")
for name in pairs (authors) do print (name) end
```

Now the order of the fields is not important. Even if some records there is no author, then we only need to change the `Entry` function :

```
function Entry (b)
if b.author then authors [b.author] = true end
end
```

Lua is not only fast to execute, but also fast to compile.

For example, the above program for making a list authors processes 1 MB of data in one tenth of a second ² . And this not by chance. Data description was one of the main applications Lua since its inception, and we pay a lot of attention to making it the compiler was fast for large programs.

12.2. Serialization

Often we need to serialize some data, that is, translate data into a stream of bytes or characters that we can write to a file or send over the network. We can represent serialized-data as Lua code in such a way that when doing this code it restores the saved values for the executing his programs.

Usually, if we want to restore the value of the global re-variable, then our block of code will be something like `varname = exp` , where *exp* is

this is the Lua code to get the value. With `varname` everything is simple, so Let's see how to write code that creates a value.

For a numeric value, the task is simple:

```
function serialize (o)
  if type (o) == "number" then
    io.write (o)
  else <other cases>
  end
end
```

When writing a number in decimal form, there is a risk of losing precision.

In Lua 5.2, you can use hexadecimal format, in order to
to avoid a similar problem:

```
if type (o) == "number" then
  io.write (string.format ("%a", o))
```

When using this format ("%a") the read value
will consist of exactly the same bits as the original one.

For a string, a naive approach would be something like the following:

```
if type (o) == "string" then
  io.write ("", o, "")
```

However, if the string contains special characters (such as tricks or `\n`), then the resulting code will no longer be a program for Lua.

You might think that this problem can be solved by changing quote type:

```
if type(o) == "string" then
  io.write("[[" .. o .. "]]")
```

However, be careful. If you try to keep something like `]] .. os.execute('rm *') .. [[` (for example, by passing this string as an address), then the resulting block of code will be:

```
varname = [[] .. os.execute('rm *') .. []]
```

As a result, you will receive an unpleasant surprise when trying to purchase honor is such "data".

The simplest way to write a string is safe to use the option `"%q"` from the `string.format` function. It surrounds the string with double quotes and safely represents double quotes and some other characters inside the string:

```
a = 'a "problematic" \'\' string'
print(string.format("%q", a)) -> "a \" problematic \'\' string"
```

Using this capability, our `serialize` function can look like this:

```
function serialize(o)
  if type(o) == "number" then
    io.write(o)
  elseif type(o) == "string" then
    io.write(string.format("%q", o))
  else <other cases>
  end
end
```

Since version 5.1 Lua offers a different way of writing strings in a safe way, using the notation `[=[...]=]` for long lines. However, this recording method is mainly intended for user-written code when we in no way want to change the character string. It is easier in auto-generated code to use `"%q"` from `string.format`.

If you nevertheless want to use a similar notation for automatically generated code, then you need to pay attention to some details. The first is that you need to pick up the correct number of equal signs. A good option is to use a number greater than that found in the original string. Insofar as strings containing a large number of equal signs are not

are rare (for example, comments separating blocks of code), then we can restrict ourselves to considering sequences of equality koks between square brackets; others sequences cannot result in an erroneous end marker strings. The second detail is that Lua always ignores symbol `\n` at the beginning of a long line; the simplest way to deal with this is the addition of the `\n` character , which will be discarded.

Listing 12.1. Outputting an arbitrary string of characters

```
function quote (s)
- find the maximum length of a sequence of equal signs
local n = -1
for w in string.gmatch (s, "[= *]") do
n = math.max (n, #w - 2) - -2 to remove ']'
end
- create a string with 'n' + 1 equal sign
local eq = string.rep ("=", n + 1)
- build a summary line
return string.format ("%s [\n% s]% s", eq, s, eq)
end
```

The `quote` function in Listing 12.1 is the result of our for previous remarks. It receives an arbitrary string as input and returns the formatted string as a long string. Call `string.gmatch` creates an iterator to iterate over all sequential of the form `] = *]` (that is, the closing square bracket, after which followed by zero or more equal signs, followed by there is another closing square bracket) on line 3 . For each occurrences are updated with the value `n` equal to the maximum number already met equal signs. After the loop, we use the function `string.rep` , to repeat the sign of equality `n + 1` times, i.e. one more than the maximum number encountered in line. Finally, the function `string.format` concludes `s` between PA-square brackets with the appropriate number of equal signs and adds extra spaces around the line and the `\n` character in the beginning of the line.

Saving tables without loops

Our next (and more challenging) challenge is to preserve tables. There are several ways to save them in accordance with

with what restrictions we impose on the structure of the table
tasy. There is no one algorithm that fits all cases.
Simple tables not only require simpler algorithms, but also
the resulting files can be visually pleasing.

Listing 12.2. Serializing tables without loops

```
function serialize (o)
if type (o) == "number" then
io.write (o)
elseif type (o) == "string" then
io.write (string.format ("% q", o))
elseif type (o) == "table" then
io.write ("{\n")
for k, v in pairs (o) do
io.write ("", k, "=")
serialize (v)
io.write ("", \n")
end
io.write ("} \n")
else
error ("cannot serialize a" .. type (o))
end
end
```

Our next attempt is shown in Listing 12.2. Despite
for its simplicity, this function does a pretty decent job.
It even handles nested tables (i.e. tables inside
other tables) as long as the table structure is a tree
(that is, there are no shared sub-tables and loops). A little visual
an improvement would be to add spaces to indent nested
tables (see Exercise 12.1).

The previous function assumes that all keys in the table are
are valid identifiers. If the table contains numeric
keys or strings that are not identifiers in Lua, then
we have a problem. A simple way to resolve it is to use
Writing the following code to write each key:

```
io.write ("["; serialize (k); io.write ("] =")
```

With this improvement, we have increased the reliability of our function.
due to the visual clarity of the resulting file. Consider
next call:

```
serialize {a = 12, b = 'Lua', key = 'another "one"'}
```

The result of this call when using the first version of the function
tion serialize will the following code:

```
{
a = 12,
b = "Lua",
key = "another \" one \"",
}
```

Compare with the result of using the second version:

```
{
["A"] = 12,
["B"] = "Lua",
["Key"] = "another \" one \"",
}
```

We can get both reliability and a beautiful view by checking in each case, whether square brackets are needed; again we will leave it improved as exercise.

Saving tables with loops

For processing tables in the general case (that is, with loops and general subtables), we need a different approach. Constructors cannot represent such tables, therefore we will not use. We need names to represent loops, so our next function will take the value as arguments to save and name. Moreover, we must keep track of the names already saved tables in order to reuse them when we discover a cycle. To do this, we will use an additional table. This table will use tables as indexes and their names as stored values.

Listing 12.3. Saving tables with loops

```
function basicSerialize (o)
if type (o) == "number" then
return tostring (o)
else -- suppose it's a string
return string.format ("% q", o)
end
end

function save (name, value, saved)
saved = saved or {} -- initial value
io.write (name, "=")
if type (value) == "number" or type (value) == "string" then
io.write (basicSerialize (value), "\ n")
elseif type (value) == "table" then
if saved [value] then -- is the value already saved?
io.write (saved [value], "\ n") -- use its name
```



```

else
saved [value] = name — save the name for next time
io.write (“{} \n”) — create a new table
for k, v in pairs (value) do — save its fields
k = basicSerialize (k)
local fname = string.format (“% s [% s]”, name, k)
save (fname, v, saved)
end
end
else
error (“cannot save a“ .. type (value))
end
end

```

The resulting code is shown in Listing 12.3. We stick with constraints that the tables we want to keep contain just numbers and strings as keys. BasicSerialize function serializes these base types. The next function, save , does em all the hard work. The parameter saved is a table that is monitors tables already saved. For example, if we build table as follows:

```

a = {x = 1, y = 2; {3,4,5}}
a [2] = a
- cycle
az = a [1]
- general subtable

```

then calling save (“a”, a) will save it like this:

```

a = {}
a [1] = {}
a [1] [1] = 3
a [1] [2] = 4
a [1] [3] = 5
a [2] = a
a [“y”] = 2
a [“x”] = 1
a [“z”] = a [1]

```

The order of these assignments can change, as it depends from traversing the table. Nevertheless, the algorithm guarantees that any the element required for building the table has already been defined .

If we want to store multiple values with common parts, then we can call the save function on the same table saved .

For example, consider the following two tables:

```

a = {{“one”, “two”}, 3}
b = {k = a [1]}

```

If we keep them independently, then the result will have no common

parts:

```
save ("a", a)
save ("b", b)
-> a = {}
-> a [1] = {}
-> a [1] [1] = "one"
-> a [1] [2] = "two"
-> a [2] = 3
-> b = {}
-> b ["k"] = {}
-> b ["k"] [1] = "one"
-> b ["k"] [2] = "two"
```

However, if we use the same saved table for both calls to `save`, then the resulting result will contain the common parts:

```
local t = {}
save ("a", a, t)
save ("b", b, t)
-> a = {}
-> a [1] = {}
-> a [1] [1] = "one"
-> a [1] [2] = "two"
-> a [2] = 3
-> b = {}
-> b ["k"] = a [1]
```

As usual, there are several other options in Lua. Among them, we can store the value without giving it a global name nor (for example, a block builds a local value and returns it), we can process functions (by constructing additional the table associating each function with its name), etc. Lua gives you have strength; you build mechanisms.

Exercises

Exercise 12.1. Modify the code in Listing 12.2 to make it equated nested tables.

(*Hint* : add an extra function parameter `serialize` containing the alignment string.)

Exercise 12.2. Modify the code in Listing 12.2 so that it used the syntax `["key"] = value` as suggested in section 12.1.

Exercise 12.3. Modify the code of the previous exercise so that to make it use the syntax `["key"] = value` only when it's necessary.

Exercise 12.4. Modify the code of the previous exercise so that so that it uses constructors whenever possible-but. For example, he should present the table `{14,15,19}` as `{14,15,19}` , not as `{[1] = 14, [2] = 15, [3] = 19}` .

(*Hint* : start by storing the values for keys 1, 2, ..., until they are *nil* . Please note what is not needed save them again when traversing the rest of the table.)

Exercise 12.5. The no-use approach calling constructors when saving tables with loops, too radical. You can save the table for more nice way, using constructors in general and then using assignments only to handle common tables and loops.

Reimplement the `save` function using this approach. Add to it everything that you have already implemented in the previous exercises.

CHAPTER 13

Metatables and metamethods

Usually, for every value in Lua, there is a completely predictable boron of operations. We can add numbers, connect strings, insert insert key-value pairs into tables, etc. However, we cannot add create tables, we cannot compare functions and we cannot call string. Unless we're using metatables.

Metatables allow you to change the behavior of a value in case when we are faced with an unexpected operation. For example, when the power of metatables we can determine how Lua should compute expression `a + b` , where `a` and `b` are tables. When Lua tries to fold two tables, then it checks if at least one of them has metatable and whether this metatable contains an `__add` field . If Lua finds this field, then it calls the corresponding value - the so-called *my metamethod* , which must be a function - to calculate

the amount.

Every value in Lua can have a metatable associated with it.

Tables and values of type `userdata` store individual values for each instance; values of other types use one general table for each type. Lua always creates new tables without metatables:

```
t = {}  
print (getmetatable (t))  
-> nil
```

We can use the `setmetatable` function to set or change the metatable for any table:

```
t1 = {}  
setmetatable (t, t1)  
print (getmetatable (t) == t1)  
-> true
```

Directly from Lua, we can set metatables only for tables; to work with metatables of values of other types so we have to use the `C` code. We will see later in chapter 21, that the string library sets up metatables for strings. All other types do not have metatables by default:

```
print (getmetatable ("hi")) -> table: 0x80772e0  
print (getmetatable ("xuxu")) -> table: 0x80772e0  
print (getmetatable (10))  
-> nil  
print (getmetatable (print)) -> nil
```

Any table can be a metatable of any value; Group related tables can share a common metatable, which defines their general behavior; the table can be metatable- to herself so that she describes her own behavior nie.

13.1. Arithmetic metamethods

In this section, we will look at a simple example for defining clarify how to use metatables. Let us use the table `ts` for representing sets with functions for calculating the connections, intersections, and so on, as shown in Listing 13.1. For in order not to litter the global namespace, we will store these functions are in the `Set` table .

Listing 13.1. Simple set implementation

```
Set = {}  
- create a new set by taking values from a given list  
function Set.new (l)
```

```

local set = {}
for _, v in ipairs (l) do set [v] = true end
return set
end
function Set.union (a, b)
local res = Set.new {}
for k in pairs (a) do res [k] = true end
for k in pairs (b) do res [k] = true end
return res
end
function Set.intersection (a, b)
local res = Set.new {}
for k in pairs (a) do
1
The main reason for this limitation is the desire to limit too much
frequent use of metatables. Experience with previous versions of Lua has shown that
such global changes often lead to unused code.
res [k] = b [k]
end
return res
end
- represent the set as a string
function Set.tostring (set)
local l = {} — a list where all elements will be placed
for e in pairs (set) do
l [#l + 1] = e
end
return “{“ .. table.concat (l, “,”) .. “}”
end
- print set
function Set.print (s)
print (Set.tostring (s))
end

```

We will now use the addition operator ('+') to calculate numbering the union of two sets. To do this, we will make it so that all tables representing sets will have one common metatable. This metatable will define how the tables should react add to the addition operator. Our first step will be to create a regular table that we will use as a metatable for sets:

```

local mt = {} — metatable for sets

```

The next step is to change the function that creates a lot of the set `Set.new` . The new version of this function will have one additional a string that sets `mt` for the generated tables as a metatable:

```

function Set.new (l) — 2nd version
local set = {}

```

```

setmetatable (set, mt)
for _, v in ipairs (l) do set [v] = true end
return set
end

```

After that, each set that we create with

`Set.new` will have the same metatable:

```

s1 = Set.new {10, 20, 30, 50}
s2 = Set.new {30, 1}
print (getmetatable (s1)) -> table: 00672B60
print (getmetatable (s2)) -> table: 00672B60

```

Finally, we'll add a metamethod to the metatable, the `__add` field, which

The swarm determines how the addition should be performed:

```
mt.__add = Set.union
```

After that, whenever Lua tries to add two

union, it will call the `Set.union` function, passing both
 and as arguments.

With the metamethod, we can use the addition operator for
 performing set union:

```

s3 = s1 + s2
Set.print (s3)
-> {1, 10, 20, 30, 50}

```

Similarly, we can define the multiplication operator for
 completing the intersection of sets:

```

mt.__mul = Set.intersection
Set.print ((s1 + s2) * s1) -> {10, 20, 30, 50}

```

For each arithmetic operator there is a corresponding
 the name of the field in the metatable. Besides `__add` and `__mul`, there is also
`__sub` (for subtraction), `__div` (for division), `__unm` (for negation),
`__mod` (for taking the remainder from division) and `__pow` (for raising the
 stump). We can also define a `__concat` field to specify the opera-
 concatenation generator.

When we add two sets, the question is what meta-
 to take the table, does not arise. However, we can write the expression,
 which involves two values with different metatables, for example
 measures as shown below:

```

s = Set.new {1,2,3}
s = s + 8

```

When looking for a metamethod, Lua takes the following steps: if y
 the first value is a metatable with a `__add` field, then Lua uses
 calls the corresponding value as a metamethod regardless
 from the second value; otherwise, if the second value has metatables
 tsu with the `__add` field, then Lua uses this value as a

there is a method; otherwise, an error occurs. In this way, the last example will call `Set.union`, just like for expressions `10 + s` and `"hello" + s`.

Lua doesn't care about mixing types, but it is important but for our application. For example, if we execute `s = s + 8`, then we get an error inside `Set.union`:

```
bad argument # 1 to 'pairs' (table expected, got number)
```

If we want to receive more accurate error messages, then we must explicitly check the types of the operands before executing the operations:

```
function Set.union (a, b)
  if getmetatable (a) ~= mt or getmetatable (b) ~= mt then
    error ("attempt to 'add' a set with a non-set value", 2)
  end
  <as before>
```

Remember that the second argument to the `error` function (2 in our case) directs the error message to where the given operation was performed called.

13.2. Comparison Methods

Metatables also allow you to make sense of operators compared previously using metamethods `__eq` (*equal*), `__lt` (*less than*) and `__le` (*less than or equal to*). No special metamethods for three remaining comparison operations: Lua translates `a ~ b` to `not (a == b)`, `a > b` in `b < a` and `a >= b` in `b <= a`.

Prior to version 4.0, Lua translated all ordering operations into one, translating `a <= b` to `not (b < a)`. However, such a translation is incorrect when yes we are dealing with *partial ordering*, that is, when not all the elements of our type are properly ordered. For instance, floating point numbers are not fully ordered on most computers because of the *NaN* (*Not a Number*) value.

In accordance with the IEEE 754 standard, NaN represents an undefined specific values such as 0/0. According to the standard, any comparable NaN including NaN must be false. It means that `NaN <= x` is always false, but `x < NaN` is also false. It follows that translating `a <= b` to `not (b < a)` is wrong in this case.

In our example with sets, we are dealing with a similar problem. An obvious (and useful) value for `<=` for sets

is an occurrence of the set: $a \leq b$ means that a is a subset property b . With this value, it is again possible that $a \leq b$ and $b < a$ are false; thus we need separate implementations for `__le` (less or equal) and `__lt` (less than):

```
mt.__le = function (a, b) -- occurrence of sets
for k in pairs (a) do
if not b[k] then return false end
end
return true
end
mt.__lt = function (a, b)
return a <= b and not (b <= a)
end
```

Finally, we can define equality of sets in terms of embedding sets:

```
mt.__eq = function (a, b)
return a <= b and b <= a
end
```

After these definitions, we are ready to compare sets:

```
s1 = Set.new {2, 4}
s2 = Set.new {4, 10, 2}
print (s1 <= s2)
-> true
print (s1 < s2)
-> true
print (s1 >= s1)
-> true
print (s1 > s1)
-> false
print (s1 == s2 * s1) -> true
```

For types that have a complete ordering, we may not define redistribute metamethod `__le`. If not, Lua uses `__lt`.

Comparison for equality also has some restrictions -

mi. If two objects have different base types or metamethods, then comparison operation for equality will return **false** without even calling metamethods. Thus, the set will always be different from the number, no matter what the metamethod returns.

13.3. Library metamethods

So far, we've seen metamethods defined in Lua itself. The virtual machine itself checks if the values contain Operation-defined metatables with corresponding metamethods.

However, since metatables are regular tables, then anyone can use them. Therefore, libraries often define have their own fields in metatables.

The `tostring` function is a typical example. As we have seen earlier, `tostring` represents tables in a fairly simple way:

```
print({})  
-> table: 0x8062ac0
```

The `print` function always calls `tostring` to format output. However, when formatting an arbitrary value `Nia` `tostring` first checks whether the values metamethod `__tostring`. If such a metamethod exists, then `tostring` calls it, passing it an object as an argument. What will return this metamethod, and will be the result of `tostring`.

In our set example, we have already defined a function for representations of the set as strings. Therefore, we only need you-put the `__tostring` field in the metatable:

```
mt.__tostring = Set.tostring
```

After that, whenever we call `print` with a set as with the same argument, `print` will call `tostring`, which in turn will calls `Set.tostring`:

```
s1 = Set.new {10, 4, 5}  
print(s1) -> {4, 5, 10}
```

The `setmetatable` and `getmetatable` functions also use meta-field, in this case to protect the metatable. Suppose you want to protect your sets so that users cannot see nor modify their metatables. If you set the `__metatable` field in a metatable, then `getmetatable` will return the value of that field, and calling `setmetatable` will throw an error:

```
mt.__metatable = "not your business"  
s1 = Set.new {}  
print(getmetatable(s1)) -> not your business  
setmetatable(s1, {})  
stdin: 1: cannot change protected metatable
```

In Lua 5.2, `pairs` and `ipairs` also have metatables, so the table can change its workaround (or add a workaround for non-table objects).

13.4. Metamets for access to the table

Metamets for arithmetic and comparison operations define behavior for situations that would otherwise lead to errors. They do not change normal behavior language. But Lua also provides a way to change the behavior of tables in two usual cases, reading and changing is existing field in the table.

Index metamethod

I said earlier that when we refer to a missing field in the table, the result is ***nil***. This is true, but this is not the whole truth. In fact, such an appeal leads to the fact that the interpretation The tator looks for the `__index` metamethod : if there is no such method, which usually happens, then ***nil*** is returned ; otherwise the result is provided by the given metamethod.

The standard example here is inheritance. Let us we want to create several tables describing windows. Each table tsa must set various parameters of the window, such as position, size, color scheme, etc. For all these parameters there is a value default and therefore we want to build windows by setting only those values that differ from the default values. Per- your choice is a constructor that fills in the missing fields. The second option is to arrange the windows in such a way so that they *inherit* any missing field from the base prototype. First, we will declare a prototype and a constructor that will creates new windows with a common metatable:

```
- create prototype with default values
prototype = {x = 0, y = 0, width = 100, height = 100}
mt = {} - create metatable
- declare a constructor function
function new (o)
  setmetatable (o, mt)
  return o
end
```

We will now define the `__index` metamethod :

```
mt.__index = function (_, key)
  return prototype[key]
end
```

After that, we will create a new window and access the missing th field:

```
w = new {x = 10, y = 20}
print (w.width) -> 100
```

Lua specifies that `w` does not have the required field, but does have metatables. with the `__index` field . Therefore, Lua calls this metamethod with an argument `w` (table) and `"width"` (missing field). Metamethod accesses this field to the prototype and returns the resulting value reading.

Using the `__index` metamethod for inheritance in Lua so common that Lua provides a simplified version.

Despite the name of the *method* , the `__index` metamethod does not have to be function: for example, it can be a table. When he is function, then Lua calls it passing the table and missing-key as arguments, as we have already seen. When this is a table, then Lua simply accesses this table. Therefore, in our in the previous example, we could simply define `__index` as follows- in a way:

```
mt.__index = prototype
```

Now when Lua searches for the `__index` metamethod , it will find prototype value , which is a table. Accordingly, Lua performs access to this table, that is, it performs an analog prototype [`"width"`] . This appeal gives the required result.

Using a table as the `__index` metamethod makes it easy to one and a quick way to implement the usual (not multiple) na-followings. The function is a more expensive option, but and provides more flexibility: we can implement multiple inheritance, caching, and more. We are we judge these forms of inheritance in chapter 16.

When we want to access a table without calling the metamethod `__index` , then we use the `rawget` function . Calling `rawget (t, i)` performs a direct access to the table `t` , that is, scheduling without using metatables. Executing directly- this call will not speed up your code (the cost of the function call will destroy all that can be won), but sometimes it turns out to be necessary, as we will see later.

`__newindex` metamethod

The `__newindex` metamethod is analogous to the `__index` metamethod , but only it works for writing values to a table. When you assign-

if you give the value to the missing field in the table, then the interpreter looks for the `__newindex` metamethod : if it exists, then the interpreter calls instead of doing the assignment. Like `__index` if metamethod is a table, then the interpreter performs the assignment for this table instead of the original one. Moreover, there is a function `rawset` that performs direct access, bypassing metamethods: `rawset(t, k, v)` writes the value `v` by key `k` to table `t`, not calling no metamethods.

Using the `__index` and `__newindex` Metamethods Together allows you to implement in Lua various rather powerful constructs such as read-only tables, tables with defaults and inheritance for object-oriented programming. In this chapter, we will see some of their applications. Object-oriented programming a separate chapter is allocated.

Tables with default values

The default value for any field in a regular table is *nil*.

It's easy to change this behavior with metatables:

```
function setDefault(t, d)
  local mt = {__index = function() return d end}
  setmetatable(t, mt)
end
tab = {x = 10, y = 20}
print(tab.x, tab.z) --> 10 nil
setDefault(tab, 0)
print(tab.x, tab.z) --> 10 0
```

After calling `setDefault`, any call to the missing field in `tab` will call its `__index` metamethod, which will return zero (value `d` for this metamethod).

The `setDefault` function creates a new closure and a new metatable. It is not good for each table that needs a default value. It can be costly if we have many tables that need default values. The metatable has a default value `nil` and «sewn up» in her metamethod, so we can not use the same metatable for all tables. So that you can use the same metatable for tables with different default values, we can remember the default value in the table itself, using a special field for this. If a

not think about possible name conflicts, then we can use call a key like “__” for our field:

```
local mt = {__index = function (t) return t.__ end}
function setDefault (t, d)
  t.__ = d
  setmetatable (t, mt)
end
```

Note that now we are creating the `mt` table only one times, outside of the `setDefault` function .

If we want to guarantee the uniqueness of the key, then this is but easy to provide. All we need is to create a new table and use it as a key:

```
local key = {} - unique key
local mt = {__index = function (t) return t[key] end}
function setDefault (t, d)
  t[key] = d
  setmetatable (t, mt)
end
```

Another way to associate a default value with each table is the use of a separate table, where the keys are the tables themselves, and the values are the default values.

However, for the correct implementation of this approach, we need special a special type of tables called *weak tables* , therefore, we will not use this approach here; we will return- See this in chapter 17.

Another option is to remember metatables, whereby we can reuse metatables corresponding to one the same default value. However, this also requires using weak tables, so we'll have to wait until Chapter 17.

Tracking table access

Both `__index` and `__newindex` work only when in the table no corresponding value. Therefore, the only way to keep track of all access to a table is to keep it empty. Thus, zom, if we want to track all access to the table, then we need create a special *proxy* table for the source table. She will- det empty with appropriate `__index` and `__newindex` metamets to track access to the table that will redirect access to the original table. Let `t` be the original table, access

to which we want to track. Then we can use the following blowing code:

```
t = {} - the source table was created somewhere
- create private access to it
local _t = t
- create a proxy
t = {}
- create a metatable
local mt = {
  __index = function (t, k)
    print ("* access to element" .. tostring (k))
    return _t [k] - access to the source table
  end,
  __newindex = function (t, k, v)
    print ("* update of element" .. tostring (k) ..
      "To" .. tostring (v))
    _t [k] = v - change the original table
  end
}
setmetatable (t, mt)
```

This code keeps track of every access to t :

```
> t [2] = "hello"
* update of element 2 to hello
> print (t [2])
* access to element 2
hello
```

If we want to be able to traverse such a table, then we you need to create a __pairs metamethod in the proxy table :

```
mt.__pairs = function ()
  return function (_, k)
    return next (_t, k)
  end
end
```

It is also possible to create something similar for __ipairs .

If we want to track access to multiple tables, then we there is no need to create a separate metatable for each of them. Inmes- then we can somehow link the proxy table with the original and use one common metatable for all proxy tables. it is similar to the task of linking a table with a default value, which we looked at earlier. For example, you can store the outcome a new table in a special field of the proxy table, using for this special key. As a result, we end up with the following code:

```
local index = {} -- create a unique key
local mt = {} - create metatable
__index = function (t, k)
```

```

print (“* access to element“ .. tostring (k))
return t [index] [k] - access to the source table
end,
__newindex = function (t, k, v)
print (“* update of element“ .. tostring (k) ..
“To“ .. tostring (v))
t [index] [k] = v - change the source table
end,
__pairs = function (t)
return function (t, k)
return next (t [index], k)
end, t
end
}
function track (t)
local proxy = {}
proxy [index] = t
setmetatable (proxy, mt)
return proxy
end

```

Now, when we want to keep track of the table `t`, all we need is
but, is to execute `t = track (t)`.

Read-only tables

It is easy to use the concept of proxy tables to create tables with pre-read-only mortar. All we need is to trigger an error every time we catch an attempt to change the table. For metame-Toda `__index` we can use the very original table instead functions, since we don't need to keep track of all reads from it; faster and more efficiently redirect such requests directly to the original table-face. This will require, however, a new metatable for each proxy-tables with an `__index` field pointing to the original table:

```

function readOnly (t)
local proxy = {}
local mt = {- create metatable
__index = t,
__newindex = function (t, k, v)
error (“attempt to update a read-only table”, 2)
end
}
setmetatable (proxy, mt)
return proxy
end

```

As an example of a read-only table, we can

Let's create a table of names of days of the week:

```
days = readOnly { "Sunday", "Monday", "Tuesday", "Wednesday",  
"Thursday", "Friday", "Saturday" }
```

```
print (days [1]) -> Sunday
```

```
days [2] = "Noday"
```

```
stdin: 1: attempt to update a read-only table
```

Exercises

Exercise 13.1. Define a `__sub` metamethod that returns the difference between the two sets. (The set `ab` is the set in all elements from `a` that are not contained in `b`.)

Exercise 13.2. Determine metamethod `__len` so that `#s` WHO-rotates the number of elements in `s`.

Exercise 13.3. Complete the implementation of proxy tables in section 13.4 with the `__ipairs` metamethod.

Exercise 13.4. Another way to implement tables, access read-only, is to use the function as the `__index` metamethod. This approach makes access to table more expensive, but the creation of such tables is cheaper, as all read-only tables can have one common metatable. Rewrite the `readOnly` function with using this approach.

Exercises

Chapter 14

Environment

Lua stores all of its global variables in a regular table, called the *global environment*. (More precisely, Lua keeps its "global" variables in several environments, but we will ignore this at first for simplicity.)

the advantage of this approach is that it simplifies the internal Lua implementation as there is no need for special structure data for storing global variables. Another advantage the fact is that we can work with this table in the same way as with any other table. To make things easier, Lua stores the environment itself in the global variable `_G`. (Yes, `_G._G` is equal

but `_G`.) For example, the following code prints the names of all global variables defined in the global environment:

```
for n in pairs (_G) do print (n) end
```

In this chapter, we will see some useful methods for working with the environment.

14.1. Global Variables with dynamic names

Usually an assignment is enough to access and set a value. a global variable. However, we often need the option metaprogramming when we want to work with global re-variable whose name is contained in another variable or computed is done in the course of work. To get the value of such a variable, many some programmers try to use something like the following snippet of code:

```
value = loadstring ("return" .. varname) ()
```

If `varname` is `x`, then as a result of concatenation we get

"Return x", which when executed will give us the desired result. Od-

Global variables with dynamic names

but this code involves creating and compiling a new block code, which is costly. You can achieve the same

go using the following code, which is more than an order of magnitude more more efficient than the previously discussed code:

```
value = _G [varname]
```

Since the environment is a regular table, you can simply access it by key (variable name). Similarly

you can also assign a value to a variable whose name computes-dynamically, using the code `_G [varname] = value`. However, be careful: some programmers are so happy about this opportunity ness that end up writing code like `_G ["a"] = _G ["var1"]`, which is just a tricky option for `a = var1`.

The generalization of the previous task is to use names fields in dynamic names such as "io.read" or "abcd".

However, if we write `_G ["io.read"]`, then we will definitely not get read field from table `io`. But we can write a `getfield` function, such that `getfield ("io.read")` will return the expected value. This

a function is a loop that starts with `_G` and then she sequentially iterates over the fields:

```
function getfield (f)
local v = _G -- start with global variable table
for w in string.gmatch (f, "[% w _] +") do
v = v [w]
end
return v
end
```

We use the `gmatch` function from the `string` library to bypass all words in `f` (a word is a sequence of letters, numbers and underscore).

The corresponding function for setting field values is more complicated. An assignment like `abcd = v` is equivalent to the following code:

```
local temp = abc
temp.d = v
```

That is, we have to extract the name without the last component and then process the last component separately. Function `setfield` You are a completes this and also creates helper tables in the path if they don't exist:

```
function setfield (f, v)
local t = _G -- we start with the table of global variables
for w, d in string.gmatch (f, "([% w _] +) (%.?)") do
if d == "." then - not the last name?
t [w] = t [w] or {} -- creates a table if it does not exist
t = t [w]
- we get the table
else
-- last name
t [w] = v
- perform the assignment
end
end
end
```

The variable `w` stores the name of the field, and possibly the following after it the point is stored in the variable `d`. If the name is not followed - there is a point, then this is the last name.

Using the previously discussed functions, the following code creates global table `t`, table `tx` and then assigns 10 `txy`:

```
setfield ("txy", 10)
print (txy) -> 10
print (getfield ("txy")) -> 10
```

14.2. Descriptions of global variables

In Lua, global variables do not need declarations. Although it is convenient but for small programs, in large programs there is only one opera-

tion that can lead to hard-to-find bugs. but

we can change this behavior if desired. Since Lua stores global variables in a regular table, then we can use use metatables to change behavior when referring to global variables.

The first approach simply tracks any calls to the missing the following keys in the global table:

```
setmetatable(_G, {  
  __newindex = function (_, n)  
    error ("attempt to write to undeclared variable" .. n, 2)  
  end,  
  __index = function (_, n)  
    error ("attempt to read undeclared variable" .. n, 2)  
  end,  
})
```

After executing this code, any attempt to address the wrong existing global variable will cause errors:

```
> print(a)  
stdin: 1: attempt to read undeclared variable a
```

However, how are we going to declare global variables? One an option is to use `rawset` which doesn't use metamethods:

```
function declare(name, initval)  
  rawset(_G, name, initval or false)  
end
```

(The **or false** construct is needed so that the global variable got a value other than **nil**.)

A simpler option is to restrict the assignments to but-global variables only inside functions, allowing piling on the outer level of the block.

To check that the assignment occurs in the main block ke, we need to use a debug library. Call `debug`.

`getinfo(2, "S")` returns a table whose `what` field says whether the function that called the metamethod is the main block,

normal function or C-function. (We will discuss `debug.getinfo` in more detail in Chapter 24.) Using this function, we can rewrite the `__newindex` metamethod as follows:

```
__newindex = function (t, n, v)
  local w = debug.getinfo (2, "S"). what
  if w ~ = "main" and w ~ = "C" then
    error ("attempt to write to undeclared variable" .. n, 2)
  end
  rawset (t, n, v)
end
```

This new version also allows assignments in C code, since usually in this code the authors know what they are doing. To check that such a variable exists, we cannot just compare it to ***nil***, because if it is ***nil***, then the children to error. Instead, we use the `rawget` function, which doesn't use metamethod:

```
if rawget (_G, var) == nil then
  - 'var' is undeclared
...
end
```

Now our approach does not allow global variables with a value ***nil***, because they will automatically be considered neobyavlenym. But this is easy to fix. All we need is an additional table containing the names of the described variables. When calling `rawget` in the metamethod, it checks against this table whether this variable is described. Similar code is shown in Listing 14.1. Now even assignments `x = nil` is enough to declare a global variable.

The cost of both solutions is extremely low. At the first decision, in normal operation, the metamethod is not called at all. At the second decision, metamethods can be called when the program is growing to a variable whose value is ***nil***.

The standard Lua package contains the `strict.lua` module, which implements checking of calls to global variables, similar to the code we have considered. A good habit is to use it when writing Lua code.

Listing 14.1. Checking global variable declarations

```
local declaredNames = {}
setmetatable (_G, {
  __newindex = function (t, n, v)
    if not declaredNames [n] then
      local w = debug.getinfo (2, "S"). what
      if w ~ = "main" and w ~ = "C" then
        error ("attempt to write to undeclared variable" ..n, 2)
      end
    end
    rawset (t, n, v)
  end
})
```

```

end
declaredNames [n] = true
end
rawset (t, n, v) - do the actual set
end,
__index = function (_, n)
if not declaredNames [n] then
error ("attempt to read undeclared variable" ..n, 2)
else
return nil
end
end,
})

```

14.3. Non-global environments

One of the problems with the environment is that it is global. Any change to it affects all parts of your program. On-example when you set metatable to control global access, your entire program must follow the appropriate policy. If you want to use a library, which uses global variables without declaring them, then you were unlucky.

In Lua, global variables don't have to be truly global ballroom. We can even say that Lua has no global roomen. This may sound strange, since from the very beginning of the book we used global variables. Obviously, Lua is very tries to create the illusion of having global variables. Come on-Let's see how Lua creates this illusion 2 .

Let's start with the concept of free names. A *free name* is not a name tied to an explicit description, that is, it is not found inside the area These are the actions of a local variable (or a **for** loop variable , or parameter) with this name. For example, both names `var1` and `var2` are free names in the next block:

```
var1 = var2 + 3
```

Unlike what was said earlier, the free name does not relate refers to a global variable (at least not directly but). Instead, Lua translates any free `var` name into `_ENV`.

`var` . Therefore, the previous block is equivalent to the following:
`_ENV.var1 = _ENV.var2 + 3`

But what is this new `_ENV` variable ? It cannot be glo-

point variable, otherwise we again return to the original sample-
leme. The compiler cheats again. I have already said that Lua will consider
Treats each block as an anonymous function. Actually Lua
compiles our source block into the following code:

```
local _ENV = <some value>
return function (...)
  _ENV.var1 = _ENV.var2 + 3
end
```

That is, Lua compiles any block of code in the presence of
limited value named `_ENV` .

Usually, when we load a block of code, the `load` function is initialized
lyses this predefined value with a reference to the global oc-
rifle. Therefore, our original block becomes equivalent to
next block:

```
local _ENV = <the global environment>
return function (...)
  _ENV.var1 = _ENV.var2 + 3
end
2
```

Note that this mechanism was one of those parts of Lua that included
Were from version 5.1 to version 5.2. The following discussion is specific to Lua
5.2 and very little applies to previous versions.

The result of all these assignments is that the `var1` field from
the global environment gets `var2` plus 3.

At first glance, this may seem a little confusing.

a way to work with the global environment. I will not argue
that this is the simplest way, but it achieves the flexibility that is difficult
but get a simpler implementation.

Before we continue, let's articulate how Lua 5.2 works.

works with global variables:

- Lua compiles any block using a value `_ENV` .
- The compiler translates any free `var` name into `_ENV.var` .
- The `load` (or `loadfile`) function initializes the `_ENV` value
a reference to the global environment.

In the end, it's not all that difficult.

Some are confused because they are trying to find some
the magic behind these rules. There is no magic here. In part
However, the first two rules are completely done by the compiler. Behind
except that `_ENV` is known to the compiler, it
is an ordinary variable. Except for compilation, `_ENV`

doesn't make any special sense in Lua 3 . Similarly translation from `var` to `_ENV.var` is just syntactic replacement without hidden meaning. In particular, after this translation, `_ENV` will refer to the `_ENV` variable , which is visible in this code snippet, is from the rules of visibility.

14.4. Using `_ENV`

In this section we will look at some of the ways to use the flexibility that the `_ENV` variable brings . Please be aware that each of these examples should be run as a separate powerful block of code. If you enter line by line in the interpreter, then each line becomes a separate block and gets its `_ENV` variable . To execute a piece of code as a separate block, you either need to run it as a file or in interactively put inside a ***do-end*** pair .

Since `_ENV` is an ordinary variable, we can assign read it and read it just like any other variable. Assigned `_ENV = nil` prohibit any access to global variables on throughout the remainder of the block. This can be useful for control which variables your code uses:

```
local print, sin = print, math.sin
_ENV = nil
print (13)
-> 13
print (sin (13)) -> 0.42016703682664
print (math.cos (13)) - error!
```

Any assignment to a free name will result in a similar error.

We can explicitly call `_ENV` to bypass the locale. variable variables:

```
a = 13
- global
local a = 12
print (a) -> 12 (local)
print (_ENV.a) -> 13 (global)
```

Of course, the main use of `_ENV` is to change ok-tool used by the code snippet. Once you have changed your environment, all calls to global variables will be used use a new table:

- change the current environment to an empty table

```
_ENV = {}
```

a = 1 - create a field in _ENV

```
print (a)
```

-> stdin: 4: attempt to call global 'print' (a nil value)

If the new environment is empty, then you lose access to all global variables including `print`. Therefore, you first need to fill it with some useful values, like the old environment:

```
a = 15
```

- create a global variable

```
_ENV = {g = _G} - change the current environment
```

```
a = 1
```

- create a field in _ENV

```
g.print (a)
```

-> 1

```
g.print (ga) -> 15
```

Now when you access the "global" `g`, you get sta

Swarm environment in which there is a `print` function.

We can rewrite the previous example using the name `_G` instead of `g`:

```
a = 15
```

- create a global variable

```
_ENV = {_G = _G} - change the current environment
```

Using `_ENV`

```
a = 1
```

- create a field in _ENV

```
_G.print (a)
```

-> 1

```
_G.print (_G.a) -> 15
```

For Lua, the name `_G` is the same name as everyone else. His from-the only characteristic is that when Lua creates a global

score table, then it assigns it to a variable named `_G`. For

Lua doesn't care about the current value of this variable. But it is usually accepted

use the same name when we refer to the global variable, as we did in the rewritten example.

Another way to populate your new environment is by inheritance

nie:

```
a = 1
```

local newgt = {} - create a new environment

```
setmetatable (newgt, {__index = _G})
```

```
_ENV = newgt
```

- install it

```
print (a) -> 1
```


In this code, the new environment inherits `print` and `a` from the old environment. `zheniya`. However, any assignment goes to the new table. Thereby there is no danger of mistakenly changing the global environment, although it can still be changed via `_G` :

- continue the previous code

```
a = 10
```

```
print (a) -> 10
```

```
print (_G.a) -> 1
```

```
_G.a = 20
```

```
print (_G.a) -> 20
```

Since `_ENV` is a regular variable, it obeys

the usual rules of visibility. In particular, the functions defined inside a block, refer to `_ENV` just like any other external

her variable:

```
_ENV = {_G = _G}
```

```
local function foo ()
```

```
_G.print (a)
```

- compiles to '`_ENV._G.print (_ENV.a)`'

```
end
```

```
a = 10
```

```
- _ENV.a
```

```
foo ()
```

```
-> 10
```

```
_ENV = {_G = _G, a = 20}
```

```
foo ()
```

```
-> 20
```

If we define a new local variable named `_ENV` , then access to free names will go through it:

```
a = 2
```

```
do
```

```
local _ENV = {print = print, a = 14}
```

```
print (a) -> 14
```

```
end
```

```
print (a) -> 2 (back to original _ENV)
```

Therefore, it is not difficult to build a function with its own (private) environment:

```
function factory (_ENV)
```

```
return function ()
```

```
return a
```

```
- “global” a
```

```
end
```

```
end
```

```
f1 = factory {a = 6}
```

```
f2 = factory {a = 7}
```

```
print (f1 ())
-> 6
print (f2 ())
-> 7
```

The `factory` function creates simple closures that return the value of local variables `a`. When the closure is created, then the visible variable `_ENV` is the parameter `_ENV` from the containing function `factory`; so the closure uses this variable to add a `stupa` for free names.

Using normal visibility rules, we can work with environments in various ways. For example, we may have several functions with a common environment for them or a function that Toraya changes the environment in common with other functions.

14.5. `_ENV` and `load`

As I mentioned, `load` usually initializes the `_ENV` value the loaded block is a pointer to the global environment. However, `load` has an optional fourth parameter that specifies the value reading for `_ENV`. (The `loadfile` function also has a similar parameter.)

As an example, let's say we have a typical configuration a file that defines the various constants and functions used washed by the program; it could be something like:

```
- file 'config.lua'
width = 200
height = 300
...
_ENV and load
```

We can load it with the following code:

```
env = {}
f = loadfile ("config.lua", "t", env)
f ()
```

All code from the config file will be executed with empty `env` environment. More importantly, all of its definitions will be named but into this environment. The config file cannot affect anything else, even by mistake. Even malicious code cannot cause thread a lot of harm. It can perform a DoS attack by wasting CPU time

and memory, but nothing else.

Sometimes you may need to execute a block multiple times, each time with a different environment table. In this case, additional the `load` argument doesn't help us. Instead, we have two option.

The first option is to use the `debug.setupvalue` function from the `debug` library. As the name implies, `setupvalue` tells us to change any incoming value (`upvalue`) of the given functions. The following code illustrates its use:

```
f = loadfile (filename)
```

```
...
```

```
env = {}
```

```
debug.setupvalue (f, 1, env)
```

The first argument when calling `setupvalue` is a function, the second is this is the index of the value, and the third is the new value. For our use use of the second argument is always one: when the function is the result of `load` or `loadfile` , Lua guarantees that there will be only one value and that value is `_ENV` .

A small disadvantage of this solution is the dependence from the `debug` library. This library breaks some standard assumptions about programs. For example `debug.setupvalue` violates Lua's visibility rules, which guarantee that the variable cannot be seen outside of its scope visibility.

Another way to run a block with different environments is

There is a slight change in the block when it is loaded. Imagine, that we add the following line to the beginning of the loaded block:

```
_ENV = ...;
```

Recall from Section 8.1 that Lua treats any block as function with variable number of arguments. Therefore, this line attaches assigns the first block argument to `_ENV` , setting it as an environment. After loading the block, we call the resulting function, passing the desired environment as the first argument. Next The following code snippet illustrates this idea using the function

`loadwithprefix` from Exercise 8.1:

```
f = loadwithprefix ("local _ENV = ...;", io.lines (filename, "* L"))
```

```
...
```

```
env = {}
```

```
f (env)
```

Exercises

Exercise 14.1. The `getfield` function that we have defined at the beginning of this chapter provides too little role, since it allows fields such as `math? sin` or `string !!! gsub`. Rewrite it to handle only one dot as a separator. (For this exercise (you may need the information in Chapter 21.)

Exercise 14.2. Explain in detail what happens next program and what its output will be.

```
local foo
do
  local _ENV = _ENV
  function foo () print (X) end
end
X = 13
_ENV = nil
foo ()
X = 0
```

Exercise 14.3. Explain in detail what happens next program and what its output will be.

```
local print = print
function foo (_ENV, a)
  print (a + b)
end
foo ({b = 14}, 12)
foo ({b = 10}, 1)
```

Exercises

Chapter 15

Modules and packages

Lua does not usually establish any conventions. Instead of this Lua provides mechanisms that are powerful enough for groups developers to implement the conventions that suit them.

However, this approach does not work well for modules. One of the main the purpose of the module system is to allow different people to share locally use the code. The lack of a common policy prevents this sharing.

Since version 5.1, Lua has defined a set of conventions for modules and packages (a package is a collection of modules). These agreements are not required

any additional features from the language; programmer

you can implement them using what we have already seen in the language:

tables, functions, metatables and environments. Programmers can

use other agreements. However, other agreements may

lead to the fact that it will not be possible to use other people's modules and your modules cannot be used in other people's programs.

From the user's point of view, a *module* is some code (in Lua or in C), which can be loaded with `require` and `co-`

which creates and returns a table. Anything the module exports

whether it be functions or tables, it defines inside this table,

which acts as a namespace.

For example, all standard libraries are modules. You can

use math library like this:

```
local m = require "math"
```

```
print (m.sin (3.14))
```

However, a separate interpreter (available as a command

strings) preloads all standard libraries with

code equivalent to the following:

```
math = require "math"
```

```
string = require "string"
```

```
...
```

This download allows us to use the normal `math.sin` notation .

The obvious advantage of using tables to implement

modules is that we can work with modules in the same way as

with tables, and use the power of Lua to do so. In most

In two languages, modules are not first class values (i.e.

they cannot be stored in variables, passed as arguments

functions, etc.), so these languages need special mechanisms

terms for every opportunity they want to offer for

modules. In Lua, you get these features for free.

For example, there are several ways to call a function from

module. The usual way is as follows:

```
local mod = require "mod"
```

```
mod.foo ()
```

User can set any local name for the module:

```
local m = require "mod"
```

```
m.foo ()
```

Alternative names can also be provided for individual

functions:

```
local m = require "mod"
```

```
local f = mod.foo
```

```
f ()
```

The nice thing about these features is that they don't require special support from the language. They only use what the language already provides.

A common complaint about `require` is that this function

This does not allow passing an argument to the loaded module. For example measures, the mathematical module could receive an argument that allows to choose between using degrees or radians:

- bad code

```
local math = require ("math", "degree")
```

The problem is that one of the main objectives `require` yav-

Avoid loading an already loaded module. As soon as possible the muzzle is loaded, it will be reused by any part of the program, who needs it. Therefore, when using the parameters there would be a problem if the same module was needed, but with other parameters:

- bad code

```
local math = require ("math", "degree")
```

- somewhere else in the same program

```
local math = require ("math", "radians")
```

In case you really want your module to be held parameters, it is better to create an explicit function to set them niya:

```
local mod = require "mod"
```

```
mod.init (0, 0)
```

If the initializing function returns the module itself, then we we can write code like the following:

```
local mod = require "mod" .init (0, 0)
```

Another option is to make the module return a function for initialization and already this function would return a table muzzle:

```
local mod = require "mod" (0, 0)
```

In any case, remember that the module is loaded only once; the module itself must resolve initializations with conflicts.

15.1. Require function

The `require` function tries to minimize the assumptions about

what is a module. For `require`, a module is just something that defines some values (such as functions or tables containing functions). Usually this code returns a table containing the functions of this module. However, since this is done by the code of the module itself, not `require`, some modules may choose to return different values or even have sideways effects.

To load a module, we simply call `require "modname"`.

The first step of `require` is to check against the `package`.

`loaded` to see if this module is already loaded. If so, then `require` returns the corresponding value. Therefore, as soon as the module is loaded, other calls requiring this module to be loaded are simply returned the same value without executing any code.

If the module has not yet been loaded, then `require` looks for a Lua file with the name

of the module. If it finds such a Lua file, then it loads it with

`loadfile`. The result of this is the function that we

called the *bootloader*. (The bootloader is a function that, when called, returns a module.)

If `require` cannot find a Lua file with a module name, then it looks for a C library with a module name. If it finds the right C library, it loads it using `package`.

`loadlib` (which we discussed in section 8.3) and looks for a function named `luaopen_modname`. In this case, the bootloader is the result of `loadlib`, that is, a `luaopen_modname` function that looks like a function in Lua.

Whether the module is a Lua file or a C

library in C, `require` now has a loader for it. For windows-

When the module is loaded, `require` calls the loader with two arguments: the name of the module and the name of the file with the loader.

(Pain-

Most modules simply ignore these arguments.)

When `require` returns some value, then `require` returns this value

and store it in the `package.loaded` table in order to

always return exactly this value for this module. If for-

the loader does not return anything, then `require` behaves the same as if

the module would return **true**. Without this clarification, subsequent calls

to `require` would execute this module again.

To force `require` to load the specified module, times, we simply erase the entry for this module from the `package.loaded`.

loaded :

`package.loaded[< modname >] = nil`

The next time you need this module, the `require` probes all the necessary work again.

Renaming a module

We usually use their original name as the module name. names, but sometimes we have to rename a module to avoid name conflict. A typical situation is loading different versions of the same module, for example for testing. Modules do not have hardcoded names inside them, so usually just rename the corresponding `.lua` file. but we cannot edit the binary library to change the name of its function `luaopen_*`. In order to maintain similar new renames, there is a little trick inside `require` : if the module name contains a minus, then `require` strips off the part of the name up to before the minus sign when creating the function name `luaopen_*`. For example measure, if the module name is `ab`, then `require` expects the corresponding the next function will be named `luaopen_b`, not `luaopen_a-b` (which wouldn't be a valid name in C anyway). therefore if we need to use two modules named `mod`, then we can Let's rename one of them to `v1-mod`, for example. When we call `require "v1-mod"`, `require` will find the renamed file `v1-mod` and inside this file will find a function named `luaopen_mod`.

Search along the path

When searching for a file in Lua, `require` uses a search path that slightly differs from the usual search paths. Typical path is a list of directories where to search for the given file. One- in ANSI C (the abstract platform that Lua runs on) there is no concept of a directory. Therefore the path used by `require` is a list of *templates*, each of which specifies its own way of converting setting the module name (the `require` argument) to the file name. More accurately, each pattern in the path is a filename containing optional

question marks. For each pattern, `require` replaces every '?' on module name and checks if there is a file with the corresponding name; if not, it moves to the next template. Patterns out of the way separated by semicolons (a character rarely used in filenames on modern operating systems). For instance, if the path is

```
?;?. lua; c: \ windows \ ?; / usr / local / lua / ? / ?. lua ,
```

then the call to `require` ("sql") will try to open the following files:

```
sql
```

```
sql.lua
```

```
c: \ windows \ sql
```

```
/usr/local/lua/sql/sql.lua
```

The `require` function uses as special characters

only semicolon (as component separator) and question mark-sign; everything else, including path delimiters and extensions files is defined by the path itself.

The path that `require` uses to find files in Lua is

always the current value of the `package.path` variable . When starting Lua it initializes this variable with the value of the next variable-environment `LUA_PATH_5_2` . If this environment variable is not installed, then Lua tries to use the environment variable named `LUA_PATH` . If both are undefined, then Lua uses the default path, set at compile time `2` . When using environment variables, Lua substitutes the default path instead of any substring ";" ... For example, if `LUA_PATH_5_2` is equal to "mydir / ?. lua ;;" then the final path will be pattern "Mydir / ?. lua" followed by the default path.

The C library search path works the same way, but

`cpath` is taken from the variable `package.cpath` (instead of `package.path`). Similarly, this variable gets its initial value

from the environment variable `LUA_CPATH_5_2` or `LUA_CPATH` . Typical the value for UNIX systems is

```
./?.so;/usr/local/lib/lua/5.2/?.so
```

Note that the path defines the file extension. Pre-

The previous example uses `.so` for all templates; in Windows t-

The pictorial template will look like the following:

```
. \ ?. dll; C: \ Program Files \ Lua502 \ dll \ ?. dll
```

The `package.searchpath` function implements all of these conventions for search for libraries. It gets the module name and path and looks for the file, following the rules described above. It returns either the name of the first

found file, or *nil* and an error message describing all the files she tried to open, as in the following example:

```
> path = “. \?. dll; C: \ Program Files \ Lua502 \ dll \ ?. dll”  
> print (package.searchpath (“X”, path))  
nil  
no file '. \ X.dll'  
no file 'C: \ Program Files \ Lua502 \ dll \ X.dll'
```

File crawlers

In reality, `require` is somewhat more complex than we have described. Finding a file in Lua and finding a C library are just two private the more general concept of a file searcher. Seeker file is just a function that takes the name of the module and returns sets a bootloader for this block, or *nil* if it cannot find any one.

The `package.searchers` array contains a list of file searchers, which `require` . When looking for a module, `require` calls each crawler in turn, passing it the module name, until

2

In Lua 5.2, the `-E` command line parameter prevents the use of variables environment and causes the compile-time path to be used.

until it finds a loader for the module. If the search ends in that is, `require` raises an error.

Using a list to control the search for a module gives more flexibility of the `require` function . For example, if you want to store thread the modules compressed into zip files, then all you need to do this is - this is to provide the appropriate finder function and add it to the list. However, more often than not, programs do not need to change `package.searchers` value . The default configuration is the Lua function finder and the C library crawler we described above, occupy the second and third positions. Before them there is a claim preload searcher.

This finder allows you to enter an arbitrary function for module load. It uses the `package.preload` table to match setting the names of the modules of the boot functions. When searching for this the crawler simply looks for the given name in the table. If he finds a function tion, then it returns it as a module loader. Otherwise it returns *nil* . This finder provides a way to handle some non-

typical cases. For example, a C library statically prelinked to Lua, can register its `luaopen_` function like this, that it will be called only when (and if) the user needs this module is being. Thus, the program does not waste time on opening a module if not in use.

By default `package.searchers` includes the fourth function that is needed for submodules. We will consider them at a time case 15.4.

15.2. Standard Approach for writing modules in Lua

The simplest way to create a module in Lua is really simple: we create a table, put all the functions we want to export into it, inside it and return this table. Listing 15.1 demonstrates this approach. Note how we define the function `inv` as a closed, simply declaring it inside the unit.

Some people don't like the terminating `return` statement. One of ways to eliminate it is to write the module table directly

```
in package.loaded :  
local M = {}  
package.loaded [...] = M  
<as before>
```

Be aware that `require` calls the loader passing the name of the module as the first argument. Therefore, the variable number expression arguments `...` results in exactly that name. After that, when the module does not return a value, then `require` will return the current value `package.loaded [modname]` (if not ***nil***). However, I prefer to return a table as it looks neater.

Another way to write a module is to define all the functions and build the table at the end, as in Listing 15.2.

What are the advantages of this approach? You don't have to start every name with `M.` or something similar; there is an explicit export list of functions; you define and use exported and internal functions exactly the same inside a module. In what are the disadvantages of this approach? Exported List of functions is at the end of the module, not at the beginning where it would be

more convenient as a quick reference; and a list to export from- is accurate, since you need to write down each name twice. (This last disadvantage can be an advantage, since it allows functions have different names outside the module and inside it, but I think that programmers rarely use this.) I personally like- This style is being developed.

However, remember that no matter how you define the module, users should be able to use it in a standard way:

```
local cpx = require "complex"
print (cpx.tostring (cpx.add (cpx.new (3,4), cpx.i)))
-> (3.5)
```

Listing 15.1. Simple module for complex numbers

```
local M = {}
function M.new (r, i) return {r = r, i = i} end
- define the constant 'i'
Mi = M.new (0, 1)
function M.add (c1, c2)
return M.new (c1.r + c2.r, c1.i + c2.i)
end
function M.sub (c1, c2)
return M.new (c1.r - c2.r, c1.i - c2.i)
end
function M.mul (c1, c2)
return M.new (c1.r * c2.r - c1.i * c2.i, c1.r * c2.i + c1.i * c2.r)
end
local function inv (c)
local n = c.r ^ 2 + c.i ^ 2
return M.new (c.r / n, -c.i / n)
end
function M.div (c1, c2)
return M.mul (c1, inv (c2))
end
function M.tostring (c)
return "(" .. c.r .. " , " .. c.i .. ")"
end
return M
```

Listing 15.2. Module with an explicit list of exported functions

```
local function new (r, i) return {r = r, i = i} end
- define the constant 'i'
local i = complex.new (0, 1)
<other functions follow the same pattern>
return {
new = new,
i = i,
add = add,
sub = sub,
```

```
mul = mul,
div = div,
tostring = tostring,
}
```

15.3. Using environments

One of the disadvantages of the considered methods for creating modules is that it is very easy to clog up the global namespace, for example just forgetting `local` in the local resource description.

Environments provide an interesting approach to creating muzzle that solves this problem. If the module has its own environment living, then not only all functions will fall into this table, but also all global variables. Therefore, we can determine everything from-covered functions as global, and they will automatically fall into co-the corresponding table. All a module needs to do is assign this table to `_ENV`. After that, when we determine we divide the `add` function, it automatically becomes `M.add`:

```
local M = {}
_ENV = M
function add (c1, c2)
return new (c1.r + c2.r, c1.i + c2.i)
end
```

Moreover, we can call other functions from this module. without any prefix. In the previous code, `add` refers to `new` from his environment, that is, he actually refers to `M.new`. This method is a good way to create modules that require buoy of very little work from a programmer. Prefixes in general Not needed. There is no difference between calling the exported and closed functions. If the programmer forgets to insert `local`, then it doesn't litter the global namespace; instead closed the function just becomes exportable.

However, I usually prefer one of the two previously considered new methods. Although they may require a little more work, the at least the code is clearer. In order not to create a global value by mistake, I just assign `_ENV` a value ***nil***. After that, any attempt to create a global value simply raises an error.

What is missing is access to other modules. After

after we changed the value of `_ENV` , we lost access to all the previous global variables. There are several ways to there is this access, each with its own pros and cons.

One option is to use inheritance:

```
local M = {}  
setmetatable (M, {__index = _G})  
_ENV = M
```

(You need to call `setmetatable` before assigning `_ENV` , so is it clear why?) When using this approach, the module receives direct access to any global variable, with a very small the price of such access. An interesting consequence of this decision is It means that your module now contains all the global variables.

For example, someone using your module can now call a standard function for computing sine with `complex`.

`math.sin (x)` . (A similar feature is also present in the Perl language.)

Another quick way to access other modules is to enter dividing a local variable containing the global environment:

```
local M = {}  
local _G = _G  
_ENV = M - or _ENV = nil
```

You should now start each global name with `_G`. but before-stupas happen a little faster as there is no use of me-there methods.

A more rigorous approach is to define as local variables of only those functions or modules that you need:

- module setting
local M = {}
- import section:

- take outside everything this module needs
local sqrt = math.sqrt
local io = io
- from this place access to the outside is impossible
_ENV = nil - or _ENV = M

This approach requires more work, but it clearly documents the dependencies of your module. It also leads to code that you- is filled a little faster than in the previously considered cases, due to for using local variables.

15.4. Submodules and packages

Lua allows for hierarchical module names using touch the point to separate the levels. For example, a module named `mod.sub` is a *submodule* of `mod`. A *package* is a complete tree of modules; it is the Lua distribution unit.

When you need a module named `mod.sub`, `require` first looks in the `package.loaded` table and then in the `package.preload` table, using the fully qualified name “`mod.sub`” as the key; in this case, exactly `ka` is the same symbol as any other.

However, when looking for a file defining this submodule, `require` re-converts point to another character, usually the system separator in the path (i.e. `'/'` for UNIX and `'\'` for Windows). After this transformation `require` looks for the resulting name just like any other name. For instance, path `'/'` is a path separator and we have the following path:

`./?.lua;/usr/local/lua/?.lua;/usr/local/lua/?.init.lua`

The `require` (“`ab`”) call will try to open the following files:

`./a/b.lua`

`/usr/local/lua/a/b.lua`

`/usr/local/lua/a/b/init.lua`

This behavior allows all modules in the package to be in useful catalog. For example, if a package contains `p` modules, `pa` and `pb`, the corresponding files can be `p / init.lua`, `p / a.lua` and `p / b.lua`, where the `p` directory is contained in the corresponding location those.

The path separator used by Lua is set at compile time. tion and can be any string (remember that Lua knows nothing about directories). For example, systems without hierarchical catalogs can

gotta use `'_'` as such a delimiter, so `require` (`"Ab"`) will look for `a_b.lua` file .

Names in C cannot contain periods, so the C library for a submodule `ab` cannot export the `luaopen_a.b` function .

In this case, `require` translates point into another character - underscore-
vanie. Thus, a C library named `ab` should call its initializing function `luaopen_a_b` . We can also use here a technique with a minus, but with a more complex result-Tom. For example, suppose we have a C library named `a` and we want to make it a submodule `mod` , then we can rename the corresponding file in `mod / va` . When calling `require "mod.va"` the `require` call will correctly find the new `mod / va` file , just like function `luaopen_a` inside it.

Also `require` has one extra crawler to load submodules in C. When it cannot find either the Lua file or the C file for submodule, this seeker again searches along the path for C, but this time seeks package name. For example, if the program wants to load a submodule `abc` , then this searcher will simply search for `a` . If he finds a library I'm fluent in C for this name, then `require` will look in that library the corresponding function, in our case `luaopen_a_b_c` . This charge the ability to place several submodules together in one a C library, each with its own initializing function.

From a Lua perspective, submodules in the same package do not have an explicit communication. Loading module `a` does not load any of its submodules; also loading `ab` doesn't load automatically `a` . End-but, when implementing the package, the developer has the right to set these links when desire. For example, module `a` may explicitly require loading as someone specific (or all) of their submodule.

Exercises

Exercise 15.1. Rewrite the code in Listing 13.1 as separate module.

Exercise 15.2. What happens when looking for a library if the path contains a fixed component (that is, a component not containing a question mark)? Could this behavior be useful?

Exercise 15.3. Write a finder that simultaneously looks for files in Lua and libraries in C. For example, the path for this searcher could be something like:

```
./?.lua;./?.so;/usr/lib/lua5.2/?.so;/usr/share/lua5.2/?.lua
```

(*Hint* : use `package.searchpath` to search the corresponding file, then try to download it, first with `loadfile` , then with `package.loadlib` .)

Exercise 15.4. What happens if you set a metatable for `package.preload` using the `__index` metamethod ? How can this be useful?

Chapter 16

Object oriented programming

A table in Lua is an object in more than one sense. Like objects, the table has a state. Like objects, a table has there is an identity (self) that does not depend on its values; in partness, two tables with the same values are different objects, an object can have different meanings at different times time. Like objects, tables have a life cycle that does not depend on who created them or where they were created. Objects have their own methods. Tables can also have their own todes as below:

```
Account = {balance = 0}
function Account.withdraw (v)
Account.balance = Account.balance - v
end
```

This definition creates a new function and stores it in the field `withdraw` of the `Account` object . Then we can call her as shown below:

```
Account.withdraw (100.00)
```

A function of this type is almost what we call a *method* . However, using the global `Account` name inside a function is bad practice. First, this feature will work only for this particular object. Secondly, even for this

object exactly as long as this object is recorded in this particular a global variable. If we change the name of the object, then `withdraw` will no longer work:

```
a, Account = Account, nil
a.withdraw (100.00) - ERROR!
```

This behavior violates the principle that every object must have its own, independent cycle of life.

A more flexible option is to use the *recipient of the opera-walkie-talkies*. To do this, our method will need an additional argument `cop` with the value of the recipient. This parameter is usually named *self* or *this* :

```
function Account.withdraw (self, v)
self.balance = self.balance - v
end
```

Now, when we call the method, we must indicate with which object it should work:

```
a1 = Account; Account = nil
...
a1.withdraw (a1, 100.00) - OK
```

When using the *self* parameter, we can use one and the same method for many objects:

```
a2 = {balance = 0, withdraw = Account.withdraw}
...
a2.withdraw (a2, 260.00)
```

This use of the *self* parameter is key in any design. an object-oriented language. Most object-oriented in languages, this mechanism is partially hidden from the programmer, therefore, this parameter does not need to be explicitly declared (although inside the method

you can still use - *self* or *this*). Lua can also hide this parameter with the *colon e operator* . We can rewrite the previous method definition as follows:

```
function Account: withdraw (v)
self.balance = self.balance - v
end
```

Then the method call will look like this:

```
a: withdraw (100.00)
```

The colon adds an extra hidden parameter in the definition method and adds an extra argument to the method call.

The colon is just syntactic sugar, albeit pre- quite comfortable; there is nothing fundamentally new here. We can

define a method when using dot syntax and call it using colon syntax, and vice versa, as long as we handle the extra parameter correctly:

```
Account = {balance = 0,  
  withdraw = function (self, v)  
    self.balance = self.balance - v  
  end  
}  
function Account: deposit (v)  
  self.balance = self.balance + v  
end  
Account:deposit (Account, 200.00)  
Account: withdraw (100.00)
```

At this point, our objects have an identity, consisting and operations on this state. They lack a class system inheritance and the ability to hide your variables (state). Let's deal with the first task first: how can we create different objects with the same behavior? Like how we can we create multiple accounts?

16.1. Classes

The class acts as a template for creating objects. Most object oriented languages offer the concept of a class. In such languages, each object is an instance of some specific class. Lua has no concept of a class; each object is defined shares his behavior and his data. However, it is not difficult at all - emulate classes in Lua, following the path of prototype languages like Self or NewtonScript. Objects have no classes in these languages. Instead each object can have a prototype that is the object in which the first object is looking for operations that it is not knows. To represent classes in such languages, we simply create an object that will only be used as a prototype for other objects (its instances). Both classes and prototypes will fall as a place to accommodate behavior common to various objects.

In Lua, we can implement prototypes using the idea of inheritance from section 13.4. More precisely, if we have two objects *a* and *b*, then

everything

what we need to do to make `b` act as a prototype for `a` is following:

```
setmetatable (a, {__index = b})
```

After that, `a` will search in `b` for all operations that it does not know.

Setting `b` as a class for `a` is actually practically same.

Let's go back to our bank account example. For co-building other accounts with behavior similar to `Account`, we will make so that these new objects will inherit their operations from `Account` when using the `__index` metamethod. As a small optimization, we we can not create separate metatables for each of the objects; instead, we'll use the `Account` table itself :

```
function Account: new (o)
o = o or {} — create a table if the user has not submitted it
setmetatable (o, self)
self.__index = self
return o
end
```

(When we call `Account: new`, `self` is equal to `Account`; so we could explicitly use `Account` instead of `self`. Odd-But using `self` is very useful for us in the following section when we introduce inheritance.) What happens when we create a new account and call its method as shown below?

```
a = Account: new {balance = 0}
a: deposit (100.00)
```

When we create a new account, `a` will have an `Account` (parameter `self` when invoking `Account: new`) as a metatable. Then, when we call `a: deposit (100.00)` we are actually calling `a.deposit (a, 100.00)`; the colon is just syntactic sugar.

However, Lua cannot find the `deposit` record in table `a`; so Lua looks for an `__index` entry in the metatable. The situation looks like in the following way:

```
getmetatable (a).__index.deposit (a, 100.00)
```

Metatable `a` is `Account` and `Account.__index` is so-the same `Account` (since the `new` method did `self.__index = self`).

Therefore, the previous expression is reduced to

```
Account.deposit (a, 100.00)
```

That is, Lua calls the original `deposit` function, but passing `a` as a parameter to `self`. Thus, the new account `a` inherited `deposit` function from `Account`. In the same way, he inherits everything

fields from `Account` .

Inheritance works not only for methods, but also for other some fields that are not in the new account. Therefore, the class can set not only methods, but also default values for fields instance. Recall that in our first definition of `Account` we provided the `balance` field with a value of 0. So if we match create an account without the initial balance value, then it will inherit this default value:

```
b = Account: new ()  
print (b.balance) -> 0
```

When we call `b`'s `deposit` method , this call will be equivalent to taped to the following code (since *self* is `b`):

```
b.balance = b.balance + v
```

Expression `b.balance` gives 0, and the method assigns an initial `b.balance` contribution . Subsequent calls to `b.balance` no longer apply. lead to a call to the corresponding metaclass, since `b` now has your `balance` field .

16.2. Inheritance

Since classes are objects, they can also receive methods from other classes. This behavior makes it easy to implement inherit (in the usual object-oriented sense le).

Let's say we have a base class `Account` :

```
Account = {balance = 0}  
function Account: new (o)  
  o = o or {}  
  setmetatable (o, self)  
  self.__index = self  
  return o  
end  
function Account: deposit (v)  
  self.balance = self.balance + v  
end  
function Account: withdraw (v)  
  if v > self.balance then error "insufficient funds" end  
  self.balance = self.balance - v  
end
```

From this class we can inherit the `SpecialAccount` class , allowing the buyer to withdraw more than is on his balance sheet. we starting with an empty class that inherits all operations from its

base class:

```
SpecialAccount = Account: new ()
```

Up to this point, SpecialAccount is just an instance

Account . However, interesting things happen next:

```
s = SpecialAccount: new {limit = 1000.00}
```

SpecialAccount inherits new from Account , like all other methods.

Today. However, this time when new is executed, its self parameter is already will refer to SpecialAccount . Therefore, the metatable s will be

SpecialAccount , whose __index is equal to SpecialAccount .

Therefore, s inherits from SpecialAccount , which in turn inherits from Account . Now if we do

```
s: deposit (100.00) ,
```

then Lua cannot find the deposit field in s , so it will look for it in SpecialAccount , there he will not find it either and will continue to search in Account , where it will find the original implementation of this method.

What makes SpecialAccount special is that we can re-define any method inherited from its parent class sa. All we need is to simply write a new method:

```
function SpecialAccount: withdraw (v)
if v - self.balance >= self: getLimit () then
error "insufficient funds"
end
self.balance = self.balance - v
end
function SpecialAccount: getLimit ()
return self.limit or 0
end
```

Now, when we call s: withdraw (200.00) , Lua will not return- in Account , since it will find a new withdraw method before in the SpecialAccount class . Since s.limit is 1000.00 (we set this field when creating s), then the program will remove it, leaving resulting in s with negative balance.

The interesting thing about objects in Lua is that you don't you need to create a new class to define the new behavior. If from- it is only necessary to change the behavior for one object, then we can re- lize this change directly in this object. For instance, if account s represents a special customer whose limit is always 10% from the current balance, then we can change only one account:

```
function s: getLimit ()
```

```
return self.balance * 0.10
end
```

After that, calling `s:withdraw(200.0)` will execute the `withdraw` method from the `SpecialAccount` class, but when `withdraw` calls `s:getLimit`, then the previously entered definition of this function will be called.

16.3. Plural inheritance

Since objects are not basic primitives, Lua has several ways to use object-oriented software programming. The approach we just saw uses the `__index` metamethod is probably the best combination of simplicity, speed and flexibility. However, there are other implementations too, which may be more suitable for some specific cases. We will now see an alternative implementation, which Toraya allows multiple inheritance in Lua.

The key in this implementation is to use the function in as the `__index` metafield. Recall that when the metatable has data table has an `__index` field, then Lua will call this function any time when it cannot find the key in the source table. In this case `__index` can search for a missing key in any number of parents.

Multiple inheritance means that a class can have more than one superclass (parent class). Therefore, we already cannot use a function like before to create a subclass. Instead, we will define a `createClass` function, which takes parent classes as arguments (see Listing 16.1). This function creates a table for the presentation new class and sets its metatable with metamethod `__index`, which implements multiple inheritance. Not looking at multiple inheritance, each created object belongs to one class, which is used to find the place to store. Therefore, the relationship between class and superclasses is different from the relationship between classes and its instances (created objects). In particular, a class cannot simultaneously be the metatable for its instances and child classes. In Listing 6.1, we use the class as the metatable for the generated

instances and create a separate table as a metatable class.

Listing 16.1. Implementing multiple inheritance

```
- look for 'k' in the list of tables 'plist'
local function search (k, plist)
  for i = 1, #plist do
    local v = plist [i] [k] - try the i-th superclass
    if v then return v end
  end
end
function createClass (...)
  local c = {} -- new class
  local parents = {...}
  - the class will search for each method in the list of its parents
  setmetatable (c, {__index = function (t, k)
    return search (k, parents)
  end})
  - prepare 'c' as a metatable of its instances
  c.__index = c
  - define a new constructor for this new class
  function c: new (o)
    o = o or {}
    setmetatable (o, c)
    return o
  end
  return c - return a new class
end
```

Let's illustrate the use of `createClass` by cabbage soup small example. Let's say we have our old `Account` class and Class Named to methods `setname` and `getname` .

```
Named = {}
function Named: getname ()
  return self.name
end
function Named: setname (n)
  self.name = n
end
```

To create a new class `NamedAccount` , which is a child with both `Account` and `Named` classes , we just call `createClass` :

```
NamedAccount = createClass (Account, Named)
```

We create and use instances of this class as before:

```
account = NamedAccount: new {name = "Paul"}
print (account: getname ()) -> Paul
```

Now let's see how the last statement works. Lua cannot find `getname` method in `account` ; so he's looking for a field

__index in the account metatable , that is, in the NamedAccount . But in NamedAccount also doesn't have a “ getname ” field , so Lua looks for a field __index in the NamedAccount metatable. Since this field contains function, Lua calls it. This function first looks for “getname” in Account and, not finding it there, looks for Named , where she finds an excellent from *nil the* value that becomes the final result.

Of course, due to the complexity of such a search, the performance for many physical inheritance differs from performance for simple that inheritance. An easy way to improve this performance is to copy inherited methods into child classes.

Using this approach, the __index metamethod would look like in the following way:

```
setmetatable (c, {__index = function (t, k)
local v = search (k, parents)
t [k] = v - save for next call
return v
end})
```

Using this technique, access to inherited methods becomes as fast as accessing local methods (except for the first call). The disadvantage is that it is difficult to change method definitions when the system is working no, since these changes do not carry over along the chain of inheritance giving.

16.4. Hiding

Many consider the possibility of hiding to be an integral part of the object-but-oriented language; the state of each object is his personal matter. In some object oriented languages, such as C ++ and Java, you can control whether the field is an object one or his method is visible from the outside. In Smalltalk, all variables are hidden, and all methods are accessible from the outside. Simula, the first object-oriented language, does not provide such protection for lei and methods.

The Lua object design we looked at earlier did not pre-delivers hiding mechanisms. This is partly a consequence of

our use of tables to represent objects. Besides
Moreover, Lua avoids redundancy and artificial limitations. If a
you don't want to access the fields inside the object, just *don't*
this .

However, another goal of Lua is flexibility, it provides
provides meta-mechanisms that allow you to emulate many of the
nosti. Although basic object design for Lua does not provide
hiding mechanisms, we can implement objects in a different way
so that you gain access control. Although this possibility
programmers use infrequently, it will be useful to learn about
her, as it reveals some interesting aspects
Lua and can be a good solution for other tasks as well.
The main idea behind alternative design is to represent each
each object using two tables: one for its state and the other
gaya - for his operations (his interface). The object is being accessed
through the second table, that is, through the operations that form its in-
terface. In order to avoid unauthorized access,
the table providing its state is not stored in another field
goy table, it is only accessible through closures within methods.
For example, to represent a bank account using this
design, we will create new objects using the following
factory functions:

```
function newAccount (initialBalance)
local self = {balance = initialBalance}
local withdraw = function (v)
self.balance = self.balance - v
end
local deposit = function (v)
self.balance = self.balance + v
end
local getBalance = function () return self.balance end
return {
withdraw = withdraw,
deposit = deposit,
getBalance = getBalance
}
end
```

The function first creates a table to store the internal
state of the object and stores it in the local variable `self` .

The function then creates methods for the object. Finally, the function is gives and returns an external object that matches the names of the methods to their implementations. The key here is that these methods don't get `self` as an optional parameter; instead they access `self` . Since the additional there is no argument, then we do not use the colon syntax for work with the object. We call their methods just as usual. functions:

```
acc1 = newAccount (100.00)
acc1.withdraw (40.00)
print (acc1.getBalance ()) -> 60
```

This design provides complete stealth for everything that is stored is stored in the `self` . After returning from the `newAccount` function there is no way to directly access this table. Although our example only stores one variable in a private table, we can store all the private parts of the object in this table face. We can also define private methods: they are like public, but we don't put them in the interface. For example, our accounts can provide an additional 10% loan with the balance above a certain value, but we do not want the user to whether they had access to the details of the calculations. We can implement this functionality as follows:

```
function newAccount (initialBalance)
local self = {
  balance = initialBalance,
  LIM = 10000.00,
}
local extra = function ()
if self.balance > self.LIM then
return self.balance * 0.10
else
return 0
end
end
local getBalance = function ()
return self.balance + extra ()
end
<as before>
```

Again there is no way to call the function directly
extra .

16.5. Single approach method

A special case of the previous approach for object-oriented The case when the object has only one method. In this case, we do not need to create an interface tab- face; we can just return this method as a view object. If this looks a little strange, let's remember the time cases 7.1, where we created iterative functions that store their co- standing as closures. An iterator storing its state is nothing is no different from an object with a single function. Another interesting case of objects with a single method is the case when this method actually performs a different tasks depending on a specific argument. Possible the implementation of such an object is shown below:

```
function newObject (value)
return function (action, v)
if action == "get" then return value
elseif action == "set" then value = v
else error ("invalid action")
end
end
end
```

Its usage is pretty simple:

```
d = newObject (0)
print (d ("get")) -> 0
d ("set", 10)
print (d ("get")) -> 10
```

This object implementation is pretty efficient. Syntax `d ("set", 10)`, although it looks strange, is only two characters long it than the traditional `d: set (10)` . Each object uses one closure, which is cheaper than one table. There is no inheritance here, but but we have complete secrecy: the only way to contact to the state of an object is to use its only th method.

Tcl / Tk uses a similar approach for its widgets. Name kind- a get in Tk denotes a function (*widget command*) that can perform various types of operations on the widget.

Exercises

Exercise 16.1. Implement the `Stack` class with `push` , `pop` , `top` and `isempty` .

Exercise 16.2. Implement class `StackQueue` as a subclass `Stack` . In addition to the inherited methods, add to this class `insertbottom` , which inserts an element at the end of the stack. (This method allows you to use objects of the given class as a queue.)

Exercise 16.3. Another way to ensure that the public is closed projects is to implement them using a proxy (see section 13.4). Each object is represented by an empty table `tsey` (proxy). The internal table sets the corresponding view between these empty tables and tables carrying the state of the object. This internal table is not available sleep guns, but the methods use it to translate their parameter self to the real table they are working with. Implement the example with the `Account` class using this approach and see its pros and cons. (There is one small problem with this approach. Try- you can find it yourself or refer to section 17.3 for suggestions her solution is proposed.)

Chapter 17

Weak tables and finalizers

Lua handles memory management. Programs create objects (tables, threads, etc.), but there is no function to destroy objects. Lua automatically destroys objects that become garbage, using *garbage collection* . This frees you from the main work with memory and, more importantly, frees most errors related to this activity, such as dangling links and memory leaks.

Using a garbage collector means Lua has no problem with cycles. You don't need any special action when using naming circular data structures; they automatically release

are given like any other data. However, sometimes even the clever the garbage collector needs your help. No garbage collector will allow you to forget about all the problems of resource management, such as external resources.

Weak tables and finalizers are mechanisms that you can be used in Lua to help the garbage collector.

Weak tables allow collection of Lua objects that are still pending are accessible to the program, while finalizers allow assembly external objects not under direct control

Lem the garbage collector. In this chapter, we will discuss both of these mechanisms.

17.1. Weak tables

The garbage collector can only collect what is guaranteed to be rubbish; he cannot guess for himself what is rubbish your opinion. A typical example is a stack implemented as an array, with a reference to the top of the stack. You know that the data is reap only from the beginning of the array to this index (top of the stack), but Lua doesn't know that. If you pop an item off the top of the stack, simply decreasing the vertex index, then the object remaining in the array is not garbage for Lua. Similarly, any object for which the referenced global variable is also not garbage for Lua even if you never use it. In both cases you (more precisely, your program) should write **nil** in the appropriate variables (or array elements) in order to avoid the appearance of indestructible objects.

However, simply removing links is not always enough. In some cases need additional interaction between your product gram and garbage collector. A typical example is the set all active objects of a certain type (for example, files) in your our program. The task seems simple: all you need is add every new object to this set. However, as soon as the object becomes part of the collection, it will never be destroyed! Even if no one refers to it, the set will still refer on him. Lua cannot know that this link should not prevent destroying that object, unless you tell Lua to do so.

Weak tables are the mechanism you use in

Lua to say that a link should not interfere
destruction of the object. A *weak reference* is such a reference to an object,
which is not counted by the garbage collector. If all links, indicate-
on an object are weak, then this object is free
and all these weak links are destroyed. Lua implements weak
links using weak tables: a *weak table* is such a
a person whose links are all weak. This means that if
the object is stored only inside a weak table, then the garbage collector
sooner or later will destroy this object.

Tables store keys and values, both of which can be object
tami of any type. Under normal conditions, the garbage collector is not
destroyed.

It also contains objects that are keys and references in an accessible
table. Both keys and values are *strong* references, i.e.
they prevent the destruction of those objects to which they point
call. In a weak table, both keys and values can be weak.

This means that there are three types of weak tables: tables with weak
strong keys, tables with weak values and completely weak
tables where both keys and values are weak. Regardless
type of table, when deleting a key or value, the entire record is deleted
is taken from the table.

The weakness of a table is specified by the `__mode` field of its metatable. Mean-
the nest of this field, when present, must be a string: if
this row is "k" , then the keys in this table are weak;
if this string is "v" , then the values in this table are weak
face; if this string is equal to "kv" , then both the keys and values in this table
face are weak. The next example, albeit artificial,
shows the behavior of weak tables:

```
a = {}  
b = {__mode = "k"}  
setmetatable(a, b) - now 'a' has weak keys  
key = {}  
- create the first key  
a[key] = 1  
key = {}  
- create the second key  
a[key] = 2  
collectgarbage() - force the garbage collector to remove garbage  
for k, v in pairs(a) do print(v) end  
-> 2
```

In this example, the second assignment `key = {}` destroys the reference to first key. Calling `collectgarbage` causes the garbage collector to delete pouring all the trash. Since there are no more references to the first key, this the key and the corresponding entry in the table are deleted. The second key is still stored in the variable `key` , so it is not deleted.

Please note that only objects can be removed from weak table. Values such as numbers and booleans, are not deleted. For example, if we insert a numeric key into table `a` (from our previous example), then the garbage collector never has it will delete. Of course, if the value corresponding to the numeric key is stored in a table with weak values, then all corresponding the entire record is removed from the table.

There is a certain subtlety with lines: although lines are deleted- by the garbage collector, in terms of implementation they differ from other objects. Other objects such as tables and threads are co- are given explicitly. For example, when Lua evaluates the expression `{}` , then it creates a new table. However, does Lua create a new line when you filling in "a" .. "b" ? What if the system already has the string "ab" ? Con- will Lua give a newline? Can the compiler generate this line before executing the program? It makes no difference: it all implementation details. From a programmer's point of view, strings are values, not objects. Therefore, just like a number or a lo- logical meaning, the row cannot be deleted from the weak table (unless the associated value is removed).

17.2. Functions with caching

A common programming technique is to obtain gain in time due to memory loss. You can speed up the function by caching its results so that when later you call the same function with the same arguments, the function can use the value stored in the cache. Imagine a server receiving requests as strings with

holding Lua code. Each time a request is received, the server executes `load` on the received line and then calls the received function. However, `load` is an expensive feature and some commands to the server can be repeated many times. Instead of constantly- call `load` every time the server receives a command like

`Closeconnection ()` , the server can remember the result of the `load` in auxiliary table.

The server checks before calling `load`, no whether there is already a value corresponding to the given string. If he can't find the corresponding value, then (and only then) the server calls `load` and stores the result in this table. We can re- lize this behavior with the following function:

```
local results = {}
function mem_loadstring (s)
local res = results [s]
if res == nil then
- no result?
res = assert (load (s)) - calculate new result
results [s] = res
- save the result
end
return res
end
```

The gains from this scheme can be very significant. but it can also cause large memory losses. Although some mandas are repeated over and over, many other teams meet just one time. Over time, the `results` table collects all the commands which the server has ever received, and the corresponding code; since time This can lead to memory exhaustion on the server. Weak tabs faces provide a simple solution to this problem. If the table `results` stores weak values, then each garbage collection cycle will remove all currently unused values (virtually all):

```
local results = {}
setmetatable (results, {__mode = "v"}) - values will be weak
function mem_loadstring (s)
<as before>
```

In fact, since indices are always strings, we we can make this table completely weak if we want to:

```
setmetatable (results, {__mode = "kv"})
```

The caching technique is also useful to ensure that the caliber of objects of a certain type. For example, let us pre-set colors as tables with `red`, `green` and `blue` fields. The simplest the color factory will create a new table every time we we turn to her:

```
function createRGB (r, g, b)
return {red = r, green = g, blue = b}
end
```

Using caching we can reuse tables

for the same colors. To create a unique key for each color, we simply connect the color components using some separator:

```
local results = {}
setmetatable (results, {__mode = "v"}) -- values will be weak
function createRGB (r, g, b)
  local key = r .. "-" .. g .. "-" .. b
  local color = results [key]
  if color == nil then
    color = {red = r, green = g, blue = b}
    results [key] = color
  end
  return color
end
```

An interesting consequence of this implementation is that the the vendor can compare colors for equality using the standard comparison operator, since two simultaneously existing the same table will always correspond to the same colors. Please note that the same color may be displayed differently. tables at different points in time, since from time to time The garbage collector will empty the `results` table . However, while this color is used, it cannot be removed from `results` . therefore if a color exists long enough to be compared with another color, its representation will also exist for just as long.

17.3. Object attributes

Another interesting use of weak tables is linking ding attributes with objects. There are an infinite number of situations ation when we may need to bind some attribute to object: names to functions, defaults to tables, sizes to arrays, etc.

When the object is a table, then we can remember the attribute in the table itself by choosing a suitable unique key. As we already seen, a simple and reliable way to create a unique key is to create give a new object (usually a table) and use it as key. However, if the object is not a table, then this approach is already

not good. Even for tables, we may need to not store attributes in the table itself. For example, we might want to do a similar thing if the attribute is private or we don't want to influence how the table gets over. In all these cases, we need a different way of connecting the use of attributes with objects.

Of course, a separate table provides an ideal way of binding attributes to objects (it is no coincidence that tables are different yes they are called *associative arrays*). We can use objects as keys, and their attributes as values. Such a table can store attributes of objects of any type, since Lua allows use of objects of any type as table keys. Moreover, attributes stored in a separate table do not affect other objects and can be private, just like the table itself.

However, this solution has a huge disadvantage: as soon as we used the object as a key in the table, it can no longer be removed by the garbage collector. Lua cannot delete an object that is used as a key. If we use the usual way to bind their names to functions, then none of these functions will ever be removed. How can you guess live, we can avoid this deficiency with weak tables. However, this time we need weak keys. Use the creation of weak keys does not prevent the garbage collector from deleting these keys, when there are no more links left. On the other hand, the table cannot be weak values; otherwise the attributes of existing objects could be deleted.

17.4. Again tables with default values

In section 13.4 we looked at how you can work with values by default non-*nil* . We showed one approach and noticed that the other two approaches require the use of weak tables, therefore we have postponed the story about them for later. Now it's time to get back to this topic. As you will see, these two approaches to implementing values by default are in fact special cases already considered well-known approaches, namely object attributes and caching. In the first approach, we use weak tables to bind the table to its default values:

```

local defaults = {}
setmetatable (defaults, {__mode = "k"})
local mt = {__index = function (t) return defaults [t] end}
function setDefault (t, d)
  defaults [t] = d
  setmetatable (t, mt)
end

```

If defaults did not use weak keys, then all tables with default values would always exist.

In the second solution, we use different metatables to define default values, but at the same time we reuse one and the same metatable when we use the same value again default. This is a typical case for caching:

```

local metas = {}
setmetatable (metas, {__mode = "v"})
function setDefault (t, d)
  local mt = metas [d]
  if mt == nil then
    mt = {__index = function () return d end}
    metas [d] = mt - remember
  end
  setmetatable (t, mt)
end

```

In this case, we use weak values in order to avoid the metatables used could be assembled by a mul-
litter.

Which of the two is the best solution? As usual this depends on the use. Both solutions have approximately the same complexity and the same speed. First solution requires several words of memory for each table with a value by default (for writing to defaults). The second solution requires no how many tens of words of memory for each unique meaning default (new table, new closure plus metas entry).

Therefore, if your application has thousands of tables with only a few different default values, then the second solution is explicitly will be better. On the other hand, if several tables have a common with the default values, you'd better prefer the first implementation.

17.5. Ephemeral tables

An interesting case occurs when in a table with weak keys

the value refers to its own key.

This case is much more common than it might seem.

A typical example is a factory that returns functions.

A factory like this takes an object and returns a function that will return this object when called:

```
function factory (o)
return function () return o end
end
```

This factory is a good candidate for caching, for in order not to create new closures when there is already a suitable the next already created closure:

```
do
local mem = {}
setmetatable (mem, {__mode = "k"})
function factory (o)
local res = mem [o]
if not res then
res = function () return o end
mem [o] = res
end
return res
end
end
```

However, there is one catch. Note that the value is (corresponding function) associated with the object located in mem , - refers to its own key (the object itself). Although the keys are weak in this table, but values are not weak

Xia. With the standard interpretation of weak tables, nothing will be removed from the caching table. Since the values are not weak, then there is always a strong reference to each function. Each the function refers to its object, that is, there is always a strong reference for each object. Therefore, these objects cannot be deleted, not looking at the use of weak keys.

However, this interpretation is not always very helpful. Pain- Most people expect the value in the table to be available only through the appropriate key. Therefore, we can consider a similar scenario as a loop case where the closure refers to an object that (via the caching table) itself references this closure.

Lua 5.2 solves this problem with ephemeral tables.

In Lua 5.2, a table with weak keys and strong values is

is an *ephemeron* table. In an ephemeral table key availability controls the availability of the corresponding key *cheniya*. Let's take a closer look at the entry (*k*, *v*) in an ephemeral table. A reference to *v* is strong only if there is a strong reference to *k* . Otherwise, the record is eventually deleted from the table, even if *v* refers (directly or indirectly) to *k* .

17.6. Finalizers

Although the purpose of the garbage collector is to remove Lua objects, it can also help the program free external resources. For these purposes, various programming languages offer mechanisms *finalizers*. *The finalizer* is a function related to *ectom*, which is called before the object is removed collection-rubbish box.

Lua implements finalizers with the `__gc` metamethod . By-look at the following example:

```
o = {x = "hi"}
setmetatable(o, {__gc = function(o) print(ox) end})
o = nil
collectgarbage() -> hi
```

In this example, we first create a table and set for her a metatable that has a `__gc` metamethod . Then we destroy- we get the only link to this table (global variable `o`) and invoke garbage collection by calling `collectgarbage` . In garbage collection time Lua detects that the given table is not available and calls its finalizer (`__gc` metamethod).

The subtle thing in Lua is marking an object for finalization. We mark an object for finalization when we set it to a metatable with a non-zero `__gc` field . If we do not mark the object, then it will not be finalized. Most of the code we write is will work, but sometimes weird cases like the following blowing:

```
o = {x = "hi"}
mt = {}
setmetatable(o, mt)
mt.__gc = function(o) print(ox) end
o = nil
collectgarbage() -> (prints nothing)
```

In this example, the metatable we are setting for `o` is not

contains the `__gc` metamethod , so the object is not marked for finalization. Even if we later add the `__gc` field to the metatable, Lua does not consider this assignment as special, so object and will not be marked. As we said, this is rarely a problem; usually the metatable does not change after it has been assigned by metatable.

If you really want to set the metamethod later, then you can choose to use any value for the `__gc` field as temporary:

```
o = {x = "hi"}
mt = {__gc = true}
setmetatable(o, mt)
mt.__gc = function(o) print(o.x) end
o = nil
collectgarbage() -> hi
```

Now, as metatable field contains the `__gc` , object `o` postponed about for finalization. There is no problem in getting give a metamethod later; Lua only calls the finalizer if it is a function.

When the garbage collector destroys multiple objects in one and in the same loop, it calls their finalizers in the reverse order the one in which the objects were marked for finalization. Consider the following example that creates a linked list of objects with finalizers:

```
mt = {__gc = function(o) print(o[1]) end}
list = nil
for i = 1, 3 do
  list = setmetatable({i, link = list}, mt)
end
list = nil
collectgarbage()
-> 3
-> 2
-> 1
```

The first object to be finalized will be object 3, which was the last marked object.

A common misconception is that linking between the destroyed objects can affect the order in where they will be finalized. For example, you might think that object 2 in the previous example must be finalized before object 1, since there is a link from 2 to 1. However, the links can form cycles. Therefore, they do not impose any order for finalization.

Another subtle point related to finalizers is *recovery*. When the finalizer is called, it receives a finalizer object to be passed as a parameter. So the object becomes alive again, at least during finalization. I call this *temporary recovery*. At run time the finalizer does not prevent it from remembering the object, for example in a global variable so that the object remains accessible after the finalizer finishes. I call it *permanent restoration*.

Recovery must be transitive. Consider the following snippet of code:

```
A = {x = "this is A"}  
B = {f = A}  
setmetatable(B, {__gc = function(o) print(ofx) end})  
A, B = nil  
collectgarbage() -> this is A
```

The finalizer for B refers to A, so A cannot be removed.

Before finalizing B and . Lua must reconstruct both A and B before by calling the finalizer.

Due to restoration, objects with finalizers have restored are carried out in two passes. The garbage collector first finds out that an object with a finalizer is unreachable (no one refers to it), then it restores that object and adds it to the queue for finalization. After executing the finalizer, Lua marks the object as finalized. The next time the garbage collector is finds out that the object is unreachable, he will destroy it. If you want to ensure that all the garbage in your program is truly its own run, then you must call `collectgarbage` twice; second call will destroy objects that were finalized during the first first call.

The finalizer for each object is executed exactly once, because Lua marks objects that have already been finalized. If the volume the object has not been deleted before the end of the program, then Lua will call it in

the very end. This feature allows you to implement in Lua an analog `atexit` functions, that is, functions that are called directly immediately before exiting the program. Anything for this what you need is to create a table with a finalizer and remember the link to it somewhere, for example in a global variable:

```
_G.AA = {__gc = function ()
```



```
- your 'atexit' code comes here
print ("finishing Lua program")
end}
setmetatable (_G.AA, _G.AA)
```

Another interesting possibility is the ability to call call a specific function every time Lua exits the loop garbage collection. Since the finalizer is called exactly once, then you need to create a new object in the finalizer to call the next blowing finalizer:

```
do
local mt = {__gc = function (o)
- whatever you want to do
print ("new cycle")
- create a new object for the next cycle
setmetatable ({}, getmetatable (o))
end}
- create the first object
setmetatable ({}, mt)
end
collectgarbage () -> new loop
collectgarbage () -> new loop
collectgarbage () -> new loop
```

Interaction of objects with finalizers and weak tables with keeps a subtle moment. The garbage collector cleans up values in a weak table before restoring while keys are being flushed after recovery. The following code snippet illustrates this behavior:

```
- table with weak keys
wk = setmetatable ({}, {__mode = "k"})
- table with weak values
wv = setmetatable ({}, {__mode = "v"})
o = {} - object
wv [1] = o; wk [o] = 10 - add to both tables
setmetatable (o, {__gc = function (o)
print (wk [o], wv [1])
end})
o = nil; collectgarbage () -> 10 nil
```

During the execution of the finalizer, it finds the object in the table wk, but not in the wv table. The rationale for this behavior is that we often store object properties in tables with weak keys (as we discussed in Section 17.3) and finalizers can be need to refer to these attributes. However, we use tab-faces with weak values to reuse existing ones objects; in this case, the finalizable objects are no longer needed.

Exercises

Exercise 17.1. Write a code to check if does Lua really use ephemeral tables. (Not for-call `collectgarbage` for garbage collection.) check your code in both Lua 5.1 and Lua 5.2.

Exercise 17.2. Consider the first example from section 17.6, creates a table with a finalizer that prints the message on call. What happens if the program ends without calling garbage collection? What happens if the program is calling `os.exit` ? What happens if the program terminates execution with error?

Exercise 17.3. Suppose you need to implement caching a table for a function that takes a string and returns string. Using a weak table will prevent deletions records as weak tables do not consider rows as objects to be deleted. How can you implement caching in this case?

Exercise 17.4. Explain the output of the following program:

```
local count = 0
local mt = {__gc = function () count = count - 1 end}
local a = {}
for i = 1, 10000 do
  count = count + 1
  a[i] = setmetatable({}, mt)
end
collectgarbage ()
print (collectgarbage "count" * 1024, count)
a = nil
collectgarbage ()
print (collectgarbage "count" * 1024, count)
collectgarbage ()
print (collectgarbage "count" * 1024, count)
```

Part III
WITH TANDARD
LIBRARIES

Chapter 18

Mathematical library

In this and the following chapters on the standard library, my goal is not to give a complete specification of each function, but to show what functionality each library. I can omit some specific options, or behavior for clarity. The main goal is to ignite your curiosity, which can then be satisfied by reading- See the Lua documentation.

The `math` library contains a standard set of mathematical functions such as trigonometric (`sin` , `cos` , `tan` , `asin` , `acos` etc.), exponentiation and logarithm (`exp` , `log` , `log10`), rounding (`floor` , `ceil`), `min` , `max` , functions for generating pseudo-tea numbers (`random` , `randomseed`) and the variables `pi` and `huge` (last it is the largest representable number, on some platforms forms can take the special value *inf*).

All trigonometric functions work with radians. You can you can use `deg` and `rad` functions to convert between degrees and radians. If you want to work with degrees, you can re-define trigonometric functions:

```
do
local sin, asin, ... = math.sin, math.asin, ...
local deg, rad = math.deg, math.rad
math.sin = function (x) return sin (rad (x)) end
math.asin = function (x) return deg (asin (x)) end
...
end
```

The `math.random` function generates pseudo-random numbers. we we can call it in three different ways. When we call her with no arguments, it returns a real pseudo-random number in the range $[0, 1)$. When we call it with a single argument, integer n , then it returns a pseudo-random integer x lying between 1 and n . Finally, we can call it with two integer- with arguments l and u , then it will return a pseudo-random integer, lying between l and u .

You can set the seed for the pseudo-random generator

numbers using the `randomseed` function ; her only clean the primary argument is "seed". Usually when starting work programs the pseudo-random number generator is initialized some fixed value. This means that every time when you run your program it generates the same a sequence of pseudo-random numbers. For debugging, it turns out is very useful, but in the game you will always get one and also. The standard trick to combat this is to use setting the current time as a "seed" using a call `math.randomseed (os.time ())` . The `os.time` function returns a number, representing the current time, usually as a number of seconds, marching from a certain date.

The `math.random` function uses the `rand` function from the standard libraries of the C language. In some implementations, returns numbers with not very good statistical properties. You can reverse-to independent distributions in search of a better generator pseudo-random numbers. (Lua standard distribution does not include into itself a similar generator due to copyright issues. It contains only code written by the Lua authors.)

Exercises

Exercise 18.1. Write a function to check if is the given number a power of two.

Exercise 18.2. Write a function to calculate the volume of the cone s by its height and angle between its generatrix and axis.

Exercise 18.3. Implement another pseudo-random generator numbers for Lua. Search the internet for a good algorithm. (You a library for bitwise operations may be needed.)

Exercise 18.4. Using the `math.random` function , write function to get pseudo-random numbers from Gaussian distribution.

Exercise 18.5. Write a function for mixing of this list. Make sure all options are equally likely.

Chapter 19

Library for bitwise operations

A constant source of complaints about Lua is the lack of bitwise operations. This absence is by no means accidental. Not so easy to reconcile bitwise operations with floating point numbers.

We can express some of the bitwise operations as arithmetic operations. For example, shifts to the left correspond to multiplication by a power of two, shifts to the right correspond to division. One-bitwise AND and OR have no such arithmetic counterparts. They are defined for binary representations of integers. Practically it is not possible to extend them to floating point operations. Even some simple operations are meaningless. What should be additional by 0.0? Should it be -1? Or 0xFFFFFFFF (what's in Lua is 4,294,967,295, which is clearly not -1)? Or maybe $2^{64} - 1$ (a number that cannot be accurately represented using a value like double)?

To avoid such problems, Lua 5.2 introduces operations using a library, not as built into the language operations. This makes it clear that these operations are not "genus" ny "for numbers in Lua, but they use a certain interpretation. for working with these numbers. Moreover, other libraries may suggest other interpretations of bitwise operations (e.g. measures using more than 32 bits).

For most of the examples in this chapter, I will use six decimal notation. I will use the word MAX to indicate values 0xFFFFFFFF (that is, $2^{32} - 1$). In the examples I will use the following additional function:

```
function printx (x)
  print (string.format ("0x% X", x))
end
```

The bitwise library in Lua 5.2 is called `bit32`. As follows from name, it works with 32-bit numbers. Since **and**, **or** and **not** are reserved words in Lua, then the corresponding the functions are named `band`, `bor` and `bnot`. For the sequence in the name The function for bitwise exclusive OR is named `bxor`:

```
printx (bit32.band (0xDF, 0xFD)) -> 0xDD
printx (bit32.bor (0xD0, 0x0D)) -> 0xDD
printx (bit32.bxor (0xD0, 0xFF)) -> 0x2F
printx (bit32.bnot (0))
-> 0xFFFFFFFF
```

The `band`, `bor` and `bxor` functions accept any number of arguments.

Tov:

```
printx (bit32.bor (0xA, 0xA0, 0xA00)) -> 0xAAA
printx (bit32.band (0xFFA, 0xFAF, 0xAFF)) -> 0xAAA
printx (bit32.bxor (0, 0xAAA, 0))
-> 0xAAA
printx (bit32.bor ())
-> 0x0
printx (bit32.band ())
-> 0xFFFFFFFF
printx (bit32.bxor ())
-> 0x0
```

(They are all commutative and associative.)

The bitwise library works with unsigned integers.

During operation, any number passed as an argument is converted to an integer in the range 0-MAX . First, the unspecified numbers are ok-swear in an unspecified way. Second, numbers out of range 0-MAX are converted to it using the modulus operation: integer n becomes $n \% (2^{32})$. This operation is equivalent to getting the binary representation of the number and then taking its least significant 32 bits.

As expected, -1 becomes MAX . You can use the following-operations to normalize a number (that is, to display it in range 0-MAX):

```
printx (bit32.bor (2 ^ 32))
-> 0x0
printx (bit32.band (-1))
-> 0xFFFFFFFF
```

Of course, in standard Lua it's easier to just do $n \% (2^{32})$.

Unless explicitly specified, all functions in the library return the result, which also lies in 0-MAX . However, you should be tricky when using the results of bitwise operations in as ordinary numbers. Sometimes Lua is compiled using other second type for numbers. In particular, some systems with limited capabilities use 32-bit numbers as numbers in Lua.

In these systems, $\text{MAX} = -1$. Moreover, some bitwise libraries thecas use different conventions for their results. By-this whenever you need to use the result of the bitwise operations as a number, be careful. Avoid comparisons: instead of $x < 0$ write `bit32.btest (x, 0x80000000)`. (We will soon see dim the `btest` function.) Use the bitwise library itself to normalizing constants:

```
if bit32.or (a, b) == bit32.or (-1) then  
  <some code>
```

The bitwise library also defines operations for shifting and bit rotation: `lshift` to shift left; `rshift` and `arshift` to shift right; `lrotate` for left rotation and `rrotate` for rotation right. Except for arithmetic shift (`arshift`), all shifts fill new bits with zeros. Arithmetic shift for fills the bits on the left with copies of its last bit.

```
printx (bit32.rshift (0xDF, 4))  
-> 0xD  
printx (bit32.lshift (0xDF, 4))  
-> 0xDF0  
printx (bit32.rshift (-1, 28))  
-> 0xF  
printx (bit32.arshift (-1, 28))  
-> 0xFFFFFFFF  
printx (bit32.lrotate (0xABCDEF01, 4)) -> 0xBCDEF01A  
printx (bit32.rrotate (0xABCDEF01, 4)) -> 0x1ABCDEF0
```

Shift or rotation by a negative number of bits shifts (rotate em) in the opposite direction. For example, shift -1 bit to the right is equivalent to shifting 1 bit to the left. The result of a shift by more than 31 bits is 0 or MAX because all original bits are gone:

```
printx (bit32.lrotate (0xABCDEF01, -4)) -> 0x1ABCDEF0  
printx (bit32.lrotate (0xABCDEF01, -36)) -> 0x1ABCDEF0  
printx (bit32.lshift (0xABCDEF01, -36)) -> 0x0  
printx (bit32.rshift (-1, 34))  
-> 0x0  
printx (bit32.arshift (-1, 34))  
-> 0xFFFFFFFF
```

In addition to these more or less standard operations, bitwise the library also provides three additional functions.

The `btest` function performs the same operation as `band` , but returns Returns the result of comparing a bitwise operation with zero:

```
print (bit32.btest (12, 1))  
-> false  
print (bit32.btest (13, 1))  
-> true
```

Another common operation is to retrieve specified bits from the number. Usually this operation involves shifting and bitwise AND; a bitwise library packs it all into one function. The call `bit32.extract (x, f, w)` returns `w` bits from `x` , on-starting from bit `f` :

```
printx (bit32.extract (0xABCDEF01, 4, 8)) -> 0xF0
```



```
printx (bit32.extract (0xABCDEF01, 20, 12)) -> 0xABC
```

```
printx (bit32.extract (0xABCDEF01, 0, 12)) -> 0xF01
```

This operation counts bits from 0 to 31. If the third argument is `w` is not specified, then it is considered equal to one:

```
printx (bit32.extract (0x0000000F, 0)) -> 0x1
```

```
printx (bit32.extract (0xF0000000, 31)) -> 0x1
```

The reverse of the `extract` operation is the `replace` operation, which replaces the given bits. The first parameter is the outcome number. The second parameter specifies the value to be inserted.

The last two parameters, `f` and `w`, have the same meaning as in `bit32`.

`extract` :

```
printx (bit32.replace (0xABCDEF01, 0x55, 4, 8)) -> 0xABCDE551
```

```
printx (bit32.replace (0xABCDEF01, 0x0, 4, 8)) -> 0xABCDE001
```

Note that for any valid values of `x`, `f` and `w` the following equality holds:

```
assert (bit32.replace (x, bit32.extract (x, f, w), f, w) == x)
```

Exercises

Exercise 19.1. Write a function to check what is this number is a power of two.

Exercise 19.2. Write a function to calculate a number single bits in binary representation of a number.

Exercise 19.3. Write a function to check if Whether the binary representation of a number is a palindrome.

Exercise 19.4. Define shift operations and bitwise AND using Lua arithmetic operations.

Exercise 19.5. Write a function that receives a string, encoded in UTF-8, and returns its first character as number. The function should return `nil` if the line does not start with a valid UTF-8 sequence.

Chapter 20

Library

for working with tables

The `table` library contains additional functions that allow for working with tables as arrays. It provides functions for inserting and removing items from the list, for sorting ki array elements and to concatenate all strings in the array.

20.1. Insert and remove functions

The `table.insert` function inserts an element at a given location in the array, shifting the rest of the elements in order to make space. On-example, if `t` is an array `{10, 20, 30}` , then after calling `table.insert(t, 1, 15)` `t` will be `{15, 10, 20, 30}` . Special (and quite common), the case is to call `insert` without specifying a insertion, then the element is inserted at the very end of the array and shift no elements happen. As an example, the following code reads input line by line, remembering all lines in the array:

```
t = {}
for line in io.lines () do
  table.insert (t, line)
end
print (#t) -> (number of lines read)
```

In Lua 5.0, this technique is fairly common. In later versions, I prefer to use `t[#t + 1] = line` in order to add a string to the array.

The `table.remove` function removes (and returns) an element from the given place of the array, shifting the next elements of the array. If, when calling, the position inside the array was not specified, then delete

The last element of the array is taken.

With these two functions it is quite easy to implement stacks, queues and double queues. We can initialize similar structures like `t = {}`. The operation of adding an item is equivalent to `table.insert(t, x)`; the operation of removing an element of equivalent

valence `table.remove(t)`. The call to `table.insert(t, 1, x)` adds element to the other end of the corresponding structure, and a call to `table.remove(t, 1)` accordingly removes an element from that end. Two the latter operations are not particularly effective as they should move all elements of the array in memory. However, since the Bible In the `table` library, these functions are implemented in C, they are not are too expensive and work well for small arrays (up to several hundred elements).

20.2. Sorting

Another useful function for working with arrays is `table.sort`; we've seen it before. It takes as arguments

an array and optionally a function for comparison. This function is takes two arguments as input and must return **true** if the first the element must come before the second. If this function is not specified, then the sort function uses the standard '`<`' operator.

Typical confusion occurs when a programmer tries to sort the indexes in the table. In the table, the indices form many property in which there is no ordering. If you want them sort, then you need to copy them into an array and sort this array. Let's take an example. May you read the input file and built a table, which for each function name co-holds the line where this function was defined: something like next:

```
lines = {  
  luaH_set = 10,  
  luaH_get = 24,  
  luaH_present = 48,  
}
```

And now you need to print these functions in alphabetical order row. If you traverse this table with the `pairs` function, then the names will appear in no particular order. You can't explicitly

sort as these names are table keys.

However, if you put them in an array, then already this array can be sorted. So you first need to create an array with these names, then sort it and only then print result:

```
a = {}  
for n in pairs (lines) do a [#a + 1] = n end  
table.sort (a)  
for _, n in ipairs (a) do print (n) end
```

Some are confused. After all, Lua arrays don't contain any what ordering (arrays are actually tables). Poet- to that we impose ordering when working with indexes that can be ordered. This is why you are better off traversing the array when help `ipairs` , not `pairs` . The first of these functions sets the order of the keys is 1, 2, 3, ..., while the second just uses arbitrary order from the table.

As a more advanced solution, we can write an iteration

A generator for traversing a table using a given key order.

The optional `f` parameter specifies this order. This sleep iterator is chala sorts the keys into a separate array, and then bypasses this array. At each step, it returns a key and the corresponding value reading from the original array:

```
function pairsByKeys (t, f)  
  local a = {}  
  for n in pairs (t) do a [#a + 1] = n end  
  table.sort (a, f)  
  local i = 0  
  return function () - iterating function  
    i = i + 1  
    return a [i], t [a [i]]  
  end  
end
```

With this iterator it is easy to print function names in alphabetical order:

```
for name, line in pairsByKeys (lines) do  
  print (name, line)  
end
```

20.3. Concatenation

We have already seen the `table.concat` function in section 11.6 . She takes on input is a list of strings and returns the result of the concatenation of all these

lines. The optional second argument specifies the delimiter string. There are also two more optional arguments that specify the dexes of the first and last strings to be concatenated. The following function is an interesting generalization of `table.concat`. It can accept nested lists of strings as input:

```
function rconcat (l)
  if type (l) ~= "table" then return l end
  local res = {}
  for i = 1, #l do
    res [i] = rconcat (l [i])
  end
  return table.concat (res)
end
```

For each item in the list, `rconcat` recursively calls itself for processing nested lists. Then it calls `table`.

`concat` to combine intermediate results.

```
print (rconcat { {"a", {"nice"}}, "and", { {"long"}, {"list"} }})
-> a nice and long list
```

Exercises

Exercise 20.1. Rewrite the `rconcat` function so that for it could be given a separator string:

```
print (rconcat ({{{ "a", "b"}, {"c"}}, "d", {}, {"e"}}, ";"))
-> a; b; c; d; e
```

Exercise 20.2. The problem with `table.sort` is that this sort is not stable (stable sort), that is, elements that the sorting function considers to be equal can change their order during the sorting process. How can you implement robust sorting in Lua?

Exercise 20.3. Write a function to check if whether the specified table is a valid sequence.

Exercises

CHAPTER 21

Library for working with strings

Immediate possibilities of working with strings interpreted Lua's are pretty limited. The program can create strings, combine them and get the length of the string. But she cannot extract substrings or examine their contents. True power to work with strings comes from her library for working with strings. A library for working with strings is available as the `string` module . Since Lua 5.1, functions are also exported as string methods (using metatables). So, a line break in capital letters can be written as `string.upper (s)` OR `s: upper ()` . Choose yourself.

21.1. Main functions for working with strings

Some functions for working with strings in the library are extremely simple: calling `string.len (s)` returns the length of string `s` . She ek- is equivalent to `#s` . Calling `string.rep (s, n)` (or `s: rep (n)`) returns string `s` repeated `n` times. You can create a 1MB string (on-example, for tests) with `string.rep ("a", 2 ^ 20)` . Call `string.lower (s)` returns a copy of the uppercase string- mi replaced by lowercase; all other characters remain unchanged. (The `string.upper` function converts lowercase letters to uppercase.) As an example, if you want to sort strings outside of the from uppercase / lowercase letters, you can use the following next piece of code:

```
table.sort (a, function (a, b)
return a: lower () <b: lower ()
end)
```

The call to `string.sub (s, i, j)` returns the substring of `s` starting at `i` -th character and ending with `j` -th (inclusive). In Lua, the first character is row has index 1. You can also use negative indices that are counted from the end of the line: index -1 referen- goes to the last character of the string, -2 to the penultimate character, and so on. So calling `string.sub (s, 1, j)` (or `s: sub (1, j)`) will return specifies the beginning of a string of length `j` ; `string.sub (s, j, -1)` (or just `s: sub (j)`) because the default for the last argument is minta is -1) returns the end of the string starting from the `j` -th character; and `string.sub (s, 2, -2)` returns a copy of the string `s` , in which the we are the first and last characters:
`s = "[in brackets]"`

```
print (s: sub (2, -2)) -> in brackets
```

Remember that strings in Lua are immutable. `String.sub` function , like any other function in Lua, it doesn't change the value of the string, but returns a new string. A common mistake is to use something like `s: sub (2, -2)` and expect this to change the value of `s` . If you want to change the value of a variable, then you must assign a new value to it:

```
s = s: sub (2, -2)
```

The `string.char` and `string.byte` functions translate between characters and their internal numeric representations. `String` function

`char` takes integers as input, converts each of them to symbol and returns a string built from all these characters. Call

`string.byte (s, i)` returns internal numeric representation

`i` -th character of string `s` ; the second argument is optional, the call to `string.`

`byte (s)` returns the internal numeric representation of the first

string character `s` . In the following examples, we assume that the characters represented by ASCII encoding:

```
print (string.char (97))
```

```
-> a
```

```
i = 99; print (string.char (i, i + 1, i + 2)) -> cde
```

```
print (string.byte ("abc"))
```

```
-> 97
```

```
print (string.byte ("abc", 2))
```

```
-> 98
```

```
print (string.byte ("abc", -1))
```

```
-> 99
```

On the last line, we used a negative index for access to the last character of the line.

Since Lua 5.1, the `string.byte` function supports the third one, optional argument. Calling `string.byte (s, i, j)` returns numerical representations of all characters at once between indices `i` and `j` (inclusive):

```
print (string.byte ("abc", 1, 2)) -> 97 98
```

The default for `j` is `i` , so the call without

the third of its argument returns the `i`- th character. Calling `{s: byte (1, -1)}`

creates a table with codes of all characters in string `s` . According to this table we can get the original string by calling `string.`

`char (table.unpack (t))` . This trick does not work for very long

lines (more than 1 MB), since Lua has a limit on the number of values rotated by the function.

The `string.format` function is a powerful tool for

formatting strings, usually for output. It returns the a matted version of its arguments (supported by free number of arguments) using the description given by its first argument, the so-called *format string*. For this line there are rules similar to those for the `printf` function from the standard C library: it consists of plain text and *pointers* that control where and how to place each argument in the resulting string. The pointer consists of a '%' character, followed by a character specifying how to format the argument: 'd' for decimal numbers, 'x' for hexadecimal numbers, 'o' for octal, 'f' for floating point numbers, 's' for lines, there are also some other options. Between '%' and sym- there may be other options that specify formatting, such as the number of decimal digits for a floating point number:

```
print (string.format ("pi =%.4f", math.pi)) -> pi = 3.1416
d = 5; m = 11; y = 1990
print (string.format ("% 02d /% 02d /% 04d", d, m, y)) -> 05/11/1990
tag, title = "h1", "a title"
print (string.format ("<% s>% s </% s>", tag, title, tag))
-> <h1> a title </h1>
```

In the first example, `%.4f` is a floating point number with three digits after the decimal point. In the second example, `% 02d` denotes a decimal number of at least two digits, if necessary walkability padded with zeros; `% 2d` without zero will pad the number spaces. For a complete description of these options, refer to the See the Lua manual, or refer to the C manual. since Lua uses the C library to do all hard work here.

21.2. Functions for work with templates

The most powerful functions in the library for working with strings are `find`, `match` and `gsub` (global substitution) functions and `gmatch` (global search). They are all *template* based. Unlike a number of other scripting languages, Lua does not use

neither POSIX syntax nor language syntax for working with templates
Perl. The main reason for this decision is size: typical
ny implementation of POSIX regular expressions takes over 4000
lines of code. This is larger than all the standard Lua libraries, taking
together. For comparison, the implementation of working with templates in Lua
is
takes less than 600 lines. Of course, the implementation of working with
templates in
Lua is inferior to a full-fledged POSIX implementation. Still working with
templates in Lua is a powerful tool and includes
some features that are difficult to relate to standard
POSIX implementations.

String.find function

The `string.find` function searches for a given pattern within a string. Pros-
The strongest case of a pattern is a word that matches
your copy. For example, the pattern 'hello' will search for the substring
“Hello” within the entire given string. When finding template `find`
returns two values: the index from which the co-
drop, and the index where the match ends. If the match is not
found, then ***nil*** is returned :

```
s = "hello world"
i, j = string.find (s, "hello")
print (i, j)
-> 1 5
print (string.sub (s, i, j))
-> hello
print (string.find (s, "world")) -> 7 11
i, j = string.find (s, "l")
print (i, j)
-> 3 3
print (string.find (s, "lll"))
-> nil
```

When the search for the template is completed successfully, we can call
`string.sub` with returned values in order to get
part of the original string that matches the pattern. For simple
templates, such a string will be the template itself.

The `string.find` function has an optional third parameter: in-
dex specifying where within the string the search should start.
This parameter is useful when we want to get all
pattern occurrences: in this case we call the search function
repeatedly, each time starting the search after the position at which

a previous match was found. As an example, the following the following code builds a table with the positions of all '\n' characters inside lines:

```
local t = {}  
- table for storing indexes  
local i = 0  
while true do  
  i = string.find (s, "\n", i + 1) - looking for the next occurrence  
  if i == nil then break end  
  t [#t + 1] = i  
end
```

Later we will see an easier way to write such loops, using the iterator `string.gmatch` .

String.match function

The `string.match` function is similar to `string.find` in the sense that it also searches for occurrences of a pattern in a string. However, instead of return the position where the pattern was found, it returns the part lines matching the pattern:

```
print (string.match ("hello world", "hello")) -> hello
```

For simple templates like 'hello' this function is meaningless.

la. It shows its power when used with difficult templates as in the following example:

```
date = "Today is 17/7/1990"  
d = string.match (date, "% d + /% d + /% d +")  
print (d) -> 7/17/1990
```

We will discuss shortly both the meaning of the pattern '% d + /% d + /% d +' and more complex use of `string.match` .

String.gsub function

The `string.gsub` function has three required parameters:

ku, pattern and replacement string. She is used to replace all occurrences of the pattern in the original string by the given string:

```
s = string.gsub ("Lua is cute", "cute", "great")
print (s)
-> Lua is great
s = string.gsub ("all lii", "l", "x")
print (s)
-> axx xii
s = string.gsub ("Lua is great", "Sol", "Sun")
print (s)
-> Lua is great
```

The optional fourth parameter limits the number of replaceable replacements:

```
s = string.gsub ("all lii", "l", "x", 1)
print (s)
-> axl lii
s = string.gsub ("all lii", "l", "x", 2)
print (s)
-> axx lii
```

The `string.gsub` function also returns as the second value the number of replacements performed. For example, in a simple way, the number of spaces in a line is

```
count = select (2, string.gsub (str, " ", ""))
```

String.gmatch function

The `string.gmatch` function returns a function that iterates over all occurrences of the pattern in a string. For example, the following example collects all words in a given string `s` :

```
words = {}
for w in string.gmatch (s, "%a+") do
words [#words + 1] = w
end
```

As we will discuss shortly, the pattern `%a +` matches the occurrence one or more letters (i.e. words). Therefore, the cycle is denotes all the words within the string, storing them in the `words` table .

The following example implements a function similar to `package`.

```
searchpath using gmatch and gsub :
function search (modname, path)
modname = string.gsub (modname, "%.", "/")
for c in string.gmatch (path, "[^:] +") do
local fname = string.gsub (c, "?", modname)
local f = io.open (fname)
if f then
f: close ()
return fname
```

```
end
end
return nil - not found
end
```

The first step is to replace all dots with a delimiter in the path, which is considered equal to '\'. (As we will see below, the point has special meaning in templates. To match the point we should write '%.'). Then the function iterates over all the components paths where all questions are replaced for each component adjectives to the module name and it is checked if such what file. If so, the function closes this file and returns it name.

21.3. Templates

You can make templates more useful with *classes* *characters*. A character class is an element in a template that can match any character from the given set. For example measures, class %d matches any digit. Therefore, one can search for date in dd / mm / yyyy format using template %d% d /% d% d /% d% d% d% d :

```
s = "Deadline is 30/05/1999, firm"
date = "%d% d /% d% d /% d% d% d% d"
print (string.sub (s, string.find (s, date))) -> 30/05/1999
```

The following table contains a list of all character classes:

```
...
All symbols
% a
Letters
% c
Control characters
% d
Numbers
% g
Printed characters other than space
% l
Lower case
% p
Punctuation symbols
% s
Whitespace characters
% u
Lower case
```

% w

Letters and numbers

% x

Hexadecimal digits

If you use the appropriate class as the class name

capital letter, then it corresponds to the complement of the class (that is, all characters outside the class). For example, '% A' matches all non-letters:

```
print (string.gsub ("hello, up-down!", "% A", "."))
```

```
-> hello..up.down. 4
```

(4 are not part of the resulting string. This is the second value.

The value returned by `gsub` is the total number of replacements performed. I will further omit this number in the following examples that print the results tat calling `gsub` .)

Some symbols, called *magic symbols* , have

a special value within a template. Magic symbols are

go

() . % + - * ? [] ^ \$

The '%' character is used to insert these characters into the pattern.

So, '%.' corresponds to the point; '%%' matches the character itself

'%'. You can use '%' like this not only with ma-

symbols, but also with any non-alphanumeric symbols

oxen. When in doubt, use '%' instead .

For a Lua parser, templates are just plain strings. They

obey the same rules as the rest of the lines. Only

functions for working with templates treat them as templates,

and only these functions use the special character meaning

'%'. To put quotes inside a template, use the same

the very tricks used to put quotes inside

other lines.

You can also create your own classes by grouping different

personal classes and symbols within square brackets. For instance,

class '['% w_]' matches alphanumeric characters and sym-

ox underlining; class '['01]' matches binary digits;

class '['% [%]]' matches square brackets. In order to

count the number of vowels in the text, you can use the following-the following code:

```
nvow = select (2, string.gsub (text, "[AEIOUaeiou]", ""))
```

You can also include ranges of sym-

oxen, writing down the first and last characters, separated by signs

lump minus. I rarely use this, as most of the ranges used are already defined; for example, '[0-9]' is the same as '% d' and '[0-9a-fA-F]' is the same as '% x'. However, if you need octal digits, then you can use '[0-7]' instead of '[01234567]', you can also get the complement of any class by prefixing the '^' character : so the pattern '[^ 0-7]' matches any character that is not an octal digit, and '[^ \n]' matches any character from-personal from '\n'. However, remember that built-in classes can be before it is easier to use the capitalized variant: '% S' is easier than '[^% s]'.

Templates can be made more useful by using the modifiers for specifying the number of repetitions and optional parts. Lua templates offer four such modifiers:

+

1 or more reps

*

0 or more reps

-

0 or shorter reps

?

Optional (0 or 1 time)

The modifier '+' matches one or more sim-oxen of a given class. It will always return the longest occurrence template. For example, the pattern '% a +' denotes one or more letters, that is the word:

```
print (string.gsub ("one, and two; and three", "% a +", "word"))
```

```
-> word, word word; word word
```

The pattern '% d +' matches one or more digits (meaning an unsigned number):

```
print (string.match ("the number 1298 is even", "% d +")) -> 1298
```

The modifier '*' is similar to '+', but it also allows null the number of occurrences of characters from this class. Used frequently to indicate optional spaces between template parts.

For example, for a pattern matching a pair of parentheses (perhaps with spaces in between), you can use the following pattern:

'% (% s *%)' : pattern '% s *' matches zero or more

spaces between brackets. (The parentheses also have a special meaning patterns, so we specify them using the '%' character .)

As another example, the pattern '[_ % a] [_ % w] *' matches identifiers inside a Lua program: starts with a space

or underscore followed by zero or more underscores and alphanumeric characters.

Like `'*'`, the modifier `'-'` also matches zero or more characters of the given class. However, instead of matches the longest sequence it matches the shortest sequence. Sometimes there is nothing in between what a difference, but usually they give different results. For instance, if you try to find an id using a template

`'[_%a] [_%a] -'`, then you will receive only the first character of the identification torus because `'[_%a] -'` matches an empty sequence.

On the other hand, let's say you want to find comments in a program me in C. Most will try to use the pattern `'/%*.*%*/'`

(that is, `"/*"` followed by any sequence of characters fishing followed by `"*/"`). However, since `'.'` Will try match as many characters as possible, then the first `"/*"` will close only with the most recent `"*/"` in the program:

```
test = "int x; /* x */ int y; /* y */"
print (string.match (test, "/%*.*%*/"))
-> /* x */ int y; /* y */
```

The pattern `'.-'` captures the least number of characters required given for the first `"/*"`, and thus gives the desired result:

```
test = "int x; /* x */ int y; /* y */"
print (string.gsub (test, "/%*.*-%*/", ""))
-> int x; int y;
```

Last modifier `'?'` matches optional with- accompanying symbol. For example, let's say we want to find the number in text, which can contain an optional character. Template `'[+ -]?%d +'` successfully copes with the job, finding such numbers like `"-12"`, `"23"` and `"+1009"`. The class `'[+ -]'` matches either the `'+'` character or the `'-'` character ; the next character `'?'` de- barks this character optional.

Unlike other systems, in Lua the modifier can be changed to character class only; you cannot group templates under one modifier sign. For example, there is no template matching an optional word (unless it consists of one symbol). Usually this limitation can be bypassed by advanced techniques that we will see at the end of this chapter.

If the pattern starts with a `'^'` character then it will match- only with the beginning of the line. Likewise if the pattern ends with the `'$'` character , it will match only the end of the string. You

you can use both of these symbols to create templates. On-example, the following test checks if a string starts with a digit:

```
if string.find (s, "^% d") then ...
```

The following test verifies that the string is a number, with no other their characters at the beginning or end:

```
if string.find (s, "^ [+ -]?% d + $") then ...
```

The characters '^' and '\$' have this meaning only when they come across are located at the beginning or at the end of the line, respectively. Otherwise, they will

fall like ordinary symbols.

Another element in the template is '% b' . We write it down

like '% b xy ' where *x* and *y* are two different characters; character *x* *will* speak falls as the opening character and *y* as the closing character . For instance, the pattern '% b (' matches the part of the string that starts with '(' and ends with ')' :

```
s = "a (enclosed (in) parentheses) line"
```

```
print (string.gsub (s, "% b ()", "")) -> a line
```

We usually use this pattern as '% b (' , '% b []' , '% b {}'

or '% b <>' , but you can use any

bye symbols.

Finally, the '% f [*char-set*]' element is a *border pattern* . is he defines the place where the next character is contained in the class *char-set* and the previous one is not:

```
s = "the anthem is the theme"
```

```
print (s: gsub ("% f [% w] the% f [% W]", "one"))
```

```
-> one anthem is one theme
```

The pattern '% f [% w]' matches the border between non-alphabetic numeric and alphanumeric characters, and the pattern '% f [% W]' matches the boundary between an alphanumeric character and a non-alphanumeric character. Therefore, the given template corresponds matches the string "the" as a whole word. note that we must write many symbols inside square side even when it's just one class.

Positions before the first and after the last are interpreted as holding the character with code 0. In the previous example, the first "the" starts with a border between a null character (not in class '[% w]') and 't' (in class '[% w]').

The border pattern was implemented in Lua 5.1, but is not documented. It only became official in Lua 5.2.

21.4. Grips

The *capture* mechanism allows the template to remember parts of the string, satisfying elements of the template for later use. You can specify the capture by capturing the parts of the template you want to capture enough, inside parentheses.

When there are captures in the template, the `string.match` function will return spreads each captured value as a separate result; others in words, it splits the string into its captured pieces.

```
pair = "name = Anna"
key, value = string.match (pair,("(% a +)% s * =% s * (% a +)")
print (key, value) -> name Anna
```

The pattern `% a +` specifies a non-empty sequence of letters; template `% s *` specifies a possibly empty sequence of spaces. By- for this in the example above, the entire pattern defines a sequence of letters, followed by a sequence of spaces, followed by there is an equal sign `'='`, followed again by the sequence spaces followed by another sequence of letters. Both their sequences of letters, their corresponding patterns are are in parentheses, so they will be captured on match.

A similar example follows:

```
date = "Today is 17/7/1990"
d, m, y = string.match (date,("(% d +) / (% d +) / (% d +)")
print (d, m, y) -> 17 7 1990
```

Inside the template, element `% d`, where *d* is a digit, matches a copy *d*-th captured line. As an example, consider the case when yes, you want to find a substring enclosed in ordinary ones inside a string or double quotes. You can try the pattern `'["] .- ["]'`, that is, a quotation mark followed by anything followed by another quote; but you will have problems with lines like "It's all right". To solve this problem, we can capture the first quote and use that to specify the second quote:

```
s = [[then he said: "it's all right"!]]
q, quotedPart = string.match (s, "([\" ']) (.-)% 1 ")
print (quotedPart)
-> it's all right
```

```
print (q)
-> “
```

The first captured value is the quote character itself, and the second the captured value is a substring between quotes (substring, satisfying `'.'`).

For another similar example, we can take a template with corresponding to long lines in Lua:

```
% [(= *)% [(.)%]% 1%]
```

It matches an opening square bracket followed by followed by zero or more equal signs, followed by blows another open square bracket followed by that anything (line itself) followed by a closing square a parenthesis followed by the same number of equal signs, followed by another closing square bracket:

```
p = “% [(= *)% [(.)%]% 1%]”
```

```
s = “a = [= [[something]]] ==]”; print (a) ”
```

```
print (string.match (s, p)) -> = [[something]] ==]
```

The first capture is a sequence of equal signs (in the example re only one character); the second captured value is the string itself.

Also the captured values can be used in the override

line in `gsub`. Like the template, the replacement string can contain the elements `%d`, which are replaced with the corresponding captured values when performing substitution. In particular, the element

`%0` matches the entire portion of the string that matches the pattern.

(Note that the `%` character in the replacement string must be written sounded like `%%`.) Another example:

```
print (string.gsub (“hello Lua!”, “% a”, “% 0-% 0”))
```

```
-> h-he-el-ll-lo-o L-Lu-ua-a!
```

The following example rearranges adjacent characters:

```
print (string.gsub (“hello Lua”, “(.) (.)”, “% 2% 1”)) -> ehll ouLa
```

As a more useful example, let's write a simple pre-format builder that receives a string with commands as input in LaTeX style and translates them into XML format:

```
\ command {some text} -> <command> some text </command>
```

If we prohibit nested commands, then the next call

`string.gsub` does the job:

```
s = [[the \ quote {task} is to \ em {change} that.]]
```

```
s = string.gsub (s, “\\ (% a +) {(.)}”, “<% 1>% 2 </% 1>”)
```

```
print (s)
```

```
-> the <quote> task </quote> is to <em> change </em> that.
```

(We'll see how to handle nested commands later.)

Another useful example is removing spaces from the beginning.

la and end of line:

```
function trim (s)
return (string.gsub (s, "^%s * (.)%s * $", "% 1"))
end
```

Pay attention to the careful use of formats. Two anchors ('^' and '\$') ensure that we get the entire string. Since ku '.' tries to pick the shortest string, then two patterns '%s *' captures all white space around the edges. Also note that since gsub returns two values, we use a circle-left brackets to discard excess (number of substitutions).

21.5. Substitutions

Instead of a string as the third argument to string.gsub, we can use a function or table. When using the function string.gsub calls the function every time it finds a substring matching the pattern; arguments of each call are the captured values are captured, and the function's return value used as a replacement string. When the third argument is a table, the string.gsub function turns into a table face using the first captured value as a key and the resulting the value from the table as a replacement string. If received from function or table value is **nil** , then for a given match no replacement is made.

As a first example, consider the execution of a simple settings - each occurrence of \$ varName is replaced with the value of the global variable varName :

```
function expand (s)
return (string.gsub (s, "$ (%w +)", _G))
end
name = "Lua"; status = "great"
print (expand ("$ name is $ status, isn't it?"))
-> Lua is great, isn't it?
```

For each match with the pattern '\$ (%w +)' (dollar sign, for followed by a variable name) the gsub function looks for a matching variable in _G , the found value replaces the occurrence of the pattern into a string. When there is no corresponding variable in the table, then replacement is not made:

```
print (expand ("$ othername is $ status, isn't it?"))
-> $ othername is great, isn't it?
```

If you are not sure if the corresponding variables have string values, then you can try `toString` to these values. In this case, as a substitute value, you you can use the function:

```
function expand (s)
return (string.gsub (s, "$ (% w +)", function (n)
return toString (_G [n])
end))
end
print (expand ("print = $ print; a = $ a"))
-> print = function: 0x8050ce0; a = nil
```

Now, for each match against the pattern `'$ (% w +)'` `gsub` is called Gets the specified function, passing the name as an argument; function of rotates the replacement value.

In the last example, we go back to the format conversion teams. We again want to transform commands from LaTeX style (`\ example {text}`) to XML style (`<example> text </example>`), but on this time we will process nested commands. Next the function uses recursion to solve our problem:

```
function toxml (s)
s = string.gsub (s, "\\ (% a +) (% b { })", function (tag, body)
body = string.sub (body, 2, -2) - remove brackets
body = toxml (body) - processing nested commands
return string.format ("<% s>% s </% s>", tag, body, tag)
end)
return s
end
print (toxml ("\\ title {The \\ bold {big} example}"))
-> <title> The <bold> big </bold> example </title>
```

URL encoding

For our next example, we will use the *coding*

The URL that HTTP uses to pass parameters in the URL.

This encoding replaces special characters (such as '=' , '&' and '+') to '% xx ' , where xx is the character hex code. After

it then replaces spaces with '+' . For example, the string "a + b = c" But det is coded as "a% 2Bb +% 3D + c" . Also the name of each parameter and its value with an equal sign between them is added to the total On the next line, the variables are separated from each other by '&' . For example, the values

```
name = "al"; query = "a + b = c"; q = "yes or no"
```

will be encoded as "name = al & query = a% 2Bb +% 3D + c & q = yes + or + no" .

Now suppose we want to decode such a URL and write each get the value into the table by its name. The next function is to completes similar decoding:

```
function unescape (s)
s = string.gsub (s, "+", "")
s = string.gsub (s, "%% (% x% x)", function (h)
return string.char (tonumber (h, 16))
end)
return s
end
```

The first operator replaces every '+' with a space. The second finds hexadecimal encoded characters and for each such character calls an anonymous function. This function converts the hexadecimal representation to a number (`tonumber` base 16) and returns the corresponding character (`string.char`). For instance:

```
print (unescape ("a% 2Bb +% 3D + c")) -> a + b = c
```

To decode `name = value` pairs, we use the function

`gmatch` . Since both name and value cannot contain characters

'&' and '=' , then we can use the pattern '[^ & =] +':

```
cgi = {}
function decode (s)
for name, value in string.gmatch (s, "([^\&=] +) = ([^\&=] +)") do
name = unescape (name)
value = unescape (value)
cgi [name] = value
end
end
```

The `gmatch` function call finds pairs of the form `name = value` . For every Doing such a pair, the iterator returns the captured values (output bracketed in the template) as the values of the `name` and `value` fields . Body the loop just calls `unescape` on both of those lines and writes matching pair into `cgi` table .

It is also easy to write down the corresponding coding. To start we will write the `escape` function ; this function encodes all special digits like '%' followed by a hexadecimal code

character (for the `format` function , the "% 02X" option is used , fetching a two-digit string), and then replaces spaces

to the '+' symbol :

```
function escape (s)
```

```

s = string.gsub (s, "[%& = + %%% c]", function (c)
return string.format ("%%% 02X", string.byte (c))
end)
s = string.gsub (s, "'", "+")
return s
end

```

The `encode` function traverses the entire table that needs to be encoded. and builds the resulting string:

```

function encode (t)
local b = {}
for k, v in pairs (t) do
b [#b + 1] = (escape (k) .. "=" .. escape (v))
end
return table.concat (b, "&")
end
t = {name = "al", query = "a + b = c", q = "yes or no"}
print (encode (t)) -> q = yes + or + no & query = a% 2Bb +% 3D + c & name = al

```

Replacing tabs

The empty capture `()` in Lua has a special meaning. Instead of in order not to capture anything (which is completely unnecessary), this template lone captures the current position within the string as a number:

```

print (string.match ("hello", "() ll ()")) -> 3 5

```

(Note that the result of this example is different from calling `string.find`, since the position of the second captured value reading comes after the found pattern.)

A nice example of using this feature is replacing tab characters with the appropriate number of spaces:

```

function expandTabs (s, tab)
tab = tab or 8 - tab size (default 8)
local corr = 0
s = string.gsub (s, "() \t", function (p)
local sp = tab - (p - 1 + corr)% tab
corr = corr - 1 + sp
return string.rep (" ", sp)
end)
return s
end

```

The `gsub` call finds all tabs within a string, grabbing their position. For each tab character, the internal `function` uses this position to compute the number of spaces it takes to get the position,

multiple of the value of `tab` : it first subtracts one for position transfer starting at zero and then adding `corr` to account previously encountered tabs (replacing each tab character affects to the positions of the following characters). Then the correction is calculated for the next tab character: minus one for the tab to be deleted plus `sp` to account for the added spaces. Finally, she returns a string with the appropriate number of spaces.

For completeness, let's look at how you can reverse this operation.

walkie-talkie, replacing spaces with tabs. At first sight

you can also use empty grips to work from position-

mi inside the line, but there is a simpler solution: on each

the eighth character, we will insert the mark inside the line. Then,

when there are spaces before this mark, we will replace the corresponding

The corresponding sequence is a tab character:

```
function unexpandTabs (s, tab)
tab = tab or 8
s = expandTabs (s)
local pat = string.rep (".", tab)
s = string.gsub (s, pat, "% 0 \ 1")
s = string.gsub (s, "+ \ 1", "\ t")
s = string.gsub (s, "\ 1", "")
return s
end
```

This function starts by replacing all existing

tab characters with spaces. She then builds an auxiliary

template and uses it to add a markup (manager

character `\ 1`) after every `tab` characters. Further, all successive

number of spaces followed by a mark are replaced with

tabs. Finally, all overlays are removed.

21.6. Tricky tricks

Templates are a very powerful tool for working with strings.

You can perform many complex operations with just a few calls to `string.gsub` . However, like any other force, it must be use carefully.

Using templates does not replace the parser. For quick solutions

ny (quick-and-dirty) you can use templates to work

with the source code, but the resulting solutions most likely won't be of high quality. As an example, let's consider

Rome template that we used for search comment in C program: `"/%*.*%*/` . If you have a line in your program, so-holding `“/ *”`, then you may get the wrong result:

```
test = [[char s [] = “a / * here”; / * a tricky string * /]]
print (string.gsub (test, “/%*.*%*/”, “<COMMENT>”))
-> char s [] = “a <COMMENT>
```

Lines with such content are quite rare, and for your personal goals, a pattern like this will likely work.

But you cannot redistribute the program with this error.

Typically, templates work quite efficiently in Lua: my an old Pentium computer only needs 0.3 seconds to would find all the words in a 4.4 MB text (850K words). But always it is better to take some precautions. It's always better to de-make the template as accurate as possible; imprecise patterns are slower accurate. A simple example is using `'(.-%) % $'` for getting the entire substring up to the first occurrence of the dollar sign. If a there is a dollar sign in the line, then everything is fine; but let's assume that there is no dollar sign in the string at all. Then the algorithm first la will try to get a substring that matches the pattern to starting from the first position within the string. Then he will move all along the line looking for a dollar sign. When the line ends, then we will get pattern mismatches *only for the first position* inside the string. Then the algorithm will do the same, starting already from the second position inside the string, etc. Thus, we get the quadratic time complexity, taking more than 4 minutes to my Pentium for a 100K character string. You can easily fix fix this situation by tying the pattern to the beginning of the line with `'^ (.-%) % $'` . When using such a binding, the execution takes it takes only one hundredth of a second.

Also, be very careful with *empty templates* , i.e. template bosoms that are satisfied by the empty string. For example, if you if you try to search for names using the `'% a *'` pattern , then you are everywhere you will find names:

```
i, j = string.find (“; $% ** # $ hello13”, “% a *”)
print (i, j) -> 1 0
```

In this example, calling `string.find` correctly finds an empty string. a sequence of letters at the beginning of a line.

You should never write a pattern that begins or begins ends with `‘ ’` , since it will be satisfied with an empty string ka. This modifier usually needs something around it, for

in order to limit it. Likewise templates that include `'.*'` are also quite tricky, as this construct can much more than you planned.

Sometimes it's easier to use Lua itself to build templates.

We have already used this technique in the function that transforms the are white in tabs. As another example, let's distribute see how we can find strings of more than 70 characters. Such a string is a sequence of 70 or more characters fishing other than `'\n'`. Single character other than `'\n'` belongs to class `'[^ \n]'`. Accordingly, we can get pattern for a long string by repeating pattern for character 70 times and by adding a pattern for zero or more of the following sym-oxen. Instead of explicitly writing out this pattern, we can co-create it with `string.rep` :

```
pattern = string.rep ("^[^ \n]", 70) .. "[^ \n] *"
```

As another example, let's say you want to do a search, an odd case-sensitive. To do this, you can replace each letter woo x in the template for the class `'[xX]'`, that is, a class that includes and lowercase and uppercase versions of the letter. We can automate this transformation using the following function:

```
function nocase(s)
s = string.gsub(s, "%a", function(c)
return "[" .. string.lower(c) .. string.upper(c) .. "]"
end)
return s
end
print(nocase("Hi there!")) -> [hH][iI][tT][hH][eE][rR][eE]!
```

Sometimes you just need to replace every occurrence of `s1` with `s2`, without considering any magic symbols. If both lines are explicitly given in the text, then you can easily add all the necessary transformations yourself. development for magic symbols, but if these are variables, then you you will need additional `gsub`s to do this job:

```
s1 = string.gsub(s1, "(%W)", "%%% 1")
s2 = string.gsub(s2, "%%%", "%%%%%")
```

In the line we are looking for, we replace all non-alphanumeric-characters, in the replacement string we replace only the `'%'` character . Another useful technique for working with templates is to completing special processing of the line before the main work. Suppose we want to convert to uppercase all letters containing inside double quotes, but inside the string itself can be `'\"'` :

follows a typical string: "This is \" great \"!".

One approach for such cases is coding input string. For example, let's replace "\" " with "\" 1 " . but if the source text already contained the "\" 1" character , then we have a lemma. An easy way to do the coding and avoid this the problem is replacing all sequences "\" x " with "\" ddd " , where *ddd* is the decimal representation of the character *x* :

```
function code (s)
return (string.gsub (s, "\" (.)", function (x)
return string.format ("\"% 03d", string.byte (x))
end))
end
```

Now any sequence "\" *ddd* " could only come from our encoding, since any "\" *ddd* " in the original line is so would be encoded. Therefore decoding is simple task:

```
function decode (s)
return (string.gsub (s, "\" (% d% d% d)", function (d)
return "\" .. string.char (tonumber (d))
end))
end
```

We can now complete our task. Since the encoded the string no longer contains "\" " , then we can safely use pattern "'.-'" :

```
s = [[follows a typical string: "This is \" great \"!".]]
s = code (s)
s = string.gsub (s, "'.-'", string.upper)
s = decode (s)
print (s) -> follows a typical string: "THIS IS \" GREAT \"!".
```

Or, writing it shorter:

```
print (decode (string.gsub (code (s), "'.-'", string.upper)))
```

21.7. Unicode

At the moment, the library for working with strings does not contain explicit support for unicode. However, it is not difficult to implement some simple tasks for working with Unicode strings encoded in UTF-8 without using additional libraries.

The primary encoding for Unicode on the Web is UTF-8. Because of her compatibility with ASCII this encoding is also very well suited for Lua. This compatibility ensures that a number of operating techniques those with strings without any modification will work with UTF-8.

UTF-8 represents each unicode character with a different number byte. For example, the character 'A' represents one byte, 65; Sim-ox Aleph, which has Unicode code 1488, is represented by a two-byte sequence 215-144. UTF-8 represents all characters from ASCII as ASCII, that is, one byte with a value less than 128.

All other characters are represented by byte sequences, where the first byte lies in the range [194, 244] and the following bytes are you are in the range [128, 191]. More precisely, the range of the first byte for two-byte sequences are [194, 223], for three-byte sequences [224, 239] and for four-byte sequences telities [240, 244]. This arrangement ensures that the a character for any character will never be found inside the after-sequence for another symbol. For example, a byte less than 128 will never occur in a multibyte sequence; it always represented by its ASCII character.

In Lua, you can read, write and store strings in UTF-8 like regular strings. String constants (literals) can also contain UTF-8 inside. (Of course, you most likely want edit your file as a UTF-8 file.) Concatenation operation is executed correctly for all strings in UTF-8. Comparison operations strings (less than, less than or equal, etc.) compare strings in UTF-8, following unicode character order.

Operating system function library and library for I / O are really just interfaces to opera-system, so their UTF-8 support depends on UTF-8 on the system itself. On Linux, for example, we can use UTF-8 for filenames, but Windows uses UTF-16. therefore to work with file names in Unicode on Windows, you will need to additional libraries or modification of standard libraries tech Lua.

Let's see how functions from the library to work with strings work with strings in UTF-8.

Functions `string.reverse` , `string.byte` , `string.char` , `string.`

`upper` and `string.lower` do not work with UTF-8 strings because each of these functions considers one character to be one byte.

The `string.format` and `string.rep` functions work without any problems.

melt with strings in UTF-8, except for the `% c` option , which is assumes that one character is one byte. `String.len` functions and `string.sub` work correctly with strings in UTF-8, but at the same time

indexes no longer refer to characters, but to bytes. Often this is exactly what you need. But we can easily count the number of characters, as we'll see shortly.

For functions for working with templates, their applicability depends on the template. Simple templates work without any problems, since the representation of one character can never occur inside

When representing another character. Character classes and sets of characters only work for ASCII characters. For example, the template

“%s” works for UTF-8 strings, but it will only match

ASCII spaces and will not match unicode spaces, so

U+00A0 as unbreakable space (U + 00A0), paragraph separator

(U + 2029) or Mongolian G + 180E.

Some templates can make good use of features

of UTF-8. For example, if you want to count the number of characters in a line

of text, then you can use the following expression:

```
# (string.gsub (s, “[\ 128- 191]”, “”))
```

In this example, `gsub` strips out the second, third and fourth bytes,

leaving one byte for each character as a result.

Similarly, the following example shows how to iterate over

all characters in a string that are in UTF-8:

```
for c in string.gmatch (s, “[\ 128- 191] *”) do
```

```
  print (c)
```

```
end
```

Listing 21.1 shows some tricks for working with UTF-8

strings in Lua. Of course, to run these examples you need

a platform where `print` supports UTF-8.

Unfortunately, Lua has nothing more to offer. Adequate

Native Unicode support requires huge tables, which is bad

because of the small size of Lua. Unicode has many features.

It is almost impossible to abstract any concept

from specific languages. Even the concept of what a symbol is is very

fuzzy, since there is no one-to-one correspondence between

Unicode-encoded characters and graphemes (that is,

characters with diacritics and “completely ignored”

symbols). Other seemingly basic concepts such as what is

a symbol are also different for different languages.

What, in my opinion, is missing in Lua is the functions for translation

between UTF-8 and unicode and validating strings in UTF-8.

Perhaps they will be included in the next version of Lua. For the rest of the

the best option would be to use an external library

like `Unicode`.

Listing 21.1. Examples of working with UTF-8 in Lua

```
local a = {}
a[#a + 1] = "Nähdään"
a[#a + 1] = "açãõ"
a[#a + 1] = "ÃØÆË"
local l = table.concat(a, ";")
print(l, # (string.gsub(l, "[\ 128- \ 191]", "")))
-> Nähdään; açãõ; ÃØÆË 18
for w in string.gmatch(l, "[^;] +") do
  print(w)
end
-> Nähdään
-> açãõ
-> ÃØÆË
for c in string.gmatch(a[3], ". [\ 128- \ 191] *") do
  print(c)
end
-> Ã
-> Ø
-> Æ
-> Ë
-> Ð
```

Exercises

Exercise 21.1. Write a `split` function that gets string and delimiter pattern and returns a sequential number of blocks separated by separator:

```
t = split("a whole new world", "")
-t = {"a", "whole", "new", "world"}
```

How does your function handle empty lines? (In particular, is the empty string an empty sequence, or sequence with one blank line?)

Exercise 21.2. The patterns `'%D'` and `'[^%d]'` are equivalent. What about the patterns `'[^%d%u]'` and `'[%D%U]'`?

Exercise 21.3. Write a function for transliteration. This function gets a string and replaces every character in this string with another character in accordance with the table given second argument. If the table maps 'a' to 'b' then the function should replace every occurrence of 'a' with 'b'. If the table maps 'a' to *false*, then the function should delete all occurrences of the 'a' character from the string.

Exercise 21.4. Write a function that reverses a string in UTF-8.

Exercise 21.5. Write a transliteration function for UTF-8.

Chapter 22

Library input / output

The I / O library provides two different models for working with files. The simple model uses the *current input* and *the current output* files, and all its operations are performed on these files. The full model uses explicit file pointers; it relies on an object-oriented approach that defines Lets all operations as methods on file pointers.

A simple model is convenient for simple things; we used her throughout the book. But it is not enough for more flexible work with files, for example, for simultaneous reading or one-temporary recording in several files at once. For this we need complete model.

22.1. Simple model input / output

A simple model performs all its operations on the two current files lami. The library uses standard input when initializing (`stdin`) as default input file and standard output (`stdout`) as the default output file. Thus, when we do something like `io.read()` that we read from standard input. We can change these current files using the functions `io.input` and `io.output` . Calling `io.input(filename)` opens the specified file to read and sets it as input file default. From now on, all input will come from this-th file until the next call to `io.input` ; `io.output` works similarly Generally, but for output. On error, both functions call mistake. If you want to explicitly handle errors, then you need complete model.

The `write` function is simpler than `read` , so we'll first look at her. The `io.write` function receives an arbitrary number of string arguments and writes them to the default output file. She pre-converts numbers to strings using standard conversion rules

vania; for complete control over this conversion use

string.format function :

```
> io.write ("sin (3) =", math.sin (3), "\n")
-> sin (3) = 0.14112000805987
> io.write (string.format ("sin (3) =%.4f \n", math.sin (3)))
-> sin (3) = 0.1411
```

Avoid code like `io.write (a..b..c)` ; call `io.write (a, b, c)` does the same thing with fewer resources, since it avoids the concatenation operation.

Use `print` for small programs or for debugging and `write` when you need full control over the output:

```
> print ("hello", "Lua"); print ("Hi")
-> hello Lua
-> Hi
> io.write ("hello", "Lua"); io.write ("Hi", "\n")
-> helloLuaHi
```

Unlike `print` , the `write` function does not add any characters like tabs or move to next string. In addition, the `write` function allows you to redirect your output, whereas `print` always uses standard output. At the same time `net`, `print` automatically applies `tostring` to its arguments; this is useful for debugging, but it can hide errors if you are not aware are thoughtful to the conclusion.

The `io.read` function reads lines from the current input file. Her arguments control what to read:

```
"* A "
```

Reads the entire file

```
"* L "
```

Reads the next line (no line feed)

```
"* L "
```

Reads the next line (with a line feed)

```
"* N "
```

Reads a number

num

Reads a string of no more than *num* characters

The `io.read ("* a")` call reads the entire current input file, starting from the current position. If we are at the end of the file or the file is empty, then the call returns an empty string.

Since Lua works efficiently with long strings, an easy way to write filters in Lua is to read the whole file into a line, perform line processing (usually using `shchi gsub`) and then write the line to the output:

```
t = io.read ("* a")
```

```

- read the entire file
t = string.gsub (t, ...)
- do the job
io.write (t)
- write file

```

As an example, the following piece of code is a law-chennaya program for encoding file content in MIME *quoted-printable* . Each non-ASCII byte is encoded as = xx , where xx is it is the hexadecimal byte value. For the integrity of the coding The equality symbol itself must also be encoded:

```

t = io.read ("* a")
t = string.gsub (t, "([\\ 128- \\ 255 =])", function (c)
return string.format ("=% 02X", string.byte (c))
end)
io.write (t)

```

The pattern used in gsub finds all bytes from 128 to 255, including the equal sign.

The io.read ("* l") call reads the next line from the current input-leg file without line feed character ('\\ n'); io.read call ("* L") is similar, but only it returns a newline character (if he was present). When we reach the end of the file, the function rotates *nil* (since there are no more lines). Pattern "* l" for function read is the default. I usually use this pattern only when it naturally processes the file line by line coy; otherwise I prefer to read the entire file at once when help "*" a" or read it in blocks, as we'll see later.

As a simple example of using this pattern, follow-

This program copies the current input to the current output by numbering with each line:

```

for count = 1, math.huge do
local line = io.read ()
if line == nil then break end
io.write (string.format ("% 6d", count), line, "\\ n")
end

```

However, in order to iterate over the entire file, line by line, it is better to use the io.lines iterator . For example, we can write Put the complete program for sorting the lines of the file as follows - in a way:

```

local lines = {}
- read lines into table 'lines'
for line in io.lines () do lines [#lines + 1] = line end
- we sort

```



```
table.sort (lines)
- write all lines
for _, l in ipairs (lines) do io.write (l, "\n") end
```

The `io.read ("* n")` call reads a number from the current input file.

This is the only case where the `read` function returns a number, not a string. When a program needs to read a lot of numbers from a file, then the absence of intermediate lines improves performance. Operation `* n` skips all spaces before the number and supports such number formats like `-3` , `+5.2` , `1000` and `-3.4e-23` . If the function is not can find a number at the current position (due to incorrect format or end of file), it returns ***nil*** .

You can call `read` by passing multiple options at once; for every

For the second argument, the function will return the corresponding value. Let u you have a file containing three numbers per line:

```
6.0 -3.23 15e12
4.3 234 1000001
```

...

Now you need to print the maximum for each line. You can

You can read all three numbers in one `read` call :

```
while true do
local n1, n2, n3 = io.read ("* n", "* n", "* n")
if not n1 then break end
print (math.max (n1, n2, n3))
end
```

Besides the standard templates, you can call `read` by passing in number *n* as argument : in this case `read` tries to read *n* characters from the input file. If she can't read one character (end of file), it returns ***nil*** ; otherwise

a string with at most *n* characters is returned . As an example

The following program demonstrates an efficient way (for Lua, of course) copy the file from `stdin` to `stdout` :

```
while true do
local block = io.read (2 ^ 13) -- 8K buffer size
if not block then break end
io.write (block)
end
```

As a separate case, `read (0)` works as an end-of-file check:

it returns an empty string if there are characters in the file, and ***nil*** if end of file reached.

22.2. Full model input / output

For more control over I / O you can use complete model. The key concept in this model is *indicating Tel file* (file handle), which is similar to the FILE * to C: he presented Lets an open file at the current location.

To open a file, use the `io.open` function , which is analogous to the `fopen` function in C. As arguments, it takes a filename and a string specifying the mode. This line can contain 'r' for reading, 'w' for writing (writing erases the previous the contents of the file) or 'a' to append to the file, also it may contain 'b' to work with binaries. Function

`open` returns a new file pointer. On error `open` returns ***nil*** as well as the error message and error code:

```
print (io.open ("non-existent-file", "r"))
-> nil non-existent-file: No such file or directory 2
print (io.open ("/ etc / passwd", "w"))
-> nil / etc / passwd: Permission denied 13
```

The interpretation of error codes is system dependent.

A typical error checking method is as follows:

```
local f = assert (io.open (filename, mode))
```

If an error occurs, the error message appears as a second

The last argument to `assert` , which prints this message.

After you open the file, you can read and write from it into it using the `read` / `write` methods . They are similar to functions `read` / `write` , but you call them as file pointer methods, using using the colon. For example, in order to open a file and honor all of it, you can use the following code:

```
local f = assert (io.open (filename, "r"))
local t = f: read ("* a")
f: close ()
```

The I / O library provides three predefined pointers to standard files in C: `io.stdin` , `io.stdout`, and

`io.stderr` . Therefore, you can send the error message directly to the appropriate standard file:

```
io.stderr: write (message)
```

You can use the complete model along with the simple model.

To get a pointer to the current input file, blows call `io.input ()` with no arguments. In order to ask the file locator as the current input file, call

`io.input (hanle)` (similar calls work for `io.output` as well).

For example, if you want to temporarily change the current input file, then you can write something like the following:

```
local temp = io.input ()
```

```
- save the current file
```

```
io.input ("newinput")
```

```
- open new current file
```

```
<process input>
```

```
io.input (): close ()
```

```
- close the current file
```

```
io.input (temp)
```

```
- restore the previous file
```

Instead of `io.read` to read from a file, we can also use

`Vat io.lines` . As we saw in the previous examples, `io.lines` returns an iterator that reads sequentially from a file.

The first argument to `io.lines` can be a file name or a specified tel per file. If a filename was passed, then `io.lines` will open the file in read-only mode and will close the file after reaching the end file. If a file pointer was passed, then `io.lines` will be used. use this file for reading; in this case `io.lines` won't close the file when it reaches the end. In the case of a call at all with no arguments, `io.lines` will read data from the current input file.

Since Lua 5.2, the `io.lines` function also accepts the same same options as `io.read` (after the first argument). As an adjunct measure the following code copies the file using `io.lines` :

```
for block in io.lines (filename, 2 ^ 13) do
```

```
io.write (block)
```

```
end
```

A small trick to increase speed

It is usually faster in Lua to read an entire file than to read a line of it. by line. However, sometimes we come across a large file (at-example, tens or even hundreds of megabytes), read which in its entirety it would be inappropriate. If you want to get the maximum performance when working with such large files, it is faster will read it in large enough blocks (for example, by 8K). To avoid a possible line break, you can

just ask to read one more line:

```
local lines, rest = f: read (BUFSIZE, "* l")
```

The `rest` variable will get the remainder of any line broken when reading a block. Then we combine the block and the resulting remainder. This way the block will always end at line boundaries.

The example in Listing 22.1 uses this technique to implement `wc`, a program that counts the number of characters, words, and lines in a file.

Note the use of `io.lines` to implement

iterations and the `"* L"` option to read a line, this is available starting at Lua 5.2.

Listing 22.1. `Wc` program

```
local BUFSIZE = 2 ^ 13 - 8K
local f = io.input (arg [1])
- open input file
local cc, lc, wc = 0, 0, 0
- counters
for lines, rest in io.lines (arg [1], BUFSIZE, "* L") do
if rest then lines = lines .. rest end
cc = cc + #lines
- count the words in the block
local _, t = string.gsub (lines, "% S +", "")
wc = wc + t
- count '\ n'
_, t = string.gsub (lines, "\ n", "\ n")
lc = lc + t
end
print (lc, wc, cc)
```

24.2. Hooks

The trap mechanism allows us to register a function that will be called upon the occurrence of certain events during program execution. There are four types of events that can trigger traps:

- a *call* event occurs when Lua calls a function;
- the *return* event occurs when the function returns;
- a *line* event occurs when Lua starts execution next line;
- the *counter* event occurs after a specified number of mand.

Lua calls hooks with a single argument, a string, the event that led to the call: `"call"` (or `"tail`

call ”), “ return ” , “ line ” or “ count ” . For line event also the second argument is passed, the new line number. To get up to additional information inside the trap should be used

`debug.getinfo` .

To register a trap, we call the function

`debug.sethook` with two or three arguments: the first argument is this is the corresponding function; the second argument is a mask string, describing exactly what events we want to track, and not the required third argument is a number that specifies how often that we want to receive counter events. In order to track to add call, return and string events, we add the letters 'c' , 'r' and 't' to the mask string. To track counter events, we simply we pass the counter as the third argument. To remove all traps just call `sethook` with no parameters.

As an example, the following code sets up a primitive a trap, which for each next run of the print run-her number:

```
debug.sethook (print, "t")
```

This call sets `print` as a hook function and sets it call only for row events. More complex hook function can use `getinfo` to add a name to the output current file:

```
function trace (event, line)
local s = debug.getinfo (2) .short_src
print (s .. ":" .. line)
end
debug.sethook (trace, "t")
```

A useful function to use in traps is

`debug.debug` . This simple function prints a prompt, reads from input and then executes the given commands. It is roughly equivalent to tape on the following code:

```
function debug1 ()
while true do
io.write ("debug>")
local line = io.read ()
if line == "cont" then break end
assert (load (line)) ()
end
end
```

When the user enters “cont” at the prompt , this the function ends. The standard implementation is very simple and you

executes commands in the global environment outside of the code being debugged.

Exercise 24.5 discusses a better implementation.

24.3. Profiling

Despite its name, the debug library is also useful for not only debugging tasks. A typical such task is profiling (obtaining information about the time spent on the execution of this or that piece of code). For profiling time, it is better to use the C-interface: the cost of the call each Lua hook is quite high and can greatly distort results. However, for simple profiling that counts many times, the Lua code is fine. In this section, we write we have the simplest profiler, which for each called the function will tell you how many times it was called during the execution the program.

The main data structure in our program will be two tables: one maps functions to their counters, and the second maps functions their names. The indexes for both of these tables will be the functions themselves can act.

```
local Counters = {}
```

```
local Names = {}
```

We can extract the function names after profiling, but we get better results if we get the names of the functions, while they are active, since in this case Lua can look at looking up the name of the function, the calling code.

Now let's define a hook function. Its task is get the called function and increment the corresponding counter, it also collects function names:

```
local function hook ()
```

```
local f = debug.getinfo (2, "f"). func
```

```
local count = Counters [f]
```

```
if count == nil then -- is the function 'f' called the first time?
```

```
Counters [f] = 1
```

```
Names [f] = debug.getinfo (2, "Sn")
```

```
else
```

```
- only increase the counter value
```

```
Counters [f] = count + 1
```

```
end
```

```
end
```

The next step is to run the program with this hook.

We will assume that the main block of the program is in the file and the name of this file is passed as an argument to the profiling program—to the box:

```
% lua profiler main-prog
```

Then the profiler can take the filename from `arg [1]` , installed hook a trap and execute the file:

```
local f = assert (loadfile (arg [1]))
```

```
debug.sethook (hook, "c") — set a hook
```

```
f ()
```

```
- execute the profiled program
```

```
debug.sethook ()
```

```
- disable the trap
```

The final step is to actually display the results. Funk-

tion `getname` Listing 24.2 outputs for each function sootvetst-name. To avoid confusion, to each name

add the place of the corresponding function in the form *file: string* .

If the function has no name, then we only print the location. If the function

is a function in C, then we use only its name (so

as she has no place). With this in mind, below is the code that I print-containing information about calls:

```
for func, count in pairs (Counters) do
```

```
print (getname (func), count)
```

```
end
```

Listing 24.2. Getting the function name

```
function getname (func)
```

```
local n = Names [func]
```

```
if n.what == "C" then
```

```
return n.name
```

```
end
```

```
local lc = string.format ("% s:% d", n.short_src, n.linedefined)
```

```
if n.what ~ = "main" and n.namewhat ~ = "" then
```

```
return string.format ("% s (% s)", lc, n.name)
```

```
else
```

```
return lc
```

```
end
```

```
end
```

If we apply our profiler to the example with the chain Markova from section 10.3, then we get something like:

```
[markov.lua]: 4 884723
```

```
write 10000
```

```
[markov.lua]: 0 1
```

```
read 31103
sub 884722
[markov.lua]: 1 (allwords) 1
[markov.lua]: 20 (prefix) 894723
find 915824
[markov.lua]: 26 (insert) 884723
random 10000
sethook 1
insert 884723
```

This shows that the anonymous function on line 4 (which is our iterator defined inside `allwords`) was called 884,723 times, the `write (io.write)` function was called 10,000 times, etc.

This profiler can be improved, for example, add a dirty output customization, improved function name printing, etc. less even this profiler is already useful and can be used Called as a basis for writing more advanced tools.

Exercises

Exercise 24.1. Why recursion in `getvarvalue` function (Listing 24.1) will it stop?

Exercise 24.2. Modify the `getvarvalue` function (Listing 24.1) to work with various coroutines (like other functions from the `debug` library).

Exercise 24.3. Write a `setvarvalue` function .

Exercise 24.4. Based on the `getvarvalue` function write `getallvars` function that returns a table with all variables that are visible at a given location (return-May the table should not include environment variables it should instead inherit them from the original arms).

Exercise 24.5. Write an improved version of `debug.debug` , which executes the given commands as if they were executed in the scope of the calling function. (*Hint* : run commands in empty environment and use the `__index` function as a metamethod `getvarvalue .`)

Exercise 24.6. Modify the previous example to you could change variables.

Exercise 24.7. Implement some of the suggested improvements for the profiler from Section 24.3.

Exercise 24.8. Write a library to work with points breakpoint. She should offer at least two functions:

setbreakpoint (function, line) -> returns handle

removebreakpoint (handle)

The breakpoint is set by a function and a line within the function.

When execution reaches a breakpoint, you should exit

name debug.debug .

(*Hint* : for the simplest implementation, use the trap strings and a hook function that checks if we hit to a breakpoint; to improve performance, we can turn on this trap only when we are inside the function we are teasing.)

Part IV

With API

CHAPTER 25

C API overview

Lua is an *embedded language* . This means Lua is not separate. package, and a library that we can link to other instructions to add Lua features to them.

You may be wondering if Lua is not a standalone program, but so far in the book, we have used Lua as a standalone program.

Rammu. The solution to this question is the Lua interpreter (executing nimy file the lua). This interpreter is a small application

ny (less than five hundred lines of code) that uses the Lua library in order to implement a separate interpretation

Lua torus. This program deals with user interaction

lem, taking files and lines and passing them to the Lua library, which

does basic work (such as running Lua code).

This ability to use the library in order to expand application possibilities - this is exactly what makes Lua *distribution widening tongue* . At the same time, the program that is used calls Lua, can register new functions in the Lua environment, thus adding features that could not be written on Lua itself. This is what makes Lua an *extensible language* . These two views of Lua (as an extension language and as a extensible language) correspond to two types of interaction between C and Lua. In the first case, C controls, and Lua is just a library. We call the C code corresponding to this type of interaction *application code* (application code). In the second case, the control is in Lua, and C is a library. In this case, the C code is called *library code* . Both of these types of code use the same API for interacting with Lua, the so-called C API. The C API is simply a collection of functions that enable C code interact with Lua ¹ . It includes functions for reading

and writing global Lua variables to call functions in Lua, to execute snippets of Lua code, to register functions in C, so that they can then be called from Lua code, etc. Practice Technically anything that Lua code can do can be done in C. using the C API.

The C API follows the C style, which differs markedly from the C API. for the Lua language. When we program in C, we follow the types data, error handling, memory allocation errors and other difficult places. Most API functions do not validate the validity of their arguments; it is your job to make sure that arguments are correct before function call ² . If you admit error, you will most likely get an error like "segmentation fault" or something like that instead of a nice error message. More Moreover, the C API emphasizes flexibility and simplicity, often at the expense of easy

use bones. Typical tasks may require several these API calls. It can be tedious, but it gives you complete control over what is happening.

As the title suggests, the purpose of this chapter is to necessary when using Lua from C. Don't try to understand now all the details of what is happening. We will dwell on this later. but

do not forget that you can always find additional information See the Lua Reference Manual. Moreover, you can find some examples of using the API in the Lua distribution itself. Department the lua Lua interpreter (lua.c) gives examples of application code, while the standard libraries (lmathlib.c , lstrlib.c etc.) provide examples of library code.

From now on, we act as C programmers.

When I talk about "you", I mean exactly you, I program on S.

An important component in the communication between Lua and C is a permanently present virtual *stack* . Almost all functions APIs work with values on this stack. All data exchange between do Lua and C go through this stack. What's more, you can also use this stack to store intermediate results.

This stack allows you to solve problems with a fundamental difference - mi between Lua and C: the first difference is that Lua has garbage collection, while C is explicit memory management; second the difference is the difference between dynamic typing in Lua and static typing in C. We will discuss the stack in more detail. in section 25.2.

25.1. First example

We'll start this overview with an example of a simple application: a stand-alone Lua interpreter. We can write a primitive interpreter

The Lua tool, as shown in Listing 25.1. lua.h header file defines the basic functions provided by Lua. It includes in functions to create a new Lua environment, to call a function tions in Lua (such as lua_pcall) to read and write global variables in the Lua environment, to register new functions, which which can be called from Lua, etc. Anything defined in the file lua.h is prefixed with _lua .

Listing 25.1. Simple standalone Lua interpreter

```
#include <stdio.h>
#include <string.h>
#include "lua.h"
#include "lauxlib.h"
#include "lualib.h"
int main (void) {
char buff [256];
```

```

int error;
lua_State * L = luaL_newstate (); /* opens Lua */
luaL_openlibs (L); /* opens standard libraries */
while (fgets (buff, sizeof (buff), stdin) != NULL) {
    error = luaL_loadstring (L, buff) || lua_pcall (L, 0, 0, 0);
    if (error) {
        fprintf (stderr, "%s\n", lua_tostring (L, -1));
        lua_pop (L, 1); /* pop the error message off the stack */
    }
}
lua_close (L);
return 0;
}

```

The `luauxlib.h` header file defines the functions provided by the *additional libraries* (auxiliary library). All definitions from that file start with `luaL_` (for example, `luaL_loadstring`). The additional library uses the basic API provided by `lua.h` to provide a higher level of abstraction, in particular the abstractions used by the standard libraries. The core API strives for economy and orthogonality, while the additional library strives for practicality for common tasks. Of course this is easy for your program too can create the necessary abstractions. Keep in mind that before the additional library does not have access to Lua internals. Everything, that it does, it does through the standard API. What does she do, maybe do your program too.

The Lua library does not define any global re-

men. It stores all its state in a dynamic structure. `lua_State`; all functions inside Lua receive a pointer to this structure round as an argument. This implementation makes Lua reentrant noisy and ready for use in multi-strand applications.

As its name suggests, the `luaL_newstate` function creates a new Lua state. When `luaL_newstate` creates a new state, it does not contains no built-in functions, not even `print`. In order to keep Lua small, all standard libraries are represented as separate packages, so you are not required to use them if you don't need them. The `luaolib.h` header file defines the functions for opening libraries. `LuaL_openlibs` function opens all standard libraries.

After creating a state and populating it with standard libraries By the way, it's time to start doing user input. For each

the line that the user enters, the program first calls `luaL_loadstring` to compile the injected code. If mistakes no, then this call returns zero and places the resulting function per stack. (Remember that we will discuss the stack in detail in the following - See the next section.) The program then calls `lua_pcall`, which pops a function off the stack and executes it in protected mode.

Like `luaL_loadstring`, `lua_pcall` returns zero if no mistakes. In case of an error, both functions post a message about an error on the stack; we will receive this message using the function `lua_tostring`, and after we print it, we remove it from stack using the `lua_pop` function.

Please note that in case of an error, the program simply prints melts the error message to standard stream for errors. Infusion how error handling in C can be quite complex, and how should be followed, often depends on the type of your application. Nucleus Lua itself does not print anything to any stream (file); it in case error simply returns an error message. Each application ny can process these messages in the most appropriate way for him way. For simplicity, we will use the following ob-an error worker who, in case of an error, prints a message about error, closes the Lua state and exits the application:

```
#include <stdarg.h>
#include <stdio.h>
#include <stdlib.h>

void error (lua_State * L, const char * fmt, ...) {
    va_list argp;
    va_start (argp, fmt);
    vfprintf (stderr, fmt, argp);
    va_end (argp);
    lua_close (L);
    exit (EXIT_FAILURE);
}
```

We'll come back to handling errors in the application code later. Since you can compile Lua both as C code and as code in C ++, `lua.h` does not include the following standard amendment, common in C code:

```
#ifdef __cplusplus
extern "C" {
#endif

...
#ifdef __cplusplus
}
#endif
```

If you are compiling Lua as C code (the most common case) and use it in C ++, you can include `lua.hpp` instead `lua.h` . It is defined as follows:

```
extern "C" {  
#include "lua.h"  
}
```

25.2. Stack

When passing values between Lua and C, we are faced with two complexities: mismatch between static and dynamic systems, manual typing and mismatch between automatic and manual memory management.

In Lua, when we write `a[k] = v` , variables can have very different types, even `a` may have a different type (due to the use of metatables). However, if we want to provide this operation in C, then any `settable` function must be of a fixed type. So dozens of functions will be needed for this simple operation (one by one functions for each combination of the types of three arguments). We can solve this problem by introducing a new type - union of new types, let's call it `lua_Value` , which can represent all values in Lua. Then we can declare `settable` as follows way:

```
void lua_settable (lua_Value a, lua_Value k, lua_Value v);
```

However, this solution has two disadvantages. First, maybe it can be very difficult to map a complex data type to other languages; we designed Lua to interoperate easily not only with C / C ++, but also with Java, Fortran, C #, etc. Secondly, Lua does garbage collection: if we store a Lua table in variable `C`, then Lua itself cannot know about this in any way and can (error sideways) decide that this table is no longer needed and delete it.

Therefore, the Lua API does not define anything like the `lua_Value` type .

Instead, it uses an abstract stack to exchange values-between Lua and C. Each slot in this stack can contain any Lua value. When we want to get a value from Lua (for example, the value of a global variable), you call Lua and it pushes a value onto the stack. When you want to pass a Lua value, then you first push the value onto the stack and then call Lua (which pops that value off the stack). We still need different

functions to push each type of C onto the stack and each type on C to remove from the stack, but we no longer have a combinatorial increasing the number of functions as before. Moreover, since this the stack lives inside Lua, then the garbage collector knows what values uses C.

Almost all functions in the API use a stack. As we already see-

In our first example, the `luaL_loadstring` function leaves its result on the stack (either as a compiled block, or as an error message); `lua_pcall` takes the called function off the stack and leaves any error message on the stack.

Lua works with the stack strictly in accordance with the LIFO principle (Last In, First Out). When you call Lua, it only changes the top of the stack. C code has more freedom; in particular, he can view any item on the stack, as well as insert and remove elements from anywhere in the stack.

Pushing items onto the stack

The API contains one function to push on the stack each type C, which can be represented in Lua: `lua_pushnil` for constants ***nil***, `lua_pushboolean` for boolean values (integers sat in C), `lua_pushnumber` (for double), `lua_pushinteger` for integers numbers, `lua_pushunsigned` for unsigned integers, `lua_pushlstring` for arbitrary strings (pointer to `char` and length) and `lua_pushstring` for regular ASCII strings:

```
void lua_pushnil (lua_State * L);
void lua_pushboolean (lua_State * L, int bool);
void lua_pushnumber (lua_State * L, lua_Number n);
void lua_pushinteger (lua_State * L, lua_Integer n);
void lua_pushunsigned (lua_State * L, lua_Unsigned n);
void lua_pushlstring (lua_State * L, const char * s, size_t len);
void lua_pushstring (lua_State * L, const char * s);
```

There are also functions for pushing functions on the stack to C and objects of type *userdata*, but we'll look at them later.

The `lua_Number` type is a numeric type in Lua. By default this is the type `double`, but it can be changed to `float` or even `long int` for different personal architectures. The `lua_Integer` type is an integer type with sign large enough to hold the size of large lines. It is usually defined as `ptrdiff_t`. The `lua_Unsigned` type is 32-bit unsigned C integer type; used by the library

for bitwise operations and various functions.

Lua strings are not null terminated, they can contain arbitrary binary data. Accordingly, they should have an explicit length. The main function for pushing a string onto the stack is `lua_pushlstring`, requiring an explicit assignment of the length of the string. For null terminated strings you can use the `lua_pushstring` function, which to calculate string length uses `strlen`. Lua never stores pointers to external lines (or to any other external object, except functions in C). For any line that needs to be stored in a thread, Lua either makes a copy or reuses an existing one. Accordingly, you can free or modify your buffer like only control will return from these functions.

When you push an item onto the stack, it is your responsibility - make sure that there is enough space on the stack for it. Remember that you are now a C programmer. When Lua starts to be executed and whenever Lua calls C, there is at least 20 free slots. (The `lua.h` header file is open - modifies this constant as `LUA_MINSTACK`.) Usually more than enough, so you usually don't have to think about it. But some tasks require more stack space, in particular if you push items onto the stack in a loop. In these cases, you can call the `lua_checkstack` function, which checks if required free space on the stack:

```
int lua_checkstack (lua_State * L, int sz);
```

Referring to elements

The API uses *indexes* to refer to items on the stack. The first element pushed onto the stack has index 1, the next one is index 2, etc. We can also refer to the elements of the stack using negative indices. In this case, -1 corresponds to an element on the top of the stack (that is, the last element pushed onto the stack), -2 matches the previous item, and so on. For example, calling `lua_tostring (L, -1)` returns the value at the top of the stack as a string.

As we will see below, there are cases when it is natural to refer to the stack, starting from the bottom of the stack (that is, using positive

dex), and there are cases when it is natural to use negative indices.

To check if an element is the value of a given type, the API offers a set of `lua_is *` functions, where `*` can be any a common type in Lua. Accordingly, there are functions `lua_isnumber`, `lua_isstring`, `lua_istable`, etc. All of these functions have the same common prototype:

```
int lua_is * (lua_State * L, int index);
```

In fact `lua_isnumber` does not check whether the knowledge reading a number, but checks if the value can be converted to number; `lua_isstring` behaves similarly: in particular, any clean lo satisfies `lua_isstring`.

There is also a `lua_type` function that returns the type of an element cop on the stack. Each type is represented by a constant defined in the `lua.h` header file: `LUA_TNIL`, `LUA_TBOOLEAN`, `LUA_TNUMBER`, `LUA_TSTRING`, `LUA_TTABLE`, `LUA_TTHREAD`, `LUA_TUSERDATA` and `LUA_TFUNCTION`. We usually use this function in conjunction with the `oparathore` switch statement. It is also useful when we need check if a value is a number or a non-cast string types.

To get values from the stack, there are `lua_to *` functions:

```
int lua_toboolean
```

```
(lua_State * L, int index);
```

```
const char * lua_tolstring (lua_State * L, int index,  
size_t * len);
```

```
lua_Number lua_tonumber (lua_State * L, int index);
```

```
lua_Integer lua_tointeger (lua_State * L, int index);
```

```
lua_Unsigned lua_tounsigned (lua_State * L, int idx);
```

`Lua_toboolean` function converts any value to boolean some value in C (0 or 1), while using the Lua rules for boolean constructs: ***nil*** and ***false*** give 0, all others - the value 1.

Any of the `lua_to *` functions can be called, even when the value is of the wrong type. `Lua_toboolean` function works for values any type; `lua_tolstring` returns NULL for non-string values cheniy. Numeric functions have no way of reporting errors ke, so they simply return zero on error. Usually

`lua_isnumber` should be called to check the type, but in Lua 5.2 we introduced the following functions:

```
lua_Number lua_tonumberx (lua_State * L, int idx, int * isnum);
```

```
lua_Integer lua_tointegerx (lua_State * L, int idx, int * isnum);
```

```
lua_Unsigned lua_tounsignedx (lua_State * L, int idx, int * isnum);
```

The `isnum` parameter returns a boolean value indicating about whether the corresponding Lua value was a number. (If you need this no value is needed, then you can as the last parameter pass `NULL` . The old `lua_to *` functions are now implemented as macros based on these functions.)

`Lua_tolstring` function returns pointer to internal a copy of the string and remembers the length of the string via `len` parameter . You do not

you can change this internal copy (the `const` specifier reminds tells you about it). Lua ensures that this pointer is valid until as long as the corresponding line is on the stack. When the function in C, a call from Lua returns, then Lua clears the stack; so never cast pointers to Lua strings outside of a function, received them.

Any string that `lua_tolstring` returns always has a null byte at the end, but it can also contain null bytes inside yourself. The actual size of the line is returned through the third argument `len` . In particular, assuming that the value at the top of the stack is this is a string, the following asserts are always true:

```
size_t l;  
const char * s = lua_tolstring (L, -1, &l); /* any Lua string */  
assert (s [l] == '\0');  
assert (strlen (s) <= l);
```

You can call `lua_tolstring` with a third parameter equal to `NULL` if you don't need the length of the string. Or you can use the macro `lua_tostring` , which is actually `lua_tolstring` with third parameter equal to `NULL` .

In order to illustrate the use of these functions ttion Listing 25.2 shows a useful helper function, which prints the contents of the stack. This function traverses the entire stack bottom to top, printing each item according to its type. Strings are printed in quotes, numbers are formatted as `'%g'` , for other values (functions, tables, etc.), only a type. (The `lua_typename` function translates a numeric value that is specifying the type, into a string.)

Listing 25.2. Printing the contents of a stack

```
static void stackDump (lua_State * L) {  
    int i;  
    int top = lua_gettop (L); /* stack depth */  
    for (i = 1; i <= top; i++) { /* repeat for each level */
```

```

int t = lua_type (L, i);
switch (t) {
case LUA_TSTRING: { /* strings */
printf ("%s", lua_tostring (L, i));
break;
}
case LUA_TBOOLEAN: { /* booleans */
printf (lua_toboolean (L, i)? "true": "false");
break;
}
case LUA_TNUMBER: { /* numbers */
printf ("%g", lua_tonumber (L, i));
break;
}
default: { /* other values */
printf ("%s", lua_typename (L, t));
break;
}
}
printf (""); /* put a separator */
}
printf ("\n"); /* end the listing */
}

```

Other stack operations

In addition to the previous functions used to exchange data between C and Lua, the API also provides the following functions to work with stack:

```

int lua_gettop (lua_State * L);
void lua_settop (lua_State * L, int index);
void lua_pushvalue (lua_State * L, int index);
void lua_remove (lua_State * L, int index);
void lua_insert (lua_State * L, int index);
void lua_replace (lua_State * L, int index);
void lua_copy (lua_State * L, int fromidx, int toidx);

```

`Lua_gettop` function returns the number of items on the stack, so is equal to the index of the element at the top of the stack. `Lua_settop` function sets the number of items on the stack. If the previous value is the top of the stack was larger, then the extra values are thrown away. Otherwise, use as missing values is *nil*. In particular, `lua_settop (L, 0)` clears the entire stack. In function `lua_settop` you can also use negative indices.

In particular, the API provides the following macro, which removes there are `n` elements from the stack :

```
#define lua_pop (L, n) lua_settop (L, - (n) - 1)
```

The `lua_pushvalue` function pushes a copy of the item with the given index; the `lua_remove` function removes an element with a given index, while shifting all other elements; `lua_insert` removes an element from the top of the stack to the given position, while shifting elements to free up space; `lua_replace` removes the value from top of the stack and sets it to the value of the element with the given index; finally, `lua_copy` copies the value at one index to the value at a different index without changing the original value. Reverse Note that the following operations do not affect a non-empty stack:

```
lua_settop (L, -1); /* set vertex to current value */  
lua_insert (L, -1); /* move element from top to top */  
lua_copy (L, x, x); /* copy the element to its position */
```

The program in Listing 25.3 uses the `stackDump` function (defined shown in Listing 25.2) to illustrate these stack operations.

Listing 25.3. An example of working with a stack

```
#include <stdio.h>  
#include "lua.h"  
#include "lauxlib.h"  
static void stackDump (lua_State * L) {  
  <same as in Listing 25.2>  
}  
int main (void) {  
  lua_State * L = luaL_newstate ();  
  lua_pushboolean (L, 1);  
  lua_pushnumber (L, 10);  
  lua_pushnil (L);  
  lua_pushstring (L, "hello");  
  stackDump (L);  
  /* true 10 nil 'hello' */  
  lua_pushvalue (L, -4); stackDump (L);  
  /* true 10 nil 'hello' true */  
  lua_replace (L, 3); stackDump (L);  
  /* true 10 true 'hello' */  
  lua_settop (L, 6); stackDump (L);  
  /* true 10 true 'hello' nil nil */  
  lua_remove (L, -3); stackDump (L);
```

```

/* true 10 true nil nil */
lua_settop (L, -5); stackDump (L);
/* true */
lua_close (L);
return 0;
}

```

25.3. C API error handling

All structures in Lua are dynamic: they grow as necessary and reduced in size when possible. This means that in Lua we are constantly faced with a possible error when memory allocation. Almost any operation over time can lead to this. Moreover, many operations can cause other many mistakes; for example, accessing a global variable can lead to a call to the `__index` metamethod, and this metamethod can call mistake. Finally, operations that allocate memory over time call the garbage collector, which calls finalizers that They can also lead to errors. In short, overwhelming most functions in Lua can lead to errors.

Instead of using error codes in its API, Lua uses exceptions for reporting errors. Unlike C++ or Java, the C language does not contain a mechanism for working with exceptions. For to get around this problem, Lua uses the `setjmp` function from C, which provides a mechanism similar to exception handling. therefore most API functions can throw an error (that is, you call `longjmp`) instead of returning a value.

When we write library code (that is, functions that will can be called from Lua), using the `longjmp` function is almost like this as convenient as using real exceptions because Lua will catch any error that occurs. When we write code for-definitions (that is, the C code that calls Lua), then we must provide a way to catch such errors.

Error handling in application code

When your application calls functions from the Lua API, then it is liable to mistakes. As we discussed, Lua usually reports these errors using the `longjmp` function. However, if there is no corresponding

the corresponding call to `setjmp`, the interpreter cannot execute and `longjmp`. In this case, any error in the API results in Lua calls a *special function* (panic function), and if the control is returned from this function, then the execution of the application will rummage. You can define your similar function with `lua_atpanic`, but there isn't much that function can do. In order to properly handle errors in your code application, you have to call your code through Lua, so it will Sets the appropriate context for catching errors (i.e. it will execute your code in the context of `setjmp`). Just like we can- Let's run Lua code in protected mode using `pcall`, we can execute C code using `lua_pcall`. More accurately, we put the C code in a function and call this function through Lua using `lua_pcall`. (We will discuss in detail how to call Lua C functions in Chapter 27.) Then our C code will run in protected mode. Even in case of memory allocation error `lua_pcall` returns the corresponding error code, leaving the the interpreter is in working order.

Error handling in library code

Lua is a *safe* language. This means that it doesn't matter what you write in Lua, no matter how wrong it is, you can always understand running the program in terms of Lua itself. Moreover, errors are also are discovered and explained in Lua terms. You can compare this is with C, where the behavior of many mis-written programs can

can only be explained in terms of the equipment used (for example, error locations are specified as command addresses). When you add a C function to Lua, you break this safety. For example, a function like `poke` that writes an arbitrary byte at an arbitrary memory address, can add This leads to a large number of errors when working with memory. You need to make sure your additions are safe for Lua and provided good error handling.

As we discussed earlier, C programs must specify your error handling with `lua_pcall`. However, when you drink

you write arbitrary functions in Lua, usually you don't need to handle make mistakes. Errors thrown by the library function will be caught either by `pcall` in Lua or by `lua_pcall` in the application code. Therefore, when a function in a C library is catches an error, it might just call `lua_error` (or whatever even better is `luaL_error`, which formats the error message and then calls `lua_error`). The `lua_error` function clears everything that needs to be cleared in Lua, and jumps back to the protected output I call, passing the error message.

Exercises

Exercise 25.1. Compile and run a simple separate Lua interpreter (Listing 25.1).

Exercise 25.2. Let's assume the stack is empty. What will be on-walk on the stack after the next call sequence

WWOW?

```
lua_pushnumber (L, 3.5);  
lua_pushstring (L, "hello");  
lua_pushnil (L);  
lua_pushvalue (L, -2);  
lua_remove (L, 1);  
lua_insert (L, -2);
```

Exercise 25.3. Use the Lua interpreter from listing ha 25.1 and the `stackDump` function (Listing 25.2) in order to check your answer to the previous exercise.

Chapter 26

Extending your applications

An important use of Lua is to use it as a *con-figurative* language. In this chapter we will show how we can use Lua to configure a program, starting with a simple example, and we will develop it to perform more and more complex tasks.

26.1. The basics

As the first task, let's look at a simple configuration rational script: your C program has a window and you want be able to set the initial window size. It is clear that for such a simple problem, there are simpler solutions than Lua uses such as environment variables or files with name-value pairs. But, even using a simple text file, you somehow need to disassemble it; so you decide to use a Lua configuration file (i.e. a text file that is a Lua program). In its simplest form, this text the file may contain the following lines:

```
- define window size
width = 200
height = 300
```

You must now use the Lua API to enable Lua to parse took this file, and then get the values of global variables `width` and `height`. The `load` function in Listing 26.1 does this. bot. This function assumes that you have already created the Lua state, similar to what we saw in the previous chapter. She calls the function clause `luaL_loadfile` to load the block from the file `fname` and then binds `lua_pcall` to run compiled block. When errors (for example, syntax errors in your config on file), these functions push the error message onto the stack and return a nonzero error code; our program then uses is `lua_tostring` with index -1 in order to get the message from the top of the stack. (We discussed the `error` function in Section 25.1.)

Listing 26.1. Getting user information from the config file

```
void load (lua_State * L, const char * fname, int * w, int * h) {
if (luaL_loadfile (L, fname) || lua_pcall (L, 0, 0, 0))
error (L, "cannot run config. file:%s", lua_tostring (L, -1));
lua_getglobal (L, "width");
lua_getglobal (L, "height");
if (! lua_isnumber (L, -2))
error (L, "'width' should be a number \n");
if (! lua_isnumber (L, -1))
error (L, "'height' should be a number \n");
* w = lua_tointeger (L, -2);
* h = lua_tointeger (L, -1);
}
```


After executing a block of code, the program needs to get the value of global variables. To do this, it calls the function twice.

`lua_getglobal`, whose parameter (except approach to others the accompanying `lua_State`) is the variable name. Everyone is like that the call pushes the corresponding value onto the stack, so the width of the window will be at position with index -2 and height at position with index -1 (at the top). (Since the stack was initially empty, which we can also index starting from the bottom of the stack, then there is to use 1 for the first value and 2 for the second. But, indexing from the top, we do not need to make any assumptions the stack is empty.) Next, our example uses the `lua_isnumber` for checking if each value is number. Then `lua_tointeger` is called, and the corresponding values are assigned.

Was Lua worth using for a similar task? As I said earlier, for such a simple task, a simple text file with two lines will be much easier than Lua. Even so, using Lua gives us some advantages. First, Lua is fully concerned with syntax for you; your config file may even contain real comments! Second, the user can already execute complex configuration with what we have. For example, the script can ask the user for some information or take value from the environment variable to select the appropriate size:

```
- configuration file
if getenv("DISPLAY") == "" then
width = 300; height = 300
else
width = 200; height = 200
end
```

Even in such simple scenarios, configuration is difficult to anticipate what users might want; but until the script defines these two variables, your C program will work unchanged.

The final reason for using Lua is that it is now easy to add new configuration options to your program; this ease creates an approach that leads to much more flexible programs.

26.2. Working with tables

Let's take this approach: now we want to set the background color for window. We will assume that the given color consists of three numbers, each of the numbers is an RGB component. Usually in C these are integer in some range for example [0, 255]. In Lua, since all numbers are floating point, we will have more it is natural to use the range [0, 1].

A naive approach would be to ask the user to ask each component in a separate global variable:

```
- configuration file
```

```
width = 200
```

```
height = 300
```

```
background_red = 0.30
```

```
background_green = 0.10
```

```
background_blue = 0
```

This approach has two drawbacks: first, it is too redundant.

and cumbersome (real programs may need dozens of colors for the background in the window, for the color in the window, for the background of the menu, etc.); and

there is no way to pre-define common colors so that the user could then simply write `background = WHITE`. To avoid these disadvantages, we can represent colors by cabbage tables:

```
background = {r = 0.30, g = 0.10, b = 0}
```

Using tables gives structure to your script; Now easy for the user (or application) to define colors for further use in the config file:

```
BLUE = {r = 0, g = 0, b = 1.0}
```

```
<other color definitions>
```

```
background = BLUE
```

To obtain these values on C, we can do the following - in a way:

```
lua_getglobal (L, "background");
```

```
if (! lua_istable (L, -1))
```

```
error (L, "'background' is not a table");
```

```
red = getcolorfield (L, "r");
```

```
green = getcolorfield (L, "g");
```

```
blue = getcolorfield (L, "b");
```

Listing 26.2. An example implementation of the `getcolorfield` function

```
#define MAX_COLOR 255
```

```
/* assume that the table is on the top of the stack */
```

```

int getcolorfield (lua_State * L, const char * key) {
int result;
lua_pushstring (L, key); /* push the key onto the stack */
lua_gettable (L, -2); /* get background [key] */
if (! lua_isnumber (L, -1))
error (L, "invalid component in background color");
result = (int) (lua_tonumber (L, -1) * MAX_COLOR);
lua_pop (L, 1); /* remove number */
return result;
}

```

We first get the value of the global variable `background` and make sure it is a table and then use `getcolorfield` to get each component.

Of course the `getcolorfield` function is not part of the API, we have to define it. Again we are faced with the problem of polymorphism: there can be many versions of `getcolorfield` function that differ key type, value type, error handling, etc. Lua API offers just one `lua_gettable` function that works for of all types. She takes the position of the table on the stack, removes the key from

stack and pushes the corresponding value onto the stack. Our function `getcolorfield`, defined in Listing 26.2, assumes that the table is at the top of the stack, so after putting the key on the stack using the `lua_pushstring` function, the table will be located at index `-2`. Before returning, `getcolorfield` pops a semi-value, leaving the stack in the same state it was in before this call.

Since indexing a table with a string key is very common, Lua 5.1 introduced a specialized version `lua_gettable` is just for this case: `lua_getfield`. Using this function, we can rewrite the following two lines:

```

lua_pushstring (L, key);
lua_gettable (L, -2); /* get background [key] */
as

```

```

lua_getfield (L, -1, key);
(Since we are not pushing a row onto the stack, the table has an index
is still -1 when lua_getfield is called .)

```

We will expand our example a little and introduce names in color for user. The user can still use the `tab-` faces for the component color setting, but you can also use Call predefined color names. To implement this

we need a color table in our C program:

```
struct ColorTable {
    char * name;
    unsigned char red, green, blue;
} colortable [] = {
    {"WHITE", MAX_COLOR, MAX_COLOR, MAX_COLOR},
    {"RED", MAX_COLOR,
    0,
    0},
    {"GREEN",
    0, MAX_COLOR,
    0},
    {"BLUE",
    0,
    0, MAX_COLOR},
    <other colors>
    {NULL, 0, 0, 0} /* terminator */
};
```

Our implementation will create global variables with color-coded names comrade and initializes these variables using color tables

Comrade The result will be the same if the user added the following lines into your script:

```
WHITE = {r = 1.0, g = 1.0, b = 1.0}
RED = {r = 1.0, g = 0, b = 0}
<other colors>
```

To set the fields of the table, we will introduce a helper function `setcolorfield`; it pushes the index and value of the field onto the stack and then calls `lua_settable` :

```
/* consider the table to be at the top of the stack */
void setcolorfield (lua_State * L, const char * index, int value) {
    lua_pushstring (L, index); /* key */
    lua_pushnumber (L, (double) value / MAX_COLOR); /* value */
    lua_settable (L, -3);
}
```

Similar to other API functions, `lua_settable` works for many different types, so it pops all of its operands off the stack. It takes the index of the table as an argument and strips off the key and value from the stack. The `setcolorfield` function assumes that before calling the tabface is at the top of the stack (index `-1`); after placing index and values on the stack, the table will be at index `-3`.

Lua 5.1 also introduced a specialized version of `lua_settable` for string keys, it's called `lua_setfield` . Using this no-function, we can rewrite `setcolorfield` as follows

at once:

```
void setcolorfield (lua_State * L, const char * index, int value) {  
    lua_pushnumber (L, (double) value / MAX_COLOR);  
    lua_setfield (L, -2, index);  
}
```

The next function, `setcolor`, defines one color. She created creates a table, sets the values of the corresponding fields and maps this table to the corresponding global variable:

```
void setcolor (lua_State * L, struct ColorTable * ct) {  
    lua_newtable (L); /* creates a table */  
    setcolorfield (L, "r", ct->red); /* table.r = ct->r */  
    setcolorfield (L, "g", ct->green); /* table.g = ct->g */  
    setcolorfield (L, "b", ct->blue); /* table.b = ct->b */  
    lua_setglobal (L, ct->name); /* 'name' = table */  
}
```

`Lua_newtable` function creates an empty table and places it on the stack; `setcolorfield` calls set the fields of this table; finally, `lua_setglobal` pops the table off the stack and uses it as a value global variable with the given name.

Using these functions, the next cycle will register all colors. the one for the config script:

```
int i = 0;  
while (colortable [i] .name != NULL)  
    setcolor (L, & colortable [i ++]);
```

Remember that the application must execute this loop before you filling in the script.

Listing 26.3. Colors as rows or tables

```
lua_getglobal (L, "background");  
if (lua_isstring (L, -1)) { /* is the value a string? */  
    const char * name = lua_tostring (L, -1); /* get string */  
    int i; /* look in the table */  
    for (i = 0; colortable [i] .name != NULL; i++) {  
        if (strcmp (colorname, colortable [i] .name) == 0)  
            break;  
    }  
    if (colortable [i] .name == NULL) /* string not found? */  
        error (L, "invalid color name (%s)", colorname);  
    else { /* use colortable [i] */  
        red = colortable [i] .red;  
        green = colortable [i] .green;  
        blue = colortable [i] .blue;  
    }  
} else if (lua_istable (L, -1)) {  
    red = getcolorfield (L, "r");  
    green = getcolorfield (L, "g");
```

```
blue = getcolorfield (L, "b");  
} else  
error (L, "invalid value for 'background'");
```

Listing 26.3 shows another option for implementing name bathroom flowers. Instead of global variables, the user can denote colors using strings, writing the settings as `background = "BLUE"`. So background can be either table, or row. With this approach, the application does not need do something before running the script. Instead, to get colors have a little more work to do. When the program is receives the value of the `background` variable, then you need to check if is this value a string, in which case look for a color in color table.

What's the best option? In C programs, using lines to indicate options is not good practice since the compiler cannot detect typos. However, in Lua the message about a typo in the name of the color will go to the person for whom it is written

this configuration. Difference between programmer and user somewhat blurry; difference between compilation error and error runtime is not that great.

With strings, the value of the `background` variable can be string `Coy` with a typo; in this case the application can add this information to the error message. The application can also compare thread strings regardless of the case of letters, so that the user can write `"white"`, `"WHITE"` or even `"White"`. Moreover, if the script is small and many errors are found, then this may be not very successful - add hundreds of colors (and create hundreds of tabs) persons and global variables) only for the user to chose several colors. With strings, you avoid this.

26.3. Lua function calls

The strength of Lua is that the configuration file can can define functions that can then be called by the application genius. For example, you can write an application to plot the function, and use Lua to plot give a function whose graph will be plotted.

The API provided way to call functions is pretty simple:
first, you push the function to be called onto the stack;
second, you put the arguments to be called on the stack as well; after that-
go use `lua_pcall` to call the function and finally remove
results from the stack.

Let, as an example, our config file contains
there is a function like the one below:

```
function f (x, y)
return (x ^ 2 * math.sin (y)) / (1 - x)
end
```

You want to compute $z = f(x, y)$ in C for given x and y . Considering
that you have already opened the Lua library and performed the configuration-
file, the function `f` in Listing 26.4 implements this call.

Listing 26.4. Calling a Lua function from C

```
/* call function 'f' defined in Lua */
double f (lua_State * L, double x, double y) {
int isnum;
double z;
/* push function and arguments onto the stack */
lua_getglobal (L, "f"); /* called function */
lua_pushnumber (L, x); /* push 1st argument onto the stack */
lua_pushnumber (L, y); /* push 2nd argument onto the stack */
/* call the function (2 arguments, 1 result) */
if (lua_pcall (L, 2, 1, 0) != LUA_OK)
error (L, "error running function 'f':%s",
lua_tostring (L, -1));
/* get the result */
z = lua_tonumberx (L, -1, & isnum);
if (! isnum)
error (L, "function 'f' must return a number");
lua_pop (L, 1); /* pop the result off the stack */
return z;
}
```

The second and third arguments to `lua_pcall` are, respectively,
the number of arguments you pass and the number of results that
you want to receive. The fourth argument is a function
tions for error handling; we will discuss this soon. As with
assignments in Lua, calling `lua_pcall` casts a valid
the number of resulting values to the number you specify; if necessary
placing on the stack walk *nil* 's or removing the extra values. Front
pushing the results onto the stack `lua_pcall` removes the function from the stack.
tion and its arguments. When a function returns multiple values,
then the first value is pushed onto the stack first; for example, if possible
three values rotate, then the first of them will have index -3 ,

and the last -1 .

If an error occurs while executing `lua_pcall` the `lua_pcall` function returns an error code; besides, it puts the error message onto the stack (but still pops function and its arguments). However, before posting a message to the `lua_pcall` stack calls the message handling function if it has been asked. To set the message processing function, use those as the last argument to the `lua_pcall` function. Zero means no processing function and the final message is the outcome of the new error message. Otherwise, this argument must be the index on the stack where the message handling function is located. In such a case, the processing function should be placed on the stack to the called function and its arguments. For normal errors, the `lua_pcall` function returns the code `LUA_ERRRUN`. Two special types of errors deserve separate codes, because they never call a function processing unit. The first type is memory allocation errors. For similar errors, `lua_pcall` always returns `LUA_ERRMEM`. Second type of errors is errors when executing the handler itself. In this case, there is no point in calling the request again. The message is lost, so `lua_pcall` returns immediately. Correction with the `LUA_ERRERR` code. Lua 5.2 highlights a third type of error: when the finalizer throws an error, `lua_pcall` returns `LUA_ERRGCM` code. This code indicates that the error is not related to the challenge itself.

Listing 26.5. Generalized function call

```
#include <stdarg.h>
void call_va (lua_State * L, const char * func,
const char * sig, ...) {
    va_list vl;
    int nargs, nres; /* number of arguments and results */
    va_start (vl, sig);
    lua_getglobal (L, func); /* push the function onto the stack */
    <put arguments on the stack (Listing 26.6)>
    nres = strlen (sig); /* number of expected results */
    if (lua_pcall (L, nargs, nres, 0) != 0) /* make a call */
        error (L, "error calling '%s':%s", func,
        lua_tostring (L, -1));
    <get the results (Listing 26.7)>
    va_end (vl);
}
```


26.4. Generalized call function

As a more complex example, we will construct a universal function to call functions in Lua using `vararg` in C. Our function, let's call it `call_va`, takes the name of the function that need to be called, a string describing the types of arguments and results, then a list of arguments, and finally a list of pointers to the variables in which we want to get the results of the call. When using this function, we can easily overwrite our previous example as follows:

```
call_va (L, "f", "dd> d", x, y, & z);
```

The line `"dd> d"` means "two arguments of type `double` and one result. The result is of type `double`. This specifier uses `'d'` for `double`, `'i'` for integers and `'s'` for strings; the `'>'` character separates arguments from results. If the function returns nothing, then the `'>'` symbol is not-required.

Listing 26.5 shows the implementation of the `call_va` function. Despite As for the general view of this function, it follows the same path as our first example: pushes a function onto the stack, pushes arguments to the stack (Listing 26.6), makes the call, and receives the results (Listing 26.7). Most of the code is pretty straightforward, but there are some subtle ty. First, it doesn't check that `func` is a function; if it is not, then `lua_pcall` will raise an error. Secondly, since it pushes an arbitrary number of arguments onto the stack, it must check Check if there is free space on the stack. Thirdly, since the functional tion can return rows, then `call_va` cannot remove results from stack. This should be done by the caller after use. result strings (or copy them elsewhere).

Listing 26.6. Generally pushing arguments

```
for (narg = 0; * sig; narg++) { /* execute for each argument */
/* check stack space */
luaL_checkstack (L, 1, "too many arguments");
switch (* sig++) {
case 'd': /* double argument */
lua_pushnumber (L, va_arg (vl, double));
break;
case 'i': /* int argument */
```

```

lua_pushinteger (L, va_arg (vl, int));
break;
case 's': /* string argument */
lua_pushstring (L, va_arg (vl, char *));
break;
case '>': /* end of arguments */
goto endargs;
default:
error (L, "invalid option (% c)", * (sig - 1));
}
}
endargs:

```

Listing 26.7. Getting call results

```

nres = -nres; /* index of the first result on the stack */
while (* sig) { /* repeat for each result */
switch (* sig++) {
case 'd': { /* double result */
int isnum;
double n = lua_tonumberx (L, nres, & isnum);
if (! isnum)
error (L, "wrong result type");
* va_arg (vl, double *) = n;
break;
}
case 'i': { /* int result */
int isnum;
int n = lua_tointegerx (L, nres, & isnum);
if (! isnum)
error (L, "wrong result type");
* va_arg (vl, int *) = n;
break;
}
case 's': { /* string result */
const char * s = lua_tostring (L, nres);
if (s == NULL)
error (L, "wrong result type");
* va_arg (vl, const char **) = s;
break;
}
default:
error (L, "invalid option (% c)", * (sig - 1));
}
nres++;
}

```

Exercises

Exercise 26.1. Write a C program that reads a Lua file that defines a function f that takes as input number and returns the function value from this number. Your the program should plot this function. (You do not- be sure to use graphics, the usual one is fine text view using '*' for graph.)

Exercise 26.2. Modify the `call_va` function (Listing 26.5) to handle boolean values.

Exercise 26.3. Let there be a program that needs keep track of several weather stations. Inside, for representation of each station, it uses a 4-byte line, and there is a config file that matches Each such line contains the URL of the corresponding station. Con- the Lua figure file must do this mapping. in several different ways:

- a set of global variables, one for each stan- tion;
- one table mapping strings to URL;
- one function that returns a URL for each line.

Discuss the pros and cons of each option, taking attention total number of stations, types of users, availability structures in the URL, etc.

Exercises

Chapter 27

Calling C from Lua

When we say that Lua can call C, it doesn't mean that Lua may call any function in C ¹. As we saw in the previous chapter, when C calls a function in Lua, you must follow the a specific protocol for passing arguments and receiving results that. Similarly, in order for Lua to call a function in C, this function must follow a specific protocol to get their arguments and return results. Moreover, so that Lua can call a function in C, we must register this function, that is, they must pass Lua its address in a certain way. When Lua calls a function in C, it uses the same the stack that C uses to call Lua code. Function in C semi-pops its arguments off the stack and pushes its results onto the stack. The important concept here is that the stack is some structure; each function has its own local stack. When Lua calls a C function, the first argument will always be have index 1 on this local stack. Even when the C code calls Lua code that calls the same (or a different) function, each of these calls will only see his personal stack with the first ar-by index 1.

27.1. Functions in C

As a first example, let's see how to implement a simplified version of a function that returns the sine of a given numbers:

```
static int l_sin (lua_State * L) {  
    double d = lua_tonumber (L, 1); /* get argument */  
    lua_pushnumber (L, sin (d)); /* push the result onto the stack */  
    1
```

There are packages that allow Lua to call any function in C, but they don't.

portable and not secure.

```
return 1; /* number of results */  
}
```

Any function registered in Lua must have one and the same prototype defined in `lua.h` file as `lua_CFunction` :

```
typedef int (* lua_CFunction) (lua_State * L);
```

From the point of view of C, the function on C receives as its unique argument is the Lua state and returns an integer equal to the number of values returned through the stack. Therefore, the functions do not need

but clear the stack before pushing your results onto it. Pos-lua itself saves the results and cleans up stack.

Before we can use this feature, we must us first to register it. We do this with `lua_pushcfunction` : it receives a pointer to a C function and creates a value of type “function” that represents this function internally ri Lua. After registration, the C function behaves like any other function inside Lua.

A quick and dirty way to check the `l_sin` function is put its code directly into our base interpreter (Listing 25.1) and add the following lines right after the call

```
luaL_openlibs :  
lua_pushcfunction (L, l_sin);  
lua_setglobal (L, “mysin”);
```

The first line pushes the function type value onto the stack, and the second assigns its value to the global variable `mysin` . After these changes, you can use this new function directly in your their Lua scripts. (In the next section we will look at more correct good ways to connect C functions to Lua.)

For a more serious sine function, we must check the type of its argument. Here we are helped by an auxiliary library. `LuaL_checknumber` function checks if whether the given argument is a number: if an error occurs, it is thrown It has a meaningful error message, otherwise it returns the number itself. The changes to our function are minimal:

```
static int l_sin (lua_State * L) {
double d = luaL_checknumber (L, 1);
lua_pushnumber (L, sin (d));
return 1; /* number of results */
}
```

With this function definition, if we call `mysin ('a')` , then we will receive the following message:

bad argument # 1 to 'mysin' (number expected, got string)

Notice how `luaL_checknumber` is automatically filled in reads the message by argument number (# 1), function name (“mysin”), expected parameter type (number) and real parameter type (string).

As a more complex example, let's write a function, which will return the contents of the given directory. Lua is not provided has such a function in its standard libraries, because ANSI C does not have a suitable function for this. Here we will consider that our function supports POSIX. Our function - let's call its `dir` in Lua and `l_dir` in C - takes as an argument a string with path to a directory and returns an array with the contents of this directory. For example, calling `dir (“/ home / lua”)` might return the following table person { “.”, “..”, “src”, “bin”, “lib” } . In case of error, the function returns *nil* and an error string. The complete code of this the function is shown in Listing 27.1. Pay attention to the use calling the `luaL_checkstring` function , which behaves similarly `luaL_checknumber` , but only for strings.

(In extreme cases, using this function may result in to a small memory leak. Three of the Lua functions it calls are may fail due to insufficient memory

ty: `lua_newtable` , `lua_pushstring` and `lua_settable` . If any of these functions will fail, then an error will be raised and complements `l_dir` be interrupted, respectively, `closedir` will not caused. As we discussed earlier, for most programs this is not is a big problem: if memory runs out, then the best what can be done is to terminate the execution of the program. However less in chapter 30 we will see another implementation of the getter function contents of the directory, which no longer contains this error.)

Listing 27.1. Function for reading directory contents

```
#include <dirent.h>
#include <errno.h>
#include <string.h>
```

```

#include "lua.h"
#include "lauxlib.h"
static int l_dir (lua_State * L) {
  DIR * dir;
  struct dirent * entry;
  int i;
  const char * path = luaL_checkstring (L, 1);
  /* open directory */
  dir = opendir (path);
  if (dir == NULL) { /* error opening directory? */
    lua_pushnil (L); /* return nil ... */
    lua_pushstring (L, strerror (errno)); /* and message */
    return 2; /* number of results */
  }
  /* create a table with the result */
  lua_newtable (L);
  i = 1;
  while ((entry = readdir (dir)) != NULL) {
    lua_pushnumber (L, i ++); /* push key */
    lua_pushstring (L, entry-> d_name); /* push value */
    lua_settable (L, -3);
  }
  closedir (dir);
  return 1; /* table is already at the top of the stack */
}

```

27.2. Continuations

With `lua_pcall` and `lua_call`, a C function called from Lua can, in turn, call Lua. Some functions from the standard Noah libraries do this: `table.sort` can call a function comparisons; `string.gsub` can call the replace function; `pcall` and `xpcall` can call functions in protected mode. If we help him that the main Lua code was itself, in turn, called from C (main program), then we get a sequence like next: C (application) calls Lua (script), which calls gives a C function (library) that calls Lua. Lua normally handles these call sequences without problems; after all, its main task is to integrate with C. However, there is a situation in which such a chain of calls can cause problems: coroutines. Every coroutine in Lua has its own stack, which contains information about pending coroutine calls. More precisely,

the stack remembers the return address, parameters and local variables every call. For function calls in Lua, the interpreter is used calls a suitable data structure to implement the stack, which is It is called *a flexible stack* (soft stack). However, for function calls in C, The interpreter must also use the C stack. After all, the address return and local variables of a function in C live on the C stack. An interpreter can easily have many flexible stacks, but the code for C only has one stack. Therefore, coroutines in Lua cannot stop execution inside the function in C: if in the call chain from `resume` to the corresponding `yield` there is a function in C, then Lua cannot save the state of this function in order to restore renew it the next time you `resume` . Let's consider the next example in Lua 5.1:

```
co = coroutine.wrap (function (a)
return pcall (function (x)
coroutine.yield (x [1])
return x [3]
end, a)
end)
print (co ({10, 3, -8, 15}))
-> false attempt to yield across metamethod / C-call boundary
```

The `pcall` call is a C function; therefore Lua cannot “freeze zit "it, because there is no ANSI C way to pause execution function in C and then continue its execution.

Lua 5.2 has dealt with this complexity with *continuations* (continuation). Lua 5.2 implements `yield` with `longjmp` , i.e. the same way it implements errors. Such a call (`longjmp`) is simply from-discards all information about functions on the C-stack, so there is no you can continue the function in C. However, the function in C `foo` can specify a function - continuation of `foo-c` that is different a C function to be called when the time is right "Continue" function `foo` . That is, the interpreter detects that it should continue executing `foo` , but since all the information mation about the `foo` was destroyed with the stack, instead of that, he is `foo-c` .

To make it clearer, let's look at an example: implementation of the `pcall` function . In Lua 5.1. this function had the following- the following code:

```
static int luaB_pcall (lua_State * L) {
int status;
luaL_checkany (L, 1); /* at least one parameter */
```

```

status = lua_pcall (L, lua_gettop (L) - 1, LUA_MULTRET, 0);
lua_pushboolean (L, (status == 0)); /* status */
lua_insert (L, 1); /* status is the first result */
return lua_gettop (L); /* return status + all results */
}

```

If the function called via `lua_pcall` is suspended if it was executed (via `yield`), it will be impossible to continue live later execution of `luaB_pcall`. Therefore, the interpreter will issue error every time we try to call `yield` inside the a protected call. Lua 5.2 implements `pcall` like shown in Listing 27.2². There are three differences from the Lua 5.1 version: firstly, the new version replaced the `lua_pcall` call with `lua_pcallk`; secondly, she has grouped everything that is done after this call in the new `finishpcall` helper function; the third difference is the `pcallcont` function, the last argument is `lua_pcallk`, which is is a continuation function.

If there are no `yield` calls, then `lua_pcallk` works exactly like `lua_pcall`. If there is a call to `yield`, then everything is completely different. If a function called by `lua_pcall` tries to call `yield`, then Lua 5.2 raises an error, just like Lua 5.1. However, when the function called `lua_pcallk`, calls `yield`, then there is no error: Lua calls `longjmp` and discards the entry for `luaB_pcall` from the C stack, but keeps This adds a reference to the `pcallcont` continuation function in the flexible stack. Later, when the interpreter finds that it should return to `luaB_pcall` (which is not possible), it calls the function instead-continued `pcallcont`.

Unlike `luaB_pcall`, `pcallcont`'s continuation function cannot can get the value returned by the `lua_pcallk` call. therefore Lua provides a special function to return status call: `lua_getctx`. When called from a normal Lua function (which in our case does not happen), `lua_getctx` returns `LUA_OK`. Kog-yes it is called from a continuation function, it returns `LUA_YIELD`. The continuation function can also be called in the case of some errors; in this case `lua_getctx` returns an error code, that is the very value that `lua_callk` would return in this case.

Listing 27.2. A `pcall` implementation with continuations

```

static int finishpcall (lua_State * L, int status) {
lua_pushboolean (L, status); /* first result (status) */
lua_insert (L, 1); /* put the first result in the first slot */
return lua_gettop (L);
}

```

```
static int pcallcont (lua_State * L) {
int status = lua_getctx (L, NULL);
return finishpcall (L, (status == LUA_YIELD));
}
2
```

The actual code is more complex than shown here, since it has some general other parts with xpcall and checking for stack overflow before placing on it boolean value.

```
static int luaB_pcall (lua_State * L) {
int status;
luaL_checkany (L, 1);
status = lua_pcallk (L, lua_gettop (L) - 2, LUA_MULTRET, 0,
0, pcallcont);
return finishpcall (L, (status == LUA_OK));
}
```

In addition to the call status, the `lua_getctx` function can also return context information. The fifth parameter in `lua_pcallk` is a free integer that can be obtained through the second parameter `lua_getctx`, which is a pointer to an integer value. This is the price the left value allows the original function to pass an arbitrary information directly to its continuation. She can also convey more information via the Lua stack. (Our example is not used takes this opportunity.)

The continuation mechanism in Lua 5.2 is awesome to support `yield`, but this is not a panacea. Some C function may need to convey too much context to your pro posts. You can use table as such an example .

`sort`, which uses the C stack for recursion, and `string.gsub`, which paradise should keep track of found substrings and a buffer for partial results. Although they can be rewritten in a way that supports `yuschim yield statement`, the gain is not worth the complexity introduced.

27.3. C modules

A module in Lua is a block of code that defines various functions. tions in Lua and remembers them in suitable places, usually tab-faces. A C module for Lua behaves the same way. Besides defining its functions in C, it must also define a special function tion, which acts as the main block in the Lua library. This the function must register all C functions from the module and remember them in appropriate places, usually in table fields.

Like the main block in Lua, this function also needs to initialize everything that needs initialization.

Lua obtains functions in C through the registration mechanism. After Once a C function is represented and stored in Lua, Lua calls her through a direct link to her address (which we transmit Lua when we register this function). In other words, Lua is not depends on the function name, package location, or visibility rules ty to call this function when it is registered.

Usually a C module has only one `extern` function, which is the function that opens the library. Everything else-functions can be closed, for example, by declaring them as `static` .

When you extend Lua with C functions, a good the idea is to organize your code as a C module, even if you only want to register one function: sooner or later (usually early) you will need other features. As usual, help The powerful library offers a helper function for this. The `luaL_newlib` macro takes a list of C functions along with their appropriate names and registers them all inside the new tab-faces. As an example, let's say we want to create a library with a function tion `l_dir` , which we defined earlier. First, we must define library functions:

```
static int l_dir (lua_State * L) {  
<as before>  
}
```

Next, we define an array with all the functions along with their names. This array contains elements of type `luaL_Reg` , which is a structure of two fields: the function name (string) and the a function factor.

```
static const struct luaL_Reg mylib [] = {  
{"Dir", l_dir},  
{NULL, NULL} /* terminator */  
};
```

In our example, there is only one function (`l_dir`), which we want to register. The last pair in the array is always `{NULL, NULL}` denoting the end of the array. Finally, we define

We write the **main** function using `luaL_newlib` :

```
int luaopen_mylib (lua_State * L) {  
luaL_newlib (L, mylib);  
return 1;  
}
```

The `luaL_newlib` call creates a new table and fills it in pairs. The name is a function from the `mylib` array. Upon returning `luaL_newlib` OS-Pushes a new table onto the stack. Function `luaopen_mylib` returns 1 in order to return this table to Lua.

After completing the library, we must link it to the interpreter. The easiest way to do this is using dynamic libraries if the interpreter is Lua supports them. In this case, you must create a dynamic library with your code (`mylib.dll` on Windows and `mylib.so` on Linux) and place it along the C-path. After these steps, you can download this library directly from Lua using `require` :

```
local mylib = require "mylib"
```

This call loads a dynamic library in Lua, finds it, dumps the `luaopen_mylib` function, registers it as a C function, and calls it, thereby opening the module. (This explains why `luaopen_mylib` must have the same prototype as any other function in C.)

When loading a dynamic library, we need to know the name of the function `luaopen_mylib` in order to find it. Will always be used a function named `luaopen_`, to which the name is attached to the module. Therefore, if your module is called `mylib`, then the function should be named `luaopen_mylib`.

If your interpreter does not support dynamic linking, then you need to rebuild Lua along with your new library. Besides this rebuilding, you also need some way of saying to the interpreter that it should open this library when creating new state. This is usually done by adding `luaopen_mylib` to the list of standard libraries that `luaL_openlib` opens in the file `linit.c`.

Exercises

Exercise 27.1. Write a summation function in C that calculates the sum of a variable number of its numeric arguments:

```
print (summation ())
```

```
-> 0
```

```
print (summation (2.3, 5.4))
```

```
-> 7.7
```

```
print (summation (2.3, 5.4, -34)) -> -26.3
```

```
print (summation (2.3, 5.4, {}))  
-> stdin: 1: bad argument # 3 to 'summation'  
(number expected, got table)
```

Exercise 27.2. Implement a function equivalent to `table.pack` from the standard library.

Exercise 27.3. Write a function that gets produced free number of parameters and returns them in reverse order:
`print (reverse (1, "hello", 20)) -> 20 hello 1`

Exercise 27.4. Write a `foreach` function that gets-takes a table and a function as input and calls this function to each key-value pair in the table:
`foreach ({x = 10, y = 20}, print)`
`-> x 10`
`-> y 20`

Exercise 27.5. Rewrite the `foreach` function from the previous th exercise so that the called function can call `yield`.

Exercise 27.6. Create a C module with all the functions from previous exercises.

Exercises

Chapter 28

Writing techniques functions in C

Both the official API and the *auxiliary* library provide provide several mechanisms to help write functions in C. In this chapter, we will look at mechanisms for working with arrays, strings and storing Lua values in C.

28.1. Working with arrays

In Lua, "array" is just a name for the table used by the spec. in a special way. We can work with arrays using those the same functions that we used to work with the table-mi, that is, `lua_settable` and `lua_gettable`. However, the API is provided by `luaL`. There are several special functions for working with arrays. One one of the advantages of using these functions is performance: often we have access to an array inside the loop of the algorithm

ma (for example, sorting), so that any increase in speed the viya in these operations can have a big impact on the final performance of the algorithm. Another plus is convenience, integer keys are common enough that would deserve special treatment.

The API provides two functions for working with arrays:

```
void lua_rawgeti (lua_State * L, int index, int key);
```

```
void lua_rawseti (lua_State * L, int index, int key);
```

The description of the `lua_rawgeti` and `lua_rawseti` functions is somewhat confusing.

is safe, since it includes two indices at once: `index` describes where the table is on the stack; `key` specifies an element in the table itself.

The `lua_rawgeti (L, t, key)` call is equivalent to the following sequence when `t` is greater than zero (otherwise it is necessary compensate for the appearance of a new element on the stack):

```
lua_pushnumber (L, key);
```

```
lua_rawget (L, t);
```

Calling `lua_rawseti (L, t, key)` (again for positive `t`) is equivalent to the following sequence:

```
lua_pushnumber (L, key);
```

```
lua_insert (L, -2); /* put 'key' below the previous value */
```

```
lua_rawset (L, t);
```

Note that both functions use direct to the table. They are faster and moreover the tables used as arrays, metamethods are rarely used.

Listing 28.1. C map function

```
int l_map (lua_State * L) {
    int i, n;
    /* 1st argument must be table (t) */
    luaL_checktype (L, 1, LUA_TTABLE);
    /* 2nd argument must be function (f) */
    luaL_checktype (L, 2, LUA_TFUNCTION);
    n = luaL_len (L, 1); /* get table size */
    for (i = 1; i <= n; i++) {
        lua_pushvalue (L, 2); /* push f */
        lua_rawgeti (L, 1, i); /* push t[i] onto the stack */
        lua_call (L, 1, 1); /* call f (t[i]) */
        lua_rawseti (L, 1, i); /* t[i] = result */
    }
    return 0; /* No results */
}
```

As an example of using these functions, Listing 28.1 re-lizuet function `map` : it applies a given function to all element the array, replacing each element with the result of the call. This the example also introduces three new functions: `luaL_checktype` , `luaL_len`

and `lua_pcall` .

The `luaL_checktype` function (from the `lua.h` file) checks that the given argument is of the given type, otherwise it is throws an error.

The `lua_len` primitive (not used in the example above) is equivalent to is valent to the `#` operator . Due to metamethods, this operator can return an object of any type, not just numbers; so `lua_len` returns its result on the stack. `LuaL_len` function (using bathroom in the example) raises an error if the length is not pure scrap, otherwise it returns the length as normal integer type.

The `lua_call` function makes an unsecured call. He is analogous `lua_pcall` is hygienic , but it passes the errors above rather than returning the code errors. When you write the main application code, you better not use `lua_call` as you want to catch any errors.

However, when you write functions, it is better to use exactly `lua_call` ; if an error occurs, then we will leave it for someone, to whom it is important.

28.2. Working with strings

When a C function receives a string argument from Lua, there is there are only two rules that need to be followed: do not remove string off the stack when working with it and never modify the string.

The situation gets more complicated when C functions need to be created string to return it to Lua. C code should worry about buffer allocation / deallocation, buffer overflow, etc. at least the Lua API provides several functions for this.

The Standard API provides assistance in two of the most common common operations: substring extraction and concatenation lines. When you select a substring remember that `lua_pushlstring` takes the length of the string as an optional argument. So if you want to pass Lua substring of string `s` with characters at positions from `i` up to `j` (inclusive), then all you have to do is:

```
lua_pushlstring (L, s + i, j - i + 1);
```

As an example, let's say you want a function that breaks down emits a string at a given delimiter (one character) and returns Creates a table with substrings. For example, calling `split` ("hi: ho:

there ",": ") should return the table { " hi ", " ho ", " there " } . Fox

Thing 28.2 shows a simple implementation of this function. She doesn't need we have additional buffers and it does not impose any restrictions on the size of the lines that it can handle. About all the buffers are taken care of by Lua itself.

Listing 28.2. Splitting a string

```
static int l_split (lua_State * L) {
  const char * s = luaL_checkstring (L, 1); /* line */
  const char * sep = luaL_checkstring (L, 2); /* delimiter */
  const char * e;
  int i = 1;
  lua_newtable (L); /* table with result */
  /* repeat for each separator */
  while ((e = strchr (s, * sep)) != NULL) {
    lua_pushlstring (L, s, e); /* substring per stack */
    lua_rawseti (L, -2, i ++); /* insert into table */
    s = e + 1; /* go beyond separator */
  }
  /* insert last substring */
  lua_pushstring (L, s);
  lua_rawseti (L, -2, i);
  return 1; /* return table */
}
```

For string concatenation, Lua provides a special function in its API called `lua_concat` . It is equivalent to operator of concatenation .. in Lua; it converts numbers to strings and when calls metamethods if necessary. Moreover, she can immediately concatenate more than two lines. Calling `lua_concat (L, n)` concatenate pops (popping from the stack) `n` values and puts the result on top stack.

Another useful function is `lua_pushfstring` :

```
const char * lua_pushfstring (lua_State * L, const char * fmt, ...);
```

It is somewhat similar to the `sprintf` function in that it creates a string `ku` by format string and additional arguments. However, in excellent from `sprintf` , you don't need to provide a buffer. Lua dynamically creates a string for you as large as needed. This function pushes the resulting string onto the stack and returns a pointer to her. You don't need to worry about buffer overflows.

Currently, this function only supports the following `1` formats :

`% s`

Insert null terminated string

`% d`

Insert integer

`% f`

Insert Lua number i.e. double

% p

Insert Pointer

% c

Insert integer as character

%%

Insert character '%'

No modifiers supported such as width or accuracy.

Both `lua_concat` and `lua_pushfstring` are useful when we want to concatenate only a few lines. However, if we need to connect many lines (or characters) together, then do it one at a time can be quite inefficient, as we saw in section 11.6.

In this case, we can use the buffers provided by auxiliary library.

In the simplest case, buffers work like two functions: one gives you a buffer of any size where you can write your string; the other converts the buffer to a string in Lua ². Listing 28.3 showing

Can't use these functions by implementing the function

`string.upper` directly from the original `lstrlib.c` file. The first step

using the buffer from the auxiliary library is

the phenomenon of a variable of type `luaL_Buffer`. The next step is

call `luaL_buffinitsize` to get a pointer to a buffer with a given

new size; then you can use this buffer to create-

this line. The last step is to call `luaL_pushresultsize`

to convert the contents of the buffer to a new Lua string on ver-

the stack bus. The size in this call is the final size of the string.

(Often, as in our example, this size is equal to the size of the buffer, but it maybe less. If you do not know the exact size of the resulting string, but you have its maximum size, then you can order larger buffer.)

Listing 28.3. `String.upper` function

```
static int str_upper (lua_State * L) {
    size_t l;
    size_t i;
    luaL_Buffer b;
    const char * s = luaL_checklstring (L, 1, & l);
    char * p = luaL_buffinitsize (L, & b, l);
    for (i = 0; i < l; i++)
        p[i] = toupper (uchar (s[i]));
    luaL_pushresultsize (& b, l);
    return 1;
}
```

```
}
```

Note that the `luaL_pushresultsize` function is not takes the Lua state as its first argument. After the initialization buffer stores a reference to the state, so we don't you need to pass it when calling other functions to work with buffers.

We can also use these buffers without knowing the maximum the length of the resulting string. Listing 28.4 shows a simplified re-
lisation of the `table.concat` function . In this function, we first call
Vai `luaL_buffinit` to initialize the buffer. Then we add

to the buffer elements one by one, in this case using the function
`luaL_addvalue` . Finally, `luaL_pushresult` releases the buffer and
Places the summary line at the top of the stack.

Listing 28.4. Simplified implementation of `table.concat`

```
static int tconcat (lua_State * L) {  
    luaL_Buffer b;  
    int i, n;  
    luaL_checktype (L, 1, LUA_TTABLE);  
    n = luaL_len (L, 1);  
    luaL_buffinit (L, & b);  
    for (i = 1; i <= n; i++) {  
        lua_rawgeti (L, 1, i); /* get row from table */  
        luaL_addvalue (b); /* add it to the buffer */  
    }  
    luaL_pushresult (& b);  
    return 1;  
}
```

The helper library provides several functions
to add values to the buffer: function `luaL_addvalue` to-
Adds the Lua string that is at the top of the stack; function
`luaL_addlstring` adds strings with the specified length; function
`luaL_addstring` adds a null terminated string and
the `luaL_addchar` function adds single characters. These functions
tions have the following prototypes:

```
void luaL_buffinit (lua_State * L, luaL_Buffer * B);  
void luaL_addvalue (luaL_Buffer * B);  
void luaL_addlstring (luaL_Buffer * B, const char * s, size_t l);  
void luaL_addstring (luaL_Buffer * B, const char * s);  
void luaL_addchar (luaL_Buffer * B, char c);  
void luaL_pushresult (luaL_Buffer * B);
```

When you use a buffer, pay attention to the following.

After initializing the buffer, it stores some auxiliary results on the Lua stack. Therefore, you cannot assume that the top of the stack will remain where it was before you put use a buffer. You can use the stack for others tasks while working with a buffer, the main thing is that calls to `push` and `pop` were balanced every time you use the buffer. Ex- the key to this rule is the `luaL_addvalue` function , which paradise assumes that the line to be added to the buffer was pushed to the top of the stack.

28.3. Saving state in C functions

Often C functions need to store some kind of non-local data, that is, data that will survive the current function call. In C we usually we use global (extern) or static variables- for this purpose. However, when you write library functions for Lua, then using global or static variables not a good solution. First, you cannot save an arbitrary Lua value in a C variable. Secondly, the library, which uses such variables will not work correctly with multiple Lua states.

A Lua function has two basic places where you can store non-lightweight local data: global variables and non-local variables.

The C API also provides two basic storage locations for non-local data: registry and values associated with the function (upvalue).

The registry is a global table that can be accessed only with a C `3` code . Typically the registry is used to storage of data that will be used by several modules. If you need to save data only for your module or functions, then you must use the values associated with function.

Registry

The registry is usually located at a *pseudo-index* , the value of which is defined divided as `LUA_REGISTRYINDEX` . The pseudo-index looks like an index

on the stack, except that the associated values are not walk on the stack. Most of the functions in the Lua API that accept take indices as arguments, they also accept pseudoin- dexes - except for those functions that change the stack, such like `lua_remove` and `lua_insert`. For example, in order to get the value associated with the "Key" in the registry, we can use make the following call:

```
lua_getfield (L, LUA_REGISTRYINDEX, "Key");
```

The registry is just a regular Lua table. Accordingly, you can for to call it, use any Lua value other than *nil*. One-

3

In fact, the registry can also be accessed from Lua using a function from the debugger. in the `debug.getregistry` sub-library.

but since all C modules share the same registry, you must be very careful in choosing the values that you will use use as keys in order to avoid possible conflicts

Comrade String keys are especially handy when you want to allow other modules access your data, since everything they need is the name. There is no completely reliable method for bunch of keys, but there are some proven approaches, such as not using common names and starting your names with the library name or something like it. (Prefixes like `lua` and `luaLib` are not good options.)

You should never use numbers as keys registry, since such keys are reserved for the *system links* (reference system). This system consists of a pair of functions in helper libraries that allow you to save readings in the table without worrying about the uniqueness of the keys.

Function

`luaL_ref` creates new links:

```
int r = luaL_ref (L, LUA_REGISTRYINDEX);
```

This call will pop the value from the top of the stack, store it in a table face at the new integer index and will return that index. By- Daubney indexes are called *links* (reference).

As the name suggests, we will use links mainly for when we need to store a Lua value inside a struct ry C. As we have already seen, we should never memorize a pointer whether to Lua strings outside of the C function that received them. More Moreover, Lua doesn't even offer pointers to other objects such as tables or functions. Therefore, we cannot refer to objects

Lua with pointers. Instead, when we need that-
these are pointers, we will create links and store them in C.
To put the value associated with the reference `r` to
stack, we just use the following piece of code:

```
lua_rawgeti (L, LUA_REGISTRYINDEX, r);
```

Finally, in order to release both meaning and reference, we
call `luaL_unref` :

```
luaL_unref (L, LUA_REGISTRYINDEX, r);
```

After that, a new call to `luaL_ref` can return the same
link.

The referencing system treats ***nil*** as a special case. When we call
`luaL_ref` for the value of ***nil*** , a new link is created, and instead
this returns the constant `LUA_REFNIL` . The next call to `sa-`

Actually does nothing:

```
luaL_unref (L, LUA_REGISTRYINDEX, LUA_REFNIL);
```

The next call pushes ***nil*** onto the stack , as expected:

```
lua_rawgeti (L, LUA_REGISTRYINDEX, LUA_REFNIL);
```

The reference system also defines the constant `LUA_NOREF`, which
is an integer other than any link. She is useful
in order to mark links as destroyed / uninitialized-
dyed.

Another reliable method for generating registry keys is
using a static variable as a key address in your

In the same code: the linker guarantees that this address is unique.

In order to use this option, you need the function

`lua_pushlightuserdata` , which pushes a value onto the Lua stack,
which is a pointer to C. The following code shows

how to save and retrieve a string from the registry using this method

Yes:

```
/* variable with unique address */
```

```
static char Key = 'k';
```

```
/* remember the line */
```

```
lua_pushlightuserdata (L, (void *) & Key); /* push address */
```

```
lua_pushstring (L, myStr); /* push value */
```

```
lua_settable (L, LUA_REGISTRYINDEX); /* registry [& Key] = myStr */
```

```
/* get string */
```

```
lua_pushlightuserdata (L, (void *) & Key); /* push address */
```

```
lua_gettable (L, LUA_REGISTRYINDEX); /* get value */
```

```
myStr = lua_tostring (L, -1); /* convert it to string */
```

We will discuss in more detail the use of the `userdata` type in
case 29.5.

In order to simplify the use of variable addresses in the

as unique keys, Lua 5.2 introduces two new functions: `lua_rawgetp` and `lua_rawsetp`. They are like `lua_rawgeti` / `lua_rawseti`, but instead of integers, they use pointers (translated to userdata) as keys. Using them, we can rewrite the previous code as follows:

```
static char Key = 'k';
/* remember the line */
lua_pushstring (L, myStr);
lua_rawsetp (L, LUA_REGISTRYINDEX, (void *) & Key);
/* get string */
lua_rawgetp (L, LUA_REGISTRYINDEX, (void *) & Key);
myStr = lua_tostring (L, -1);
```

Function related values

While the registry offers global variables, the mechanism values associated with functions offers an analogue of static-variables-in C, which are visible only inside a separate function. Each time when you create a new C function in Lua you can bind with it any number of similar values; every such value is a Lua value. Then, when the function is called, it accesses any of these values freely using pseudo-indices.

We call this relationship functions in C with their values *z closures* (closure). The C closure is analogous to the Lua closure for the C language. In particular, you can create various closures-using the same function code, but different associated values *cheniya*.

As a simple example, let's write a function

`newCounter` in C. This function is a factory: it returns

Creates a new counting function each time it is called. Although everyone is like that

functions have the same C code, each of them stores its own own counter. This factory function looks like this way:

```
static int counter (lua_State * L); /* forward declaration */
int newCounter (lua_State * L) {
    lua_pushinteger (L, 0);
    lua_pushcclosure (L, & counter, 1);
    return 1;
}
```

```
}
```

The main function here is `lua_pushcclosure`, which creates a new closure. Its second argument is the base function (in the example this is `counter`), and the third argument is a number associated values (in the example, this is 1). Before creating a new job we have to put the initial values for the related values on the stack. In our example, we put 0 as the initial value for the only associated value. As expected, `lua_pushcclosure` leaves a new closure on the stack, so the closure is already ready to return as the result of `newCounter`. Let's now take a look at the definition of the `counter` function :

```
static int counter (lua_State * L) {  
  int val = lua_tointeger (L, lua_upvalueindex (1));  
  lua_pushinteger (L, ++ val); /* new value */  
  lua_pushvalue (L, -1); /* duplicate it */  
  lua_replace (L, lua_upvalueindex (1)); /* update value */  
  return 1; /* return new value */  
}
```

The key element here is the `lua_upvalueindex` macro, which returns the pseudo-index of the associated value. In part No. 8 expression `lua_upvalueindex (1)` returns the pseudo-index of the first associated value of the current function. This pseudo-index looks like any index on the stack, only it's not on the stack. Therefore, the call to `lua_tointeger` returns the current value first (and only) associated value as a number. Then the function pushes the new value `++ val` onto the stack, makes a copy of it, and uses one of the copies to replace the associated value. Finally, it returns another copy as its value.

As a more complex example, we will implement tuples using associated values. A *tuple* is something like a constant Yanny records with anonymous fields; you can get specific new field by index or you can get all fields at once. In our implementations, we will represent tuples as functions that remember their values in related values. When the function called with a numeric argument, then it returns the specific field. When called with no arguments, it returns all fields. Next The following code demonstrates the use of tuples:

```
x = tuple.new (10, "hi", {}, 3)  
print (x (1)) -> 10
```



```
print (x (2)) -> hi
print (x ()) -> 10 hi table: 0x8087878 3
```

In C, we represent all tuples using the same the `t_tuple` function, shown in Listing 28.5. Since we can let's call a tuple with a numeric argument or without arguments, the `t_tuple` function uses `luaL_optint` to reading an optional argument. `LuaL_optint` function is like `luaL_checkint`, but if the argument is absent, then it simply returns the specified default value (in the example it is 0).

Listing 28.5. Implementing tuples

```
int t_tuple (lua_State * L) {
    int op = luaL_optint (L, 1, 0);
    if (op == 0) { /* no arguments? */
        int i;
        /* push each associated value onto the stack */
        for (i = 1; ! lua_isnone (L, lua_upvalueindex (i)); i++)
            lua_pushvalue (L, lua_upvalueindex (i));
        return i - 1; /* number of values on the stack */
    }
    else { /* get the 'op' field */
        luaL_argcheck (L, 0 < op, 1, "index out of range");
        if (lua_isnone (L, lua_upvalueindex (op)))
            return 0; /* no field */
        lua_pushvalue (L, lua_upvalueindex (op));
        return 1;
    }
}

int t_new (lua_State * L) {
    lua_pushcclosure (L, t_tuple, lua_gettop (L));
    return 1;
}

static const struct luaL_Reg tuplelib [] = {
    {"New", t_new},
    {NULL, NULL}
};

int luaopen_tuple (lua_State * L) {
    luaL_newlib (L, tuplelib);
    return 1;
}
```

When we refer to a non-existent associated value, the result is a pseudo-value of type `LUA_TNONE`. (When we accessing the value above the top of the stack, then we also get pseudo-value of type `LUA_TNONE`.) Therefore, our function `t_tuple` uses `lua_isnone` to check if there is a corresponding value

reading. However, we should never call `lua_upvalueindex` with a negative index, so we must check this when yes the index is provided by the user. `LuaL_argcheck` function checks for any given value, throwing an error if needed availability.

Function for creating `t_new` tuples (also in Listing 28.5) trivial: since all of her arguments are already on the stack, she just binds `lua_pushcclosure` to create a circuit using their arguments as bound values. Finally, the `tupplelib` array and the `luaopen_tuple` function (also in Listing 28.5) are standard code to create a tuple library with a single function `new` .

The values associated with the function are used by multiple functions

Quite often we need to give access to multiple values or variables to all functions of this module. Although we can use call the registry for this purpose, we can also use values related to functions.

Unlike Lua closures, C closures cannot be shared. associated values. Each closure has its own independent related values. However, we can make it so that values of several functions will point to the same table, so this table becomes the environment where everything these functions can store general data.

Lua 5.2 has a feature that makes the task of separating connections easier. value between all library functions. We are open Get C libraries with `luaL_newlib` . Lua implements this function with the following macro:

```
#define luaL_newlib (L, l) \
(luaL_newlibtable (L, l), luaL_setfuncs (L, l, 0))
```

The `luaL_newlibtable` macro simply creates a table for the library. theca. (We could have used `lua_newtable` as well , but this macro uses `lua_createtable` to create a table with pre-emitted small size, optimal for the number of functions in this library lioteke.) Function `luaL_setfuncs` adds features from the list `l` to a new table at the top of the stack.

We are interested in the third parameter of the `luaL_setfuncs` function here . It tells how many related values the functions will have. libraries. Initial values for these related values must be on the stack, as with `lua_pushc-closure` . Thus, to create a library where functions will be are supposed to have a common table as the only associated value, we we can use the following code:

```
/* create a table for the library ('lib' is a list of its functions) */
luaL_newlibtable (L, lib);
/* create shared value */
lua_newtable (L);
/* add functions from the 'lib' list to the new library, so */
/* they will all have this table as their associated value */
luaL_setfuncs (L, lib, 1)
```

The last call also pops the table off the stack, leaving just a new library.

Exercises

Exercise 28.1. Write a C function `filter` . She gets-takes a list and a function as input and returns all elements from a given list for which the function returns true value:

```
t = filter ({1, 3, 20, -4, 5}, function (x) return x < 5 end)
- t = {1, 3, -4}
```

Exercise 28.2. Modify the `l_split` function (from Listing 28.2) so that it can work with lines containing well-left byte. (Apart from other changes, it must also use `memchr` instead of `strchr` .)

Exercise 28.3. Implement the `transliterate` function (exercise 21.3) on C.

Exercise 28.4. Implement a library with modified functionality it `transliterate` so that replacement of the table is not transmitted as argument, but is stored by the library itself. Your library should provide the following features:

`lib.settrans (table)` - set the replacement table

`lib.gettrans ()` - return replacement table

`lib.transliterate (s)` - translate 's' using the current table

Use a registry to store the replacement table.

Exercise 28.5. Repeat the previous exercise using specifying the associated value to store the table.

Exercise 28.6. Do you think it is a good design to store the replacement table as part of the library state, rather than passing to use it as a parameter?

Exercises

Chapter 29

Asked user-defined types in C

In the previous chapter, we saw how to extend Lua with new functions written in C. Now we will see how to extend Lua with the new types defined in C. We'll start with a little example; throughout this chapter, we will expand it with help of metamethods and other possibilities.

Our example will be pretty simple: an array of logical (boolean) values. Such a simple structure was chosen because with it does not involve any complex algorithms and we can completely concentrate on API. Nevertheless, this example is still useful. Of course in Lua we can use tables to implement arrays of boolean values. But in implementation on With we will use one bit for each element, that is we only need about 3% of the memory that would be needed for the corresponding table.

For our implementation, we need the following definitions:

```
#include <limits.h>
#define BITS_PER_WORD (CHAR_BIT * sizeof (unsigned int))
#define I_WORD (i) ((unsigned int) (i) / BITS_PER_WORD)
#define I_BIT (i) (1 << ((unsigned int) (i) % BITS_PER_WORD))
```

The `BITS_PER_WORD` constant is the number of bits in an unsigned string. scrap number. Macro `I_WORD` evaluates a word that contains a bit at the given index, and the macro `I_BIT` calculates the bit mask for corresponding bit.

We will represent our arrays with the following structures:

```
typedef struct NumArray {
    int size;
    unsigned int values [1]; /* mutable part */
} NumArray;
```

We are declaring a `values` array of size 1, since C89 is not allows you to declare arrays with size 0; in fact we will allocate the required number of elements when allocating memory for our array. The following expression calculates the total size for our bitmap with `n` elements:

```
sizeof (NumArray) + I_WORD (n - 1) * sizeof (unsigned int)
```

(We subtract one from `n`, since in our structure we already allocated space for one element.)

29.1. User data (userdata)

Our first task is to represent the structure `NumArray` in Lua. Lua provides a special base type for this: *userdata*. This type simply corresponds to an area, memory in so we can store anything without any specific operations.

The `lua_newuserdata` function allocates a block of memory for a given time measure, pushes the corresponding Lua value onto the stack and returns dedicated block address:

```
void * lua_newuserdata (lua_State * L, size_t size);
```

If for some reason you need to allocate a block of memory otherwise, you can easily create the corresponding object Lua with pointer size and remember there pointer to allocated block. We will discuss this technique in Chapter 30.

Using the `lua_newuserdata` function, the function to create new out arrays of boolean values looks like this:

```
static int newarray (lua_State * L) {  
    int i;  
    size_t nbytes;  
    NumArray * a;  
    int n = luaL_checkint (L, 1);  
    luaL_argcheck (L, n >= 1, 1, "invalid size");  
    nbytes = sizeof (NumArray) + I_WORD (n - 1) * sizeof (unsigned int);  
    a = (NumArray *) lua_newuserdata (L, nbytes);  
    a-> size = n;  
    for (i = 0; i <= I_WORD (n - 1); i++)
```

```

a-> values [i] = 0; /* initialize the array */
return 1; /* the new object is already on the stack */
}

```

Once the `newarray` function is registered in Lua, we can create new arrays using expressions like: `a = array.`

`new (1000) .`

In order to write a value to our array, we will use

Call expressions of the form: `array.set (a, index, value) .` Later we

see how you can use metatables to support more

the traditional syntax is `a [index] = value .` In both cases

the function that writes the element to the array is the same. we

we assume that, as is customary in Lua, indices start at 1:

```

static int setarray (lua_State * L) {
    NumArray * a = (NumArray *) lua_touserdata (L, 1);
    int index = luaL_checkint (L, 2) - 1;
    luaL_argcheck (L, a != NULL, 1, "'array' expected");
    luaL_argcheck (L, 0 <= index && index <a-> size, 2,
        "Index out of range");
    luaL_checkany (L, 3);
    if (lua_toboolean (L, 3))
        a-> values [I_WORD (index)] |= I_BIT (index); /* set a bit */
    else
        a-> values [I_WORD (index)] &= ~ I_BIT (index); /* remove a bit */
    return 0;
}

```

Since Lua can be used as a boolean

to call any value, then we use `luaL_checkany` to

to make sure there is some value for this parameter.

If we call `setarray` with incorrect arguments, then we will

we get the corresponding error messages:

```
array.set (0, 11, 0)
```

```
-> stdin: 1: bad argument # 1 to 'set' ('array' expected)
```

```
array.set (a, 1)
```

```
-> stdin: 1: bad argument # 3 to 'set' (value expected)
```

The following function returns the value at the given index:

```

static int getarray (lua_State * L) {
    NumArray * a = (NumArray *) lua_touserdata (L, 1);
    int index = luaL_checkint (L, 2) - 1;
    luaL_argcheck (L, a != NULL, 1, "'array' expected");
    luaL_argcheck (L, 0 <= index && index <a-> size, 2,
        "Index out of range");
    lua_pushboolean (L, a-> values [I_WORD (index)] & I_BIT (index));
    return 1;
}

```

We will define a separate function to return array size:

```
static int getsize (lua_State * L) {  
    NumArray * a = (NumArray *) lua_touserdata (L, 1);  
    luaL_argcheck (L, a != NULL, 1, "'array' expected");  
    lua_pushinteger (L, a-> size);  
    return 1;  
}
```

Finally, we need additional code to initialize our library:

```
static const struct luaL_Reg arraylib [] = {  
    {"New", newarray},  
    {"Set", setarray},  
    {"Get", getarray},  
    {"Size", getsize},  
    {NULL, NULL}
```

```
};
int luaopen_array (lua_State * L) {
  luaL_newlib (L, arraylib);
  return 1;
}
```

Again we use the `luaL_newlib` function from the helper libraries. She creates a table and populates it with name-function pairs. set by the array `arraylib` .

After opening the library, we are ready to use our new type in Lua:

```
a = array.new (1000)
print (a)
-> userdata: 0x8064d48
print (array.size (a))
-> 1000
for i = 1, 1000 do
  array.set (a, i, i% 5 == 0)
end
print (array.get (a, 10))
-> true
```

29.2. Metatables

Our current implementation has a big security problem.

thu. Let's say the user writes something like `array.set (io.`

`stdin, 1, false)` . The `io.stdin` value is an object of type `userdata` with

a pointer to `FILE` . Due to the fact that this is also a `userdata` value ,

then `array.set` will take it as a valid argument; as a result

we will most likely get a write to an arbitrary memory location (if

we are very lucky, we will only receive a message about the recording by

invalid index). This behavior is unacceptable for any

fight the Lua library. No matter how you use the library

edema, we shouldn't write something to an arbitrary memory address

(or cause the entire application to crash).

The usual way to distinguish one type of `userdata` object from another is

Go is to define a unique metatable for this type. Each

time when we create an object of type `userdata` , we expose it to

the corresponding metatable; every time we get an object

of type `userdata` , we check that it has the correct metatable.

Since Lua code cannot change the metatable for objects

type `userdata` , we are guaranteed to be fine.

We also need a place to store this metatable.

so that we can refer to it when creating new objects and checks whether an object of type `userdata` has the type we need. Like us said earlier, there are two options for storing the metatable: in the registry or as a bound value for functions in a library. In Lua, when-

It is easy to register each new C type in the registry using the type name as the index and the metatable as its corresponding values. As with any other indexes in the registry, we should choose the type name carefully to avoid possible conflicts. In our example, we will use the name

“`LuaBook.array`” .

As usual, the helper library provides us with the functionality. These new helper functions are as follows -

other functions:

```
int luaL_newmetatable (lua_State * L, const char * tname);
void luaL_getmetatable (lua_State * L, const char * tname);
void * luaL_checkudata (lua_State * L, int index,
const char * tname);
```

`LuaL_newmetatable` function creates a new table (which will be our metatable), puts it on the top of the stack and links a table with a given name in the registry. `LuaL_getmetatable` function returns the metatable associated with `tname` in the registry. At the same time nets, the `luaL_checkudata` function checks that the object at the given place on the stack is an object of type `userdata` with a metatable, with corresponding to the given name. It raises an error if the object that other metatable (or it doesn't exist) or is it not an object of type `userdata` ; otherwise, it returns the address of the object.

Now we can start our implementation. The first step would be changing the function that opens our library. New ver-

This should create a metatable for our arrays:

```
int luaopen_array (lua_State * L) {
luaL_newmetatable (L, “LuaBook.array”);
luaL_newlib (L, arraylib);
return 1;
}
```

The next step is to change the `newarray` function to be all at once so that it sets the metatable for the created mas-

Sivov:

```
static int newarray (lua_State * L) {
<as before>
```

```

luaL_getmetatable (L, "LuaBook.array");
lua_setmetatable (L, -2);
return 1; /* the new object is already on the stack */
}

```

The `lua_setmetatable` function pops a table off the stack and sets casts it as a metatable for an object on the stack by request given index. In our case, this object is the created an object of type `userdata` .

Finally, the `setarray` , `getarray` and `getsize` functions need to know did they actually get a valid array as their th first argument. To simplify this task, we will define next macro:

```

#define checkarray (L) \
(NumArray *) luaL_checkudata (L, 1, "LuaBook.array")

```

Using this macro, the new `getsize` implementation becomes very simple:

```

static int getsize (lua_State * L) {
    NumArray * a = checkarray (L);
    luaL_pushinteger (L, a-> size);
    return 1;
}

```

Since `setarray` and `getarray` share common code to check index as our second argument, we will place the common parts in the next blowing function:

```

static unsigned int * getindex (lua_State * L,
    unsigned int * mask) {
    NumArray * a = checkarray (L);
    int index = luaL_checkint (L, 2) - 1;
    luaL_argcheck (L, 0 <= index && index <a-> size, 2,
        "Index out of range");
    /* return the address of the element */
    * mask = I_BIT (index);
    return & a-> values [I_WORD (index)];
}

```

Following are the resulting implementations of `setarray` and `getarray` :

```

static int setarray (lua_State * L) {
    unsigned int mask;
    unsigned int * entry = getindex (L, & mask);
    luaL_checkany (L, 3);
    if (lua_toboolean (L, 3))
        * entry | = mask;
    else
        * entry & = ~ mask;
    return 0;
}

static int getarray (lua_State * L) {

```

```

unsigned int mask;
unsigned int * entry = getindex (L, & mask);
lua_pushboolean (L, * entry & mask);
return 1;
}

```

Now if you try to do something like `array`.

`get (io.stdin, 10)` , then you will receive a corresponding message about error:

```
error: bad argument # 1 to 'get' ('array' expected)
```

29.3. Object- oriented access

Our next step will be to convert our new type to an object so that we can work with it using object-oriented tied syntax as shown below:

```

a = array.new (1000)
print (a: size ())
-> 1000
a: set (10, true)
print (a: get (10))
-> true

```

Recall that `a: size ()` is the same as `a.size (a)` . So- we must make it so that `a.size` returns our function `getsize` . The key mechanism here is the `__index` metamethod . For tables, Lua calls this metamethod when it cannot find values for the given key. For objects of type `Lua userdata`, you- calls it every time it is accessed, since such objects have there are no keys yet.

Let's assume we ran the following code:

```

local metaarray = getmetatable (array.new (1))
metaarray .__ index = metaarray
metaarray.set = array.set
metaarray.get = array.get
metaarray.size = array.size

```

In the first line, we create an array just to get read its metatable, which we write to `metaarray` . (We are not we can set the metatable of an object of type `userdata` from Lua, but we can get it.) Then we set the `metaarray .__ index`

equal to `metatable` . Then when we execute `a.size` , Lua cannot find the key `size` in object `a` , since it is an object of type `userdata` . So Lua tries to get this value from the `__index` field metatable `a` , which is the same as the `metatable` itself . But `metatable.size` is `array.size` , so `a.size(a)` returns `array.size(a)` , which is what we wanted.

Of course we can do the same in C. We can do even better: now that arrays are objects with their own operations, we no longer need to have these operations in the `array` table . The only function from our library, which we have to pass out is the `new` function to create new arrays. All other operations will be available only as methods. The C code can register them itself.

The `getsize` , `getarray` and `setarray` operations will not change compared to with our previous approach. All that will change is how we will register them. To do this, we need to change the code that opens the library. First, we need two separate lists functions: one for regular functions and one for methods.

```
static const struct luaL_Reg arraylib_f [] = {
    {"New", newarray},
    {NULL, NULL}
};
static const struct luaL_Reg arraylib_m [] = {
    {"Set", setarray},
    {"Get", getarray},
    {"Size", getsize},
    {NULL, NULL}
};
```

The new version of the opening function `luaopen_array` should match create a metatable, assign it to its own `__index` field , register all methods and create and populate the `array` table :

```
int luaopen_array (lua_State * L) {
    luaL_newmetatable (L, "LuaBook.array");
    /* metatable.__index = metatable */
    lua_pushvalue (L, -1); /* create a copy of the metatable */
    lua_setfield (L, -2, "__index");
    luaL_setfuncs (L, arraylib_m, 0);
    luaL_newlib (L, arraylib_f);
    return 1;
}
```

Here we again use `luaL_setfuncs` to write put the functions from `arraylib_m` into the metatable located on

top of the stack. We then use `luaL_newlib` to create a new table and register functions from `arraylib_f` (actually just the `new` function).

As a finishing touch, we'll add the `__tostring` method to our type so that `print(a)` prints "array" and the size of the array in parentheses, something like "array (1000)". The corresponding function is shown below:

```
int array2string (lua_State * L) {
    NumArray * a = checkarray (L);
    lua_pushfstring (L, "array (%d)", a->size);
    return 1;
}
```

The `lua_pushfstring` call builds the string and leaves it at the top stack. We also have to add `array2string` to the list of `arraylib_m`, in order to include it in the corresponding metatable-`tsu`:

```
static const struct luaL_Reg arraylib_m [] = {
    {"__ToString", array2string},
    <other methods>
};
```

29.4. Access as usual array

An alternative to the object-oriented way of writing is the usual way of working with arrays. Instead of writing `a:get(i)`, we can just write `a[i]`. In our example, this is pretty easy. `do` as our `setarray` and `getarray` functions are already semi-give their arguments in the order in which they should be passed to use the appropriate metamethods. A quick fix would be to define these metamethods right in your Lua code:

```
local metaarray = getmetatable (array.new (1))
metaarray.__index = array.get
metaarray.__newindex = array.set
metaarray.__len = array.size
```

(We have to execute this code for our original implementation arrays, without modifications for the object-oriented syntax

sis.) That's all we need to use the standard syntax:

```
a = array.new (1000)
a [10] = true - 'setarray'
print (a [10]) - 'getarray' -> true
print (#a) - 'getsize' -> 1000
```

If we want this, then we can register these metamethods right in the C code. To do this, we must again change our initializing function:

```
static const struct luaL_Reg arraylib_f [] = {
    {"New", newarray},
    {NULL, NULL}
};

static const struct luaL_Reg arraylib_m [] = {
    {"__Newindex", setarray},
    {"__Index", getarray},
    {"__Len", getsize},
    {"__ToString", array2string},
    {NULL, NULL}
};

int luaopen_array (lua_State * L) {
    luaL_newmetatable (L, "LuaBook.array");
    luaL_setfuncs (L, arraylib_m, 0);
    luaL_newlib (L, arraylib_f);
    return 1;
}
```

In this version, we again have only one `new` function visible to everyone . All other functions are available only as metamethods for corresponding corresponding operations.

29.5. Light objects like userdata (light userdata)

The type of objects we have used so far is called *full userdata*. Lua offers another type of *userdata* object called *light*, - *light userdata*.

Such objects are just a pointer in C (i.e. value of type `void *`). It is a value, not an object; we do not create them (just like we don't create numbers). In order to place such object on the stack, we call `lua_pushlightuserdata`:

```
void lua_pushlightuserdata (lua_State * L, void * p);
```

Despite the common name, full and lightweight objects like *userdata* is actually quite different. Light objects are not buffers, just pointers. They don't have metatables. Like numbers, they not managed by the garbage collector.

Sometimes we use the lightweight option as a cheap alternative. full-fledged objects of type *userdata*. However, this is not their typical using. First, light objects do not have metatables, so we cannot find out their type. Secondly, despite its name full-fledged *userdata* objects are pretty cheap. They add They have very little overhead compared to calling `malloc`.

The real use of light objects comes from equality.

A fully-fledged *userdata* object is equal only to itself. Lay down cue object is just a pointer. And as such he is equal to any other object of type *userdata* representing that the same pointer. This way we can use the lungs objects of type *userdata* so that C objects are inside Lua.

We have already seen the typical use of lightweight objects as a key whose in the registry (see section 28.3). There the equality of light objects was extremely important. Every time we push a light object onto the stack with `lua_pushlightuserdata`, we get the same value Lua and, accordingly, the same entry in the registry.

Another typical scenario is the need to obtain a full-fledged `userdata` object at its address in C. Let's say we organize the connection between Lua and the window system. Then we can use full-fledged `userdata` objects for presentation windows. Each such object contains or the entire structure, representing a window, or just a pointer to a structure created by the system. When an event occurs inside the window (for example, clicking a button mouse), the system calls the appropriate handler that identifies quoting the window at its address. In order to pass a Lua handler, we have to find an object of type `userdata` representing the given window. In order to find it, we can use the table, where indices are light objects containing window addresses, and the values are full-fledged objects of the `userdata` type, setting the appropriate windows. If we have a window address, we push it onto the stack as a light object of type `userdata` and use its as an index on the table. (Most likely this table should have weak values. Otherwise, they will never be collected by the collector. garbage.)

Exercises

Exercise 29.1. Modify the `setarray` implementation so that it took only boolean values as input.

Exercise 29.2. We can consider a boolean array as a set of integers (indices that correspond to are true values in the array). Add built-in array functions that compute union and intersection of two arrays. These functions should receive to input two arrays and return a new array without changing input arrays.

Exercise 29.3. Change the implementation of the `__tostring` metamethod so that it shows the full contents of the array by some either way. Use buffers (see Section 28.2) to co-building the summary line.

Exercise 29.4. Based on the example with boolean arrays, re-
Alize the C library for working with arrays of integers
sat down.

Chapter 30

Resource management

In our implementation of boolean arrays from the previous chapter, we do not worry about resource management. These arrays use only memory. Each `userdata` object that represents a mass, has its own block of memory, which is managed by Lua. When the array becomes garbage (that is, no one stores references to it), Lua will collect it in time and free the occupied memory.

However, life is not always so easy. Sometimes the object needs other resources other than memory such as file descriptors, pointers to windows, etc. (often these resources are also memory, but it is controlled by other part of the system). In such cases, when the object becomes garbage, it is necessary to somehow release these resources.

As we saw in section 17.6, Lua provides a finalizer `__gc` metamethod. To show the use of this metamethod in C, we implement two libraries in C, providing access to external resources. The first example is another realization of the function for crawling the contents of the directory. The second (and more complex) example is using the *Expat* library to parse XML files.

30.1. Directory iterator

In section 27.1, we implemented a `dir` function that returned a table with all files from a given directory. Our new realization will return an iterator that returns a new file each time call. Using this implementation, we can iterate over the content of the directory using a loop as shown below:

```
for fname in dir.open(".") do
  print(fname)
end
```

In order to iterate over the contents of the directory in C, we need

structure `DIR` . These structures are created by calling `opendir` and are destroyed by calling `closedir` . Our previous the implementation of the `dir` function kept this structure as local to belt and released when getting the name of the last file. Our new implementation cannot store DIRs in local variables-`noah`, since this structure will be needed for a whole series of calls. Moreover, we cannot destroy it upon receiving the name nor the last file, as the program may prematurely get out of the loop, in which case we'll never get to the last file. Therefore, in order to ensure that this structure will be always freed, we need to store its address in an object like `userdata` and use the `__gc` metamethod to free this structures.

Despite its central role in our implementation, this an object representing a directory does not have to be den from Lua. The `dir` function returns an iterating function; this is whatever Lua sees. The directory can be the associated value of this iterating function. In this case, the iterating function will be have direct access to this structure, but Lua code to it does not have access (and he does not need it).

In total, we need three functions in C. First, we need `dir.open`. `open` is a function that Lua calls to create iterators; it should create a `DIR` structure and an iteration function closure with by this structure (as a bound value). Second, we need iterating function. Third, we need the `__gc` metamethod , which releases the created `DIR` structure . As usual, we also you will need a function for initial setup, such as creating and initializing the metatable for the directory.

Let's start our code with the `dir.open` function shown in figure 30.1. The important point is that this function should create a `userdata` object before opening the directory. Otherwise, if he will open the directory first and then calling `lua_newuserdata` will result in error when working with memory, then a memory leak occurs, since nobody will release the created structure. With the right order the `DIR` structure , as soon as it is created, is immediately linked to the `userdata` ; whatever happens after, the `__gc` metamethod over time it will release this structure.

The next function is `dir.iter` (Listing 30.2), the iterator itself. Its code is pretty simple. It gets the address of the `DIR` structure from the

associated

value with it and calls `readdir` to get the next values.

The `dir_gc` function (also in Listing 30.2) is the `__gc` metamethod. It frees the generated DIR structure, but you need to be careful: since the `userdata` object is created before opening the directory, but even if `opendir` returns an error, then the `userdata` object is all will be created equal. Therefore, we need to check what is, what close.

The last function in Listing 30.2 is `luaopen_dir`, a function that which opens our library.

There is one subtlety in the complete example. At first it may seem that the function `dir_gc` needs to know if its argument is valid is a directory. Otherwise, the user can call its with a different type of `userdata` object (like file), which will result to a serious error. However, a Lua program has no way to Refer to this function: it is stored as a metatable of catalogs, which, in turn, are stored as associated with the iterating value function. Lua programs can't access to objects of this type.

Listing 30.1. Dir.open function

```
#include <dirent.h>
#include <errno.h>
#include <string.h>
#include "lua.h"
#include "lauxlib.h"
/* declare an iteration function */
static int dir_iter (lua_State * L);
static int l_dir (lua_State * L) {
    const char * path = luaL_checkstring (L, 1);
    /* create a userdata object to store the address of the DIR */
    DIR ** d = (DIR **) lua_newuserdata (L, sizeof (DIR *));
    /* set metatable */
    luaL_getmetatable (L, "LuaBook.dir");
    lua_setmetatable (L, -2);
    /* trying to open the directory */
    * d = opendir (path);
    if (* d == NULL) /* error opening directory? */
        luaL_error (L, "cannot open% s:% s", path, strerror (errno));
    /* create and return an iteration function;
    its associated value is the userdata object,
    already on the stack */
}
```

```

lua_pushcclosure (L, dir_iter, 1);
return 1;
}

```

Listing 30.2. Other functions in the dir library

```

static int dir_iter (lua_State * L) {
DIR * d = * (DIR **) lua_touserdata (L, lua_upvalueindex (1));
struct dirent * entry;
if ((entry = readdir (d)) != NULL) {
lua_pushstring (L, entry-> d_name);
return 1;
}
else return 0; /* no more values */
}

static int dir_gc (lua_State * L) {
DIR * d = * (DIR **) lua_touserdata (L, 1);
if (d) closedir (d);
return 0;
}

static const struct luaL_Reg dirlib [] = {
{"Open", l_dir},
{NULL, NULL}
};

int luaopen_dir (lua_State * L) {
luaL_newmetatable (L, "LuaBook.dir");
/* set field __gc */
lua_pushcfunction (L, dir_gc);
lua_setfield (L, -2, "__gc");
/* create library */
luaL_newlib (L, dirlib);
return 1;
}

```

30.2. XML parser

We now turn to a simplified library implementation for links between Lua and the Expat library, which we will call `l_xp`. Expat is an open source XML 1.0 parser written in C. It implements SAX, that is, a *simple API for XML* (simple API for XML). SAX is event-driven API. This means that the SAX parser is clean the XML document melts and tells the application as it reads that it finds using user - defined *functions-process-*

chikov (callback). For example, if we want Expat to parse line like “<tag cap =” 5 ”> hi </tag>” , then it will create three events: start event when it reads the line “<tag cap =” 5 ”>” ; event text when it reads "hi" , and the end of element event when it reads “</tag>” melts . Each of these events triggers a corresponding handler in the application.

We won't cover the entire Expat library here. We focus we focus only on those parts that show newer methods of interaction Modeling with Lua. While Expat handles over a dozen different events, we will consider only those three events that we saw in previous example (start of element, end of element and text) ¹ .

The part of the Expat API that we need is pretty small. In-First, we need functions to create and destroy the parser:

```
XML_Parser XML_ParserCreate (const char * encoding);  
void XML_ParserFree (XML_Parser p);
```

The encoding argument is optional, we will pass it instead to be NULL.

Once we have a parser, we must register our handler functions:

```
void XML_SetElementHandler (XML_Parser p,  
XML_StartElementHandler start,  
XML_EndElementHandler end);  
void XML_SetCharacterDataHandler (XML_Parser p,  
XML_CharacterDataHandler hndl);
```

The first function sets handlers for start and end events element. The second function sets the handler for the text.

All handlers receive a non-

which is a pointer. The element start handler is also named tag and its attributes:

```
typedef void (* XML_StartElementHandler) (void * uData,  
const char * name,  
const char ** atts);
```

Attributes are passed as a NULL terminated array of strings , where each pair of consecutive lines contains an attribute and its value nie. The end-of-element handler only receives one extra nth element - tag name:

```
typedef void (* XML_EndElementHandler) (void * uData,  
const char * name);
```

Finally, the text processor receives as an additional parameter is the text itself. The line of text is not null terminated, and the length is explicitly passed for it:

```
typedef void (* XML_CharacterDataHandler) (void * uData,
const char * s,
int len);
```

In order to pass text to Expat for parsing, we use the following function:

```
int XML_Parse (XML_Parser p, const char * s, int len, int isLast);
```

Expat receives the document to be parsed hourly cham through successive calls to XML_Parse . The last argument such a call to isLast tells Expat whether the passed chunk was last in the document. Note that each snippet the text does not have to be terminated with a null byte, we explicitly we pass its length. XML_Parse function returns zero in case errors. (Expat also provides functions for getting information error messages, but for simplicity, we will not consider them here. vat.)

The last function we need from Expat is the function ttion that allows you to specify the pointer that will be passed handlers:

```
void XML_SetUserData (XML_Parser p, void * uData);
```

Now let's see how we can use this bib-

library in Lua. The first approach is the simplest: let's just yes-

Let's access all these functions from Lua. More successful will be

adapt this functionality for Lua. For example, since

Lua is an atypical language (more precisely, a language without strong typing), then we

no need for different functions for each type of handler. More

Moreover, we can avoid registering handlers altogether. Inmes-

then we will create a parser, pass it a table of handlers,

each with a matching key. For example, if we want to print

structure of the document, then we can use the following table

zu handlers:

```
local count = 0
callbacks = {
  StartElement = function (parser, tagname)
    io.write ("+", string.rep (" ", count), tagname, "\n")
    count = count + 1
  end,
  EndElement = function (parser, tagname)
    count = count - 1
    io.write ("-", string.rep (" ", count), tagname, "\n")
  end,
}
```

If we give the input the string “<to> <yes /> </to>” , then these bots will generate the following output:

```
+ to
+ yes
- yes
- to
```

With such an API, we do not need functions to work with the handler-mi. We work with them directly in the handler table.

Thus, the entire API will consist of only three functions: one for creating parsers, one for processing a piece of text and one to destroy the parser. In fact, we are implementing two last function as parser methods. As a result, we come to the following - typical usage of our API:

```
local lxp = require "lxp"
p = lxp.new (callbacks)
- create a new parser
for l in io.lines () do
- process input lines
assert (p: parse (l))
- parse the string
assert (p: parse ("\ n"))
- add '\ n'
end
assert (p: parse ())
- complete the document
p: close ()
```

Let's now turn to the implementation. Our first decision

This will be how we will represent our parser in Lua. Quite it is natural to use an object of type `userdata` for this , but what do we need to put inside it? At least we need the parser itself and the handler table. We cannot remember the table zu inside an object of type `userdata` (or inside a C structure), but Lua allows each object of type `userdata` to have a *custom* *a* user value, which can be any Lua 2 table . we should also remember the Lua state into the parser object as all that the Expat handler receives is the parser itself, and in order to call Lua, we need this state. Therefore, we will use define the following parser:

```
#include <stdlib.h>
#include "expat.h"
#include "lua.h"
#include "luaXlib.h"
typedef struct lxp_userdata {
XML_Parser parser; /* corresponding Expat parser */
```

```
lua_State * L;
} lxp_userdata;
2
```

In Lua 5.1, the environment of the userdata object acts as a user value.
niya.

Our next step is to create a function that creates no parsers, `lxp_make_parser`. Its code is shown in Listing 30.3. This the function consists of four important steps:

- Its first step follows a standard pattern: first create a userdata object; then it is initialized appropriately values, and finally, a metatable is assigned to it. faces. The reason for this initialization is as follows: if in any error occurs during initialization, you must dimo that the finalizer (metamethod `__gc`) finds our data holistic.
- At step 2, the function creates an Expat parser, stores it in userdata object and checks for errors.
- Step 3 checks that the first argument of the function is valid a table (table of handlers) is stored, and assigns it as a custom value for the userdata object.
- The last step initializes the Expat parser. Our userdata is the object is given as a pointer to be passed to be included in handlers, handler functions are also set. Note that these handlers are the same for everyone parsers; after all, in C one cannot dynamically construct function. Instead, fixed functions using a table of handlers, decide which Lua functions to follow call.

Listing 30.3. Function for creating XML parsers

```
/* descriptions of handler functions */
static void f_StartElement (void * ud,
const char * name,
const char ** atts);
static void f_CharData (void * ud, const char * s,
int len);
static void f_EndElement (void * ud, const char * name);
static int lxp_make_parser (lua_State * L) {
XML_Parser p;
/* (1) create parser object */
lxp_userdata * xpu = (lxp_userdata *)
lua_newuserdata (L,
sizeof (lxp_userdata));
/* initialize it in case of error */
```



```

xpu-> parser = NULL;
/* set a metatable for it */
luaL_getmetatable (L, "Expat");
lua_setmetatable (L, -2);
/* (2) create Expat parser */
p = xpu-> parser = XML_ParserCreate (NULL);
if (! p)
luaL_error (L, "XML_ParserCreate failed");
/* (3) check and save the handler table */
luaL_checktype (L, 1, LUA_TTABLE);
lua_pushvalue (L, 1); /* push the table onto the stack */
lua_setuservalue (L, -2);
/* (4) configure the Expat parser */
XML_SetUserData (p, xpu);
XML_SetElementHandler (p, f_StartElement,
f_EndElement);
XML_SetCharacterDataHandler (p, f_CharData);
return 1;
}

```

Listing 30.4. Function for parsing a piece of text

```

static int lxp_parse (lua_State * L) {
int status;
size_t len;
const char * s;
lxp_userdata * xpu;
/* get and check the first argument */
xpu = (lxp_userdata *) luaL_checkudata (L, 1, "Expat");
/* check that it is not closed */
luaL_argcheck (L, xpu-> parser != NULL, 1, "parser is closed");
/* get second argument (string) */
s = luaL_optlstring (L, 2, NULL, & len);
/* put the handler table at index 3 on the stack */
lua_settop (L, 2);
lua_getuservalue (L, 1);
xpu-> L = L; /* set Lua state */
/* call Expat to parse the string */
status = XML_Parse (xpu-> parser, s, (int) len, s == NULL);
/* return error code */
lua_pushboolean (L, status);
return 1;
}

```

The next step is the method for parsing the text `lxp_parse` (Listing 30.4), which parses a chunk of XML data. He semi-takes two arguments: a parser (*self* in the method) and an optional fragment XML. When called with no data, it tells Expat that more

there are no parts.

When `lxp_parse` calls `XML_Parse` it will call handlers for those elements that it finds in the transferred text fragment that. These handlers will need access to the handler table, so `lxp_parse` pushes this table onto the stack at index 3 (immediately after parameters). There is one caveat to the `XML_Parse` call : remember, that the last argument to this function tells Expat whether the last piece of text transferred. When we call `parse` with no arguments, `s` will be `NULL` , and this last argument will be met the true value.

Now let's turn our attention to the handler functions

`f_StartElement` , `f_EndElement` and `f_CharData` . All these functions are have the same structure: each of them checks if there is handler table for this event, and if such a handler is present, then prepares the arguments and then you-calls this handler.

Let's take a look at the `f_CharData` handler in listing ge 30.5. Its code is pretty simple. The handler receives the structure `lxp_userdata` as its first argument, since we called `XML_SetUserData` when we created our parser. After receiving co-Lua state handler can refer to handler table on the stack at index 3, given by `lxp_parse` , and the parser itself by index 1. Then it calls the corresponding Lua handler (when present) with two arguments: parser and character data nym (string).

Listing 30.5. Character data processor

```
static void f_CharData (void * ud, const char * s, int len) {
    lxp_userdata * xpu = (lxp_userdata *) ud;
    lua_State * L = xpu->L;
    /* get handler */
    lua_getfield (L, 3, "CharacterData");
    if (lua_isnil (L, -1)) { /* no handler? */
        lua_pop (L, 1);
        return;
    }
    lua_pushvalue (L, 1); /* push the parser ('self') onto the stack */
    lua_pushlstring (L, s, len); /* push a line onto the stack */
    lua_call (L, 2, 0); /* call handler */
}
```

The `f_EndElement` handler is quite similar to `f_CharData` ; education See Listing 30.6. It also calls the appropriate Lua bot with two arguments - a parser and a tag name (again

a string, this time terminated with a null byte).

Listing 30.6. End of element handler

```
static void f_EndElement (void * ud, const char * name) {
    lxp_userdata * xpu = (lxp_userdata *) ud;
    lua_State * L = xpu->L;
    lua_getfield (L, 3, "EndElement");
    if (lua_isnil (L, -1)) {/ * no handler? */
    lua_pop (L, 1);
    return;
    }
    lua_pushvalue (L, 1); / * push the parser ('self') onto the stack */
    lua_pushstring (L, name); / * push the tag onto the stack */
    lua_call (L, 2, 0); / * call handler */
}
```

Listing 30.7 shows the final handler, `f_StartElement`.

It calls Lua with three arguments: parser, tag name, and list of attributes. This handler is slightly more complex than the others, because how long it is necessary to translate the list of attributes into Lua. He used suggests a completely natural translation, building a table that compares Specifies attribute names and their values. For example, for a short tag, shown below

```
<to method = "post" priority = "high">
```

the following attribute table is generated:

```
{method = "post", priority = "high"}
```

Listing 30.7. Element start handler

```
static void f_StartElement (void * ud,
    const char * name,
    const char ** atts) {
    lxp_userdata * xpu = (lxp_userdata *) ud;
    lua_State * L = xpu->L;
    lua_getfield (L, 3, "StartElement");
    if (lua_isnil (L, -1)) {/ * no handler? */
    lua_pop (L, 1);
    return;
    }
    lua_pushvalue (L, 1); / * push the parser ('self') onto the stack */
    lua_pushstring (L, name); / * push the tag name onto the stack */
    / * create and populate the attribute table */
    lua_newtable (L);
    for (; * atts; atts += 2) {
        lua_pushstring (L, * (atts + 1));
```

```

lua_setfield (L, -2, * atts); /* table [* atts] = * (atts + 1) * /
}
lua_call (L, 3, 0); /* call handler */
}

```

The final method for parsers is `close`, shown in listing 30.8. When we close the parser, we must release all of its resources, namely the Expat structure. Remember that due to errors in the parser may not have this structure. Pay attention—how we keep the parser consistent across as we close it, so there will be no problem, if we try to close it again or when the garbage collector finalizes it. This ensures that each parser over time will free its resources even if the programmer hasn't closed it.

Listing 30.8. Method for closing an XML parser

```

static int lxp_close (lua_State * L) {
    lxp_userdata * xpu =
    (lxp_userdata *) luaL_checkudata (L, 1, "Expat");
    /* free the Expat parser (if any) */
    if (xpu-> parser)
        XML_ParserFree (xpu-> parser);
    xpu-> parser = NULL; /* if we close it again */
    return 0;
}

```

The final step is shown in Listing 30.9: it shows the functionality `luaopen_lxp`, which opens up the library, combining together all previously discussed functions. We use the same scheme here as we used for object-oriented boolean array in section 29.3: we create a metatable, set it the `__index` field on it and put all methods inside it. For this—Then we need a list with all the parser methods (`lxp_meths`). We also need a list of functions of this library (`lxp_funcs`). how and is accepted in object-oriented libraries, this list is composed of just one function that creates new parsers.

Listing 30.9. Initializing code for lxp library

```

static const struct luaL_Reg lxp_meths [] = {
    {"Parse", lxp_parse},
    {"Close", lxp_close},
    {"__Gc", lxp_close},
}

```

```

{NULL, NULL}
};
static const struct luaL_Reg lxp_funcs [] = {
{"New", lxp_make_parser},
{NULL, NULL}
};
int luaopen_lxp (lua_State * L) {
/* create metatable */
luaL_newmetatable (L, "Expat");
/* metatable.__index = metatable */
lua_pushvalue (L, -1);
lua_setfield (L, -2, "__index");
/* register methods */
luaL_setfuncs (L, lxp_meths, 0);
/* register functions (lxp.new only) */
luaL_newlib (L, lxp_funcs);
return 1;
}

```

Exercises

Exercise 30.1. Modify the `dir_iter` function so that it would close the `DIR` structure when it reaches the end directory. With this change, the program does not need to wait garbage collection to release a larger resource not needed.

(When you close the directory, you must set the address, written in the `userdata` object to `NULL` to inform the parser that the directory is already closed. Also the `dir_iter` function before using the directory should check that it is not closed.)

Exercise 30.2. In the example with the `lxp` library, the handler starts with an element gets a table with the attributes of the element. In this table the order in which the elements were given internally is already lost. How can you convey this information to the handler?

Exercise 30.3. In the example with the `lxp` library, we used whether a custom value for linking the table handlers with the corresponding `userdata` object, pre-installing the parser. This choice created a small problem,

since what the C handlers get is the structure `lxp_userdata`, and this structure does not provide direct access to `stupa` to this table. We solved this problem by co-storing the handler table at a given place on the stack in parsing time of each fragment.

Another solution could be to link the handler table with a `userdata` object using links (section 28.3): we create a link to the handler table and remember this link (integer) in the `lxp_userdata` structure. Implement this option. Don't forget to release the link on close parser.

CHAPTER 31

Threads and states

Lua does not support true multi-threading, that is, it displaces threads sharing shared memory. There are two reasons for this.

The first reason is that such support was not provided.

is ANSI C and therefore there is no portable way to implement this support in Lua. The second and more serious reason is that we don't think multi-threading is a good idea for Lua.

Multi-threading was developed for low-level production programming. Synchronization mechanisms like semaphores and monitors have been proposed for operating systems (and experienced programmers), not for applications. It is extremely difficult to find and fix bugs related to multi-threading and some of them can lead to security holes. Also multi-threading can lead to serious performance problems due to the need for synchronization at a number of critical points in the program, such as memory allocation.

Multi-threading problems arise from the combination of displacement threads and shared memory, so we can avoid

them, either without using displacing threads, or without using storage space. Lua offers support for both.

Lua threads (also known as coroutines) are not preemptive and therefore avoid the problems associated with unpredictable by switching the threads. Lua states have no shared memory, so form a good basis for parallel computing. In this chapter we will consider both of these options.

31.1. Numerous threads

A *thread* is the essence of a coroutine in Lua. We are considering a coroutine as a thread and a user-friendly interface, or we can consider a thread as a coroutine with a low-level API.

From a C perspective, it can be helpful to think of a thread as a ke - what the thread really is from the point of view of implementation.

Each stack stores information about the current calls to the thread, as well as the same parameters and local variables of each call. Others in words, the stack contains all the information a thread needs to continuation of its implementation. Therefore, many threads mean many independent stacks.

When we call most of the functions from the Lua-C API, then these functions work with a specific stack. For example `lua_pushnumber` must push a number onto a specific stack; also `lua_pcall` needs a stack to call. How does Lua know which stack to follow use? The secret is that the `lua_State` type , the first the argument of all these functions, is not only a state Lua, but also a thread inside that state. (Many believe that this type should be called `lua_Thread` .)

When you create a new state, Lua automatically creates thread inside this state and returns `lua_State` representing that thread. This *main thread* is destroyed along with the state, when you call `lua_close` .

You can create other threads within the state with

`lua_newthread` :

`lua_State * lua_newthread (lua_State * L);`

This function returns a pointer to `lua_State` representing the new thread, and also pushes the new thread onto the stack as a value of type `thread` . For example, after executing the operator

```
L1 = lua_newthread (L);
```

we will have two strands, L1 and L, both referring internally to the same the same state of Lua. Each thread has its own stack. New thread L1 starts from an empty stack; old thread L has new thread on top of stack:

```
printf ("%d \n", lua_gettop (L1)); -> 0  
printf ("%s \n", luaL_typename (L, -1)); -> thread
```

With the exception of the main yarn, the yarns can be collected by the picker. a lump of garbage like any other Lua object. When you create a new thread, then it is pushed onto the stack, which ensures that this thread is not rubbish. You should never use a thread that

not tied to state. (The main thread is tied from the beginning, so you don't have to worry about it.) Any Lua API call can destroy an unattached thread, even a call using this self-wash the thread. For example, let's look at the following snippet:

```
lua_State * L1 = lua_newthread (L);  
lua_pop (L, 1); /* L1 is now garbage for Lua */  
lua_pushstring (L1, "hello");
```

Calling `lua_pushstring` can invoke the garbage collector and collect L1 (resulting in an application error) even though L1 is still in use enjoys. To avoid this, always keep links to the threads you are using, for example on the anchored thread stack or in the registry.

As soon as we have a new thread, we can immediately start use it like the main thread. We can put values onto its stack and pop values from its stack, we can use it to function calls, etc. For example, the following code calls `f(5)` on the new thread and then puts the result on the old thread:

```
lua_getglobal (L1, "f"); /* consider that there is a global 'f' */  
lua_pushinteger (L1, 5);  
lua_call (L1, 1, 1);  
lua_xmove (L1, L, 1);
```

The `lua_xmove` function moves a Lua value between two stacks-the same state. Calling type `lua_xmove (F, T, n)` from below meth `n` elements from the stack `F` and places them on a stack `T`.

However, we do not need a new thread for these purposes; we can light to use the main thread. The main purpose of using non-how many threads is the implementation of coroutines so that we can suspend their execution and resume it again. For this

we need the `lua_resume` function :

```
int lua_resume (lua_State * L, lua_State * from, int narg);
```


To start the execution of the coroutine, we use `lua_resume` just like we use `lua_pcall` : we put the function on stack, push its arguments onto the stack and call `lua_resume` , re-giving `nargs` the number of arguments. (The `from` parameter is the thread that makes the call.) This is very similar to `lua_pcall` , however there are three differences. First, `lua_resume` does not contain a parameter for the number desired results; it always returns all the results you-called function. Secondly, it has no parameter for the handler errors; the error does not unwind the stack, so you can later explore. Third, if a function suspends its execution, (with `yield`), then `lua_resume` returns a special code `LUA_YIELD` and leaves the thread in such a state that we can resume later.

When `lua_resume` returns `LUA_YIELD` , the visible portion of the thread contains only the values passed to `yield` . `lua_gettop` will return a number of these values. In order to transfer these values to another thread, we can use `lua_xmove` .

To continue execution of a suspended thread, we call `lua_resume` again . In this case, Lua assumes that everything knows the values on the stack are the values returned by `yield` . For example, if you don't touch the stack thread in between returns eat from the previous `lua_resume` and the next `lua_resume` , then `yield` will return exactly the values with which it was called.

We usually run a coroutine with a Lua function as body. This Lua function can call other functions, and any from these functions can call `yield` , ending the call to `lua_resume` .

For example, consider the following definitions:

```
function foo (x) coroutine.yield (10, x) end
function foo1 (x) foo (x + 1); return 3 end
```

We will now execute the following C code:

```
lua_State * L1 = lua_newthread (L);
lua_getglobal (L1, "foo1");
lua_pushinteger (L1, 20);
lua_resume (L1, L, 1);
```

Call `lua_resume` return `LUA_YIELD` , to inform, that the thread is suspended. At this point, the `L1` stack contains the values passed to `yield` :

```
printf ("%d\n", lua_gettop (L1)); -> 2
printf ("%d\n", lua_tointeger (L1, 1)); -> 10
printf ("%d\n", lua_tointeger (L1, 2)); -> 21
```

When we call `lua_resume` again , the thread will continue executing

from where it left off (call to `yield`). From there `foo` will return the control
Lenie `foo1`, and she, in turn, returns management `lua_resume`:

```
lua_resume (L1, L, 0);  
printf ("%d \n", lua_gettop (L1)); -> 1  
printf ("%d \n", lua_tointeger (L1, 1)); -> 3
```

This second call to `lua_resume` will return `LUA_OK` which means ok
minimal return.

Coroutines can also call C functions that can
call other Lua functions. We have already discussed how to use
continue to allow these functions in Lua

call `yield` (section 27.2). A C function itself can also call
`yield`. In this case, you must provide a continuation function,
which will be called when execution continues. On the C following-
This function plays the role of `yield`:

```
int lua_yieldk (lua_State * L, int nresults, int ctx,  
lua_CFunction k);
```

We must always use this function in a return statement
the one as shown below:

```
static int myCfunction (lua_State * L) {  
...  
return lua_yieldk (L, nresults, ctx, k);  
}
```

This call immediately suspends the currently running
program. The `nresults` parameter is the number of values on the stack,
which should be returned to the corresponding `lua_resume`; `ctx` is
the context to be passed on to the continuation and `k` is a function-
continuation. When the coroutine continues execution, the control
the extension goes to the continuation function `k`. After calling `lua_yieldk`
the `myCfunction` cannot do anything else; she must
delegate all further work to your continuation.

Let's look at a hypothetical example. Let us want
write a function that reads some data by calling
`yield` when the data is not ready. We can write this function in
With the following:

```
int prim_read (lua_State * L) {  
if (nothing_to_read ())
```

```

return lua_yieldk (L, 0, 0, & prim_read);
lua_pushstring (L, read_some_data ());
return 1;
}

```

If the function has any data, then it reads and returns them. Otherwise, it calls `lua_yieldk`. When the thread is continues to execute, it will call the continuation function. In that in the example, the continuation function is `prim_read` itself, so the thread will call it over and over again to read data. (This template, when the calling `lua_yieldk` function is itself position is not uncommon.)

If a C function has nothing to do after calling `lua_yieldk`, then it can call `lua_yieldk` without a continuation function or using `lua_yield` macro :

```

return lua_yield (L, nres);

```

After this call, when the thread continues its execution, raids the function that called `myCfunction`.

31.2. Lua states

Each call to `luaL_newstate` (Or `lua_newstate`, as we will see in chapter 32) creates a new Lua state. Different states of Lua do not depend on each other in any way. And they don't have any common data. This means that no matter what happens in one state in Lua, it cannot "harm" another state in any way. So- this means that different states of Lua cannot be communicate; for this we have to use a special code on C. For example, if we have two states `L1` and `L2`, then the following command - push to the top of the `L2` stack the value from the top of the stack in `L1`:

```

lua_pushstring (L2, lua_tostring (L1, -1));

```

Since the data must pass through C, different co-states in Lua can exchange only those types

data that is representable in C, such as strings and numbers. Other types,

for example, tables to be transferred must be serialized.

In systems that offer multi-threading, an interesting architectural solution would be to create on a separate state Lua for each thread. As a result, we get threads that behave like processes in UNIX, that is, we have parallelism without shared (shared) memory. In this section, we will build a prototype applications using this approach. For this implementation I will use POSIX threads (`pthread`). Since I am using only the most basic features, it will be easy to transfer this code to other multi-threaded systems.

The system we want to build is very simple. Its purpose is to show the use of multiple Lua states in a context of multi-threading. After it is ready, you yourself can add additional features to it. We will call our `lproc` library . It offers just four functions:

- `lproc.start (chunk)` starts a new process to execute the given block of code (chunk). The library implements the process in Lua as a C thread and associated Lua state.
- `lproc.send (channel, val1, val2, ...)` sends the given values (which must be strings) to a given channel, identified by its name (string).
- `lproc.receive (channel)` receives values from the given channel.
- `lproc.exit ()` ends the process. This function is only needed in the main process. If this process ends without calling `lproc.exit` , then the entire program terminates without waiting for other processes to finish.

The library identifies channels using strings and using calls them to match the sender of the recipient. Operation can send any number of string values that are returned by the corresponding receive operation. All interaction is synchronous: the process sending the message feeds into pipe, waits until there is a process reading from this pipe while the process reading from the pipe is also waiting until there is a process sending to it.

The `lproc` library , like its interface, is pretty simple. It is using two doubly linked ring lists, one for processes waiting to send a message, and another waiting to send processes waiting to receive a message. Also used one

a mutex to control access to both of these lists. Everyone has it process has its own *condition variable* . When the process wants to send a message to the channel, it looks in the waiting list the process that is waiting for this particular channel. If he finds such a process, then it removes it from the waiting list, transfers values from itself to the found process and signals the mental processes. Otherwise it inserts itself into the list sending and waiting for its conditional variable. Getting co-communication behaves similarly.

The main element of the implementation is a structure that represents the following process:

```
#include <pthread.h>
#include "lua.h"
typedef struct Proc {
    lua_State * L;
    pthread_t thread;
    pthread_cond_t cond;
    const char * channel;
    struct Proc * previous, * next;
} Proc;
```

The first two fields represent the Lua state used by process, and the corresponding C thread performing this process. Other fields are used only when the process has to wait corresponding `send / receive` . The third field `cond` is a conditional the variable that the thread uses to wait; four- The `th` field is the channel the process is waiting for; and the last two fields, `previous` and `next` are used to connect a process in a double-link list.

The following code declares two lists of waiting processes and their associated mutex:

```
static Proc * waitsend = NULL;
static Proc * waitreceive = NULL;
static pthread_mutex_t kernel_access = PTHREAD_MUTEX_INITIALIZER;
```

Each process needs a corresponding `Proc` structure , and it needs access to it whenever his body calls `send` or `receive`. The only parameter these functions receive is the corresponding Lua state, so each process has to remember its `Proc` structure inside its Lua state.

In our implementation, each Lua state stores a corresponding

the Proc structure as an object of type `userdata` associated with the key `"_SELF"`. The `getself` helper function returns state

Proc corresponding to the given state:

```
static Proc * getself (lua_State * L) {
    Proc * p;
    lua_getfield (L, LUA_REGISTRYINDEX, "_SELF");
    p = (Proc *) lua_touserdata (L, -1);
    lua_pop (L, 1);
    return p;
}
```

The next function, `movevalues`, transfers values from send-process to receiving:

```
static void movevalues (lua_State * send, lua_State * rec) {
    int n = lua_gettop (send);
    int i;
    for (i = 2; i <= n; i++) /* transfer values to the receiver */
        lua_pushstring (rec, lua_tostring (send, i));
}
```

It transfers to the receiver all values from the sender's stack, except for the first value, which is a channel.

Listing 31.1 defines a `searchmatch` function that bypasses pending list looking for a process waiting on the given channel

Lua states

la. If the function finds such a channel, then it removes it from the list and returns it, otherwise it returns `NULL`.

Listing 31.1. Function to find a process waiting for a given channel

```
static Proc * searchmatch (const char * channel, Proc ** list) {
    Proc * node = * list;
    if (node == NULL) return NULL; /* the list is empty? */
    do {
        if (strcmp (channel, node-> channel) == 0) { /* found? */
            /* remove a node from the list */
            if (* list == node) /* is this the first item in the list? */
                * list = (node-> next == node)? NULL: node-> next;
            node-> previous-> next = node-> next;
            node-> next-> previous = node-> previous;
            return node;
        }
        node = node-> next;
    } while (node != * list);
    return NULL; /* not found */
}
```

The last helper function defined in listing 31.2, called when the process cannot find the friend it needs

gogo process. In this case, the process connects itself to the end of the the corresponding list and waits until another process finds and will not wake him up. (The loop around `pthread_cond_wait` protects against random wakes that are possible on POSIX threads.) When process wakes up another process, then it sets the `channel` field to the awakened process to `NULL`. So if `p->channel` is not equal `NULL`, it means that no other process woke up this process, so you have to wait further.

Listing 31.2. Function for adding a process to the waiting list

```
static void waitonlist (lua_State * L, const char * channel,
Proc ** list) {
    Proc * p = getself (L);
    /* connect yourself to the end of the list */
    if (* list == NULL) { /* is the list empty? */
        * list = p;
        p->previous = p->next = p;
    }
    else {
        p->previous = (* list) -> previous;
        p->next = * list;
        p->previous->next = p->next->previous = p;
    }
    p->channel = channel;
    do { /* expects a conditional variable */
        pthread_cond_wait (& p->cond, & kernel_access);
    } while (p->channel);
}
```

Now with these helper functions we can write `send` and `receieve` (Listing 31.3). `Send` function starts with checking the channel. Then it closes the mutex and looks for the matching the recipient. If she finds him, then she carries her values to this recipient, marks the recipient as ready to receive fullness and wakes him up. Otherwise, she waits herself. Upon completion Upon doing this, it opens a mutex and returns to Lua. The `receieve` function is similar, but it must return everything received. values.

Listing 31.3. Functions for sending and receiving messages

```
static int ll_send (lua_State * L) {
    Proc * p;
    const char * channel = luaL_checkstring (L, 1);
    pthread_mutex_lock (& kernel_access);
    p = searchmatch (channel, & waitreceieve);
    if (p) { /* found a matching recipient? */
        movevalues (L, p->L); /* transfer values to recipient */
    }
}
```

```

p-> channel = NULL; /* mark recipient */
pthread_cond_signal (& p-> cond); /* wake him up */
}
else
waitonlist (L, channel, & waitsend);
pthread_mutex_unlock (& kernel_access);
return 0;
}
static int ll_receive (lua_State * L) {
Proc * p;
const char * channel = luaL_checkstring (L, 1);
lua_settop (L, 1);
pthread_mutex_lock (& kernel_access);
p = searchmatch (channel, & waitsend);
if (p) { /* found a matching recipient? */
movevalues (p-> L, L); /* transfer values to recipient */
p-> channel = NULL; /* mark recipient */
pthread_cond_signal (& p-> cond); /* wake him up */
}
else
waitonlist (L, channel, & waitreceive);
pthread_mutex_unlock (& kernel_access);
/* return all values from the stack except the pipe */

```

Lua states

```

return lua_gettop (L) - 1;
}

```

Now let's see how to create new processes. New

The process needs a new POSIX thread, and the new thread needs a body to execution. We will define this body later; here is the totype:

```

static void * ll_thread (void * arg);

```

Listing 31.4. Function for creating a new process

```

static int ll_start (lua_State * L) {
pthread_t thread;
const char * chunk = luaL_checkstring (L, 1);
lua_State * L1 = luaL_newstate ();
if (L1 == NULL)
luaL_error (L, "unable to create new state");
if (luaL_loadstring (L1, chunk) != 0)
luaL_error (L, "error starting thread:%s",
lua_tostring (L1, -1));
if (pthread_create (& thread, NULL, ll_thread, L1) != 0)
luaL_error (L, "unable to create new thread");
pthread_detach (thread);
return 0;
}

```



```
}
```

To create and launch a new process, the system needs to create a new lua state, start new thread, compile passed block, call it and finally release its resources. Original thread performs the first three tasks and the new thread does the rest. (For to simplify error handling, the system starts a new thread after how she successfully compiled the given block.)

The `ll_start` function creates a new process (Listing 31.4). This the function creates a new Lua L1 state and compiles the given block in this new state. In case of an error, she informs initial states of L . Then she creates a new thread (with `pthread_create`) with the body `ll_thread` , passing the new state to L1 as a body argument. The `pthread_detach` call tells the system that we not expecting a definitive answer from this thread.

Listing 31.5. A body for new threads

```
int luaopen_lproc (lua_State * L);
static void * ll_thread (void * arg) {
    lua_State * L = (lua_State *) arg;
    luaL_openlibs (L); /* open standard libraries */
    luaL_requiref (L, "lproc", luaopen_lproc, 1);
    lua_pop (L, 1);
    if (lua_pcall (L, 0, 0, 0) != 0) /* call main chunk */
        fprintf (stderr, "thread error:%s", lua_tostring (L, -1));
    pthread_cond_destroy (& getself (L) -> cond);
    lua_close (L);
    return NULL;
}
```

The body of each new thread is the `ll_thread` function (the Thing 31.5). It gets its Lua state (created by `ll_start`) with an already compiled block on the stack. A new thread opens the standard Lua libraries, opens the `lproc` library and then calls has its own block. At the end, it releases its conditional variable (which Thoraya was created `luaopen_lproc`) and closes his fortune Lua.

Note the use of `luaL_require` for that- would open `lproc` 1 . This function is somewhat equivalent to `require` , but instead of searching for the bootloader, it uses the specified function (in our this is `luaopen_lproc`) to open the library. Pos-

before calling the opening function `luaL_requiref` registers a result in the `package.loaded` table. If its last parameter is true, then it also registers the library in the corresponding global variable (in our case `lproc`).

The last function in our module, `exit`, is very simple:

```
static int ll_exit (lua_State * L) {  
    pthread_exit (NULL);  
    return 0;  
}
```

Only the main process needs to call this function when it will complete execution in order not to interrupt immediately completing the entire program.

Our final step is to determine the opening functions for the `lproc` module. This function is `luaopen_lproc` (see Thing 31.6) must register module functions, but it also should create and initialize the `Proc` structure of the current process.

As I said earlier, this process definition in Lua is very simple. There are an infinite number of improvements you can make. Here I want to briefly discuss some of them.

The first obvious improvement would be to replace linear search a process waiting on the specified channel. A beautiful alternative would be using a hash table to find the channel and using creation of independent waiting lists for each channel.

Another improvement relates to the efficiency of process creation. Creating a new state in Lua is a very fast operation.

However, opening all standard libraries is no longer so fast, and most processes will likely not need all the standard new libraries. We can avoid the price associated with opening libraries, by pre-registering libraries, as we discussed or in section 15.1. When using this approach, instead of calling `luaL_requiref` for each standard library, we simply

We place the function that opens the library into the `package.preload`. If the process calls `require "lib"`, then and only then `require` will call the appropriate function in order to cover the library. Function `registerlib` (listing 31.7) performs this registration.

Listing 31.6. Opening function for `lproc` module

```
static const struct luaL_reg ll_funcs [] = {  
    {"Start", ll_start},
```

```

{"Send", ll_send},
{"Receive", ll_receive},
{"Exit", ll_exit},
{NULL, NULL}
};
int luaopen_lproc (lua_State * L) {
/* create your own control block */
Proc * self = (Proc *) lua_newuserdata (L, sizeof (Proc));
lua_setfield (L, LUA_REGISTRYINDEX, "_SELF");
self-> L = L;
self-> thread = pthread_self ();
self-> channel = NULL;
pthread_cond_init (& self-> cond, NULL);
luaL_register (L, "lproc", ll_funcs); /* open library */
return 1;
}

```

Listing 31.7. Registration of libraries on request

```

static void registerlib (lua_State * L, const char * name,
lua_CFunction f) {
lua_getglobal (L, "package");
lua_getfield (L, -1, "preload"); /* get 'package.preload' */
lua_pushcfunction (L, f);
lua_setfield (L, -2, name); /* package.preload [name] = f */
lua_pop (L, 2); /* pop 'package' and 'preload' off the stack */
static void openlibs (lua_State * L) {
luaL_requiref (L, "_G", luaopen_base, 1);
luaL_requiref (L, "package", luaopen_package, 1);
lua_pop (L, 2); /* remove the results of previous calls */
registerlib (L, "io", luaopen_io);
registerlib (L, "os", luaopen_os);
registerlib (L, "table", luaopen_table);
registerlib (L, "string", luaopen_string);
registerlib (L, "math", luaopen_math);
registerlib (L, "debug", luaopen_debug);
}

```

It is always a good idea to open up the main library. You also need a package library, otherwise you won't be able to use `require` to load other libraries. (You don't even get the `package.preload` table.) All other libraries are optional. So instead of calling `luaL_openlibs` we will substitute our own `openlibs` function (also shown in Listing 31.7) when creating new states. When the process will need any of these libraries, it will explicitly require it, and `require` will call the corresponding `luaopen_*` function. Other improvements include primitives for communication. For example, it would be helpful to set limits on how long

`lproc.send` and `lproc.receive` may wait. In particular, the limit of expected a denier of zero will make these functions non-blocking. In threads POSIX we can implement this functionality with `pthread_cond_timedwait` .

Exercises

Exercise 31.1. As we have seen, if a function calls `lua_yield` (version without continuation function), control passed to the function that called it when the thread is again will continue its execution. What are the meanings of the caller will the function get how the results of this call?

Exercise 31.2. Modify the `lproc` library so that it could send other basic types such as booleans values and numbers. (*Hint* : you only need to change function `movevalues` .)

Exercise 31.3. In the `lproc` library, implement a non-blocking the `send` function .

Exercises

Chapter 32

Memory management

Lua dynamically allocates all of its data structures. All these structures grow dynamically as needed and over time decrease change their size or disappear.

Lua is strict about its memory usage. When we force we hide the state of Lua, then Lua explicitly frees all its memory. Moreover, all objects inside Lua are subject to garbage collection: not only to tables and rows, but also functions, threads and modules (since they are actually tables).

The way Lua manages memory is comfortable for most applications. However, for some applications you may need - adaptation, for example, to work in a limited space memory space or to reduce garbage collector delays to a minimum. Lua allows such adaptations to be done right away on two levels. At the bottom level, we can define a function to use used to allocate memory. At a higher level, we can let's set some parameters to control the garbage collector or we can even take direct control of the garbage collector. In this chapter, we'll cover both of these options.

32.1. Function to highlight memory

The Lua core does not assume anything about memory allocation. To highlight it does not call `malloc` or `realloc`. Instead of this it performs all its allocation and deallocation of memory through single *out a function* (allocation function), which user must provide when creating Lua state.

The `luaL_newstate` function we used to create Lua states is a helper function that creates a Lua state with a default highlighting function. This function defaults to standard `malloc-realloc-free` from the C standard library, which should be enough for common applications. However, it is very easy to gain control over memory allocation, creating your state using the function

`lua_newstate` :

```
lua_State * luaL_newstate (lua_Alloc f, void * ud);
```

This function takes two arguments: the highlighting function and *the user data* (user data). The state created by this way, performs all allocation and deallocation of memory when using function calls `f`. (Even the structure `lua_State` `vyde-` is done with `f`.)

The type of the `lua_Alloc` allocating function is defined as follows. at once:

```
typedef void * (* lua_Alloc) (void * ud,  
void * ptr,  
size_t osize,  
size_t nsize);
```

The first parameter is the user data that we will provide delivered `lua_newstate`; the second parameter is the address of the block that we want to free or resize it; the third parameter is the original size of this block, and the fourth parameter is the request desired block size.

Lua guarantees that if `ptr` is not `NULL` then it was previously issued linen with `osize` size .

Lua uses `NULL` for zero-sized blocks. When `nsz` is zero, then the function must free the block at the address `ptr` and verify a `NULL` string that corresponds to the requested block size. When `ptr` is `NULL`, the function should allocate and return the block a given size; if she cannot allocate a block of a given time measure, then it must return `NULL`. If `ptr` is both `NULL` and `nsz` is zero, then the function does nothing and returns `NULL`. Finally, when `ptr` is both non-`NULL` and `nsz` is non-zero, the function the tion should re-allocate this block (like `realloc`) and return a new one address (which may be the same as the original address, or may differ from him). Again, in case of an error, the function should return `NULL`. Lua assumes that an allocating function is always successful. works when the new size is less than or equal to the old size. (Lua shrinks some structures during garbage collection and not in able to handle errors correctly at this time.)

Function for allocating memory

The standard highlighting function used by `luaL_newstate` is looks like this (taken from `lauxlib.c` file):

```
void * l_alloc (void * ud, void * ptr, size_t osize, size_t nsz) {
if (nsz == 0) {
free (ptr);
return NULL;
}
else
return realloc (ptr, nsz);
}
```

She believes that `free (NULL)` does nothing and that the call `realloc (NULL, size)` is equivalent to `malloc (size)`. This guarantees- xia ANSI C.

You can get an emitting function for a given state-

lua with `lua_getallocf`:

```
lua_Alloc lua_getallocf (lua_State * L, void ** ud);
```

If `ud` is not `NULL`, then the function will set `*ud` to user values. data used for this emitting function

tion. You can change the allocating function for Lua state by calling `lua_setallocf`:

```
void lua_setallocf (lua_State * L, lua_Alloc f, void * ud);
```

Keep in mind that the new highlighting function must be in able to free blocks allocated by the old function. More often the whole new highlighting function is just a wrapper over the old function, for example, to track secretions or synchronize

heap access control.

Internally, Lua does not cache free blocks for reuse.

education. It assumes that the highlighting function does it,

many good memory allocation functions do this. Lua is not

tries to minimize memory fragmentation. Research on

show fragmentation is more the result of bad design

memory allocation than program behavior; good features for

memory allocations do not create strong fragmentation.

It is quite difficult to make a good highlighting function, but

sometimes you can try it. For example Lua gives you

the old size of any block when it is freed or changed

its size. Accordingly, a specialized emitting function

it is not necessary to store information about the block size somewhere, thus

reducing the amount of memory required for each block.

Another case where you can improve memory allocation is

case of multi-strand systems. Such systems usually require syn-

timing to allocate memory, since they use global

ny resource (memory). However, accessing Lua state must also

be synchronized - or better yet, limited to just one

thread, as in our `lproc` implementation in Chapter 31. So if each

before Lua state will allocate memory from its own memory pool-

ty, you can remove the explicit synchronization requirement.

32.2. Garbage collector

Prior to version 5.0, Lua used a simple *mark-* type garbage *collector*

and-sweep . This garbage collector is sometimes called *collector-stop-*

vi-world . This means that from time to time Lua will stop interpreting

run the main program to complete a complete build cycle

garbage. Each such cycle consists of three phases: *mark* ,

clean and *sweep* .

Lua starts the mark phase by marking it as live.

the root set that includes all objects to which Lua

has direct access: registry and main thread. Any object that

is stored in a living object, is reachable by the program and therefore also

tosses about as if alive. The tagging phase ends when all the

Objects are marked as live.

Before starting the sweeping phase, Lua performs a cleanup phase,

which is related to finalizers and weak tables. First of all, it traverses all objects marked for finalization, looking for unmarked objects. These objects are marked as alive (re-bans) and placed on a separate list for use on finalization stage. Second, Lua bypasses its weak tables and removes from them all elements where either the key or the value itself is not marked.

The sweeping phase bypasses all Lua objects. (In order for this it was possible that Lua kept all created objects with a coherent list ke.) If the object is not marked as alive, then it is deleted. Otherwise Lua unchecks it to prepare for the next loop.

During this phase, Lua also invokes object finalizers, which some were collected during the cleaning phase.

Since version 5.1, Lua uses an incremental collector garbage. This collector follows the same steps as the old one, but for this he does not need to "stop the world." Instead, he worked

Garbage collector

works with the interpreter. Every time the interpreter allocates some memory, the garbage collector executes small step. This means that while the garbage collector is running works, the interpreter can change the visibility of the object. For to ensure that the garbage collector works correctly, some ry operations in the interpreter have special barriers, which rye detect dangerous changes and correct the markings accordingly corresponding objects.